

Squeak by Example

Andrew P. Black Stéphane Ducasse

Oscar Nierstrasz Damien Pollet

with Damien Cassou and Marcus Denker

Version of 2007-09-13

This book is available as a free download from SqueakByExample.org, hosted by the Institute of Computer Science and Applied Mathematics of the University of Bern, Switzerland.

Copyright © 2007 by Andrew P. Black, Stéphane Ducasse, Oscar Nierstrasz and Damien Pollet.

The contents of this book are protected under Creative Commons Attribution-ShareAlike 3.0 Unported license.

You are free:

to Share — to copy, distribute and transmit the work

to Remix — to adapt the work

Under the following conditions:

Attribution. You must attribute the work in the manner specified by the author or licensor (but not in any way that suggests that they endorse you or your use of the work).

Share Alike. If you alter, transform, or build upon this work, you may distribute the resulting work only under the same, similar or a compatible license.

- For any reuse or distribution, you must make clear to others the license terms of this work. The best way to do this is with a link to this web page: creativecommons.org/licenses/by-sa/3.0/
- Any of the above conditions can be waived if you get permission from the copyright holder.
- Nothing in this license impairs or restricts the author's moral rights.



Your fair dealing and other rights are in no way affected by the above. This is a human-readable summary of the Legal Code (the full license): creativecommons.org/licenses/by-sa/3.0/legalcode

Published by Square Bracket Associates, Switzerland. SquareBracketAssociates.org

ISBN 978-3-9523341-0-2

First Edition, September, 2007.

Contents

	Preface	ix
I	Getting Started	
1	A Quick Tour of Squeak	3
1.1	Getting Started	3
1.2	The World Menu	7
1.3	Saving, Quitting and Restarting a Squeak Session	9
1.4	Workspaces and Transcripts	11
1.5	Keyboard shortcuts	13
1.6	SqueakMap	15
1.7	The System Browser	17
1.8	Finding Classes	18
1.9	Finding Methods	20
1.10	Defining a new Method.	22
1.11	Chapter Summary	26
2	A First Application	29
2.1	The Quinto Game	29
2.2	Creating a new class Category	30
2.3	Defining the Class SBECcell	31
2.4	Adding Methods to a Class	32
2.5	Inspecting an Object	34
2.6	Defining the Class SBEGame	36

2.7	Organizing Methods into Protocols	39
2.8	Let's Try Our Code	42
2.9	Saving and Sharing Smalltalk Code	45
2.10	Chapter Summary	49
3	Syntax in a Nutshell	51
3.1	Syntactic Elements	51
3.2	Pseudo-variables	54
3.3	Message Sends	55
3.4	Method Syntax	56
3.5	Block Syntax.	57
3.6	Conditionals and Loops in a Nutshell	59
3.7	Primitives and Pragmas	60
3.8	Chapter Summary	61
4	Understanding Message Syntax	63
4.1	Identifying Messages	63
4.2	Three Kinds of Messages	65
4.3	Message Composition	68
4.4	Hints for Identifying Keyword Messages	74
4.5	Expression Sequences	76
4.6	Cascaded Messages	76
4.7	Chapter Summary	76
II	Developing in Squeak	
5	The Smalltalk Object Model	81
5.1	The Rules of the Model	81
5.2	Everything is an Object	81
5.3	Every object is an instance of a class	82
5.4	Every class has a superclass	89
5.5	Everything Happens by Message Sending	93
5.6	Method lookup follows the inheritance chain	94
5.7	Shared Variables	101

5.8	Chapter Summary	106
6	The Squeak Programming Environment	109
6.1	Overview	110
6.2	The System Browser	111
6.3	Monticello	125
6.4	The Inspector and Explorer	132
6.5	The Debugger	135
6.6	The Process Browser	144
6.7	Finding Methods	145
6.8	Change Sets and the Change Sorter	145
6.9	The File List Browser	149
6.10	In Smalltalk, You Can't Lose Code	151
6.11	Chapter Summary	152
7	SUnit	155
7.1	Introduction	155
7.2	Why Testing is Important	156
7.3	What makes a Good Test?	157
7.4	SUnit by Example	158
7.5	The SUnit Cook Book	162
7.6	The SUnit Framework	164
7.7	Advanced Features of SUnit	166
7.8	The Implementation of SUnit.	168
7.9	Some Advice on Testing	171
7.10	Chapter Summary	172
8	Basic Classes	175
8.1	Object	175
8.2	Numbers	184
8.3	Characters	188
8.4	Strings	189
8.5	Booleans	190
8.6	Chapter Summary	192

9	Collections	193
9.1	Introduction	193
9.2	The Varieties of Collection	195
9.3	Implementations of Collections	198
9.4	Examples of Key Classes	199
9.5	Collection Iterators	209
9.6	Some Hints for using Collections	213
9.7	Chapter Summary	214
10	Streams	217
10.1	Two Sequences of Elements	217
10.2	Streams vs. Collections	218
10.3	Streaming over Collections	219
10.4	Using Streams for File Access	227
10.5	Chapter Summary	230
11	Morphic	231
11.1	The History of Morphic	231
11.2	Manipulating morphs	232
11.3	Composing Morphs	233
11.4	Creating and Drawing Your Own Morphs	234
11.5	Mouse and Keyboard Interaction	236
11.6	Interactors	239
11.7	Drag-and-drop	239
11.8	A Complete Example	241
11.9	More About the Canvas	245
11.10	Chapter Summary	247
III	Advanced Squeak	
12	Classes and Metaclasses	251
12.1	Rules for Classes and Metaclasses	251
12.2	Revisiting the Smalltalk Object Model	252
12.3	Every Class is an Instance of a Metaclass	254

12.4	The Metaclass Hierarchy Parallels the Class Hierarchy . . .	255
12.5	Every Metaclass Inherits from Class and Behavior	257
12.6	Every metaclass is an instance of Metaclass	260
12.7	The Metaclass of Metaclass is an Instance of Metaclass . . .	260
12.8	Chapter Summary	262

IV Appendices

A	Frequently Asked Questions	267
A.1	Getting started	267
A.2	Collections	267
A.3	Browsing the system	268
A.4	Using Monticello and SqueakSource	269
A.5	Tools	270
A.6	Regular Expressions and Parsing	271
	Bibliography	273
	Index	274

Preface

What is Squeak?

Squeak is a modern, open source, fully-featured implementation of the Smalltalk programming language and environment.

Squeak is highly portable — even its virtual machine is written entirely in Smalltalk, making it easy to debug, analyze, and change. Squeak is the vehicle for a wide range of innovative projects from multimedia applications and educational platforms to commercial web development environments.

Who should read this book?

This book presents the various aspects of Squeak, starting with the basics, and proceeding to more advanced topics.

This book will not teach you how to program. The reader should have some familiarity with programming languages. Some background with object-oriented programming would be helpful.

This book will introduce the Squeak programming environment, the language and the associated tools. You will be exposed to common idioms and practices, but the focus is on the technology, not on object-oriented design. Wherever possible, we will show you lots of examples. (We have been inspired by Alec Sharp's excellent book on Smalltalk¹.)

There are numerous other books on Smalltalk freely available on the web but none of these focuses specifically on Squeak. See for example: stephane.ducasse.free.fr/FreeBooks.html

¹Alec Sharp, *Smalltalk by Example*. McGraw-Hill, 1997 (URL: <http://stephane.ducasse.free.fr/FreeBooks/ByExample/>).

A word of advice

Do not be frustrated by parts of Smalltalk that you do not immediately understand. You do not have to know everything! Alan Knight expresses this principle as follows²:

Try not to care. Beginning Smalltalk programmers often have trouble because they think they need to understand all the details of how a thing works before they can use it. This means it takes quite a while before they can master Transcript show: 'Hello World'. One of the great leaps in OO is to be able to answer the question "How does this work?" with "I don't care".

An Open Book

This book is an open book in the following senses:

- The content of this book is released under the Creative Commons Attribution-ShareAlike (by-sa) license. In short, you are allowed to freely share and adapt this book, as long as you respect the conditions of the license available at the following URL: creativecommons.org/licenses/by-sa/3.0/.
- This book just describes the core of Squeak. Ideally we would like to encourage others to contribute chapters on the parts of Squeak that we have not described. If you would like to participate in this effort, please contact us. We would like to see this book grow!

For more details, visit the book's web site, SqueakByExample.org, hosted by the Institute of Computer Science and Applied Mathematics of the University of Bern, Switzerland.

The Squeak community

The Squeak community is friendly and active. Here is a short list of resources that you may find useful:

²www.surfscranton.com/architecture/KnightsPrinciples.htm

- www.squeak.org is the main web site of Squeak. (Do not confuse it with www.squeakland.org which is dedicated to the eToy environment built on top of Squeak but whose audience is elementary school teachers.)
- www.squeaksource.com is the equivalent of SourceForge for Squeak projects.
- wiki.squeak.org/squeak is a wiki with up-to-date information about Squeak.

About mailing-lists. There are a lot of mailing-lists and sometimes they can be just a little bit too active. If you do not want to get flooded by mail but would still like to participate we suggest you to use news.gmane.org or www.nabble.com/Squeak-f14152.html to browse the lists.

You can find the complete list of Squeak mailing-lists at lists.squeakfoundation.org/mailman/listinfo.

- Note that *Squeak-dev* refers to the developers' mailing-list, which can be browsed here:
news.gmane.org/gmane.comp.lang.smalltalk.squeak.general
- *Newbies* refers to a friendly mailing-list for beginners where any question can be asked:
news.gmane.org/gmane.comp.lang.smalltalk.squeak.beginners
(There is so much to learn that we are all beginners in some aspect of Squeak!)

IRC. Have a question that you need answered quickly? Would you like to meet with other squeakers around the world? A great place to participate in longer-term discussions is the IRC channel on the “#squeak” channel at irc.freenode.net. Stop by and say “Hi!”

Other sites. There are several websites supporting the Squeak community today in various ways. Here are some of them:

- people.squeakfoundation.org is the site of SqueakPeople, which is a kind of “advogato.org” for squeakers. It offers articles, diaries and an interesting trust metric system.
- planet.squeak.org is the site of PlanetSqueak which is an RSS aggregator. It is good place to get a flood of squeaky things. This includes the latest blog entries from developers and others who have an interest in Squeak.

- www.frapppr.com/squeak is a site that tracks Squeak users around the world.


Examples and Exercises


We make use of two special conventions in this book.

We have tried to provide as many examples as possible. In particular, there are many examples that show a fragment of code which can be evaluated. We use the symbol \longrightarrow to indicate the result that you obtain when you select an expression and `print it`:

`3 + 4` \longrightarrow `7` *"if you select 3+4 and 'print it', you will see 7"*

In case you want to play in Squeak with these code snippets, you can download a plain text file with all the example code from the book's web site: SqueakByExample.org.

The second convention that we use is to display the icon  to indicate when there is something for you to do:

 *Go ahead and read the next chapter!*

Acknowledgments

We would like to thank Hilaire Fernandes and Serge Stinckwich who allowed us to translate parts of their columns on Smalltalk, and Damien Cas-sou for contributing the chapter on streams. We also thank Tim Rowledge for the use of the Squeak logo, and Frederica Nierstrasz for the original cover art.

We especially thank Lukas Renggli and Orla Greevy for their comments on drafts of the first release.

We thank the University of Bern, Switzerland, for graciously supporting this open-source project and for hosting the web site of this book.


Part I

Getting Started

Chapter 1

A Quick Tour of Squeak

In this chapter we will give you a high-level tour of Squeak to help you get comfortable with the environment. There will be plenty of opportunities to try things out, so it would be a good idea if you have a computer handy when you read this chapter.

We will use this icon:  to mark places in the text where you should try something out in Squeak. In particular, you will fire up Squeak, learn about the different ways of interacting with the system, and discover some of the basic tools. You will also learn how to define a new method, create an object and send it messages.

1.1 Getting Started

Squeak is available as a free download from www.squeak.org. There are three parts that you will need to download, consisting of four files (see Figure 1.1).



Figure 1.1: The Squeak download files.

1. The *virtual machine* (VM) is the only part of the system that is different for each operating system and processor. Pre-compiled virtual

machines are available for all the major computing environments. In Figure 1.1 we see the VM for the Mac is called *Squeak 3.8.15beta1U.app*.

2. The *sources* file contains the source code for all of the parts of Squeak that don't change very frequently. In Figure 1.1 it is called *SqueakV39.sources*. Note that the file *SqueakV39.sources* is only for versions 3.9 and later of Squeak. For earlier versions, use a *sources* file corresponding to the main version *e.g.*, *SqueakV3.sources* for versions of Squeak from 3.0 up to 3.8.
3. The current *system image* is a snapshot of a running Squeak system, frozen in time. It consists of two files: an *.image* file, which contains the state of all of the objects in the system (including classes and methods, since they are objects too), and a *.changes* file, which contains a log of all of the changes to the source code of the system. In Figure 1.1, we see that we have grabbed the *Squeak3.9-final-7067* image and *sources* files.



Download and install Squeak on your computer.

The version of Squeak that we have used for developing the examples in this book is *Squeak-dev*, available from <http://damien.cassou.free.fr/squeak-dev>. We chose this image because it has a wider variety of programming tools pre-installed, and because additional packages can be loaded with a single click. Most of the introductory material in this book will work with any version, so if you already have one installed, you may as well continue to use it. However, if you notice differences between the appearance or behavior of your system and what is described here, do not be surprised. On the other hand, if you are about to download Squeak for the first time, you may as well grab the *Squeak-dev* image.

As you work in Squeak, the image and changes files are modified, so you need to make sure that they are writable. Always keep these two files together. Never edit them directly with a text editor, as Squeak uses them to store the objects you work with and to log the changes you make to the source code. It is a good idea to keep a backup copy of the downloaded image and changes files so you can always start from a fresh image and reload your code.

The *sources* file and the VM can be read-only—they can be shared between different users. All of these files can be placed in the same directory, but it is also possible to put the Virtual Machine and sources file in separate directory where everyone has access to them. Do whatever works best for your style of working and your operating system.

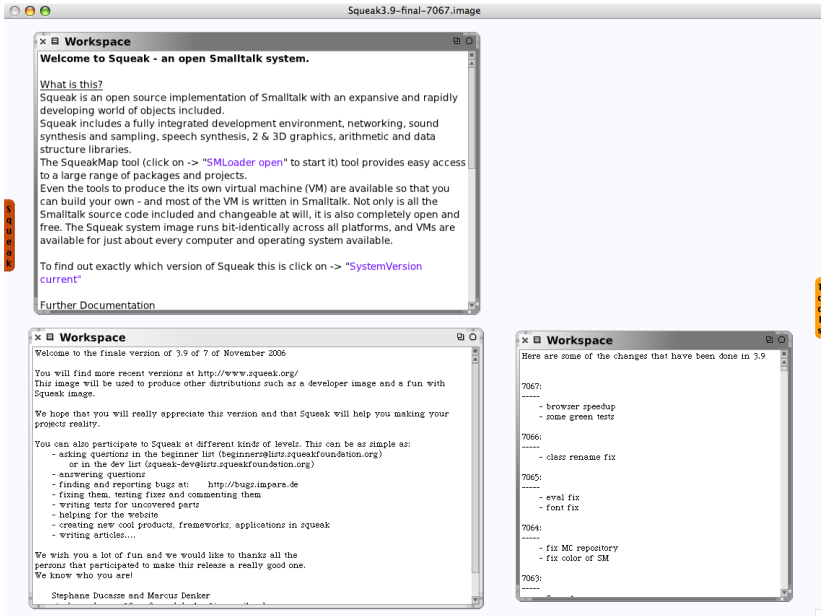





Figure 1.2: A fresh Squeak image.

Launching. To start Squeak, do whatever your operating system expects: drag the *.image* file onto the icon of the virtual machine, or double-click the *.image* file, or at the command line type the name of the virtual machine followed by the path to the *.image* file. (When you have multiple VMs installed on your machine the operating system may not automatically pick the right one; in this case it is safer to drag and drop the image onto the virtual machine, or to use the command line.)

Once Squeak is running, you should see a single large window, possibly containing some open workspace windows (see Figure 1.2), and it's not obvious how to proceed! You will notice that there is no menu bar, or at least not a useful one. Instead, Squeak makes heavy use of context-dependent pop-up menus.

 *Start Squeak.* You can dismiss any open workspaces by clicking on the X in the top left corner of the workspace window. You can collapse the windows (so that they can be expanded again later) by clicking on the  in the top-right corner.

First Interaction. A good place to get started is the world menu shown in Figure 1.3 (a).

 Click with the mouse on the background of the main window to show the world menu, then choose `open ... ▸ workspace` to create a new workspace.

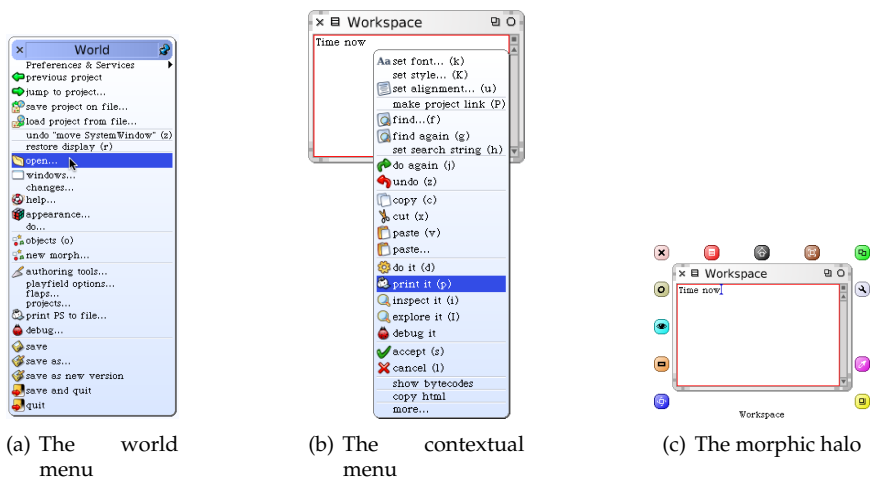


Figure 1.3: The world menu (brought up by the red mouse button), a contextual menu (yellow mouse button), and a morphic halo (blue mouse button).


Squeak was originally designed for a computer with a three button mouse. If your mouse has fewer than three buttons, you will have to press extra keys while clicking the mouse to simulate the extra buttons. A two-button mouse works quite well with Squeak, but if you have only a single-button mouse, you should seriously consider buying a two-button mouse with a clickable scroll wheel: it will make working with Squeak much more pleasant.

Squeak avoids terms like “left mouse click” because different computers, mice, keyboards and personal configurations mean that different users will need to press different physical buttons to achieve the same effect. Instead, the mouse buttons are labeled with colors. The mouse button that you pressed to get the “World” menu is called the *red button*; it is most often used for selecting items in lists, selecting text, and selecting menu items. When you start using Squeak, it can be surprisingly helpful to actually label your mouse, as shown in Figure 1.4.





Figure 1.4: The author’s mouse. Clicking the scroll wheel activates the blue button.

The *yellow button* is the next most used button; it is used to bring up a contextual menu, that is, a menu that offers different sets of actions depending on where the mouse is pointing; see Figure 1.3 (b).


 *Type Time now in the workspace. Now click the yellow button in the workspace. Select print it.*

Finally, there is the *blue button*, which is used to activate the “morphic halo”, an array of handles that are used to perform operations on the on-screen objects themselves, such as rotating them or resizing them; see Figure 1.3 (c). If you let the mouse linger over a handle, a help balloon will explain its function.


 *Click the blue button on the workspace. Grab the  handle near the bottom left corner and drag it to rotate the workspace.*

We recommend that right-handed people configure their mouse to put the red button on the left side of their mouse, the yellow button on the right, and use a clickable scroll wheel, if one is available, for the blue button. If you don’t have a clickable scroll wheel, then you can get the Mophic halo by holding down the alt or option key while clicking the red button. If you are using a Macintosh without a second mouse button, you can simulate one by holding down the ⌘ key while clicking the mouse. However, if you are going to be using Squeak at all often, we recommend investing in a mouse with at least two buttons.

You can configure your mouse to work the way you want by using the preferences of your operating system and mouse driver. Squeak has some preferences for customising the mouse and the meta keys on your keyboard. You can find the preference browser in the `open` item of the `World` menu. In the preference browser, the `general` category contains an option `swapMouseButtons` that switches the yellow and blue functions (see Figure 1.5). The `keyboard` category has options to duplicate the various command keys.

 *Open the preference browser and find the `swapMouseButtons` option using the search box.*

1.2 The World Menu

 *Red-click again on the Squeak background.*

You will see the `World` menu again. Most Squeak menus are not modal; you can leave them on the screen for as long as you wish by clicking the

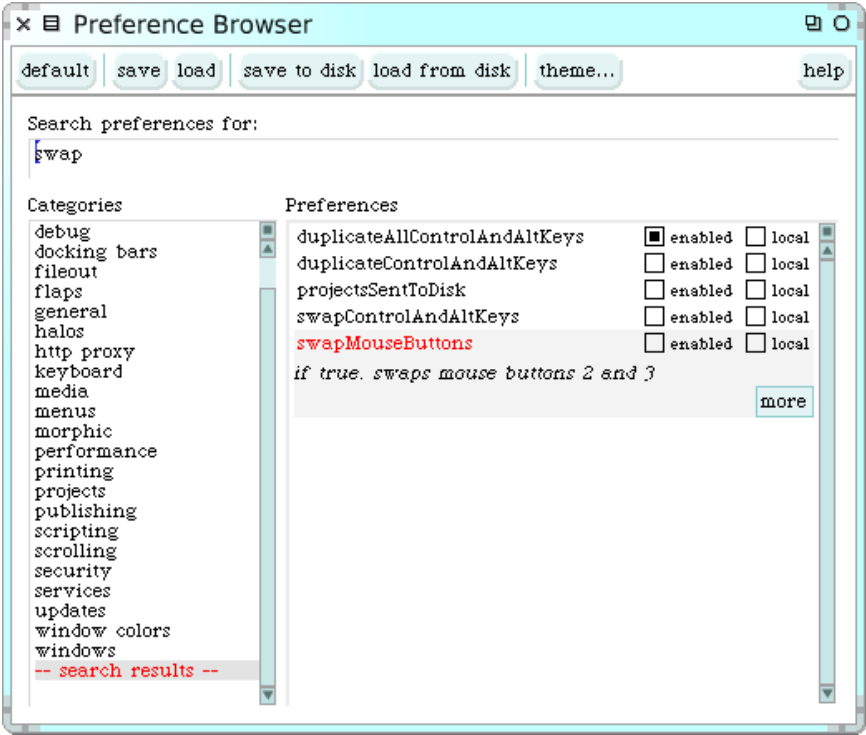



Figure 1.5: The Preference Browser.

push pin icon in the top-right corner. Do this. Also, notice that menus appear when you click the mouse, but do not disappear when you release it; they stay visible until you make a selection, or until you click outside of the menu. You can even move the menu around by grabbing its title bar.

The world menu provides you a simple means to access many of the tools that Squeak offers.

 Have a closer look at the `world>open ...` menu.

You will see a list of several of the core tools in Squeak, including the system browser (one of many available class browsers) and the workspace. We will encounter most of them in the coming chapters.





Figure 1.6: The `open ...` dialogue of the world menu.

1.3 Saving, Quitting and Restarting a Squeak Session

 Bring up the world menu. Now select `new morph ...` and navigate to `from alphabetical list > A-C > BlobMorph`. You now have a blob “in hand”. Put the blob down (by red-clicking) somewhere. It will start to move.







Figure 1.7: An instance of BlobMorph. Figure 1.8: The `save as ...` dialogue.

 Select `World > save as ...`, and enter the name “SBE”. Now click on the `Accept(s)` button. Now select `World > save and quit`.

Now if you go to the location where the original image and changes files were, you will find two new files called “SBE.image” and “SBE.changes” that represent the working state of the Squeak image at the moment before you told Squeak to `save and quit`. If you wish, you can move these two files anywhere that you like on your disk, but if you do so you may (depending on your operating system) need to also move, copy or link to the virtual machine and the `.source` file.

 *Start up Squeak from the newly created “SBE.image” file.*

Now you should find yourself in precisely the state you were when you quit Squeak. The blob is there again and it continues to move from where it was when you left it.

When you start Squeak for the first time, the Squeak virtual machine loads the image file that you provide. This file contains a snapshot of a large number of objects, including a vast amount of pre-existing code and a large number of programming tools (all of which are objects). As you work with Squeak, you will send messages to these objects, you will create new objects, and some of these objects will die and their memory will be reclaimed (*i.e.*, garbage-collected).

When you quit Squeak, you will normally save a snapshot that contains all of your objects. If you save normally, you will overwrite your old image file with the new snapshot. Alternatively you may save the image under a new name, as we just did.

In addition to the *.image* file, there is also a *.changes* file. This file contains a log of all the changes to the source code that you have made using the standard tools. Most of the time you do not need to worry about this file at all. As we shall see, however, the *.changes* file can be very useful for recovering from errors, or replaying lost changes. More about this later!

The image that you have been working with is a descendant of the original Smalltalk-80 image created in the late 1970s. Some of these objects have been around for decades!

You might think that the image is the key mechanism for storing and managing software projects, but you would be wrong. As we shall see very soon, there are much better tools for managing code and sharing software developed by teams. Images are very useful, but you should learn to be very cavalier about creating and throwing away images, since tools like Monticello offer much better ways to manage versions and share code amongst developers.

 *Blue-click on the Blob.*

You will see a collection of colored dots that are collectively called the blob’s morphic halo. Each dot is called a *handle*. Click in the pink handle containing the cross; the blob should go away. (You may have to try several times as the Blob squirms around and tries to get away from the mouse.)

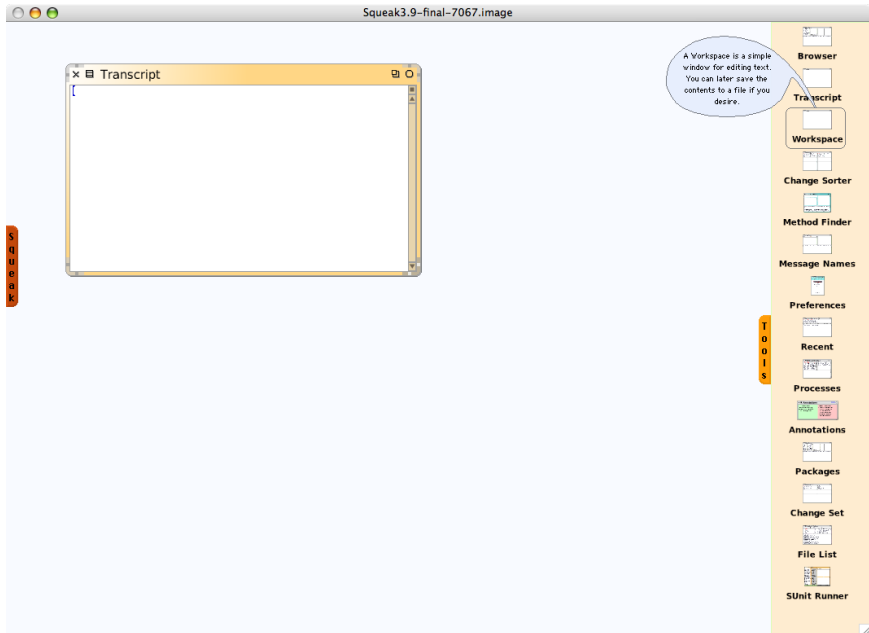




Figure 1.9: The Squeak *Tools* flap.

1.4 Workspaces and Transcripts

 *Close all open windows. Click on the **Tools** tab at the far right of the Squeak environment to open the Tools flap.*

You will see icons for some of the key tools in Squeak (Figure 1.9). Drag out a transcript and a workspace.

 *Position and resize the transcript and workspace windows so that the workspace just overlaps the transcript.*


You can resize windows either by dragging one of the corners, or by blue-clicking the window to bring up the morphic handles, and dragging the yellow, bottom right handle.

At any time only one window is active; it is in front and has its label highlighted. The mouse cursor must be in the window in which you wish to type.

The transcript is an object that is often used for logging system messages. It is a kind of “system console”. Note that the transcript is terribly slow, so

if you keep it open and write to it certain operations can become 10 times slower. In addition the transcript is not thread-safe so you may experience strange problems if multiple objects write concurrently to the transcript.

Workspaces are useful for typing snippets of Smalltalk code that you would like to experiment with. You can also use workspaces simply for typing arbitrarily text that you would like to remember, such as to-do lists or instructions for anyone who will use your image. Workspaces are often used to hold documentation about a captured image, as is the case with the standard image that we downloaded earlier (see Figure 1.2).

 *Type the following text into the workspace:*

Transcript show: 'hello world'; cr.

Try double-clicking in the workspace at various points in the text you have just typed. Notice how an entire word, entire string, or the whole text is selected, depending on where you click.

 *Select the text you have typed and yellow-click. Select **do it (d)**.*

Notice how the text “hello world” appears in the transcript window (Figure 1.10). Do it again. (The **(d)** in the menu item **do it (d)** tells you that the keyboard shortcut to *do it* is CMD–d. More on this in the next section!)

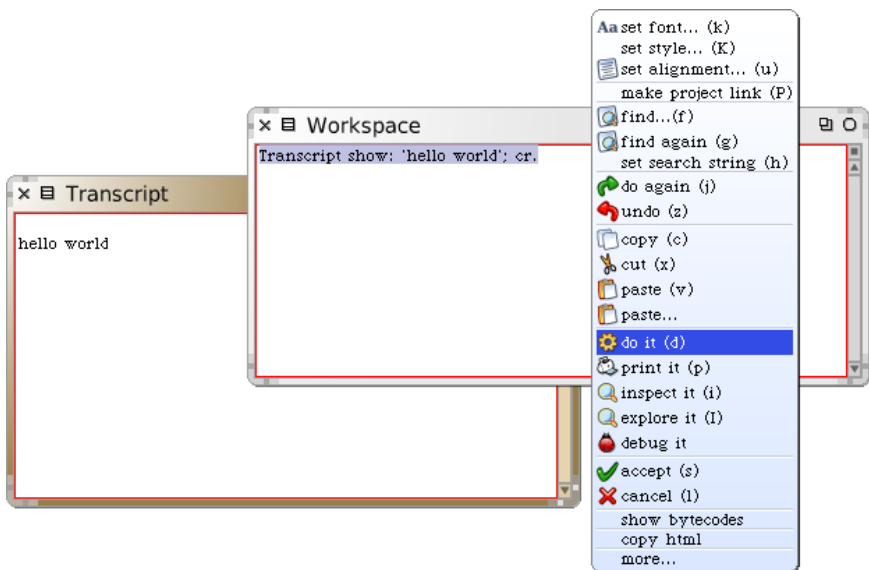



Figure 1.10: “Doing” an expression

You have just evaluated your first Smalltalk expression! You just sent the message `show: 'hello world'` to the Transcript object, followed by the message `cr` (carriage return). The Transcript then decided what to do with this message, that is, it looked up its *methods* for handling `show:` and `cr` messages and reacted appropriately.


If you talk to Smalltalkers for a while, you will quickly notice that they generally do not use expressions like “call an operation” or “invoke a method”, but instead they will say “send a message”. This reflects the idea that objects are responsible for their own actions, so you should never tell an object what to do — instead you ask it politely by sending it a message. The object selects the method for responding to your message, not you!

1.5 Keyboard shortcuts


If you want to evaluate an expression, you do not always have to bring up the yellow-button menu. Instead, you can directly use keyboard shortcuts. These are the parenthesized expressions in the menu. Depending on your platform, you may have to press one of the modifier keys (control, ALT, command, or meta). (We will indicate these generically as `CMD-key`.)

 *Evaluate the expression in the workspace again, but using the keyboard shortcut: `CMD-d`.*

In addition to `do it`, you will have noticed `print it`, `inspect it` and `explore it`. Let's have a quick look at each of these.

 *Type the expression `3 + 4` into the workspace. Now `do it` with the keyboard shortcut.*

Do not be surprised if you saw nothing happen! What you just did is send the message `+` with argument `4` to the number `3`. Normally the result `7` will have been computed and returned to you, but since the workspace did not know what to do with this answer, it simply threw the answer away. If you want to see the result, you should `print it` instead. `print it` actually compiles the expression, executes it, sends the message `printString` to the result, and displays the resulting string.

 *Select `3+4` and `print it` (`CMD-p`).*

This time we see the result we expect (Figure 1.11).

`3 + 4` \longrightarrow `7`

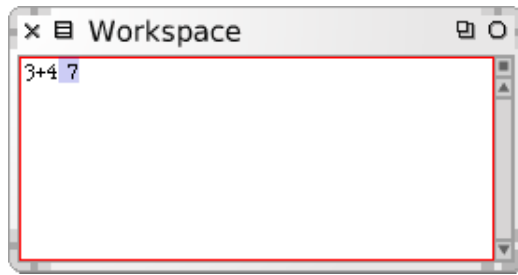



Figure 1.11: “Print it” rather than “do it”.

We use the notation \longrightarrow as a convention in this book to indicate that a particular Squeak expression yields a given result when you **print it**.

 Delete the highlighted text “7” (Squeak should have selected it for you, so you can just press the delete key). Select `3+4` again and this time **inspect it** (CMD-*i*).

Now you should see a new window, called an *inspector*, with the heading `SmallInteger: 7` (Figure 1.12). The inspector is an extremely useful tool that will allow you to browse and interact with any object in the system. The title tells us that 7 is an instance of the class `SmallInteger`. The left panel allows us to browse the instance variables of an object, which are shown in the right panel. The bottom panel can be used to write expressions to send messages to the object.

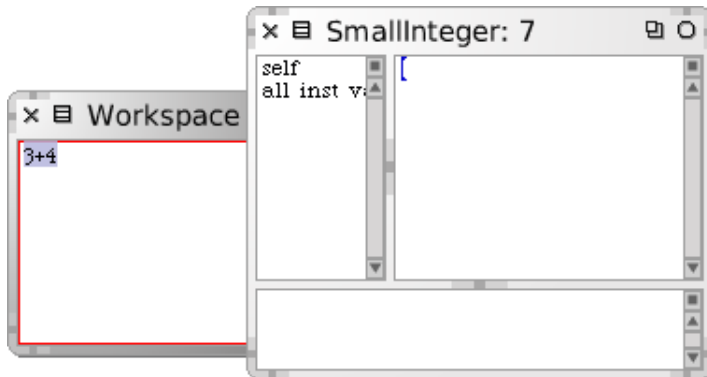




Figure 1.12: Inspecting an object.

 Type `self squared` in the bottom panel of the inspector on 7 and **print it**.

 Close the inspector. Type the expression `Object` in a workspace and this time explore it (CMD-I, uppercased i).

This time you should see a window labelled `Object` containing the text `root: Object`. Click on the triangle to open it up (Figure 1.13).

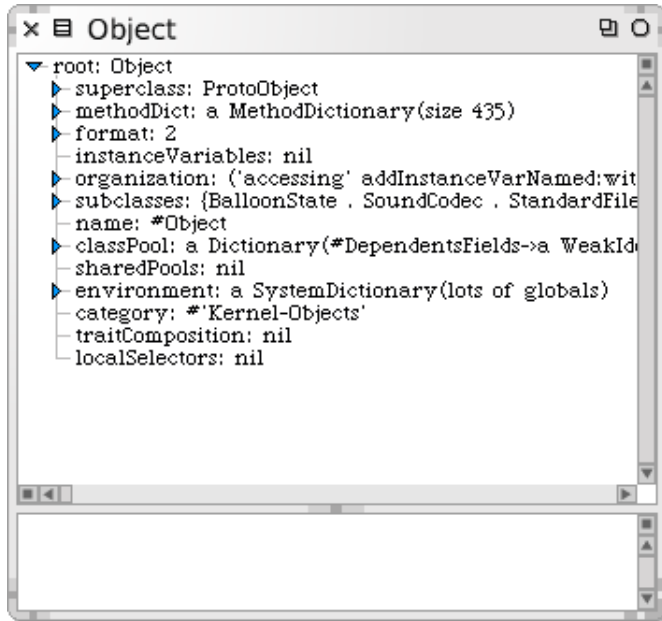


Figure 1.13: Exploring an object.

The explorer is similar to the inspector, but it offers a tree view of a complex object. In this case the object we are looking at is the `Object` class. We can see directly all the information stored in this class, and we can easily navigate to all its parts.

1.6 SqueakMap

SqueakMap is a web-based catalog of “packages” — applications and libraries — that you can download to your image. The packages are hosted in many different places in the world and maintained by many different people. Some of them may only work with a specific version of Squeak.

 Open `World > open... > SqueakMap package loader`.

You will need an Internet connection for this to work. After some time,

the SqueakMap loader window should appear (Figure 1.14). On the left side is a very long list of packages. The field in the top-left corner is a search pane that can be used to find what you want in the list. Type “Sokoban” in the search pane and hit the return key. Clicking on the right-pointing triangle by the name of a package reveals a list of available versions. When a package or a version is selected, information about it appears in the right-hand pane. Navigate to the latest version of Sokoban. With the mouse in the list pane, use the yellow-button menu to **install** the selected package. (If Squeak complains that it is not sure this version of the game will work in your image, just say “yes” and go ahead.) Notice that once a package has been installed, it is marked with an asterisk in the list in the SqueakMap package loader.

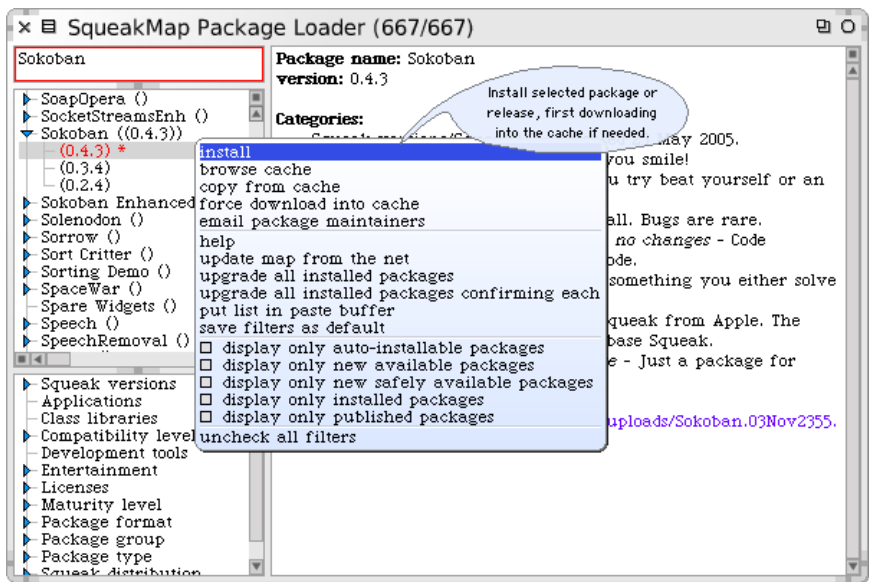




Figure 1.14: Using SqueakMap to install the Sokoban game.

 After installing this package, start up Sokoban by evaluating `SokobanMorph random openInWorld` in a workspace.

The bottom-left pane in the SqueakMap package loader provides various ways to filter the list of packages. You can choose to see only those packages that are compatible with a particular version of Squeak, or only games, and so on.

1.7 The System Browser

The system browser is one of the key tools used for programming. As we shall see, there are several interesting class browsers available for Squeak, but this is the basic one you will find in any image.

 Open a browser by selecting **World > open ... > class browser**, or by dragging a Browser from the Tools flap.

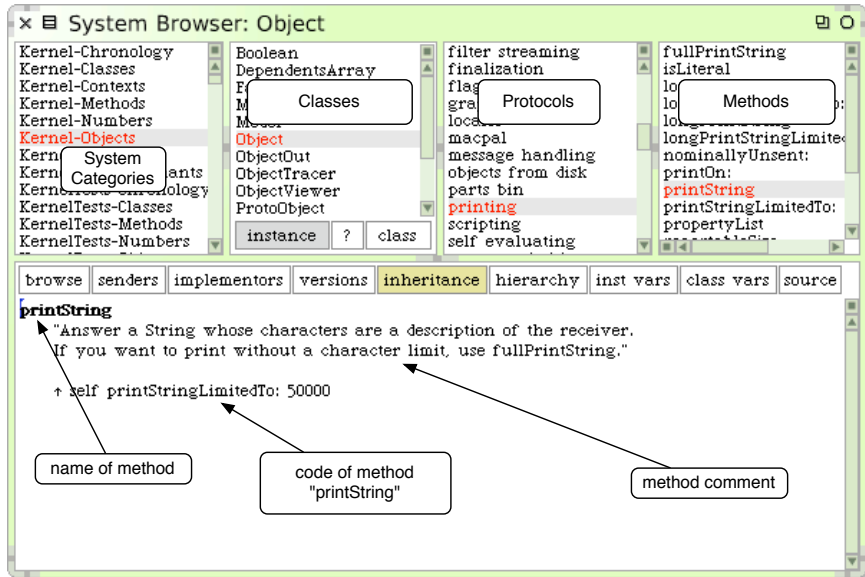



Figure 1.15: The system browser showing the `printString` method of class `Object`.

We can see a system browser in Figure 1.15. The title indicates that we are browsing the class `Object`.


When the browser first opens, all panes are empty but the leftmost one. This first pane lists all known *system categories*, which are groups of related classes.

 Click on the category `Kernel-Objects`.


This causes the second pane to show a list of all of the classes in the selected category.

 Select the class `Object`.

Now the remaining two panes will be filled with text. The third pane displays the *protocols* of the currently selected class. These are convenient groupings of related methods. If no protocol is selected you should see all methods in the fourth pane.

 *Select the printing protocol.*

You may have to scroll down to find it. Now you will see in the fourth pane only methods related to printing.


 *Select the printString method.*

Now we see in the bottom pane the source code of the printString method, shared by all objects in the system (except those that override it).


1.8 Finding Classes

There are several ways to find a class in Squeak. The first, as we have just seen above, is to know (or guess) what category it is in, and to navigate to it using the browser.


A second way is to send the browse message to the class, asking it to open a browser on itself. Suppose we want to browse the class Boolean.

 *Type Boolean browse into a workspace and do it.*

A browser will open on the Boolean class (Figure 1.16). There is also a keyboard shortcut CMD-b (browse) that you can use in any tool where you find a class name.

 *Use the keyboard shortcut to browse the class Boolean.*

Notice that when the Boolean class is selected but no protocol or method is selected, two panes rather than one appear below the four panes at the top (Figure 1.16). The middle pane is the *class definition*. This is nothing more than an ordinary Smalltalk message that is sent to the parent class, asking it to create a subclass. Here we see that the class Object is being asked to create a subclass named Boolean with no instance variables, class variables or “pool dictionaries”, and to put the class Boolean in the *Kernel-Objects* category.

The bottom pane shows the *class comment* — a piece of plain text describing the class. If you click on the  at the bottom of the class pane, you can see the class comment in a dedicated pane.

If you would like to explore Squeak’s inheritance hierarchy, the *hierarchy browser* can help you. This can be useful if you are looking for an unknown

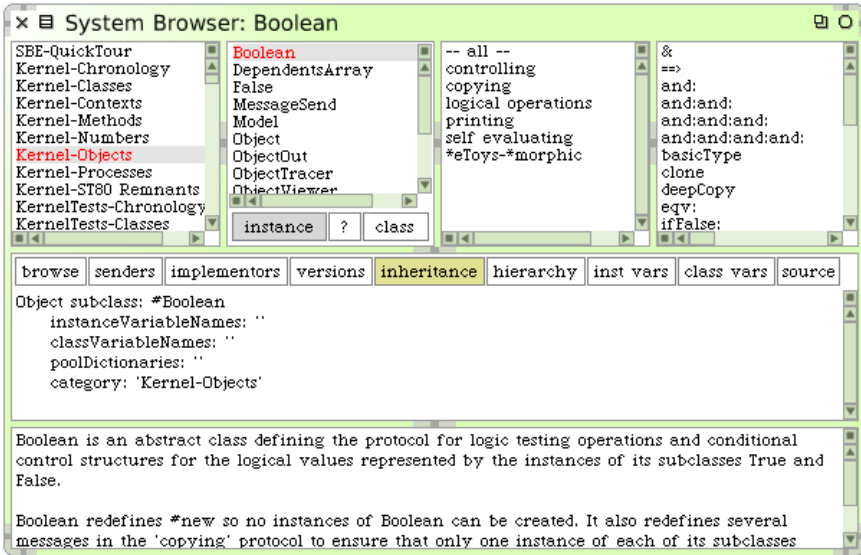




Figure 1.16: The system browser showing the definition of class Boolean.

subclass or superclass of a known class. The hierarchy browser is like the system browser, except that the list of classes is arranged as an indented tree mirroring the inheritance hierarchy.

 Click on **hierarchy** in the browser while the class Boolean is selected.

This will open a hierarchy browser showing the superclasses and subclasses of Boolean. Navigate to the immediate superclass and subclasses of Boolean.

Often, the fastest way to find a class is to search for it by name. For example, suppose that you are looking for some unknown class that represents dates and times.

 Put the mouse in the system category pane of the system browser and type **CMD-f**, or select **find class ... (f)** from the yellow-button menu. Type "time" in the dialog box and accept it.

You will be presented with a list of classes whose names contain "time" (see Figure 1.17). Choose one, say, Time, and the browser will show it, along with a class comment that suggests other classes that might be useful. If you want to browse one of the others, select its name (in any text pane), and type **CMD-b**.

Note that if you type the complete (and correctly capitalized) name of a class in the find dialog, the browser will go directly to that class without

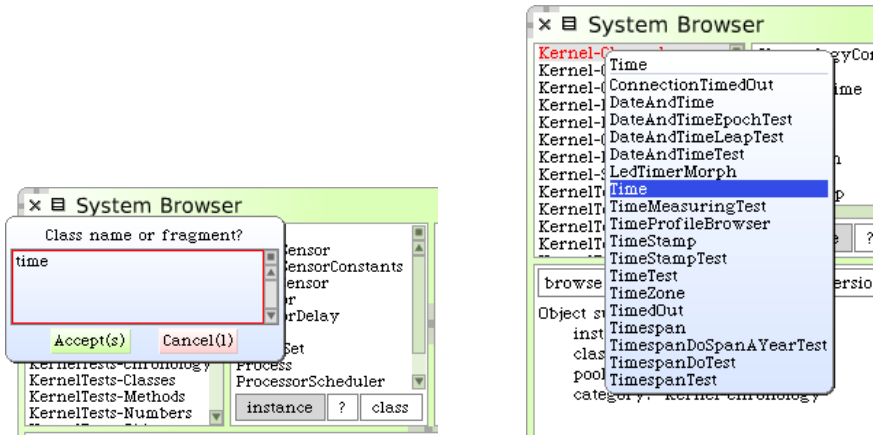



Figure 1.17: Searching for a class by name.

showing you the list of options.


1.9 Finding Methods

Sometimes you can guess the name of a method, or at least part of the name of a method, more easily than the name of a class. For example, if you are interested in the current time, you might expect that there would be a method called “now”, or containing “now” as a substring. But where might it be? The *method finder* can help you.

 Drag the **method finder** icon out of the Tools flap. Type “now” in the top left pane, and **accept** it (or just press the RETURN key).

The method finder will display a list of all the method names that contain the substring “now”. To scroll to now itself, type “n”; this trick works in all scrolling windows. Select “now” and the right-hand pane shows you the three classes that define a method with this name, as shown in Figure 1.18. Selecting any one of them will open a browser on it.

At other times you may have a good idea that a method exists, but will have no idea what it might be called. The method finder can still help! For example, suppose that you would like to find a method that turns a string into upper case, for example, it would translate ‘eureka’ into ‘EUREKA’.

 Type ‘eureka’. ‘EUREKA’ into the method finder, as shown in Figure 1.19.

The method finder will suggest a method that does what you want.

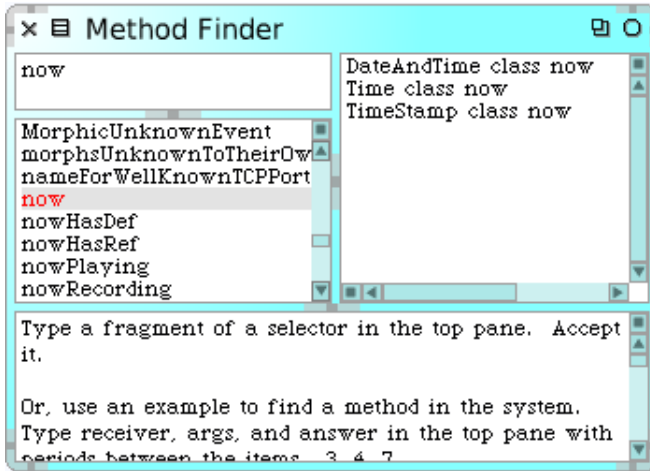


Figure 1.18: The method finder showing three classes that define a method named now.

A star at the beginning of a line in the right pane of the method finder indicates that this method is the one that was actually used to obtain the requested result. So, the star in front of `String asUppercase` lets us know that the method `asUppercase` defined on the class `String` was executed and returned the result we wanted. The methods that do not have a star are just the other methods that have the same name as the ones that returned the expected result. So `Character asUppercase` was not executed on our example, because 'eureka' is not a `Character` object.

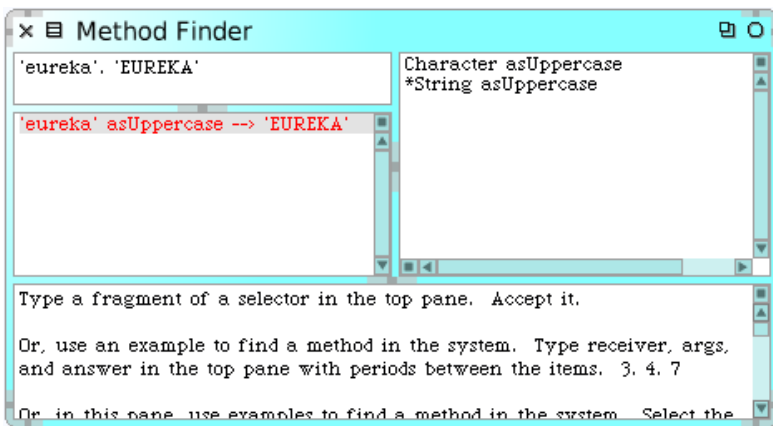


Figure 1.19: Finding a method by example.

You can also use the method finder for methods with arguments; for example, if you are looking for a method that will find the greatest common factor of two integers, you might try 25. 35. 5 as an example. You can also give the method finder multiple examples to narrow the search space; the help text in the bottom pane explains more.

1.10 Defining a new Method

The advent of Test Driven Development¹ has changed the way that we write code. The idea behind Test Driven Development, also called Behavior Driven Development or just TDD, is that we write a test that defines the desired behavior of our code *before* we write the code itself. Only then do we write the code that satisfies the test.

Suppose that our assignment is to write a method that “says something loudly and with emphasis”. What exactly could that mean? What would be a good name for such a method? How can we make sure that programmers who may have to maintain our method in the future have an unambiguous description of what it should do? We can answer all of these questions by giving an example:


When we send the message shout to the string “Don’t panic” the result should be “DON’T PANIC!”.

To make this example into something that the system can use, we turn it into a test method:

Method 1.1: A test for a shout method

```
testShout
  self assert: ('Don't panic' shout = 'DON'T PANIC!')
```

How do we create a new method in Squeak? First, we have to decide which class the method should belong to. In this case, the shout method that we are testing will go in class String, so the corresponding test will, by convention, go in a class called StringTest.

 Open a browser on the class StringTest, and select an appropriate protocol for our method, in this case **tests - converting**, as shown in Figure 1.20. The highlighted text in the bottom pane is a template that reminds you what a Smalltalk method looks like. Delete this and enter the code from method 1.1.

¹Kent Beck, *Test Driven Development: By Example*. Addison-Wesley, 2003, ISBN 0-321-14653-0.

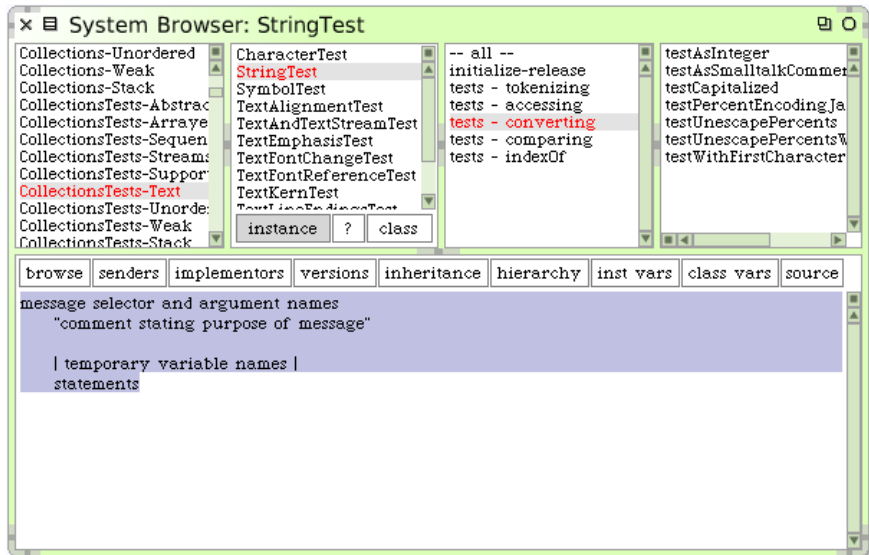




Figure 1.20: The new method template in class StringTest.

Once you have typed the text into the browser, notice that the bottom pane is outlined in red. This is a reminder that the pane contains unsaved changes. So select **accept (s)** from the yellow-button menu in the bottom pane, or just type **CMD-s**, to compile and save your method.

Because there is as yet no method called **shout**, the browser will ask you to confirm that this is the name that you really want—and it will suggest some other names that you might have intended (Figure 1.21). This can be quite useful if you have merely made a typing mistake, but in this case, we really *do* mean **shout**, since that is the method we are about to create, so we have to confirm this by selecting the first option from the menu of choices, as shown in Figure 1.21.

 *Run your newly created test: open the SUnit TestRunner, either by dragging it from the Tools flap, or by selecting **World > open... > Test Runner**.*

The leftmost two panes are a bit like the top panes in the system browser. The left pane contains a list of system categories, but it's restricted to those categories that contain test classes.

 *Select **CollectionsTests-Text** and the pane to the right will show all of the test classes in that category, which includes the class **StringTest**. The names of the classes are already selected, so click **Run Selected** to run all these tests.*

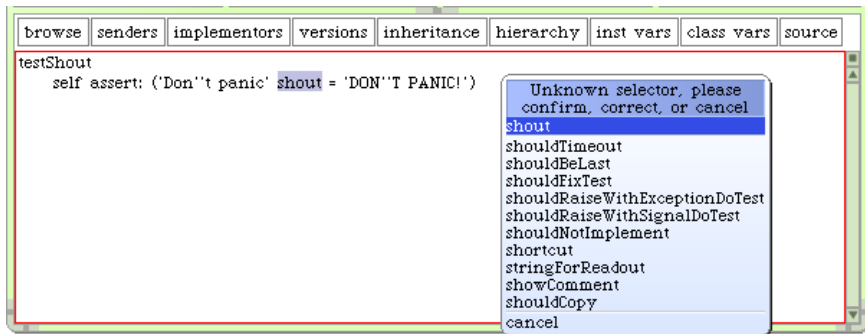


Figure 1.21: Accepting the testShout method class StringTest.

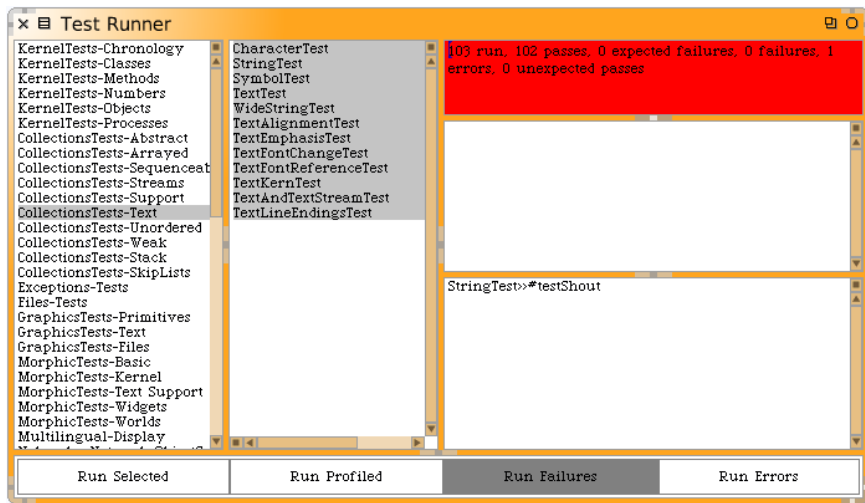


Figure 1.22: Running the String tests.

You should see a message like that shown in Figure 1.22, which indicates that there was an error in running the tests. The list of tests that gave rise to errors is shown in the bottom right pane; as you can see, `StringTest>>#testShout` is the culprit. (Note that `StringTest>>#testShout` is the Smalltalk way of identifying the testShout method of the StringTest class.) If you click on that line of text, the erroneous test will run again, this time in such a way that you see the error happen: "MessageNotUnderstood: ByteString>shout".

The window that opens with the error message is the Smalltalk debugger (see Figure 1.23). We will look at the debugger and how to use it in Chapter 6.

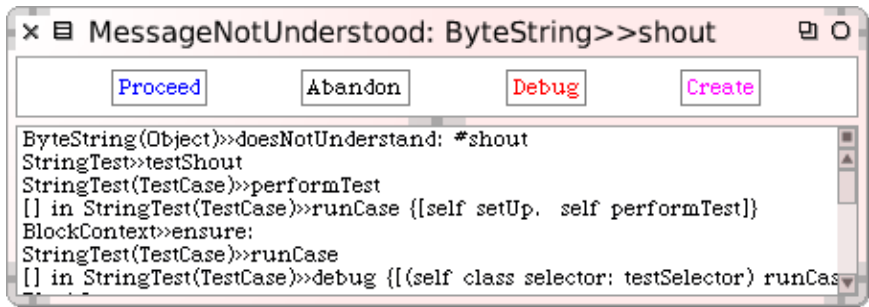



Figure 1.23: The (pre-)debugger.

The error is, of course, exactly what we expected: running the test generates an error because we haven't yet written a method that tells strings how to shout. Nevertheless, it's good practice to make sure that the test fails because this confirms that we have set up the testing machinery correctly and that the new test is actually being run. Once you have seen the error, you can **Abandon** the running test, which will close the debugger window. Note that often with Smalltalk you can define the missing method using the **Create** button, edit the newly-created method in the debugger, and then **Proceed** with the test.

Now let's define the method that will make the test succeed!

 Select class `String` in the system browser, select the `converting` protocol, type the text in method 1.2 over the method creation template, and **accept** it.


(Note: to get a `↑`, type `^`).

Method 1.2: The shout method

```
shout
  ↑ self asUppercase, '!'
```

The comma is the string concatenation operation, so the body of this method appends an exclamation mark to an upper-case version of whatever `String` object the shout message was sent to. The `↑` tells Squeak that the expression that follows is the answer to be returned from the method, in this case the new concatenated string.

Does this method work? Let's run the tests and see.

 Click on **Run Selected** again in the test runner, and this time you should see a green bar and text indicating that all of the tests ran with no failures and no errors.

When you get to a green bar², it's a good idea to save your work and take a break. So do that right now!

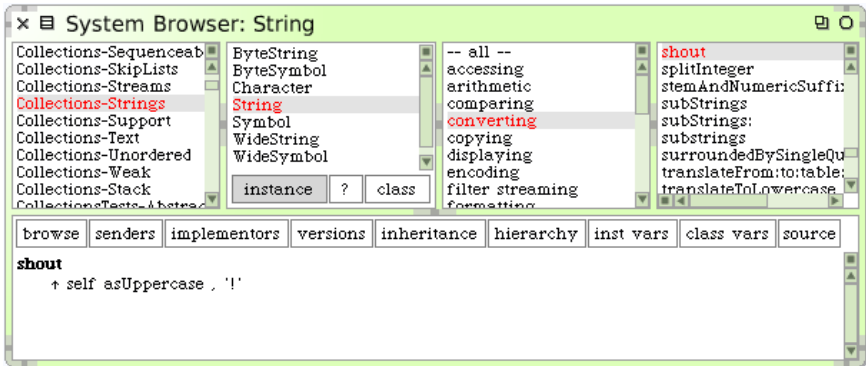


Figure 1.24: The shout method defined on class String.

1.11 Chapter Summary

This chapter has introduced you to the Squeak environment and shown you how to use some of the major tools, such as the system browser, the method finder, and the test runner. You have also seen a little of Squeak's syntax, even though you may not understand it all yet.

- A running Squeak system consists of a *virtual machine*, a *sources* file, and *image* and *changes* files. Only these last two change, as they record a snapshot of the running system.
- When you restore a Squeak image, you will find yourself in exactly the same state — with the same running objects — that you had when you last saved that image.
- Squeak is designed to work with a three-button mouse. The buttons are known as the *red*, the *yellow* and the *blue* buttons. If you don't have a three-button mouse, you can use modifier keys to obtain the same effect.
- Use the red button on the Squeak background to bring up the *World menu* and launch various tools. You can also launch tools from the *Tools* flap at the right of the Squeak screen.

²Actually, you might not get a green bar since some images contains tests for bugs that need to be fixed. Don't worry about this. Squeak is constantly evolving.

- A *workspace* is a tool for writing and evaluating snippets of code. You can also use it to store arbitrary text.
- You can use keyboard shortcuts on text in the workspace, or any other tool, to evaluate code. The most important of these are `do it` (CMD-d), `print it` (CMD-p), `inspect it` (CMD-i) and `explore it` (CMD-I).
- SqueakMap is a tool for loading useful packages from the internet.
- The *system browser* is the main tool for browsing Squeak code, and for developing new code.
- The *test runner* is a tool for running unit tests. It also supports Test Driven Development.

Chapter 2

A First Application

In this chapter, we will develop a simple game: Quinto. Along the way we will demonstrate most of the tools that Squeak programmers use to construct and debug their programs, and show how programs are exchanged with other developers. We will see the system browser, the object inspector, the debugger and the Monticello package browser. Development in Smalltalk is efficient: you will find that you spend far more time actually writing code and far less managing the development process. This is partly because the Smalltalk language is very simple, and partly because the tools that make up the programming environment are very well integrated with the language.

2.1 The Quinto Game




Figure 2.1: The Quinto game board. The user has just clicked the mouse as shown by the cursor.


To show you how to use Squeak's programming tools, we will build a simple game called *Quinto*. The game board is shown in Figure 2.1; it

consists of rectangular array of light yellow *cells*. When you click on one of the cells with the mouse, the four surrounding cells turn blue. Click again, and they toggle back to light yellow. The object of the game is to turn blue as many cells as possible.

The Quinto game shown in Figure 2.1 is made up of two kinds of objects: the game board itself, and 100 individual cell objects. The Squeak code to implement the game will contain two classes: one for the game and one for the cells. We will now show you how to define these classes using the Squeak programming tools.

2.2 Creating a new class Category

We have already seen the system browser in Chapter 1, where we learned how to navigate to classes and methods, and saw how to define new methods. Now we will see how to create system categories and classes.

 Open a system browser and yellow-click in the category pane. Select `add item ...`.

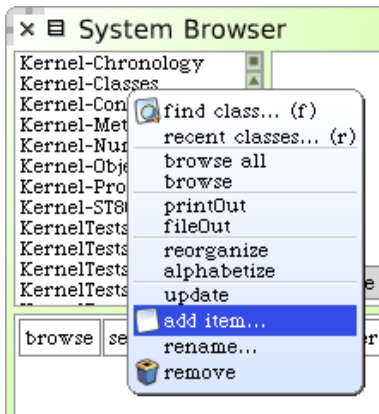


Figure 2.2: Adding a system cate-

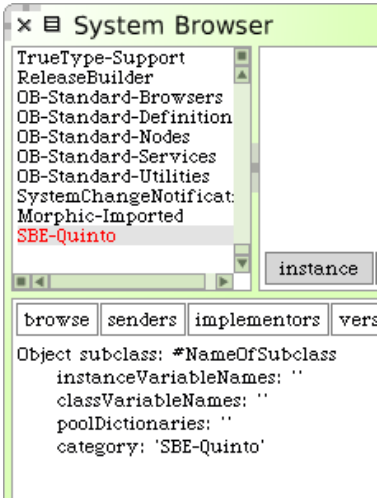


Figure 2.3: The class-creation Tem-

Type the name of the new category (we will use *SBE-Quinto*) in the dialog box and click `accept` (or just press the return key); the new category is created, and positioned at the end of the category list. If you selected an ex-


isting category first, then the new category will be positioned immediately ahead of the selected one.

2.3 Defining the Class SBECeIl

As yet there are of course no classes in the new category. However, the main editing pane displays a template to make it easy to create a new class (see Figure 2.3).

This template shows us Smalltalk expression that sends a message to a class called `Object`, asking it to create a subclass called `NameOfSubClass`. The new class has no variables, and should belong to the category *SBE-Quinto*.

We simply modify the template to create the class that we really want.

 *Modify the class creation template as follows:*

- Replace `Object` by `SimpleSwitchMorph`.
- Replace `NameOfSubClass` by `SBECeIl`.
- Add `mouseAction` to the list of instance variables.


The result should look like class 2.1.

Class 2.1: *Defining the class SBECeIl*

```
SimpleSwitchMorph subclass: #SBECeIl
instanceVariableNames: 'mouseAction'
classVariableNames: ''
poolDictionaries: ''
category: 'SBE-Quinto'
```

This new definition consists of a Smalltalk expression that sends a message to the existing class `SimpleSwitchMorph`, asking it to create a subclass called `SBECeIl`. (Actually, since `SBECeIl` does not exist yet, we passed as argument the *symbol* `#SBECeIl` which stands for the name of the class to create.) We also tell it that instances of the new class should have a `mouseAction` instance variable, which we will use to define what action the cell should take if the mouse should click over it.

At this point you still have not created anything. Note that the border of the class template pane has changed to red (Figure 2.4). This means that there are *unsaved changes*. To actually send this message, you must **accept** it.

 *Accept the new class definition.*

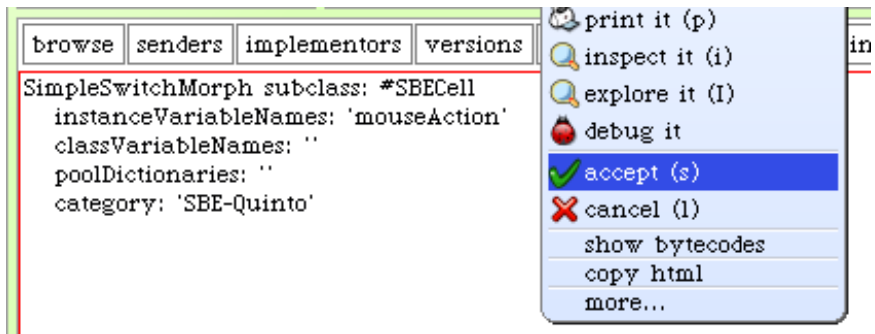



Figure 2.4: The class-creation Template.


Either yellow-click and select `accept`, or use the shortcut CMD-s (for “save”). The message will be sent to `SimpleSwitchMorph`, which will cause the new class to be compiled.

Once the class definition is accepted, the class will be created and appear in the classes pane of the browser (Figure 2.5). The editing pane now shows the class definition, and a small pane below it will remind you to write a few words describing the purpose of the class. This is called a *class comment*, and it is quite important to write one that will give other programmers a high-level overview of the purpose of this class. Smalltalkers put a very high value on the readability of their code, and detailed comments in methods are unusual: the philosophy is that the code should speak for itself. (If it doesn’t, you should refactor it until it does!) A class comment need not contain a detailed description of the class, but a few words describing its overall purpose are vital if programmers who come after you are to know whether to spend time looking at this class.

 *Type a class comment for SBECell and accept it; you can always improve it later.*

2.4 Adding Methods to a Class

Now let’s add some methods to our class.

 *Select the protocol --all-- in the protocol pane.*

You will see a template for method creation in the editing pane. Select it, and replace it by the text of method 2.2.




Figure 2.5: The newly-created class SBECCell


Method 2.2: *Initializing instances of SBECCell*

```

1 initialize
2   super initialize.
3   self label: ".
4   self borderWidth: 2.
5   bounds := 0@0 corner: 16@16.
6   offColor := Color paleYellow.
7   onColor := Color paleBlue darker.
8   self useSquareCorners.
9   self turnOff

```

Note that the characters " on line 3 are two separate single quotes with nothing between them, not a double quote! " denotes the empty string.

 **Accept** *this method definition.*

What does the above code do? We won't go into all of the details here (that's what the rest of the book is for!), but we will give you a quick preview. Let's take it line by line.

Notice that the method is called `initialize`. The name is very significant! By convention, if a class defines a method named `initialize`, it will be called right after the object is created. So, when we evaluate `SBECCell new`, the message `initialize` will be sent automatically to this newly created object. Initialize

methods are used to set up the state of objects, typically to set their instance variables; this is exactly what we are doing here.

The first thing that this method does (line 2) is to execute the initialize method of its superclass, SimpleSwitchMorph. The idea here is that any inherited state will be properly initialized by the initialize method of the superclass. It is always a good idea to initialize inherited state by sending super initialize before doing anything else; we don't know exactly what SimpleSwitchMorph's initialize method will do, and we don't care, but it's a fair bet that it will set up some instance variables to hold reasonable default values, so we had better call it, or we risk starting in an unclean state.

The rest of the method sets up the state of this object. Sending self label: "", for example, sets the label of this object to the empty string.


The expression 0@0 corner: 16@16 probably needs some explanation. 0@0 represents a Point object with x and y coordinates both set to 0. In fact, 0@0 sends the message @ to the number 0 with argument 0. The effect will be that the number 0 will ask the Point class to create a new instance with coordinates (0,0). Now we send this newly created point the message corner: 16@16, which causes it to create a Rectangle with corners 0@0 and 16@16. This newly created rectangle will be assigned to the bounds variable, inherited from the superclass.

Note that the origin of the Squeak screen is the *top left*, and the y coordinate increases *downwards*.


The rest of the method should be self-explanatory. Part of the art of writing good Smalltalk code is to pick good method names so that Smalltalk code can be read like a kind of pidgin English. You should be able to imagine the object talking to itself and saying "Self use square corners!", "Self turn off!".

2.5 Inspecting an Object

You can test the effect of the code you have written by creating a new SBCell object and inspecting it.

 *Open a workspace. Type the expression SBCell new and inspect it.*

The left-hand pane of the inspector shows a list of instance variables; if you select one (try bounds), the value of the instance variable is shown in the right pane. You can also use the inspector to change the value of an instance variable.

 *Change the value of bounds to 0@0 corner: 50@50 and accept it.*




Figure 2.6: The inspector used to examine a SBECeCell object.

The bottom pane of the inspector is a mini-workspace. It’s useful because in this workspace the pseudo-variable `self` is bound to the object being inspected.

 *Type the text `self openInWorld` in the bottom pane and `do it`.*

The cell should appear at the top left-hand corner of the screen, indeed, exactly where its bounds say that it should appear. Blue-click on the cell to bring up the morphic halo. Move the cell with the brown (next to top-right) handle and resize it with the yellow (bottom-right) handle. Notice how the bounds reported by the inspector also change.

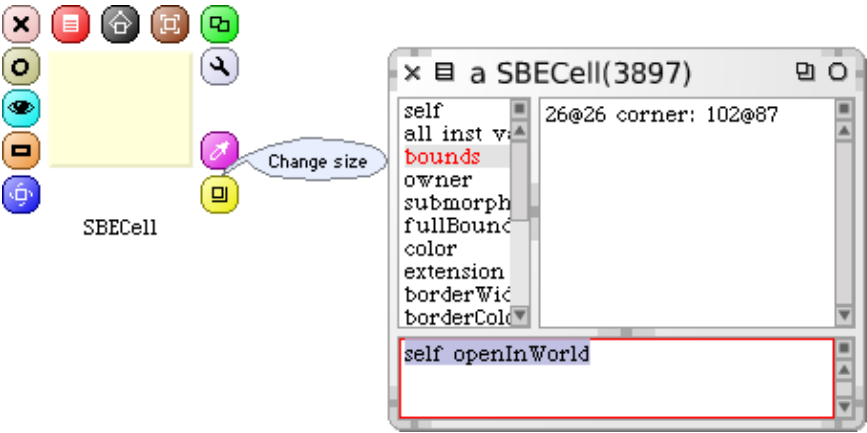



Figure 2.7: Resizing the cell.

 Delete the cell by clicking on the x in the pink handle.

2.6 Defining the Class SBEGame

Now let's create the other class that we need for the game, which we will call SBEGame.

 Make the class definition template visible in the browser main window.


(Do this by clicking twice on the name of the already-selected class category, or by displaying the definition of SBECe ll again (by clicking the `instance` button.) Edit the code so that it reads as follows, and `accept` it.

Class 2.3: Defining the SBEGame class

```
BorderedMorph subclass: #SBEGame
  instanceVariableNames: "
  classVariableNames: "
  poolDictionaries: "
  category: 'SBE-Quinto'
```

Here we subclass BorderedMorph; Morph is the superclass of all of the graphical shapes in Squeak, and (surprise!) a BorderedMorph is a Morph with a border. We could also insert the names of the instance variables between the quotes on the second line, but for now, let's just leave that list empty.

Now let's define an initialize method for SBEGame.

 Type the following into the browser as a method for SBEGame and try to `accept` it:

Method 2.4: Initializing the game

```
1 initialize
2   | sampleCell width height n |
3   super initialize.
4   n := self cellsPerSide.
5   sampleCell := SBECe ll new.
6   width := sampleCell width.
7   height := sampleCell height.
8   self bounds: (5@5 extent: ((width*n) @(height*n)) + (2 * self borderWidth)).
9   cells := Matrix new: n tabulate: [ :i j | self newCe llAt: i at: j ].
```

Squeak will complain that it doesn't know the meaning of some of the terms. Squeak tells you that it doesn't know of a message `cellsPerSide`, and suggests a number of corrections, in case it was a spelling mistake.

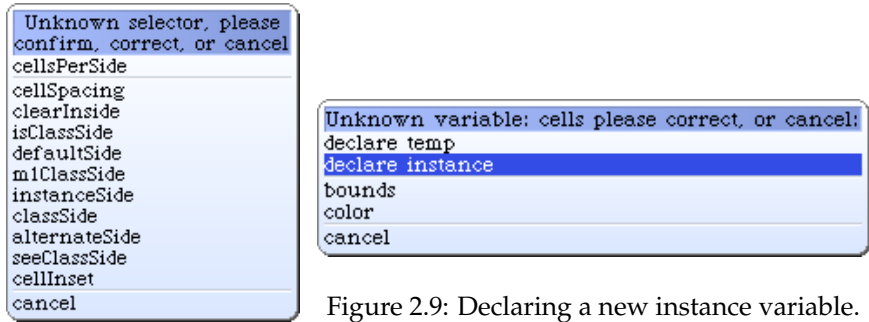




Figure 2.9: Declaring a new instance variable.

Figure 2.8: Squeak detecting an unknown selector.

But `cellsPerSide` is not a mistake — it is just a method that we haven’t yet defined — we will do so in a minute or two.

 So just select the first item from the menu, which confirms that we really meant `cellsPerSide`.

Next, Squeak will complain that it doesn’t know the meaning of `cells`. It offers you a number of ways of fixing this.

 Choose `declare instance` because we want `cells` to be an instance variable.

Finally, Squeak will complain about the message `newCellAt:at:` sent on the last line; this is also not a mistake, so confirm that message too.

If you now look at the class definition once again (which you can do by clicking on the `instance` button), you will see that the browser has modified it to include the instance variable `cells`.

Let’s look at this `initialize` method. The line `| sampleCell width height n |` declares 4 temporary variables. They are called temporary variables because their scope and lifetime are limited to this method. Temporary variables with explanatory names are helpful in making code more readable. Smalltalk has no special syntax to distinguish constants and variables, and in fact all four of these “variables” are really constants. Lines 4–7 define these constants.

How big should our game board be? Big enough to hold some integral number of cells, and big enough to draw a border around them. How many cells is the right number? 5? 10? 100? We don’t know yet, and if we did, we would probably change our minds later. So we delegate the responsibility for knowing that number to another method, which we will call `cellsPerSide`, and which we will write in a minute or two. It’s because we are sending the

cellsPerSide message before we define a method with that name that Squeak asked us to “confirm, correct, or cancel” when we accepted the method body for initialize. Don’t be put off by this: it is actually good practice to write in terms of other methods that we haven’t yet defined. Why? Well, it wasn’t until we started writing the initialize method that that we realized that we needed it, and at that point, we can give it a meaningful name, and move on, without interrupting our flow.

The fourth line uses this method: the Smalltalk self cellsPerSide sends the message cellsPerSide to self, i.e., to this very object. The response, which will be the number of cells per side of the game board, is assigned to n.

The next three lines create a new SBECeIl object, and assign its width and height to the appropriate temporary variables.

Line 8 sets the bounds of the new object. Without worrying too much about the details just yet, just believe us that the expression in parentheses creates a square with its origin (*i.e.*, its top-left corner) at the point (5,5) and its bottom-right corner far enough away to allow space for the right number of cells.

The last line sets the SBEGame object’s instance variable cells to a newly created Matrix with the right number of rows and columns. We do this by sending the message new:tabulate: to the Matrix class (classes are objects too, so we can send them messages). We know that new:tabulate: takes two arguments because it has two colons (:) in its name. The arguments go right after the colons. If you are used to languages that put all of the arguments together inside parentheses, this may seem weird at first. Don’t panic, it’s only syntax! It turns out to be a very good syntax because the name of the method can be used to explain the roles of the arguments. For example, it is pretty clear that Matrix rows: 5 columns: 2 has 5 rows and 2 columns, and not 2 rows and 5 columns.

Matrix new: n tabulate: [:i j | self newCellAt: i at: j] creates a new $n \times n$ matrix and initializes its elements. The initial value of each element will depend on its coordinates. The $(i,j)^{\text{th}}$ element will be initialized to the result of evaluating self newCellAt: i at: j.

That’s initialize. When you accept this message body, you might want to take the opportunity to pretty-up the formatting. You don’t have to do this by hand: from the the yellow-button menu select `more ... ▷ prettyprint`, and the browser will do it for you. You have to `accept` again after you have pretty-printed a method, or of course you can cancel (CMD-1 — that’s a lower-case letter L) if you don’t like the result. Alternatively, you can set up the browser to use the pretty-printer automatically whenever it shows you code: use the the right-most button in the button bar to adjust the view.

If you find yourself using `more ...` a lot, it’s useful to know that you can

hold down the SHIFT key when you click to directly bring up the **more ...** menu.

2.7 Organizing Methods into Protocols

Before we define any more methods, let's take a quick look at the third pane at the top of the browser. In the same way that the first pane of the browser lets us categorize classes so we are not overwhelmed by a very long list of class names in the second pane, so the third pane lets us categorize methods so that we are not overwhelmed by a very long list of method names in the fourth pane. These categories of methods are called "protocols".

If there are only a few methods in a class, the extra level of hierarchy provided by protocols is not really necessary. This is why the browser also offers us the *--all--* virtual protocol, which, you will not be surprised to learn, contains all of the methods in the class.

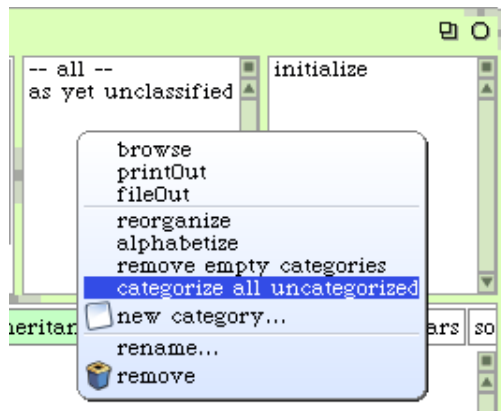



Figure 2.10: Categorize all uncategorized methods.

If you have followed along with this example, the third pane may well contain the protocol *as yet unclassified*.

 Select the yellow button menu item **categorize all uncategorized** to fix this, and move the *initialize* methods to a new protocol called *initialization*.

How does Squeak know that this is the right protocol? Well, in general Squeak can't know, but in this case there is also an *initialize* method in a superclass, and Squeak assumes that our *initialize* method should go in the same category as the one that it overrides.

You may find that Squeak has already put your initialize method into the *initialization* protocol. If so, it's probably because you have loaded a package called AutomaticMethodCategorizer into your image.

A typographic convention. Smalltalkers frequently use the notation ">>" to identify the class to which a method belongs, so, for example, the `cellsPerSide` method in class `SBEGame` would be referred to as `SBEGame >> cellsPerSide`. To indicate that this is *not* Smalltalk syntax, we will use the special symbol » instead, so this method will appear in the text as `SBEGame»cellsPerSide`.

From now on, when we show a method in this book, we will write the name of the method in this form. Of course, when you actually type the code into the browser, you don't have to type the class name or the »; instead, you just make sure that the appropriate class is selected in the class pane.

Now let's define the other two methods that are used by the `SBEGame»initialize` method. Both of them can go in the *initialization* protocol.

Method 2.5: A constant method.

`SBEGame»cellsPerSide`

"The number of cells along each side of the game"

↑ 10

This method could hardly be simpler: it answers the constant 10. One advantage of representing constants as methods is that if the program evolves so that the constant then depends on some other features, the method can be changed to calculate this value.

Method 2.6: An initialization helper method

`SBEGame»newCellAt: i at: j`

"Create a cell for position (i,j) and add it to my on-screen representation at the appropriate screen position. Answer the new cell"

| c origin |

c := SBECell new.

origin := self innerBounds origin.

self addMorph: c.

c position: ((i - 1) * c width) @ ((j - 1) * c height) + origin.

c mouseAction: [self toggleNeighboursOfCellAt: i at: j].

 Add the methods `SBEGame»cellsPerSide` and `SBEGame»newCellAt:at:`.


Confirm the spelling of the new selectors `toggleNeighboursOfCellAt:at:` and `mouseAction:`.

Method 2.6 answers a new SBECcell, specialized to position (i, j) in the Matrix of cells. The last line defines the new cell’s mouseAction to be the block [self toggleNeighboursOfCellAt: i at: j]. In effect, this defines the callback behaviour to perform when the mouse is clicked. The corresponding method also needs to be defined.

Method 2.7: *The callback method*

```
SBEGame>toggleNeighboursOfCellAt: i at: j
(i > 1) ifTrue: [ (cells at: i - 1 at: j ) toggleState].
(i < self cellsPerSide) ifTrue: [ (cells at: i + 1 at: j) toggleState].
(j > 1) ifTrue: [ (cells at: i at: j - 1) toggleState].
(j < self cellsPerSide) ifTrue: [ (cells at: i at: j + 1) toggleState].
```

Method 2.7 toggles the state of the four cells to the north, south, west and east of cell (i, j). The only complication is that the board is finite, so we have to make sure that a neighboring cell exists before we toggle its state.

 Place this method in a new protocol called game logic.

To move the method, you can simply click on its name and drag it to the newly-created protocol (Figure 2.11).

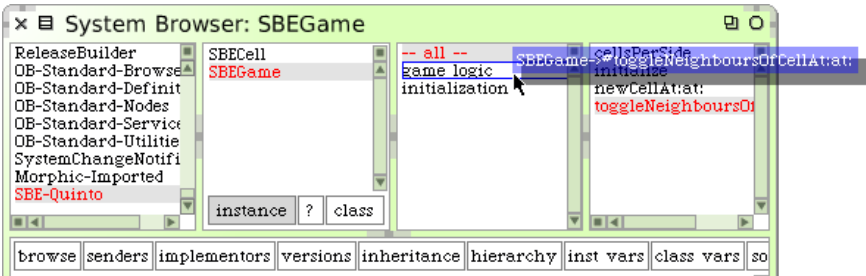


Figure 2.11: Drag a method to a protocol.

To complete the Quinto game, we need to define two more methods in class SBECcell to handle mouse events.

Method 2.8: *A typical setter method*


```
SBECcell>mouseAction: aBlock
↑ mouseAction := aBlock
```

Method 2.8 does nothing more than set the cell’s mouseAction variable to the argument, and then answers the new value. Any method that *changes* the value of an instance variable in this way is called a *setter method*; a

method that *answers* the current value of an instance variable is called a *getter method*.

If you are used to getters and setters in other programming languages, you might expect these methods to be called `setmouseAction` and `getmouseAction`. The Smalltalk convention is different. A getter always has the same name as the variable it gets, and a setter is named similarly, but with a trailing “:”, hence `mouseAction` and `mouseAction:`.


Collectively, setters and getters are called *accessor methods*, and by convention they should be placed in the *accessing* protocol. In Smalltalk, *all* instance variables are private to the object that owns them, so the only way for another object to read or write those variables in the Smalltalk language is through accessor methods like this one¹.

 Go to the class `SBCell`, define `SBCell>mouseAction:` and put it in the accessing protocol.

Finally, we need to define a method `mouseUp:`; this will be called automatically by the GUI framework if the mouse button is released while the mouse is over this cell on the screen.

Method 2.9: An event handler

```
SBCell>mouseUp: anEvent
    mouseAction value
```

 Add the method `SBCell>mouseUp:` and then `categorize all uncategorized methods`.

What this method does is to send the message `value` to the object stored in the instance variable `mouseAction`. Recall that in `SBGame>newCellAt: i at: j` we assigned the following code fragment to `mouseAction`:

```
[self toggleNeighboursOfCellAt: i at: j]
```

Sending the `value` message causes this code fragment to be evaluated, and consequently the state of the cells will toggle.

2.8 Let's Try Our Code

That's it: the Quinto game is complete!

If you have followed all of the steps, you should be able to play the game, consisting of just 2 classes and 7 methods.

¹In fact, the instance variables can be accessed in subclasses too.

 In a workspace, type `SBEGame new openInWorld and do it`.


The game will open, and you should be able to click on the cells and see how it works.

Well, so much for theory... When you click on a cell, a *notifier* window called the `PreDebugWindow` window appears with an error message! As depicted in Figure 2.12, it says `MessageNotUnderstood: SBEGame>>toggleState`.




Figure 2.12: There is a bug in our game when a cell is clicked!

What happened? To find out, let's use one of Smalltalk's more powerful tools: the debugger.

 Click on the `debug` button in the notifier window.

The debugger will appear. In the upper part of the debugger window you can see the execution stack, showing all the active methods; selecting any one of them will show, in the middle pane, the Smalltalk code being executed in that method, with the part that triggered the error highlighted.

 Click on the line labelled `SBEGame>>toggleNeighboursOfCellAt:at:` (near the top).

The debugger will show you the execution context within this method where the error occurred (Figure 2.13).

At the bottom of the debugger are two small inspector windows. On the left, you can inspect the object that is the receiver of the message that caused the selected method to execute, so you can look here to see the values of the instance variables. On the right you can inspect an object that represents the currently executing method itself, so you can look here to see the values of the method's parameters and temporary variables.

Using the debugger, you can execute code step by step, inspect objects in parameters and local variables, evaluate code just as you can in a workspace,





Figure 2.13: The debugger, with the method `toggleNeighboursOfCell:at:` selected.

and, most surprisingly to those used to other debuggers, change the code while it is being debugged! Some Smalltalkers program in the debugger almost all the time, rather than in the browser. The advantage of this is that you see the method that you are writing as it will be executed, with real parameters in the actual execution context.

In this case we can see in the first line of the top panel that the `toggleState` message has been sent to an instance of `SBEGame`, while it should clearly have been an instance of `SBECell`. The problem is most likely with the initialization of the cells matrix. Browsing the code of `SBEGame>>initialize` shows that `cells` is filled with the return values of `newCellAt:at:`, but when we look at that method, we see that there is no return statement there! By default, a method returns `self`, which in the case of `newCellAt:at:` is indeed an instance of `SBEGame`.

 Close the debugger window. Add the expression “`↑ c`” to the end of the method `SBEGame>>newCellAt:at:` so that it returns `c`. (See method 2.10.)

Method 2.10: Fixing the bug.

`SBEGame>>newCellAt: i at: j`

“Create a cell for position (i,j) and add it to my on-screen representation at the appropriate screen position. Answer the new cell”
| c origin |



```

c := SBECCell new.
origin := self innerBounds origin.
self addMorph: c.
c position: ((i - 1) * c width) @ ((j - 1) * c height) + origin.
c mouseAction: [self toggleNeighboursOfCellAt: i at: j].
↑ c

```

Recall from Chapter 1 that the construct to return a value from a method in Smalltalk is `↑`, which you obtain by typing `^`.

Often, you can fix the code directly in the debugger window and click **Proceed** to continue running the application. In our case, because the bug was in the initialization of an object, rather than in the method that failed, the easiest thing to do is to close the debugger window, destroy the running instance of the game (with the halo), and create a new one.


 *Do: SBEGame new openInWorld again.*

Now the game should work properly.


2.9 Saving and Sharing Smalltalk Code

Now that you have the Quinto game working, you probably want to save it somewhere so that you can share it with your friends. Of course, you can save your whole Squeak image, and show off your first program by running it, but your friends probably have their own code in their images, and don't want to give that up to use your image. What you need is a way of getting source code out of your Squeak image so that other programmers can bring it into theirs.

The simplest way of doing this is by *filing out* the code. The yellow-button menu in the System Categories pane will give you the option to file out the whole of category *SBE-Quinto*. The resulting file is more or less human readable, but is really intended for computers, not humans. You can email this file to your friends, and they can file it into their own Squeak images using the file list browser.

 *Yellow-click on the SBE-Quinto category and **fileOut** the contents.*

You should now find a file called "SBE-Quinto.st" in the same folder on disk where your image is saved. Have a look at this file with a text editor.

 *Open a fresh Squeak image and use the File List tool to **file in** the SBE-Quinto.st fileout. Verify that the game now works in the new image.*

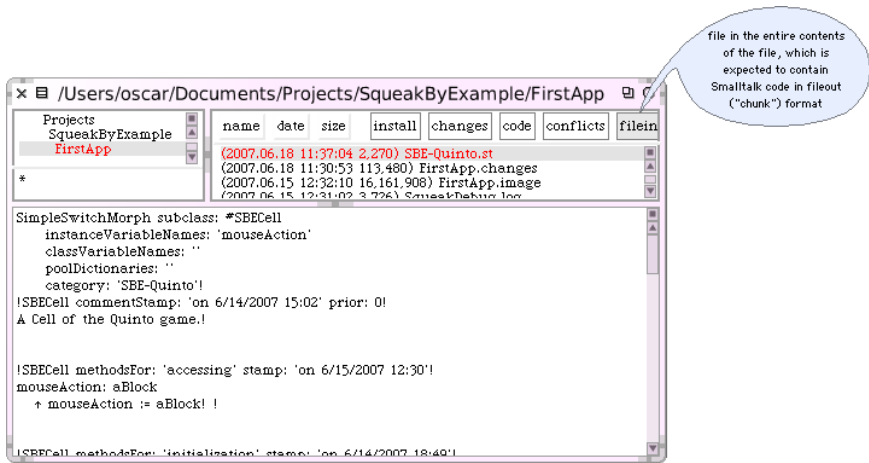


Figure 2.14: Filing in Squeak source code.

Monticello Packages

Although fileouts are a convenient way of making a snapshot of the code you have written, they are decidedly “old school”. Just as most open-source projects find it much more convenient to maintain their code in a repository using CVS² or Subversion³, so Squeak programmers find it more convenient to manage their code using Monticello packages. These packages are represented as files with names ending in `.mcz`; they are actually zip-compressed bundles that contain the complete code of your package.

Using the Monticello package browser, you can save packages to repositories on various types of server, including FTP and HTTP servers; you can also just write the packages to a repository in a local file system directory. A copy of your package is also always cached on your local hard-disk in the *package-cache* folder. Monticello lets you save multiple versions of your program, merge versions, go back to an old version, and browse the differences between versions. In fact, it supports the same sort of operations that you are used to if you share your work using CVS or Subversion.


A good trick is to always develop in the same folder. This way you get a copy of all the code that you publish on squeaksource on your local machine. You can then backup and browse at will.

You can also send a `.mcz` file by email. The recipient will have to place it

²www.nongnu.org/cvs

³subversion.tigris.org

in her *package-cache* folder; she will then be able to use Monticello to browse and load it.

 *Open the Monticello browser by selecting World ▸ open ... ▸ Monticello browser.*

In the right-hand pane of the browser (see Figure 2.15) is a list of Monticello repositories, which will include all of the repositories from which code has been loaded into the image that you are using.




Figure 2.15: The Monticello browser.


At the top of the list in the Monticello browser is a repository in a local directory called the *package cache*, which caches copies of the packages that you have loaded or published over the network. This local cache is really handy because it lets you keep your own local history; it also allows you to work in places where you do not have internet access, or where access is slow enough that you do not want to save to a remote repository very frequently.

Saving and Loading Code with Monticello.

On the left-hand side of the Monticello browser is a list of packages that have a version loaded into the image; packages that have been modified since they were loaded are marked with an asterisk. (These are sometimes referred to as dirty packages.) If you select a package, the list of repositories is restricted to just those repositories that contain a copy of the selected package.


What is a package? For now, you can think of a package as a group of class and method categories that share the same prefix. Since we put all of

the code for the Quinto game into the class category called *SBE-Quinto*, we can refer to it as the SBE-Quinto package.

 Add the SBE-Quinto package to your Monticello browser using the +Package button.


SqueakSource: a SourceForge for Squeak.

We think that the best way to save your code and share it is to create an account for your project on a SqueakSource⁴ server. SqueakSource is like SourceForge⁵: it is a web front-end to a HTTP Monticello server that lets you manage your projects. There is a public SqueakSource server at <http://www.squeaksource.com>, and a copy of the code related to this book is stored there at <http://www.squeaksource.com/SqueakByExample.html>. You can look at this project with a web browser, but it's a lot more productive to do so from inside Squeak, using a special-purpose tool called the Monticello browser, which lets you manage your packages.

 Open a web browser to www.squeaksource.com. Create an account for yourself and then create (i.e., “register”) a project.

SqueakSource will show you the information that you should use when adding a repository using the Monticello browser.

Once your project has been created on SqueakSource, you have to tell your Squeak system to use it.

 With the SBE-Quinto package selected, click the +Repository button in the Monticello browser.

You will see a list of the different types of Repository that are available; to add a SqueakSource repository select HTTP. You will be presented with a dialog in which you can provide the necessary information about the server. You should copy the presented template to identify your SqueakSource project, paste it into Monticello and supply your initials and password:

```

MCHttpRepository
location: 'http://www.squeaksource.com/YourProject'
user: 'yourInitials'
password: 'yourPassword'

```

If you provide empty initials and password strings, you can still load the project, but you will not be able to update it:

⁴www.squeaksource.com

⁵sourceforge.net

MCHttpRepository

location: 'http://www.squeaksource.com/SqueakByExample'

user: "

password: "

Once you have accepted this template, your new repository should be listed on the right-hand side of the Monticello browser.

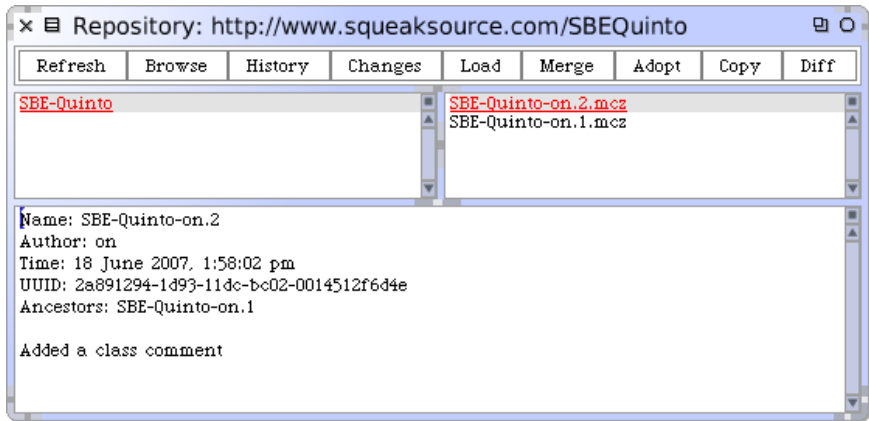



Figure 2.16: Browsing a Monticello Repository

 Click on the **Save** button to save a first version of your Quinto game on SqueakSource.

To load a package into your image, you must first select a particular version. You can do this in the repository browser, which you can open with using the **Open** button or the yellow-button menu. Once you have selected a version, you can load it onto your image.

 Open the SBE-Quinto repository you have just saved.

Monticello has many more capabilities, which will be discussed in depth in Chapter 6. You can also look at the on-line documentation for Monticello at <http://www.wiresong.ca/Monticello/>.

2.10 Chapter Summary

In this chapter you have seen how to create categories, classes and methods. You have seen how to use the system browser, the inspector, the debugger and the Monticello browser.

- Categories are groups of related classes.
- A new class is created by sending a message to its superclass.
- Protocols are groups of related methods.
- A new method is created or modified by editing its definition in the browser and then *accepting* the changes.
- The inspector offers a simple, general-purpose GUI for inspecting and interacting with arbitrary objects.
- The system browser detects usage of undeclared methods and variables, and offers possible corrections.
- The initialize method is automatically executed after an object is created in Squeak. You can put any initialization code there.
- The debugger provides a high-level GUI to inspect and modify the state of a running program.
- You can share source code *filing out* a category.
- A better way to share code is to use Monticello to manage an external repository, for example defined as a SqueakSource project.

Chapter 3

Syntax in a Nutshell

Squeak, like most modern Smalltalk dialects, adopts a syntax very close to that of Smalltalk-80. The syntax is designed so that program text can be read aloud as though it were a kind of pidgin English:

```
(Smalltalk includes: Class) ifTrue: [ Transcript show: Class superclass ]
```

Squeak's syntax is minimal. Essentially there is syntax only for *sending messages* (i.e., expressions) and *declaring methods*. Expressions are built up from a very small number of primitive elements. There are only 6 keywords, and there is no syntax for control structures or declaring new classes. Instead, nearly everything is achieved by sending messages to objects. For instance, instead of an if-then-else control structure, Smalltalk sends messages like `ifTrue:` to Boolean objects. New (sub-)classes are created by sending a message to their superclass.

3.1 Syntactic Elements

Expressions are composed of the following building blocks: (i) six reserved keywords, or *pseudo-variables*: `self`, `super`, `nil`, `true`, `false`, and `thisContext`, (ii) constant expressions for *literal objects* including numbers, characters, strings, symbols and arrays, (iii) variable declarations, (iv) assignments, (v) block closures, and (vi) messages.

We can see examples of the various syntactic elements in Table 3.1.

Local variables `startPoint` is a variable name, or identifier. By convention, identifiers are composed of words in “camelCase” (i.e., each word except the first starting with an upper case letter). The first letter of an

Syntax	What it represents
startPoint	a variable name
Transcript	a global variable name
self	pseudo-variable
1	decimal integer
2r101	binary integer
1.5	floating point number
2.4e7	exponential notation
\$a	the character 'a'
'Hello'	the string "Hello"
#Hello	the symbol #Hello
#{1 2 3}	a literal array
{1. 2. 1+2}	a dynamic array
"a comment"	a comment
x y	declaration of variables x and y
x := 1	assign 1 to x
[x + y]	a block that evaluates to x+y
<primitive: 1>	virtual machine primitive or annotation
3 factorial	unary message
3+4	binary messages
2 raisedTo: 6 modulo: 10	keyword message
↑ true	return the value true
Transcript show: 'hello'. Transcript cr	expression separator (.)
Transcript show: 'hello'; cr	message cascade (;)

Table 3.1: Squeak Syntax in a Nutshell

instance variable, method or block argument, or temporary variable must be lower case. This indicates to the reader that the variable has a private scope.

Shared variables Identifiers that start with upper case letters are global variables, class variables, pool dictionaries or class names. Transcript is a global variable, an instance of the class TranscriptStream.

The receiver. self is a keyword that refers to the object inside which the current method is executing. We call it “the receiver” because this object will normally have received the message that caused the method to execute. self is called a “pseudo-variable” since we cannot assign to it.

Integers. In addition to ordinary decimal integers like 42, Squeak also

provides a radix notation. 2r101 is 101 in radix 2 (*i.e.*, binary), which is equal to decimal 5.

Floating point numbers can be specified with their base-ten exponent: 2.4e7 is 2.4×10^7 .

Characters. A dollar sign introduces a literal character: \$a is the literal for 'a'. Instances of non-printing characters can be obtained by sending appropriately named messages to the Character class, such as Character space and Character tab.

Strings. Single quotes are used to define a literal string. If you want a string with a quote inside, just double the quote, as in 'G"day'.

Symbols are like Strings, in that they contain a sequence of characters. However, unlike a string, a literal symbol is guaranteed to be globally unique. There is only one Symbol object #Hello but there may be multiple String objects with the value 'Hello'.

Compile-time arrays are defined by #(), surrounding space-separated literals. Everything within the parentheses must be a compile-time constant. For example, #(27 #(true false) abc) is a literal array of three elements: the integer 27, the compile-time array containing the two booleans, and the symbol #abc.

Run-time arrays. Curly braces {} define a (dynamic) array at run-time. Elements are expressions separated by periods. So { 1. 2. 1+2 } defines an array with elements 1, 2, and the result of evaluating 1+2. (The curly-brace notation is peculiar to the Squeak dialect of Smalltalk! In other Smalltalks you must build up dynamic arrays explicitly.)

Comments are enclosed in double quotes. "hello" is a comment, not a string, and is ignored by the Squeak compiler. Comments may span multiple lines.

Local variable definitions. Vertical bars || enclose the declaration of one or more local variables in a method (and also in a block).

Assignment. := assigns an object to a variable. Sometimes you will see ← used instead. Unfortunately, since this is not an ASCII character, it will appear as an underscore unless you are using a special font. So, x := 1 is the same as x ← 1 or x _ 1. You should use := since the other representations have been deprecated since Squeak 3.9.

Blocks. Square brackets [] define a block, also known as a block closure or a lexical closure, which is a first-class object representing a function. As we shall see, blocks may take arguments and can have local variables.

Primitives. <primitive: ...> denotes an invocation of a virtual machine primitive. (<primitive: 1> is the VM primitive for `SmallInteger`»+.) Any code following the primitive is executed only if the primitive fails. Since Squeak 3.9, the same syntax is also used for method annotations.

Unary messages consist of a single word (like `factorial`) sent to a receiver (like 3).

Binary messages are operators (like `+`) sent to a receiver and taking a single argument. In `3+4`, the receiver is 3 and the argument is 4.

Keyword messages consist of multiple keywords (like `raisedTo: modulo:`), each ending with a colon and taking a single argument. In the expression `2 raisedTo: 6 modulo: 10`, the *message selector* `raisedTo:modulo:` takes the two arguments 6 and 10, one following each colon. We send the message to the receiver 2.

Method return. `↑` is used to *return* a value from a method. (You must type `^` to obtain the `↑` character.)

Sequences of statements. A period or full-stop (`.`) is the *statement separator*. Putting a period between two expressions turns them into independent statements.

Cascades. Semicolons can be used to send a *cascade* of messages to a single receiver. In `Transcript show: 'hello';` or we first send the keyword message `show: 'hello'` to the receiver `Transcript`, and then we send the same receiver the unary message `cr`.

The classes `Number`, `Character`, `String` and `Boolean` are described in more detail in Chapter 8.

3.2 Pseudo-variables

In Smalltalk, there are 6 reserved keywords, or *pseudo-variables*: `nil`, `true`, `false`, `self`, `super`, and `thisContext`. They are called pseudo-variables because they are predefined and cannot be assigned to. `true`, `false`, and `nil` are constants while the values of `self`, `super`, and `thisContext` vary dynamically as code is executed.

`true` and `false` are the unique instances of the Boolean classes `True` and `False`. See Chapter 8 for more details.

`self` always refers to the receiver of the currently executing method. `super` also refers to the receiver of the current method, but when you send a message to `super`, the method-lookup changes so that it starts from the

superclass of the class containing the method that uses super. For further details see Chapter 5.

nil is the undefined object. It is the unique instance of the class UndefinedObject. Instance variables (and Class variables) are initialized to nil.

thisContext is a pseudo-variable that represents the top frame of the run-time stack. In other words, it represents the currently executing MethodContext or BlockContext. thisContext is normally not of interest to most programmers, but it is essential for implementing development tools like the debugger and it is also used to implement exception handling and continuations.

3.3 Message Sends

There are three kinds of messages in Squeak.

1. *Unary* messages take no argument. 1 factorial sends the message factorial to the object 1.
2. *Binary* messages take exactly one argument. 1 + 2 sends the message + with argument 2 to the object 1.
3. *Keyword* messages take an arbitrary number of arguments. 2 raisedTo: 6 modulo: 10 sends the message consisting of the message selector raisedTo:modulo: and the arguments 6 and 10 to the object 2.

Unary message selectors consist of alphanumeric characters, and start with a lower case letter.

Binary message selectors consist of one or more characters from the following set:

+ - / \ * ~ < > = @ % | & ? ,

Keyword message selectors consist of a series of alphanumeric keywords, where each keyword starts with a lower-case letter and ends with a colon.

Unary messages have the highest precedence, then binary messages, and finally keyword messages, so:

2 raisedTo: 1 + 3 factorial \longrightarrow 128

(First we send `factorial` to 3, then we send `+ 6` to 1, and finally we send `raisedTo: 7` to 2.) Recall that we use the notation *expression* \longrightarrow *result* to show the result of evaluating an expression.

Precedence aside, evaluation is strictly from left to right, so

$$1 + 2 * 3 \longrightarrow 9$$

not 7. Parentheses must be used to alter the order of evaluation:

$$1 + (2 * 3) \longrightarrow 7$$

Message sends may be composed with periods and semi-colons. A period-separated sequence of expressions causes each expression in the series to be evaluated as a *statement*, one after the other.

```
Transcript cr.  
Transcript show: 'hello world'.  
Transcript cr
```

This will send `cr` to the `Transcript` object, then send it `show: 'hello world'`, and finally send it another `cr`.

When a series of messages is being sent to the *same* receiver, then this can be expressed more succinctly as a *cascade*. The receiver is specified just once, and the sequence of messages is separated by semi-colons:

```
Transcript cr;  
  show: 'hello world';  
  cr
```

This has precisely the same effect as the previous example.

3.4 Method Syntax

Whereas expressions may be evaluated anywhere in Squeak (for example, in a workspace, in a debugger, or in a browser), methods are normally defined in a browser window, or in the debugger. (Methods can also be filed in from an external medium, but this is not the usual way to program in Squeak.)

Programs are developed one method at a time, in the context of a given class. (A class is defined by sending a message to an existing class, asking it to create a subclass, so there is no special syntax required for defining classes.)

Here is the method `lineCount` in the class `String`. (The usual convention is to refer to methods as `ClassName»methodName`, so we call this method `String»lineCount`.)

Method 3.1: *Line count*

```
String»lineCount
  "Answer the number of lines represented by the receiver,
  where every cr adds one line."
  | cr count |
  cr := Character cr.
  count := 1 min: self size.
  self do:
    [:c | c == cr ifTrue: [count := count + 1]].
  ↑ count
```

Syntactically, a method consists of:

1. the method pattern, containing the name (*i.e.*, `lineCount`) and any arguments (none in this example);
2. comments (these may occur anywhere, but the convention is to put one at the top that explains what the method does);
3. declarations of local variables (*i.e.*, `cr` and `count`); and
4. any number of expressions separated by dots; here there are four.

The evaluation of any expression preceded by a `↑` (typed as `^`) will cause the method to exit at that point, returning the value of that expression. A method that terminates without explicitly returning some expression will implicitly return `self`.

Arguments and local variables should always start with lower case letters. Names starting with upper-case letters are assumed to be global variables. Class names, like `Character`, for example, are simply global variables referring to the object representing that class.

3.5 Block Syntax

Blocks provide a mechanism to defer the evaluation of expressions. A block is essentially an anonymous function. A block is evaluated by sending it the message value. The block answers the value of the last expression in its body, unless there is an explicit return (with `↑`), in which case it does not answer any value.

```
[ 1 + 2 ] value  →  3
```

Blocks may take parameters, each of which is declared with a leading colon. A vertical bar separates the parameter declaration(s) from the body of the block. To evaluate a block with one parameter, you must send it the message `value:` with one argument. A two-parameter block must be sent `value:value:`, and so on, up to 4 arguments.

```
[ :x | 1 + x ] value: 2  →  3
[ :x :y | x + y ] value: 1 value: 2  →  3
```

If you have a block with more than four parameters, you must use `valueWithArguments:` and pass the arguments in an array. (A block with a large number of parameters is often a sign of a design problem.)

Blocks may also declare local variables, which are surrounded by vertical bars, just like local variable declarations in a method. Locals are declared after any arguments:

```
[ :x :y || z | z := x + y. z ] value: 1 value: 2  →  3
```

Blocks are actually lexical *closures*, since they can refer to variables of the surrounding environment. The following block refers the variable `x` of its enclosing environment:

```
| x |
x := 1.
[ :y | x + y ] value: 2  →  3
```

Blocks are instances of the class `BlockContext`. This means that they are objects, so they can be assigned to variables and passed as arguments just like any other object.

Caveat: In the current version (3.9), Squeak does not actually support true block-closures, since block arguments are actually simulated as temporary variables of the enclosing method. There is a new compiler that supports full block closures, but it is still being worked on and is not used by default. In somewhat obscure situations this problem can lead to naming conflicts. This situation arises because Squeak is based on an early implementation of Smalltalk. If you encounter this problem, have a look at the senders of the method `fixTemps`, or load the *Closure Compiler*.

3.6 Conditionals and Loops in a Nutshell

Smalltalk offers no special syntax for control constructs. Instead, these are typically expressed by sending messages to booleans, numbers and collections, with blocks as arguments.

Conditionals are expressed by sending one of the messages `ifTrue:`, `ifFalse:` or `ifTrue:ifFalse:` to the result of a boolean expression. See Chapter 8 for more about booleans.

```
(17 * 13 > 220)
  ifTrue: [ 'bigger' ]
  ifFalse: [ 'smaller' ]  →  'bigger'
```

Loops are typically expressed by sending messages to blocks, integers or collections. Since the exit condition for a loop may be repeatedly evaluated, it should be a block rather than a boolean value. Here is an example of a very procedural loop:

```
n := 1.
[ n < 1000 ] whileTrue: [ n := n*2 ].
n  →  1024
```

`whileFalse:` reverses the exit condition.

```
n := 1.
[ n > 1000 ] whileFalse: [ n := n*2 ].
n  →  1024
```

`timesRepeat:` offers a simple way to implement a fixed iteration:

```
n := 1.
10 timesRepeat: [ n := n*2 ].
n  →  1024
```

We can also send the message `to:do:` to an number which then acts as the initial value of a loop counter. The two arguments are the upper bound, and a block that takes the current value of the loop counter as its argument:

```
n := 0.
1 to: 10 do: [ :counter | n := n + counter ].
n  →  55
```

High-Order Iterators. Collections comprise a large number of different classes, many of which support the same protocol. The most important

messages for iterating over collections include `do:`, `collect:`, `select:`, `reject:`, `detect:` and `inject:into:`. These messages define high-level iterators that allow one to write very compact code.

An Interval is a collection that lets one iterate over a sequence of numbers from the starting point to the end. `1 to: 10` represents the interval from 1 to 10. Since it is a collection, we can send the message `do:` to it. The argument is a block that is evaluated for each element of the collection.

```
n := 0.
(1 to: 10) do: [ :element | n := n + element ].
n  → 55
```

`collect:` builds a new collection of the same size, transforming each element.

```
(1 to: 10) collect: [ :each | each * each ]  →  #(1 4 9 16 25 36 49 64 81 100)
```

`select:` and `reject:` build new collections, each containing a subset of the elements satisfying (or not) the boolean block condition. `detect:` returns the first element satisfying the condition. Don't forget that strings are also collections, so you can iterate over all the characters.

```
'hello there' select: [ :char | char isVowel ]  →  'eooo'
'hello there' reject: [ :char | char isVowel ]  →  'hll thr'
'hello there' detect: [ :char | char isVowel ]  →  $e
```

Finally, you should be aware that collections also support a functional-style *fold* operator in the `inject:into:` method. This lets you generate a cumulative result using an expression that starts with a seed value and injects each element of the collection. Sums and products are typical examples.

```
(1 to: 10) inject: 0 into: [ :sum :each | sum + each ]  →  55
```

This is equivalent to `0+1+2+3+4+5+6+7+8+9+10`.

More about collections and streams can be found in Chapter 9 and Chapter 10.

3.7 Primitives and Pragmas

In Smalltalk everything is an object, and everything happens by sending messages. Nevertheless, at certain points we hit rock bottom. Certain objects can only get work done by invoking virtual machine primitives.

For example, the following are all implemented as primitives: memory allocation (`new`, `new:`), bit manipulation (`bitAnd:`, `bitOr:`, `bitShift:`), pointer and integer arithmetic (`+`, `-`, `<`, `>`, `*`, `/`, `=`, `==...`), and array access (`at:`, `at:put:`).

Primitives are invoked with the syntax `<primitive: aNumber>`. A method that invokes such a primitive may also include Smalltalk code, which will be evaluated *only* if the primitive fails.

Here we see the code for `SmallInteger>+`. If the primitive fails, the expression `super + aNumber` will be evaluated and returned.

Method 3.2: A *primitive method*

`+ aNumber`

"Primitive. Add the receiver to the argument and answer with the result if it is a SmallInteger. Fail if the argument or the result is not a SmallInteger Essential No Lookup. See Object documentation whatIsAPrimitive."

`<primitive: 1>`

`↑ super + aNumber`

Since Squeak 3.9, the angle bracket syntax is also used for method annotations called pragmas.

3.8 Chapter Summary

- Squeak has (only) six reserved identifiers also called *pseudo-variables*: `true`, `false`, `nil`, `self`, `super`, and `thisContext`.
- There are five kinds of literal objects: numbers (`5`, `2.5`, `1.9e15`, `2r111`), characters (`$a`), strings (`'hello'`), symbols (`#hello`), and arrays (`#('hello' #hello)`)
- Strings are delimited by single quotes, comments by double quotes. To get a quote inside a string, double it.
- Unlike strings, symbols are guaranteed to be globally unique.
- Use `#(...)` to define a literal array. Use `{ ... }` to define a dynamic array. Note that `#(1 + 2) size` \longrightarrow `3`, but `{ 1 + 2 } size` \longrightarrow `1`
- There are three kinds of messages: *unary* (e.g., `1 asString`, `Array new`), *binary* (e.g., `3 + 4`, `'hi'`, `' there'`), and *keyword* (e.g., `'hi' at: 2 put: $o`)
- A *cascaded* message send is a sequence of messages sent to the same target, separated by semi-colons: `OrderedCollection new add: #calvin; add: #hobbes; size` \longrightarrow `2`

- Local variables are declared with vertical bars. Use `:=` for assignment; `←` or `_` will work as well but are deprecated since Squeak 3.9. `|x| x:=1`
- Expressions consist of message sends, cascades and assignments, possibly grouped with parentheses. *Statements* are expressions separated by periods.
- Block closures are expressions enclosed in square brackets. Blocks may take arguments and can contain temporary variables. The expressions in the block are not evaluated until you send the block a `value...` message with the correct number of arguments.
`[:x | x + 2] value: 4 → 6.`
- There is no dedicated syntax for control constructs, just messages that conditionally evaluate blocks.
 (Smalltalk includes: `Class`) `ifTrue: [Transcript show: Class superclass]`

Chapter 4

Understanding Message Syntax

Although Smalltalk's message syntax is extremely simple, it is unconventional and can take some getting used to. This chapter offers some guidance to help you get acclimatized to this special syntax for sending messages. If you already feel comfortable with the syntax, you may choose to skip this chapter, or come back to it later.

4.1 Identifying Messages

In Smalltalk, except for the syntactic elements listed in Chapter 3 (`:=`, `↑`, `.`, `;`, `#`, `()`, `{}`, `[]`, `:`, `|`), everything is a message send. As in C++, you can define operators like `+` for your own classes, but all operators have the same precedence. Moreover, you cannot change the arity of a method: `-` is always a binary message, and there is no way to have a unary form with a different overloading.

In Smalltalk the order in which messages are sent is determined by the kind of message. There are just three kinds of messages: *unary*, *binary*, and *keyword* messages. Unary messages are always sent first, then binary messages and finally keyword ones. As in most languages, parentheses can be used to change the order of evaluation. These rules make Smalltalk code as easy to read as possible. And most of the time you do not have to think about the rules.

As most computation in Smalltalk is done by message passing, correctly identifying messages is crucial. The following terminology will help us:

- A message is composed of the message *selector* and the optional message arguments.
- A message is sent to a *receiver*.
- The combination of a message and its receiver is called a *message send* as shown in Figure 4.1.

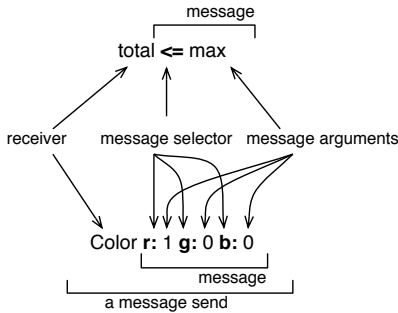


Figure 4.1: Two messages composed of a receiver, a method selector, and a set of arguments.




Figure 4.2: aMorph color: Color yellow is composed of two expressions: Color yellow and aMorph color: Color yellow.

A message is always sent to a receiver, which can be a single literal or a variable or the result of evaluating another expression.

In order to help you identify the receiver of a message, we will underline it for you. We will also surround each expression with an ellipse and number expressions starting from the first one that will be sent to help you see the order in which messages are sent.

Figure 4.2 represents two message sends, Color yellow and aMorph color: Color yellow, hence there are two ellipses. The expression Color yellow is sent first so its ellipse is numbered 1. There are two receivers: aMorph which receives the message color: ... and Color which receives the message yellow. Both receivers are underlined.

A receiver can be the first element of a message, such as 100 in the expression 100 + 200 or Color in the expression Color yellow. However, a receiver can also be the result of other messages. For example in the message Pen new go: 100, the receiver of the message go: 100 is the object returned by

Expression	Message type	Result
Color yellow	unary	Creates a color.
aPen go: 100.	keyword	The receiving pen moves forward 100 pixels.
100 + 20	binary	The number 100 receives the message + with the number 20.
Browser open	unary	Opens a new browser.
Pen new go: 100	unary and keyword	A pen is created and moved 100 pixels.
aPen go: 100 + 20	keyword and binary	The receiving pen moves forward 120 pixels.

Table 4.1: Examples of Messages

the expression `Pen new`. In all the cases, a message is sent to an object called the *receiver* which may be the result of another message send.

Table 4.1 shows several examples of messages. You should note that not all messages have arguments. Unary messages like `open` do not have arguments. Single keyword and binary messages like `go: 100` and `+ 20` each have one argument. There are also simple messages and composed ones. `Color yellow` and `100 + 20` are simple: a message is sent to an object, while the expression `aPen go: 100 + 20` is composed of two messages: `+ 20` is sent to `100` and `go:` is sent to `aPen` with the argument being the result of the first message. A receiver can be an expression which returns an object. In `Pen new go: 100`, the message `go: 100` is sent to the object that results from the evaluation of the expression `Pen new`.

4.2 Three Kinds of Messages

Smalltalk defines a few simple rules to determine the order in which the messages are sent. These rules are based on the distinction between 3 different kinds of messages:

- *Unary messages* are messages that are sent to an object without any other information. For example in `3 factorial`, `factorial` is a unary message.
- *Binary messages* are messages consisting of operators (often arithmetic). They are binary because they always involve only two objects: the receiver and the argument object. For example in `10 + 20`, `+` is a binary

message sent to the receiver 10 with argument 20.

- *Keyword messages* are messages consisting of one or more keywords, each ending with a colon (:) and taking an argument. For example in anArray at: 1 put: 10, the keyword at: takes the argument 1 and the keyword put: takes the argument 10.

Unary Messages

Unary messages are messages that do not require any argument. They follow the syntactic template: receiver *messageName*. The selector is simply made up of a succession of characters not containing : (*e.g.*, factorial, open, class).

89 sin	→	0.860069405812453
3 sqrt	→	1.732050807568877
Float pi	→	3.141592653589793
'blop' size	→	4
true not	→	false
Object class	→	Object class <i>"The class of Object is Object class (!)"</i>

Unary messages are messages that do not require any argument.
They follow the syntactic template: receiver **selector**

Binary Messages

Binary messages are messages that require exactly one argument *and* whose selector consists of a sequence of one or more characters from the set: +, -, *, /, &, =, >, |, <, ~, and @. Note that -- is not allowed for parsing reasons.

100@100	→	100@100 <i>"creates a Point object"</i>
3 + 4	→	7
10 - 1	→	9
4 <= 3	→	false
(4/3) * 3 = 4	→	true <i>"equality is just a binary message, and Fractions are exact"</i>
(3/4) == (3/4)	→	false <i>"two equal Fractions are not the same object"</i>

Binary messages are messages that require exactly one argument *and* whose selector is composed of a sequence of characters from: +, -, *, /, &, =, >, |, <, ~, and @. -- is not possible.

They follow the syntactic template: receiver **selector** argument

Keyword messages

Keyword messages are messages that require one or more arguments and whose selector consists of one or more keywords each ending in :. Keyword messages follow the syntactic template: receiver **selectorWordOne**: argumentOne **wordTwo**: argumentTwo

Each keyword takes an argument. Hence r:g:b: is a method with three arguments, playFileNameed: and at: are methods with one argument, and at:put: is a method with two arguments. To create an instance of the class Color one can use the method r:g:b: as in Color r: 1 g: 0 b: 0, which creates the color red. Note that the colons are part of the selector.

In Java or C++, the Smalltalk method invocation Color r: 1 g: 0 b: 0 would be written Color.rgb(1,0,0).

1 to: 10	→	(1 to: 10) "create an interval"
Color r: 1 g: 0 b: 0	→	Color red "create a new color"
12 between: 8 and: 15	→	true

nums := Array newFrom: (1 to: 5).

nums at: 1 put: 6.

nums → #(6 2 3 4 5)

Keyword based messages are messages that require one or more arguments. Their selector consists of one or more keywords each ending in a colon (:). They follow the syntactic template:

receiver **selectorWordOne**: argumentOne **wordTwo**: argumentTwo

4.3 Message Composition

The three kinds of messages each have different precedence, which allows them to be composed in an elegant way.

1. Unary messages are always sent first, then binary messages and finally keyword messages.
2. Messages in parentheses are sent prior to any kind of messages.
3. Messages of the same kind are evaluated from left to right.

These rules lead to a very natural reading order. Now if you want to be sure that your messages are sent in the order that you want you can always put more parentheses as shown in Figure 4.3. In this figure, the message `yellow` is an unary message and the message `color:` a keyword message, therefore the expression `Color yellow` is sent first. However as expressions in parentheses are sent first putting (unnecessary) parentheses around `Color yellow` just emphasizes that it will be sent first. The rest of the section illustrates each of these points.




Figure 4.3: Unary messages are sent first so `Color yellow` is sent. This returns a color object which is passed as argument of the message `aPen color:`.

Unary > Binary > Keywords

Unary messages are sent first, then binary messages, and finally keyword messages. We also say that unary messages have a higher priority over the other kinds of messages.

Rule One. Unary messages are sent first, then binary messages, and finally keyword based messages.
Unary > Binary > Keyword

As these examples show, Smalltalk's syntax rules generally ensure that expressions can be read in a natural way:

1000 factorial / 999 factorial	→	1000
2 raisedTo: 1 + 3 factorial	→	128

Unfortunately the rules are a bit too simplistic for arithmetic expressions, so you need to introduce parentheses whenever you want to impose a priority over binary operators:

1 + 2 * 3	→	9
1 + (2 * 3)	→	7

The following example, which is a bit more complex (!), offers a nice illustration that even complicated Smalltalk expressions can be read in a natural way:

[aClass aClass methodDict keys select: [:aMethod (aClass>>aMethod) isAbstract]]	
value: Boolean	→ an IdentitySet(#or: # #and: #& #ifTrue: #ifTrue:ifFalse: #ifFalse: #not #ifFalse:ifTrue:)

Here we want to know which methods of the Boolean class are abstract. We ask some argument class, aClass, for the keys of its method dictionary, and select those methods of that class that are abstract. Then we bind the argument aClass to the concrete value Boolean. We need parentheses only to send the binary message >>, which selects a method from a class, before sending the unary message isAbstract to that method. The result shows us which methods must be implemented by Boolean’s concrete subclasses True and False.

Example. In the message aPen color: Color yellow, there is one *unary* message yellow sent to the class Color and a *keyword* message color: sent to aPen. Unary messages are sent first so the expression Color yellow is sent (1). This returns a color object which is passed as argument of the message aPen color: aColor (2) as shown in example 4.1. Figure 4.3 shows graphically how messages are sent.

Example 4.1: <i>Decomposing the evaluation of</i> aPen color: Color yellow		
	aPen color: Color yellow	
(1)	Color yellow	"unary message is sent first"
	→ aColor	
(2)	aPen color: aColor	"keyword message is sent next"

Example. In the message aPen go: 100 + 20, there is a *binary* message + 20 and a *keyword* message go:. Binary messages are sent prior to keyword

messages so 100 + 20 is sent first (1): the message + 20 is sent to the object 100 and returns the number 120. Then the message aPen go: 120 is sent with 120 as argument (2). Example 4.2 shows how the expression is evaluated.

Example 4.2: Decomposing aPen go: 100 + 20

aPen go: 100 + 20		
(1)	100 + 20	"binary message first"
	→ 120	
(2)	aPen go: 120	"then keyword message"




Figure 4.4: Unary messages are sent first so Color yellow is sent. This returns a color object which is passed as argument of the message aPen color:.




Figure 4.5: Decomposing Pen new go: 100 + 20

Example. As an exercise we let you decompose the evaluation of the message Pen new go: 100 + 20 which is composed of one unary, one keyword and one binary message (see Figure 4.5).

Parentheses First

Rule Two. Parenthesised messages are sent prior to other messages.

(Msg) > Unary > Binary > Keyword

1.5 tan rounded asString = (((1.5 tan) rounded) asString)	→	true	"parentheses not needed here"
3 + 4 factorial	→	27	"(not 5040)"
(3 + 4) factorial	→	5040	

Here we need the parentheses to force sending lowMajorScaleOn: before play.

(FMSound lowMajorScaleOn: FMSound clarinet) play
"(1) send the message clarinet to the FMSound class to create a clarinet sound.
(2) send this sound to FMSound as argument to the lowMajorScaleOn: keyword message.
(3) play the resulting sound."

Example. The message (65@325 extent: 134 @ 100) center returns the center of a rectangle whose top left point is (65, 325) and whose size is 134×100. Example 4.3 shows how the message is decomposed and sent. First the message between parentheses is sent: it contains two binary messages 65@325 and 134@100 that are sent first and return points, and a keyword message extent: which is then sent and returns a rectangle. Finally the unary message center is sent to the rectangle and a point is returned. Evaluating the message without parentheses would lead to an error because the object 100 does not understand the message center.

Example 4.3: Example of Parentheses.

(65 @ 325 extent: 134 @ 100) center		
(1)	65@325	"binary"
	→ aPoint	
(2)	134@100	"binary"
	→ anotherPoint	
(3)	aPoint extent: anotherPoint	"keyword"
	→ aRectangle	
(4)	aRectangle center	"unary"
	→ 132@375	

From Left to Right

Now we know how messages of different kinds or priorities are handled. The final question to be addressed is how messages with the same priority are sent. They are sent from the left to the right. Note that you already saw this behavior in example 4.3 where the two point creation messages (@) were sent first.

Rule Three. When the messages are of the same kind, the order of evaluation is from left to right.

Example. In the expression Pen new down all messages are unary messages, so the leftmost one, Pen new, is sent first. This returns a newly created pen to which the second message down is sent, as shown in Figure 4.6.





Figure 4.6: Decomposing Pen new down



Arithmetic Inconsistencies

The message composition rules are simple but they result in inconsistency for the evaluation of arithmetic expressions expressed in terms of binary messages. Here we see the common situations where extra parentheses are needed.

3 + 4 * 5	→	35	"(not 23) Binary messages sent from left to right"
3 + (4 * 5)	→	23	
1 + 1/3	→	(2/3)	"and not 4/3"
1 + (1/3)	→	(4/3)	
1/3 + 2/3	→	(7/9)	"and not 1"
(1/3) + (2/3)	→	1	

Example. In the expression 20 + 2 * 5, there are only binary messages + and *. However in Smalltalk there is no specific priority for the operations + and *. They are just binary messages, hence * does not have priority over +. Here the leftmost message + is sent first (1) and then the * is sent to the result as shown in example 4.4.

Example 4.4: Decomposing 20 + 2 * 5

*"As there is no priority among binary messages, the leftmost message + is evaluated first even if by the rules of arithmetic the * should be sent first."*

- 20 + 2 * 5
- (1) 20 + 2 → 22
- (2) 22 * 5 → 110

As shown in example 4.4 the result of this expression is not 30 but 110. This result is perhaps unexpected but follows directly from the rules used

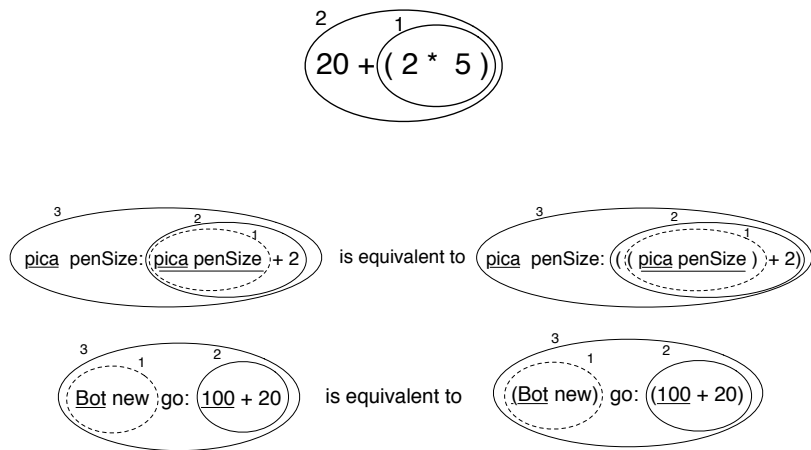


Figure 4.7: Equivalent messages using parentheses.

to send messages. This is somehow the price to pay for the simplicity of the Smalltalk model. To get the correct result, we should use parentheses. When messages are enclosed in parentheses, they are evaluated first. Hence the expression `20 + (2 * 5)` returns the result as shown in example 4.5.

Example 4.5: *Decomposing 20 + (2 * 5)*

*"The messages surrounded by parentheses are evaluated first therefore * is sent prior to + which produces the correct behavior."*

	20 + (2 * 5)	
(1)	(2 * 5)	→ 10
(2)	20 + 10	→ 30

In Smalltalk, arithmetic operators such as + and * do not have different priority. + and * are just binary messages, therefore * does not have priority over +. Use parentheses to obtain the desired result.

Note that the first rule stating that unary messages are sent prior to binary and keyword messages avoids the need to put explicit parentheses around them. Table 4.8 shows expressions written following the rules and equivalent expressions if the rules would not exist. Both expressions result in the same effect or return the same value.

Implicit precedence	Explicitly parenthesized equivalent
aPen color: Color yellow	aPen color: (Color yellow)
aPen go: 100 + 20	aPen go: (100 + 20)
aPen penSize: aPen penSize + 2	aPen penSize: ((aPen penSize) + 2)
2 factorial + 4	(2 factorial) + 4

Figure 4.8: Expressions and their fully parenthesized equivalents

4.4 Hints for Identifying Keyword Messages

Often beginners have problems understanding when they need to add parentheses. Let’s see how keywords messages are recognized by the compiler.

Parentheses or not?

The characters [], and (,) delimit distinct areas. Within such an area, a keyword message is the longest sequence of words terminated by : that is not cut by the characters ., or ;. When the characters [], and (,) surround some words with colons, these words participate in the keyword message *local* to the area defined.

In this example, there are two distinct keyword messages: rotatedBy:magnify:smoothing: and at:put:.

```
aDict
at: (rotatingForm
    rotateBy: angle
    magnify: 2
    smoothing: 1)
put: 3
```

The characters [], and (,) delimit distinct areas. Within such an area, a keyword message is the longest sequence of words terminated by : that is not cut by the characters ., or ;. When the characters [], and (,) surround some words with colons, these words participate in the keyword message local to the area defined.

Hints. If you have problems with these precedence rules, you may start simply by putting parentheses whenever you want to distinguish two messages having the same precedence.

The following expression does not require parentheses because the expression `x isNil` is unary hence is sent prior to the keyword message `ifTrue:.`

```
(x isNil)
  ifTrue:[...]
```

The following expression requires parentheses because the messages `includes:` and `ifTrue:` are both keyword messages.

```
ord := OrderedCollection new.
(ord includes: $a)
  ifTrue:[...]
```

Without parentheses the unknown message `includes:ifTrue:` would be sent to the collection!

When to use [] or ()

You may also have problems understanding when to use square brackets rather than parentheses. The basic principle is that you should use `[]` when you do not know how many times, potentially zero, an expression should be evaluated. `[expression]` will create a block closure (*i.e.*, an object) from *expression*, which may be evaluated any number of times (possibly zero), depending on the context.

Hence the conditional branches of `ifTrue:` or `ifTrue:ifFalse:` require blocks. Following the same principle both the receiver and the argument of a `whileTrue:` message require the use of square brackets since we do not know how many times either the receiver or the argument should be evaluated.

Parentheses, on the other hand, only affect the order of sending messages. No object is created, so in *(expression)*, the *expression* will *always* be evaluated exactly once (assuming the sounding code is evaluated once).

<code>[x isReady] whileTrue: [y doSomething]</code>	<i>"both the receiver and the argument must be blocks"</i>
<code>4 timesRepeat: [Beeper beep]</code>	<i>"the argument is evaluated more than once, so must be a block"</i>
<code>(x isReady) ifTrue: [y doSomething]</code>	<i>"receiver is evaluated once, so is not a block"</i>

4.5 Expression Sequences

Expressions (*i.e.*, messages sends, assignments ...) separated by periods are evaluated in sequence. Note that there is no period between a variable definition and the following expression. The value of a sequence is the value of the last expression. The values returned by all the expressions except the last one are ignored. Note that the period is a separator and not a terminator. Therefore a final period is optional.

```
| box |
box := 20@30 corner: 60@90.
box containsPoint: 40@50    →    true
```

4.6 Cascaded Messages

Smalltalk offers a way to send multiple messages to the same receiver using a semicolon (;). This is called the *cascade* in Smalltalk jargon.

Expression Msg1 ; Msg2

```
Transcript show: 'Squeak
is '.
```

```
Transcript show: 'fun '.    is equivalent to:
Transcript cr.
```

```
Transcript
show: 'Squeak is';
show: 'fun ';
cr
```

Note that the object receiving the cascaded messages can itself be the result of a message send. In fact the receiver of all the cascaded messages is the receiver of the first message involved in a cascade. In the following example, the first cascaded message is `setX:setY` since it is followed by a cascade. The receiver of the cascaded message `setX:setY` is the newly created point resulting from the evaluation of `Point new`, and *not* `Point`. The subsequent message `isZero` is sent to that same receiver.

```
Point new setX: 25 setY: 35; isZero    →    false
```

4.7 Chapter Summary

- A message is always sent to an object named the *receiver* which may be the result of other message sends.

- Unary messages are messages that do not require any argument. They are of the form of receiver **selector**.
- Binary messages are messages that involve two objects, the receiver and another object *and* whose selector is composed of one or two characters from the following list: +, -, *, /, |, &, =, >, <, ~, and @. They are of the form: receiver **selector** argument
- Keyword messages are messages that involve more than one object and that contain at least one colon character (:). They are of the form: receiver **selectorWordOne:** argumentOne **wordTwo:** argumentTwo
- **Rule One.** Unary messages are sent first, then binary messages, and finally keyword messages.
- **Rule Two.** Messages in parentheses are sent before any others.
- **Rule Three.** When the messages are of the same kind, the order of evaluation is from left to right.
- In Smalltalk, traditional arithmetic operators such as + and * have the same priority. + and * are just binary messages, therefore * does not have priority over +. You must use parentheses to obtain a different result.

Part II

Developing in Squeak

Chapter 5

The Smalltalk Object Model

Smalltalk's programming model is simple and uniform: everything is an object, and objects communicate only by sending each other messages. However, this simplicity and uniformity can be a source of difficulty for programmers used to other languages. In this chapter we present the core concepts of the Smalltalk object model; in particular we discuss the consequences of representing classes as objects.

5.1 The Rules of the Model

The Smalltalk object model is based on a set of simple rules that are applied *uniformly*. The rules are as follows:

Rule 1. Everything is an object.

Rule 2. Every object is an instance of a class.

Rule 3. Every class has a superclass.

Rule 4. Everything happens by message sends.

Rule 5. Method lookup follows the inheritance chain.

Let us look at each of these rules in some detail.

5.2 Everything is an Object

The mantra “everything is an object” is highly contagious. After only a short while working with Smalltalk, you will start to be surprised at how

this rule simplifies everything you do. Integers, for example, are truly objects, so you can send messages to them, just as you do to any other object.

3 + 4	→	7	"send '+ 4' to 3, yielding 7"
20 factorial	→	2432902008176640000	"send factorial, yielding a big number"

The representation of 20 factorial is certainly different from the representation of 7, but because they are both objects, none of the code — not even the implementation of factorial — needs to now about this.

Perhaps the most fundamental consequence of this rule is the following:

Classes are objects too.

Furthermore, classes are not second-class objects: they are really first-class objects that you can send messages to, inspect, and so on. This means that Squeak is a truly reflective system, which gives a great deal of expressive power to developers.

Deep in the implementation of Smalltalk, there are three different kinds of objects. There are (1) ordinary objects with instance variables that are passed by references, there are (2) *small integers* that are passed by value, and there are (3) indexable objects like arrays that hold a contiguous portion of memory. The beauty of Smalltalk is that you normally don't need to care about the differences between these three kinds of object.

5.3 Every object is an instance of a class

Every object has a class; you can find out which by sending it the message `class`.

1 class	→	SmallInteger
20 factorial class	→	LargePositiveInteger
'hello' class	→	ByteString
#(1 2 3) class	→	Array
(4@5) class	→	Point
Object new class	→	Object

A class defines the *structure* of its instances via instance variables, and the *behavior* of its instances via methods. Each method has a name, called its *selector*, which is unique within the class.

Since *classes are objects*, and *every object is an instance of a class*, it follows that classes must also be instances of classes. A class whose instances are

classes is called a *metaclass*. Whenever you create a class, the system automatically creates a metaclass for you. The metaclass defines the structure and behavior of the class that is its instance. 99% of the time you will not need to think about metaclasses, and may happily ignore them. (We will have a closer look at metaclasses in Chapter 12.)

Instance variables

Instance variables in Smalltalk are private to the *instance* itself. This is in contrast to Java and C++, which allow instance variables (also known as “fields” or “member variables”) to be accessed by any other instance that happens to be of the same class. We say that the *encapsulation boundary* of objects in Java and C++ is the class, whereas in Smalltalk it is the instance.

In Smalltalk, two instances of the same class cannot access each other’s instance variables unless the class defines “accessor methods”. There is no language syntax that provides direct access to the instance variables of any other object. (Actually, a mechanism called reflection does provide a way to ask another object for the values of its instance variables; meta-programming is intended for writing tools like the object inspector, whose sole purpose is to look inside other objects.)

Instance variables can be accessed by name in any of the instance methods of the class that defines them, and also in the methods defined in its subclasses. This means that Smalltalk instance variables are similar to *protected* variables in C++ and Java. However, we prefer to say that they are private, because it is considered bad style in Smalltalk to access an instance variable directly from a subclass.

Example

Method `Point>>dist:` (method 5.1) computes the distance between the receiver and another point. The instance variables `x` and `y` of the receiver are accessed directly by the method body. However, the instance variables of the other point must be accessed by sending it the messages `x` and `y`.

Method 5.1: *the distance between two points*

`Point>>dist:` `aPoint`

"Answer the distance between aPoint and the receiver."

| `dx dy` |

`dx := aPoint x - x.`

`dy := aPoint y - y.`

↑ `((dx * dx) + (dy * dy)) sqrt`

1@1 dist: 4@5 \longrightarrow 5.0

The key reason to prefer instance-based encapsulation to class-based encapsulation is that it enables different implementations of the same abstraction to coexist. For example, method `point>dist:`, need not know or care whether the argument `aPoint` is an instance of the same class as the receiver. The argument object might be represented in polar coordinates, or as a record in a database, or on another computer in a distributed system; as long as it can respond to the messages `x` and `y`, the code in method 5.1 will still work.

Methods

All methods are public.¹ Methods are grouped into protocols that indicate their intent. Some common protocol names have been established by convention, for example, *accessing* for all accessor methods, and *initialization* for establishing a consistent initial state for the object. The protocol *private* is sometimes used to group methods that should not be seen from outside. Nothing, however, prevents you from sending a message that is implemented by such a “private” method.

Methods can access all instance variables of the object. Some Smalltalk developers prefer to access instance variables only through accessors. This practice has some value, but it also clutters the interface of your classes, and worse, exposes private state to the world.

The Instance Side and the Class Side

Since classes are objects, they can have their own instance variables and their own methods. We call these *class instance variables* and *class methods*, but they are really no different from ordinary instance variables and methods: class instance variables are just instance variables defined by a metaclass, and class methods are just methods defined by a metaclass.

A class and its metaclass are two separate classes, even though the former is an instance of the latter. However, this is largely irrelevant to you as a programmer: you are concerned with defining the behavior of your objects and the classes that create them.

For this reason, the browser helps you to browse both class and meta-class as if they were a single thing with two “sides”: the “” and the “”, as

¹Well, almost all. In Squeak, methods whose selectors start with the string `pvt` are private: a `pvt` message can be sent *only* to self. However, `pvt` methods are not used very much.

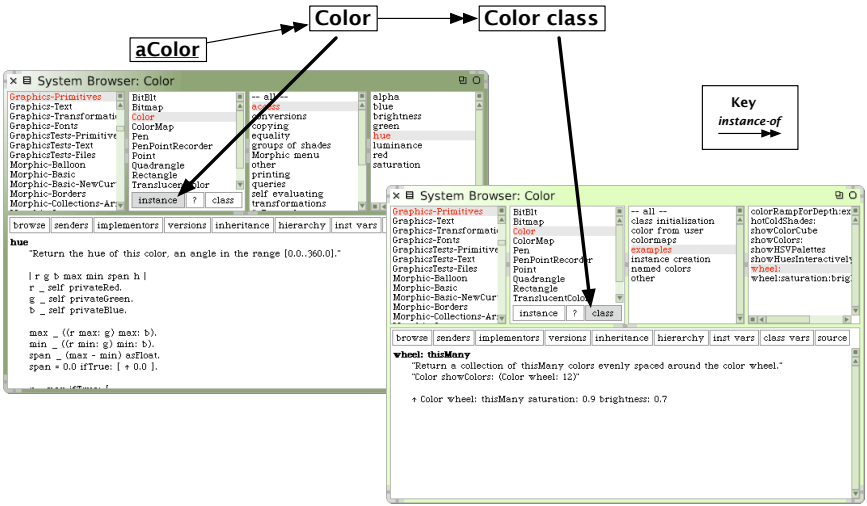


Figure 5.1: Browsing a class and its metaclass.

shown in Figure 5.1. Clicking on the `instance` button browses the class `Color`, *i.e.*, you browse the methods that are executed when messages are sent to an instance of `Color`, like the color `blue`. Pressing the `class` button browses the class `Color class`, *i.e.*, you see the methods that will be executed when messages are sent to the class `Color` itself. For example, `Color blue` sends the message `blue` to the class `Color`. You will therefore find the method `blue` defined on the class side of `Color`, not on the instance side.

<code>aColor := Color blue.</code>	<i>"Class side method blue"</i>
<code>aColor</code> → <code>Color blue</code>	
<code>aColor red</code> → <code>0.0</code>	<i>"Instance side accessor method red"</i>
<code>aColor blue</code> → <code>1.0</code>	<i>"Instance side accessor method blue"</i>


You define a class by filling in the template proposed on the instance side. When you accept this template, the system creates not just the class that you defined, but also the corresponding metaclass. You can browse the metaclass by clicking on the `class` button. The only part of the metaclass creation template that makes sense for you to edit directly is the list of instance variable names.

Once a class has been created, clicking the `instance` button lets you edit and browse the methods that will be possessed by instances of that class (and of its subclasses). For example, we can see in Figure 5.1 that the method `hue` is defined on instances of the class `Color`. In contrast, the `class` button lets you browse and edit the metaclass (in this case `Color class`).

Class Methods

Class methods can be quite useful; browse Color class for some good examples. You will see that there are two kinds of method defined on a class: those that create instances of the class, like Color class»blue and those that perform a utility function, like Color class»showColorCube. This is typical, although you will occasionally find class methods used in other ways.

It is convenient to place utility methods on the class side because they can be executed without having to create any additional objects first. Indeed, many of them will contain a comment designed to make it easy to execute them.

 Browse method Color class»showColorCube, double-click just inside the quotes on the comment "Color showColorCube" and type CMD-d.

You will see the effect of executing this method. (Select World ▸ restore display (r) to undo the effects.)

For those familiar with Java and C++, class methods may seem similar to static methods. However, the uniformity of Smalltalk means that they are somewhat different: whereas Java static methods are really just statically-resolved procedures, Smalltalk class methods are dynamically-dispatched methods. This means that inheritance, overriding and super-sends work for class methods in Smalltalk, whereas they don't work for static methods in Java.

Class Instance Variables

With ordinary instance variables, all the instances of a class have the same set of variable names, and the instances of its subclasses inherit those names; however, each instance has its own private set of values. The story is exactly the same with class instance variables: each class has its own private class instance variables. A subclass will inherit those class instance variables, *but it has its own private copies of those variables*. Just as objects don't share instance variables, neither do classes and their subclasses share class instance variables.

You could use a class instance variable called count to keep track of how many instances you create of a given class. However, any subclass would have its own count variable, so subclass instances would be counted separately.

Example: class instance variables are not shared with subclasses. Suppose we define classes Dog and Hyena, where Hyena inherits the class instance

variable count from Dog.

Class 5.2: Dogs and Hyenas

```
Object subclass: #Dog
  instanceVariableNames: "
  classVariableNames: "
  poolDictionaries: "
  category: 'SBE-CIV'
```

```
Dog class
  instanceVariableNames: 'count'
```

```
Dog subclass: #Hyena
  instanceVariableNames: "
  classVariableNames: "
  poolDictionaries: "
  category: 'SBE-CIV'
```

Now suppose we define class methods for Dog to initialize its count to 0, and to increment it when new instances are created:

Method 5.3: Keeping count of new dogs

```
Dog class»initialize
  super initialize.
  count := 0.
```

```
Dog class»new
  count := count + 1.
  ↑ super new
```

```
Dog class»count
  ↑ count
```

Now when we create a new Dog its count is incremented, and so is that of every Hyena, but they are counted separately:

```
Dog initialize.
Hyena initialize.
Dog count      → 0
Hyena count    → 0
Dog new.
Dog count      → 1
Dog new.
Dog count      → 2
Hyena new.
Hyena count    → 1
```

Note also that class instance variables are private to a class in exactly the same way that instance variables are private to the instance. Since classes and their instances are different objects, this has the following immediate consequences:

A class does not have access to the instance variables of its own instances.

An instance of a class does not have access to the class instance variables of its class.

For this reason, instance initialization methods must always be defined on the instance side — the class side has no access to instance variables, so cannot initialize them! All that the class can do is to send initialization messages, which may use using accessors, to newly created instances.

Similarly, instances can only access class instance variables indirectly, by sending accessor messages to their class.

Java has nothing equivalent to class instance variables. Java and C++ static variables are more like Smalltalk class variables, which we will discuss in Section 5.7: all of the subclasses and all of their instances share the same static variable.

Example: Defining a Singleton. The Singleton pattern² provides a typical example of the use of class instance variables and class methods. Imagine that we would like to implement a class `WebServer` and use the Singleton pattern to ensure that it has only one instance.

Clicking on the `instance` button in the browser, we define the class `WebServer` as follows (class 5.4).

Class 5.4: *A singleton class*

```
Object subclass: #WebServer
  instanceVariableNames: 'sessions'
  classVariableNames: ''
  poolDictionaries: ''
  category: 'Web'
```

Then, clicking on the `class` button, we add the instance variable `uniqueInstance` to the class side.

²Sherman R. Alpert, Kyle Brown and Bobby Woolf, *The Design Patterns Smalltalk Companion*. Addison Wesley, 1998, ISBN 0-201-18462-1.

Class 5.5: *The class side of the singleton class*

```
WebServer class
  instanceVariableNames: 'uniqueInstance'
```

The consequence of this is that the class `WebServer` now has another instance variable, in addition to the variables that it inherits, such as `superclass` and `methodDict`.

We can now define a class method named `uniqueInstance` as shown in method 5.6. This method first checks whether `uniqueInstance` has been initialized. If it has not, the method creates an instance and assigns it to the class instance variable `uniqueInstance`. Finally the value of `uniqueInstance` is returned. Since `uniqueInstance` is a class instance variable, this method can directly access it.

Method 5.6: *uniqueInstance (on the class side)*

```
WebServer class>uniqueInstance
  uniqueInstance ifNil: [uniqueInstance := self new].
  ↑ uniqueInstance
```

The first time that `WebServer uniqueInstance` is executed, an instance of the class `WebServer` will be created and assigned to the `uniqueInstance` variable. The next time, the previously created instance will be returned instead of creating a new one.


Note that the instance creation code inside the conditional in method 5.6 is written as `self new` and not as `WebServer new`. What is the difference? Since the `uniqueInstance` method is defined in `WebServer class`, you might think that they were the same. And indeed, until someone creates a subclass of `WebServer`, they are the same. But suppose that `ReliableWebServer` is a subclass of `WebServer`, and inherits the `uniqueInstance` method. We would clearly expect `ReliableWebServer uniqueInstance` to answer a `ReliableWebServer`. Using `self` ensures that this will happen, since it will be bound to the respective class. Note also that `WebServer` and `ReliableWebServer` will each have their own class instance variable called `uniqueInstance`. These two variables will of course have different values.

5.4 Every class has a superclass

Each class in Smalltalk inherits its behavior and the description of its structure from a single *superclass*. This means that Smalltalk has single inheritance.

SmallInteger superclass	→	Integer
Integer superclass	→	Number
Number superclass	→	Magnitude
Magnitude superclass	→	Object
Object superclass	→	ProtoObject
ProtoObject superclass	→	nil

Traditionally the root of the Smalltalk inheritance hierarchy is the class `Object` (since everything is an object). In Squeak, the root is actually a class called `ProtoObject`, but you will normally not pay any attention to this class. `ProtoObject` encapsulates the minimal set of messages that all objects *must* have. However, most classes inherit from `Object`, which defines many additional messages that almost all objects ought to understand and respond to. Unless you have a very good reason to do otherwise, when creating application classes you should normally subclass `Object`, or one of its subclasses.

 A new class is normally created by sending the message `subclass: instanceVariableNames: ...` to an existing class. There are a few other methods to create classes. Have a look at the protocol `Kernel-Classes > Class > subclass` creation to see what they are.

Although Squeak does not provide multiple inheritance, since version 3.9 it has incorporated a mechanism called *traits* for sharing behavior across unrelated classes. Traits are collections of methods that can be reused by multiple classes that are not related by inheritance. Using traits allows one to share code between different classes without duplicating code.

Abstract Methods and Abstract Classes

An abstract class is a class that exists to be subclassed, rather than to be instantiated. An abstract class is usually incomplete, in the sense that it does not define all of the methods that it uses. The “missing” methods — those that the other methods assume, but which are not themselves defined — are called abstract methods.

Smalltalk has no dedicated syntax to specify that a method or a class is abstract. By convention, the body of an abstract method consists of the expression `self subclassResponsibility`. This is known as a “marker method”, and indicates that subclasses have the responsibility to define a concrete version of the method. `self subclassResponsibility` methods should always be overridden, and thus should never be executed. If you forget to override one, and it is executed, an exception will be raised.

A class is considered abstract if one of its methods is abstract. Nothing actually prevents you from creating an instance of an abstract class; everything will work until an abstract method is invoked.

Example: the class **Magnitude**.

Magnitude is an abstract class that helps us to define objects that can be compared to each other. Subclasses of Magnitude should implement the methods `<`, `=` and `hash`. Using such messages Magnitude defines other methods such as `>`, `>=`, `<=`, `max:`, `min:`, `between:` and `and:` and others for comparing objects. Such methods are inherited by subclasses. The method `<` is abstract and defined as shown in method 5.7.

Method 5.7: Magnitude»<

Magnitude»< aMagnitude

"Answer whether the receiver is less than the argument."

↑self subclassResponsibility

By contrast, the method `>=` is concrete; it is defined in terms of `<`:

Method 5.8: Magnitude»>=

>= aMagnitude

"Answer whether the receiver is greater than or equal to the argument."

↑(self < aMagnitude) not

The same is true of the other comparison methods.

Character is a subclass of Magnitude; it overrides the `subclassResponsibility` method for `<` with its own version of `<` (see method 5.9). Character also defines methods `=` and `hash`; it inherits from Magnitude the methods `>=`, `<=`, `~=` and others.

Method 5.9: Character»<

Character»< aCharacter

"Answer true if the receiver's value < aCharacter's value."

↑self asciiValue < aCharacter asciiValue

Traits

A *trait* is a collection of methods that can be included in the behavior of a class without the need for inheritance. This makes it easy for classes to have a unique superclass, yet still share useful methods with otherwise unrelated classes.

To define a new trait, simply replace the subclass creation template by a message to the class `Trait`.

Class 5.10: *Defining a new trait*

```
Trait named: #TAuthor
  uses: { }
  category: 'SBE-Quinto'
```

Here we define the trait `TAuthor` in the category *SBE-Quinto*. This trait does not *use* any other existing traits. In general we can specify a *trait composition expression* of other traits to use as part of the `uses:` keyword argument. Here we simply provide an empty array.

Traits may contain methods, but no instance variables. Suppose we would like to be able to add an `author` method to various classes, independent of where they occur in the hierarchy. We might do this as follows:

Method 5.11: *An author method*

```
TAuthor>>author
  "Returns author initials"
  ↑ 'on'  "oscar nierstrasz"
```

Now we can use this trait in a class that already has its own superclass, for instance the `SBEGame` class that we defined in Chapter 2. We simply modify the class creation template for `SBEGame` to include a `uses:` keyword argument that specifies that `TAuthor` should be used.

Class 5.12: *Using a trait*

```
BorderedMorph subclass: #SBEGame
  uses: TAuthor
  instanceVariableNames: 'cells'
  classVariableNames: "
  poolDictionaries: "
  category: 'SBE-Quinto'
```

If we now instantiate `SBEGame`, it will respond to the `author` message as expected.

```
SBEGame new author  →  'on'
```

Trait composition expressions may combine multiple traits using the `+` operator. In case of conflicts (*i.e.*, if multiple traits define methods with the same name), these conflicts can be resolved by explicitly removing these methods (with `-`), or by redefining these methods in the class or trait that you are defining. It is also possible to *alias* methods (with `@`), providing a new name for them.

Traits are used in the system kernel. One good example is the class Behavior.

Class 5.13: Behavior defined using traits

```
Object subclass: #Behavior
  uses: TPureBehavior @ {#basicAddTraitSelector:withMethod:->
    #addTraitSelector:withMethod:}
  instanceVariableNames: 'superclass methodDict format'
  classVariableNames: 'ObsoleteSubclasses'
  poolDictionaries: ''
  category: 'Kernel-Classes'
```

Here we see that the method `basicAddTraitSelector:withMethod:` defined in the trait `TPureBehavior` has been renamed to `addTraitSelector:withMethod:`. Support for traits is currently being added to the browsers.

5.5 Everything Happens by Message Sending

This rule captures the essence of programming in Smalltalk.

In procedural programming, the choice of which piece of code to execute when a procedure is called is made by the caller. The caller chooses the procedure or function to execute *statically*, by name.

In object-oriented programming, we do *not* “call methods”: we “send messages.” The choice of terminology is significant. Each object has its own responsibilities. We do not *tell* an object what to do by applying some procedure to it. Instead, we politely *ask* an object to do something for us by sending it a message. The message is *not* a piece of code: it is nothing but a name and a list of arguments. The receiver then decides how to respond by selecting its own *method* for doing what was asked. Since different objects may have different methods for responding to the same message, the method must be chosen *dynamically*, when the message is received.

3 + 4	→	7	"send message + with argument 4 to integer 3"
(1@2) + 4	→	5@6	"send message + with argument 4 to point (1@2)"

As a consequence, we can send the *same message* to different objects, each of which may have *its own method* for responding to the message. We do not tell the `SmallInteger 3` or the `Point 1@2` how to respond to the message `+ 4`. Each has its own method for answering this message send, and responds accordingly.

One of the consequences of Smalltalk’s model of message sending is that it encourages a style in which objects tend to have very small methods and

delegate tasks to other objects, rather than implementing huge, procedural methods that assume too much responsibility. Joseph Pelrine expresses this principle succinctly as follows:

Don't do anything that you can push off onto someone else.

Many object-oriented languages provide both static and dynamic operations for objects; in Smalltalk there are only dynamic message sends. Instead of providing static class operations, for instance, classes are objects and we simply send messages to classes.

Nearly everything in Smalltalk happens by message sends. At some point action must take place:

- *Variable declarations* are not message sends. In fact, variable declarations are not even executable. Declaring a variable just causes space to be allocated for an object reference.
- *Assignments* are not message sends. An assignment to a variable causes that variable name to be freshly bound in the scope of its definition.
- *Returns* are not message sends. A return simply causes the computed result to be returned to the sender.
- *Primitives* are not message sends. They are implemented in the virtual machine.

Other than these few exceptions, pretty much everything else does truly happen by sending messages. In particular, since there are no “public fields” in Smalltalk, the only way to update an instance variable of another object is to send it a message asking that it update its own field. Of course, providing setter and getter methods for all the instance variables of an object does is not good object-oriented style. Joseph Pelrine also states this very nicely:

Don't let anyone else play with your data.

5.6 Method lookup follows the inheritance chain

What exactly happens when an object receives a message?

The process is quite simple: the class of the receiver looks up the method to use to handle the message. If this class does not have a method, it asks its superclass, and so on, up the inheritance chain. When the method is found, the arguments are bound to the parameters of the method, and the virtual machine executes it.

It is essentially as simple as this. Nevertheless there are a few questions that need some care to answer:

- *What happens when a method does not explicitly return a value?*
- *What happens when a class reimplements a superclass method?*
- *What is the difference between `self` and `super` sends?*
- *What happens when no method is found?*

The rules for method lookup that we present here are conceptual: virtual machine implementors use all kinds of tricks and optimizations to speed-up method lookup. That's their job, but you should never be able to detect that they are doing something different from our rules.

First let us look at the basic lookup strategy, and then consider these further questions.

Method lookup

Suppose we create an instance of `EllipseMorph`.

```
anEllipse := EllipseMorph new.
```

If we now send this object the message `defaultColor`, we get the result `Color yellow`:

```
anEllipse defaultColor  →  Color yellow
```

The class `EllipseMorph` implements `defaultColor`, so the appropriate method is found immediately.

Method 5.14: *A locally implemented method*

```
EllipseMorph»defaultColor
  "answer the default color/fill style for the receiver"
  ↑ Color yellow
```

In contrast, if we send the message `openInWorld` to an `Ellipse`, the method is not immediately found, since the class `EllipseMorph` does not implement

openInWorld. The search therefore continues in the superclass, BorderedMorph, and so on, until an openInWorld method is found in the class Morph (see Figure 5.2).

Method 5.15: *An inherited method*

Morph»openInWorld

"Add this morph to the world. If in MVC, then provide a Morphic window for it."

self couldOpenInMorphic

ifTrue: [self openInWorld: self currentWorld]

ifFalse: [self openInMVC]

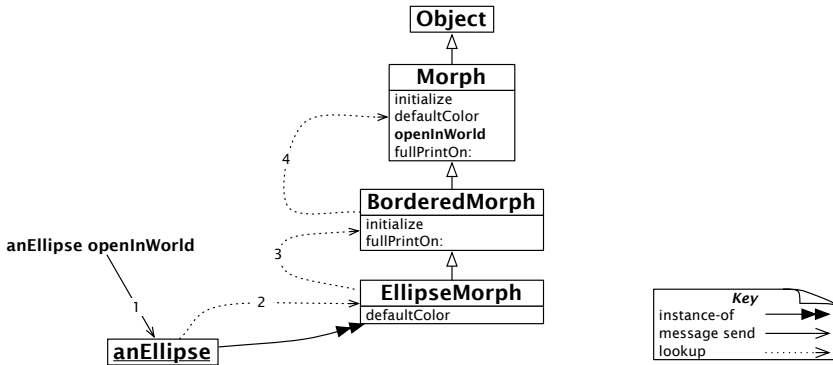


Figure 5.2: Method lookup follows the inheritance hierarchy.

Returning self

Notice that EllipseMorph»defaultColor (method 5.14) explicitly returns Color yellow whereas Morph»openInWorld (method 5.15) does not appear to return anything.

Actually a method *always* answers a message with a value — which is, of course, an object. The answer may be defined by the ↑ construct in the method, but if execution reaches the end of the method without executing a ↑, the method still answers a value: it answers the object that received the message. We usually say that the method “answers self”, because in Smalltalk the pseudo-variable self represents the receiver of the message, rather like this in Java.

This suggests that method 5.15 is equivalent to method 5.16:

Method 5.16: *Explicitly returning self*

```

Morph»openInWorld
  "Add this morph to the world. If in MVC,
  then provide a Morphic window for it."
  self couldOpenInMorphic
    ifTrue: [self openInWorld: self currentWorld]
    ifFalse: [self openInMVC].
  ↑ self    "Don't do this unless you mean it"

```

Why is writing `↑ self` explicitly not a good thing to do? Well, when you return something explicitly, you are communicating that you are returning something of interest to the sender. When you explicitly return `self`, you are saying that you expect the sender to use the returned value. This is not the case here, so it is best not to explicitly return `self`.

This is a common idiom in Smalltalk, which Kent Beck refers to as “Interesting return value”³:

Return a value only when you intend for the sender to use the value.

Overriding and Extension

If we look again at the `EllipseMorph` class hierarchy in Figure 5.2, we see that the classes `Morph` and `EllipseMorph` both implement `defaultColor`. In fact, if we open a new morph (`Morph new openInWorld`) we see that we get a blue morph, whereas an ellipse will be yellow by default.

We say that `EllipseMorph` *overrides* the `defaultColor` method that it inherits from `Morph`. The inherited method no longer exists from the point of view of an `Ellipse`.

Sometimes we do not want to override inherited methods, but rather *extend* them with some new functionality, that is, we would like to be able to invoke the overridden method *in addition to* the new functionality we are defining in the subclass. In Smalltalk, as in many object-oriented languages that support single inheritance, this can be done with the help of `super` sends.

The most important application of this mechanism is in the `initialize` method. Whenever a new instance of a class is initialized, it is critical to also initialize any inherited instance variables. However, the knowledge of how to do this is already captured in the `initialize` methods of each of

³Kent Beck, *Smalltalk Best Practice Patterns*. Prentice-Hall, 1997.

the superclass in the inheritance chain. The subclass has no business even trying to initialize inherited instance variables!

It is therefore good practice whenever implementing an initialize method to send super initialize before performing any further initialization:

Method 5.17: *Super initialize*

```
BorderedMorph>initialize
  "initialize the state of the receiver"
  super initialize.
  self borderInitialize
```

An initialize method should always start by sending super initialize.

Self sends and super sends

We need super sends to compose inherited behavior that would otherwise be overridden. The usual way to compose methods, whether inherited or not, however, is by means of self sends.

How do self sends differ from super sends? Like self, super represents the receiver of the message. The only thing that changes is the method lookup. Instead of lookup starting in the class of the receiver, it starts in the superclass of the class of the method where the super send occurs.

Note that super is *not* the superclass! It is a common and natural mistake to think this. It is also a mistake to think that lookup starts in the superclass of the receiver. We shall see with the following example precisely how this works.

Consider the message `initString`, which we can send to any morph:

```
anEllipse initString  →  '(EllipseMorph newBounds: (0@0 corner: 50@40) color:
  Color yellow) setBorderWidth: 1 borderColor: Color black'
```

The return value is a string that can be evaluated to recreate the morph.

How exactly is this result obtained through a combination of self and super sends? First, `anEllipse initString` will cause the method `initString` to be found in the class `Morph`, as shown in Figure 5.3.

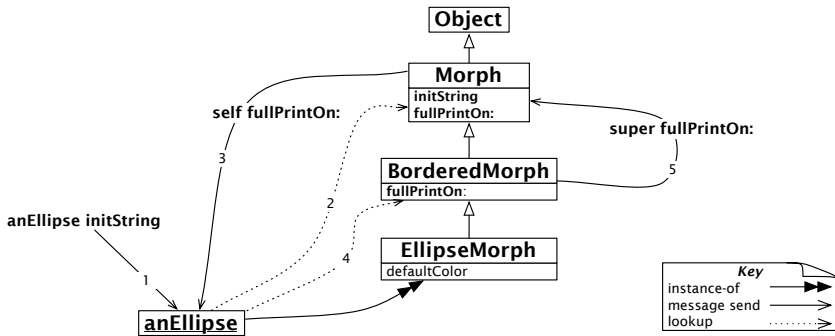


Figure 5.3: self and super sends

Method 5.18: A self send

```

Morph»initString
  ↑ String streamContents: [:s | self fullPrintOn: s]

```

The method `Morph»initString` performs a self send of `fullPrintOn:`. This causes a second lookup to take place, starting in the class `EllipseMorph`, and finding `fullPrintOn:` in `BorderedMorph` (see Figure 5.3 once again). What is critical to notice is that the self send causes the method lookup to start again in the class of the receiver, namely the class of `anEllipse`.

A self send triggers a *dynamic* method lookup starting in the class of the receiver.

Method 5.19: Combining super and self sends

```

BorderedMorph»fullPrintOn: aStream
  aStream nextPutAll: '('.
  super fullPrintOn: aStream.
  aStream nextPutAll: ')' setBorderWidth: '; print: borderWidth;
  nextPutAll: ' borderColor: ', (self colorString: borderColor)

```

At this point, `BorderedMorph»fullPrintOn:` does a super send to extend the `fullPrintOn:` behavior it inherits from its superclass. Because this is a super send, the lookup now starts from the superclass of the class where the super send occurs, namely in `Morph`. We then immediately find and evaluate `Morph»fullPrintOn:`.

Note that the super lookup did not start in the superclass of the receiver.

This would have caused lookup to start from BorderedMorph, resulting in an infinite loop!

A super send triggers a *static* method lookup starting in the superclass of the class of the method performing the super send.

If you think carefully about super send and Figure 5.3, you will realize that super bindings are static: all that matters is the class in which the text of the super send is found. By contrast, the meaning of self is dynamic: it always represents the receiver of the currently executing message. This means that *all* messages sent to self are looked-up by starting in the receiver's class.

Message not Understood

What happens if the method we are looking for is not found?

Suppose we send the message foo to our ellipse. First the normal method lookup would go through the inheritance chain all the way up to Object (or rather ProtoObject) looking for this method. When this method is not found, the virtual machine will cause the object to send self doesNotUnderstand: #foo. (See Figure 5.4.)

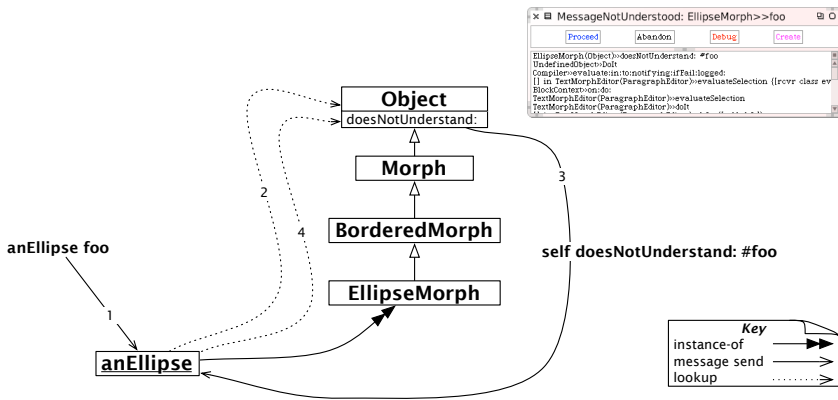


Figure 5.4: Message foo is not understood

Now, this is a perfectly ordinary, dynamic message send, so the lookup starts again from the class EllipseMorph, but this time searching for the method doesNotUnderstand:. As it turns out, Object implements

doesNotUnderstand:. This method will create a new MessageNotUnderstood object which is capable of starting a Debugger in the current execution context.

Why do we take this convoluted path to handle such an obvious error? Well, this offers developers an easy way to intercept such errors and take alternative action. One could easily override the method doesNotUnderstand: in any subclass of Object and provide a different way of handling the error.

In fact, this can be an easy way to implement automatic delegation of messages from one object to another. A Delegator object could simply delegate all messages it does not understand to another object whose responsibility it is to handle them, or raise an error itself!

5.7 Shared Variables

Now we will look at an aspect of Smalltalk that is not so easily covered by our five rules: shared variables.

Smalltalk provides three kinds of shared variables: (1) *globally* shared variables; (2) variables shared between instances and classes (*class variables*), and (3) variables shared amongst a group of classes (*pool variables*). The names of all of these shared variables start with a capital letter, to warn us that they are indeed shared between multiple objects.

Global Variables.

In Squeak, all global variables are stored in a namespace called Smalltalk, which is implemented as an instance of the class SystemDictionary. Global variables are accessible everywhere. Every class is named by a global variable; in addition, a few globals are used to name special or commonly useful objects.

The variable Transcript names an instance of TranscriptStream, a stream that writes to a scrolling window. The following code displays some information and then goes to the next line in the Transcript.

```
Transcript show: 'Squeak is fun and powerful' ; cr
```

Before you **do it**, open a transcript by dragging one from the *Tools* flap.

HINT Writing to the Transcript is slow, especially when the transcript window is open. So, if you experience some sluggishness and are writing to the Transcript, think about collapsing it.

Other useful Global variables.

- Smalltalk is the instance of SystemDictionary that defines all of the globals—including Smalltalk itself. The keys to this dictionary are the symbols that name the global objects in Smalltalk code. So, for example,

Smalltalk at: #Boolean → Boolean

Since Smalltalk is itself a global variable,

Smalltalk at: #Smalltalk → a SystemDictionary(lots of globals)}

and

(Smalltalk at: #Smalltalk) == Smalltalk → true

- Sensor is an instance of EventSensor, and represents input to Squeak. For example, Sensor keyboard answers the next character input on the keyboard, and Sensor leftShiftDown answers true if the left shift key is being held down, while Sensor mousePoint answers a Point indicating the current mouse location.
- World is an instance of PasteUpMorph that represents the screen. World bounds answers a rectangle that defines the whole screen space; all Morphs on the screen are submorphs of World.
- ActiveHand is the current instance of HandMorph, the graphical representation of the cursor. ActiveHand's submorphs hold anything being dragged by the mouse.
- Undeclared is another dictionary—it contains all the undeclared variables. If you write a method that references an undeclared variable, the browser will normally prompt you to declare it, for example, as a global or as an instance variable of the class. However, if you later delete the declaration, the code will then reference an undeclared variable. Inspecting Undeclared can sometimes help explain strange behaviour!
- SystemOrganization is an instance of SystemOrganizer: it records the organization of classes into packages. More precisely, it categorizes the *names* of classes, so

SystemOrganization categoryOfElement: #Magnitude → #'Kernel-Numbers'

Current practice is to strictly limit the use of global variables; it is usually better to use class instance variables or class variables, and to provide class methods to access them. Indeed, if Squeak were to be implemented from scratch today, most of the global variables that are not classes would be replaced by singletons.

The usual way to define a global is just to `do it` on an assignment to a capitalized but undeclared identifier. The parser will then offer to declare the global for you. If you want to define a global programatically, just execute Smalltalk at: `#AGlobalName put: nil`. To remove it, execute Smalltalk `removeKey: #AGlobalName`.

Class Variables

Sometimes we need to share some data amongst all the instances of a class and the class itself. This is possible using *class variables*. The term class variable indicates that the lifetime of the variable is the same as that of the class. However, what the term does not convey is that these variables are shared amongst all the instances of a class as well as the class itself, as shown in Figure 5.5. Indeed, a better name would have been *shared variables* since this expresses more clearly their role, and also warns of the danger of using them, particularly if they are modified.

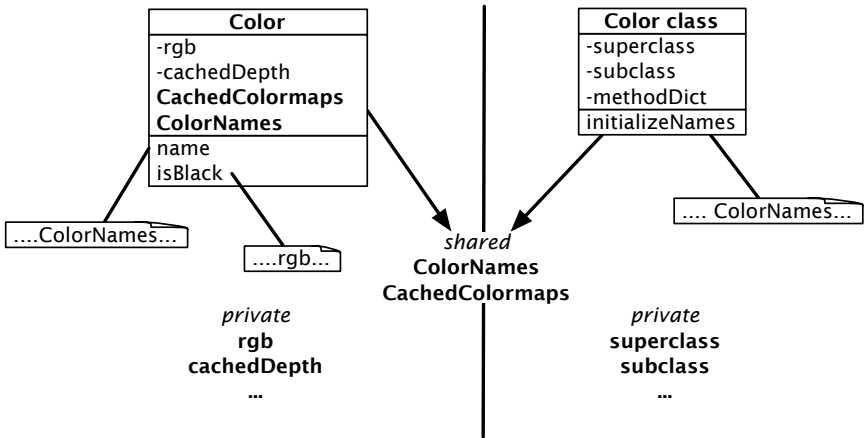


Figure 5.5: Instance and class methods accessing different variables.

In Figure 5.5 we see that `rgb` and `cachedDepth` are instance variables of `Color`, hence only accessible to instances of `Color`. We also see that `superclass`, `subclass`, `methodDict` and so on are class instance variables, *i.e.*, instance

variables only accessible to the Color class.

But we can also see something new: `ColorNames` and `CachedColormaps` are *class variables* defined for `Color`. The capitalization of these variables gives us a hint that they are shared. In fact, not only may all instances of `Color` access these shared variables, but also the `Color` class itself, *and any of its subclasses*. Both instance methods and class methods can access these shared variables.

A class variable is declared in the class definition template. For example, the class `Color` defines a large number of class variables to speed up color creation; its definition is shown below (class 5.20).

Class 5.20: *Color and its class variables*

Object subclass: `#Color`

```
instanceVariableNames: 'rgb cachedDepth cachedBitPattern'
classVariableNames: 'Black Blue BlueShift Brown CachedColormaps ColorChart
ColorNames ComponentMask ComponentMax Cyan DarkGray Gray
GrayToIndexMap Green GreenShift HalfComponentMask HighLightBitmaps
IndexedColors LightBlue LightBrown LightCyan LightGray LightGreen
LightMagenta LightOrange LightRed LightYellow Magenta MaskingMap Orange
PaleBlue PaleBuff PaleGreen PaleMagenta PaleOrange PalePeach PaleRed
PaleTan PaleYellow PureBlue PureCyan PureGreen PureMagenta PureRed
PureYellow RandomStream Red RedShift TranslucentPatterns Transparent
VeryDarkGray VeryLightGray VeryPaleRed VeryVeryDarkGray
VeryVeryLightGray White Yellow'
poolDictionaries: ''
category: 'Graphics-Primitives'
```

The class variable `ColorNames` is an array containing the name of frequently-used colors. This array is shared by all the instances of `Color` and its subclass `TranslucentColor`. It is accessible from all the instance and class methods.

`ColorNames` is initialized once in `Color class>>initializeNames`, but it is accessed from instances of `Color`. The method `Color>>name` uses the variable to find the name of a color. Since most colors do not have names, it was thought inappropriate to add an instance variable name to every color.

Class initialization

The presence of class variables raises the question: how do we initialize them? One solution is lazy initialization. This can be done by introducing an accessor method which, when executed, initializes the variable if it has not yet been initialized. This implies that we must use the accessor all the time and never use the class variable directly. This furthermore imposes the

cost of the accessor send and the initialization test. It also arguably defeats the point of using a class variable, since in fact it is no longer shared.

Method 5.21: *Color class»colorNames*

```
Color class»colorNames
  ColorNames ifNil: [self initializeNames].
  ↑ ColorNames
```

Another solution is to override the class method initialize.

Method 5.22: *Color class»initialize*

```
Color class»initialize
  ...
  self initializeNames
```

If you adopt this solution, you need to remember to invoke the initialize method after you define it, *e.g.*, by evaluating `Color initialize`. Although class side initialize methods are executed automatically when code is loaded into memory, they are *not* executed automatically when they are first typed into the browser and compiled, or when they are edited and re-compiled.

Pool Variables

Pool variables are variables that are shared between several classes that may not be related by inheritance. Pool variables were originally stored in pool dictionaries; now they should be defined as class variables of dedicated classes (subclasses of `SharedPool`). Our advice is to avoid them; you will need them only in rare and specific circumstances. Our goal here is therefore to explain pool variables just enough so that you can understand them when you are reading code.

A class that accesses a pool variable must mention the pool in its class definition. For example, the class `Text` indicates that it is using the pool dictionary `TextConstants`, which contains all the text constants such as `CR` and `LF`. This dictionary has a key `#CR` that is bound to the value `Character cr`, *i.e.*, the carriage return character.

Class 5.23: *Pool dictionaries in the Text class*

```
ArrayedCollection subclass: #Text
  instanceVariableNames: 'string runs'
  classVariableNames: ''
  poolDictionaries: 'TextConstants'
  category: 'Collections-Text'
```

This allows methods of the class `Text` to access the keys of the dictionary in the method body *directly*, *i.e.*, by using variable syntax rather than an explicit dictionary lookup. For example, we can write the following method.

Method 5.24: *Text»testCR*

```
Text»testCR
  ↑ CR == Character cr
```

Once again, we recommend that you avoid the use of pool variables and pool dictionaries.

5.8 Chapter Summary

The object model of Squeak is both simple and uniform. Everything is an object, and pretty much everything happens by message sends.

- Everything is an object. Primitive entities like integers are objects, but also classes are first-class objects.
- Every object is an instance of a class. Classes define the structure of their instances via *private* instance variables and the behavior of their instances via *public* methods. Each class is the unique instance of its metaclass. Class variables are private variables shared by the class and all the instances of the class. Classes cannot directly access instance variables of their instances, and instances cannot access class variables of their class. Accessors must be defined if this is needed.
- Every class has a superclass. The root of the single inheritance hierarchy is `ProtoObject`. Classes you define, however, should normally inherit from `Object` or its subclasses. There is no syntax for defining abstract classes. An abstract class is simply a class with an abstract method — one whose implementation consists of the expression `self subclassResponsibility`. Although Squeak supports only single inheritance, it is easy to share implementations of methods by packaging them as *traits*.
- Everything happens by message sends. We do not “call methods”, we “send messages”. The receiver then chooses its own method for responding to the message.
- Method lookup follows the inheritance chain; `self` sends are dynamic and start the method lookup again in the class of the receiver, whereas

super sends are static, and start in the superclass of class in which the super send is written.

- There are three kinds of shared variables. Global variables are accessible everywhere in the system. Class variables are shared between a class, its subclasses and its instances. Pool variables are shared between a select set of classes. You should avoid shared variables as much as possible.

Chapter 6

The Squeak Programming Environment

The goal of this chapter is to show you how to develop programs in the Squeak programming environment. You have already seen how to define methods and classes using the system browser; this chapter will show you more of the features of the system browser, and introduce you to some of the other browsers.

Of course, very occasionally you may find that your program does not work as you expect. Squeak has an excellent debugger, but like most powerful tools, it can be confusing on first use. We will walk you through a debugging session and demonstrate some of the features of the debugger.

One of the unique features of Smalltalk is that while you are programming, you are living in a world of live objects, not in a world of static program text. This makes it possible to get very rapid feedback while programming, which makes you more productive. There are two tools that let you look at, and indeed change, live objects: the *inspector* and the *explorer*.

The consequence of programming in a world of live objects rather than with files and a text editor is that you have to do something explicit to export your program from your Smalltalk image. The old way of doing this, also supported by all Smalltalk dialects, is by creating a *fileout* or a *change set*, which are essentially encoded text files that can be imported into another system. The new way of doing this in Squeak is to upload your code to a versioned repository on a server. This is done using a tool called Monticello, and is a much more powerful and effective way to work, especially when working in a team.

Finally, you may find a bug in Squeak as you work; we explain how to

report bugs, and how to submit bug fixes.

6.1 Overview

Smalltalk and modern graphical interfaces were developed together. Even before the first public release of Smalltalk in 1983, Smalltalk had a self-hosting graphical development environment, and all Smalltalk development was taking place in it. Lets start by looking at the main tools in Squeak, all of which can be dragged out of the *Tools* flap in the *Squeak-dev* image (see Section 1.1). Depending on your personal settings, the *Tools* flap can be opened by mousing over or clicking on the orange tab on the right-hand edge of Squeak's main window.

- The **Browser** is the central development tool. You will use it to create, define, and organize your classes and methods. Using it you can also navigate through all the library classes: unlike other environments where the source code is stored in separate files, in Smalltalk all classes and methods are contained in the image.
- The **Message Names** tool is used to look at all of the methods with a particular selector, or with a selector containing a substring.
- The **Method Finder** tool will also let you find methods, but according to what they *do* as well as what they are called.
- The **Monticello Browser** is the starting point for loading code from, and saving code in, Monticello packages.
- The **Process Browser** provides a view on all of the processes (threads) executing in Smalltalk.
- The **Test Runner** lets you run and debug SUnit tests, and is described in Chapter 7.
- The **Transcript** is a window on the Transcript output stream, which is useful for writing log messages and has already been described in Section 1.4.
- The **Workspace** is a window into which you can type input. It can be used for any purpose, but is most often used for typing Smalltalk expressions and executing them as **do** its. The use of the workspace was also illustrated in Section 1.4.

The **Debugger** has an obvious role, but you will discover that it has a more central place compared to debuggers for other programming languages, because in Smalltalk you can *program* in the debugger. The debugger is not launched from a menu or from the *Tools* flap; it is normally entered by running a failing test, by typing `CMD-.` to interrupt a running process, or by inserting a self halt expression in code.

6.2 The System Browser

There are actually several browsers in Squeak: the standard system browser, the package browser, the Omnibrowser, and the Refactoring Browser. We will take a look at the standard system browser first, since the others are variations on it. Figure 6.1 shows the browser as it appears when you first drag it out of the *Tools* flap.

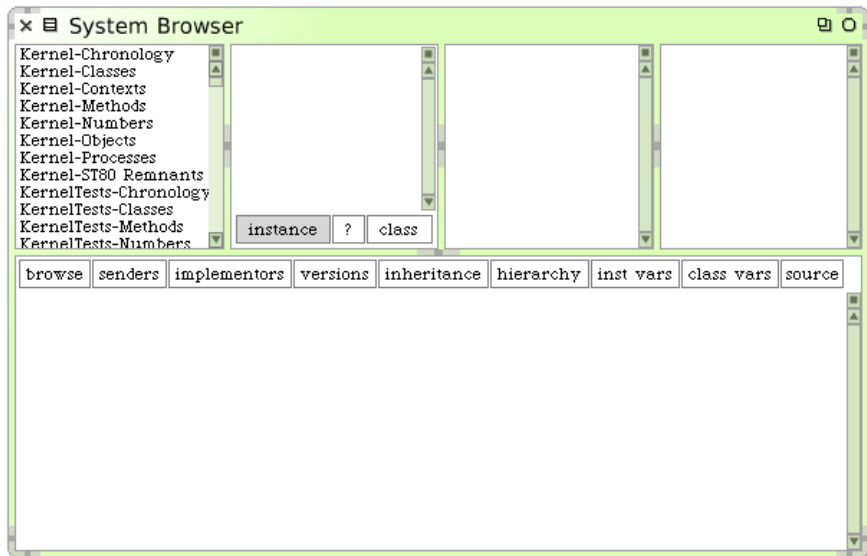


Figure 6.1: The System Browser

The four small panes at the top of the browser represent a hierarchic view of the methods in the system, much in the same way as the NeXTstep *File Viewer* and the Mac OS X *Finder* in column mode provide a view of the files on the disk. The leftmost pane lists *categories* of classes; select one (say *Kernel-Objects*) and the pane immediately to the right will then show all of the classes in that category.

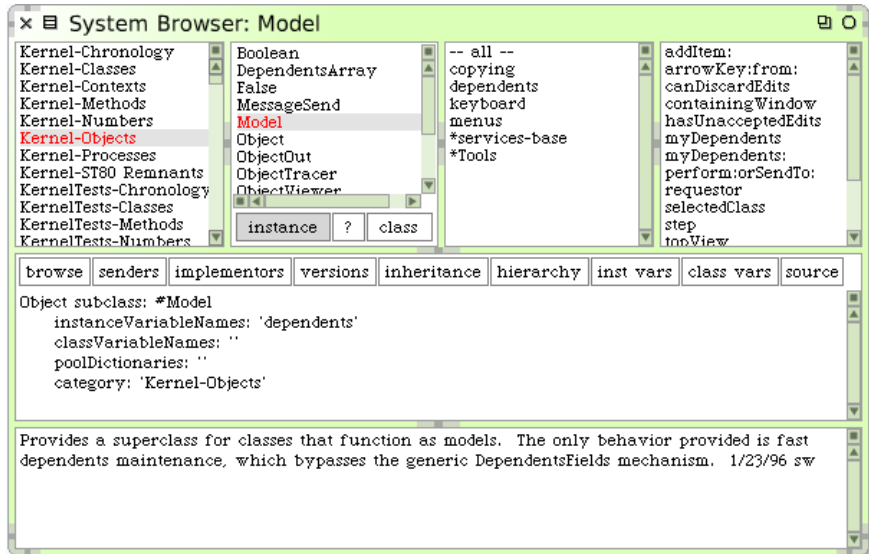


Figure 6.2: System Browser with the class Model selected

Similarly, if you select one of the classes in the second pane, say, `Model` (see Figure 6.2), the third pane will show all of the *protocols* defined for that class, as well as a virtual protocol `--all--`, which is selected by default. Protocols are a way of categorizing methods; they make it easier to find and think about the behaviour of a class by breaking it up into smaller, conceptually coherent pieces. The fourth pane shows the names of all of the methods defined in the selected protocol. If you then select a method name, the source code of the corresponding method appears in the large pane at the bottom of the browser, where you can view it, edit it, and save the edited version. If you select class `Model`, protocol `dependents` and the method `myDependents`, the browser should look like Figure 6.3.

Unlike directories in the Mac OS X *Finder*, the four top panes of the browser are not quite equal. Whereas classes and methods are part of the Smalltalk language, system categories and message protocols are not: they are a convenience introduced by the browser to limit the amount of information that needs to be shown in each pane. For example, if there were no protocols, the browser would have to show a list of all of the methods in the selected class; for many classes this list would be too large to navigate conveniently.

Because of this, the way that you create a new category or a new protocol is different from the way that you create a new class or a new method. To

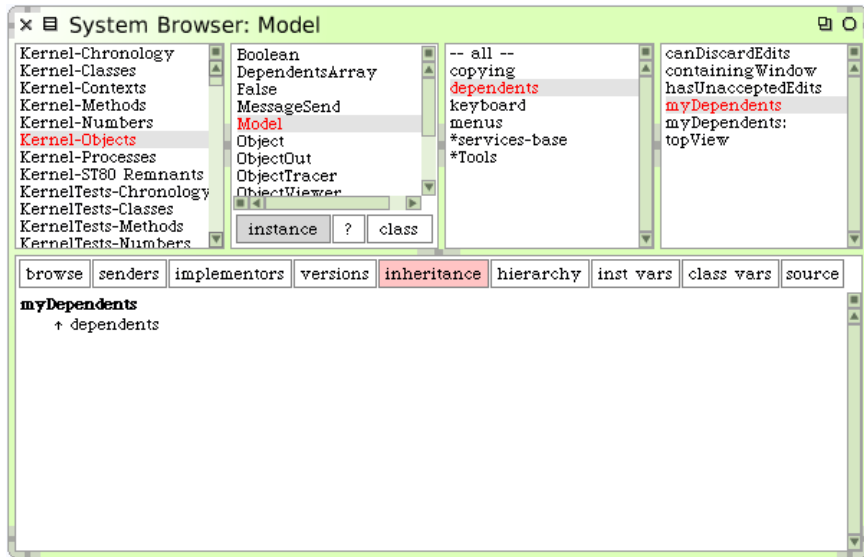


Figure 6.3: System Browser showing the myDependents method in class Model

create a new category, select **new category** from the yellow button menu in the category pane; to create a new protocol, select **new protocol** from the yellow button menu in the protocol pane. Enter the name of the new thing in the dialog, and you are done: there is nothing more to a category or a protocol than its name and its contents.

In contrast, to create a new class or a new method, you will actually have to write some Smalltalk code. If you deselect the currently selected class, so that no class is selected at all, the main browser pane will display a class creation template (Figure 6.4). You create a new class by editing this template: replace **Object** by the name of the existing class of which you wish to create a subclass, replace **NameOfSubclass** by the name that you would like to give to your new subclass, and fill in the instance variable names if you know them. The category for the new class is by default the currently selected category, but you can change this too if you like. If you already have the browser focussed on the class that you wish to subclass, you can get the same template with slightly different initialization by using the yellow button menu in the class pane, and selecting **more ... ▸ subclass template**. You can also just edit the definition of an existing class, changing the class name to something new. In all cases, when you accept the new definition, the new class (the one whose name follows the #) is created (as is the corresponding metaclass). Creating a class also creates a global variable that




Figure 6.4: System Browser showing the class-creation template

references the class, which is why you can refer to all of the existing classes by using their names.

Can you see why the name of the new class has to appear as a Symbol (*i.e.*, prefixed with #) in the class creation template, but after the class is created, code can refer to the class by using the name as an identifier (*i.e.*, without the #)?

The process of creating a new method is similar. First select the class in which you want the method to live, and then select a protocol. The browser will display a method-creation template, as shown in Figure 6.5, which you can fill-in or edit.

The Button Bar

The system browser provides several tools for exploring and analysing code. Those tools are most simply accessed from the horizontal button bar in the middle of the browser window. The buttons are labeled `browse`, `senders`, `implementors`... Figure 6.5 shows the complete set.

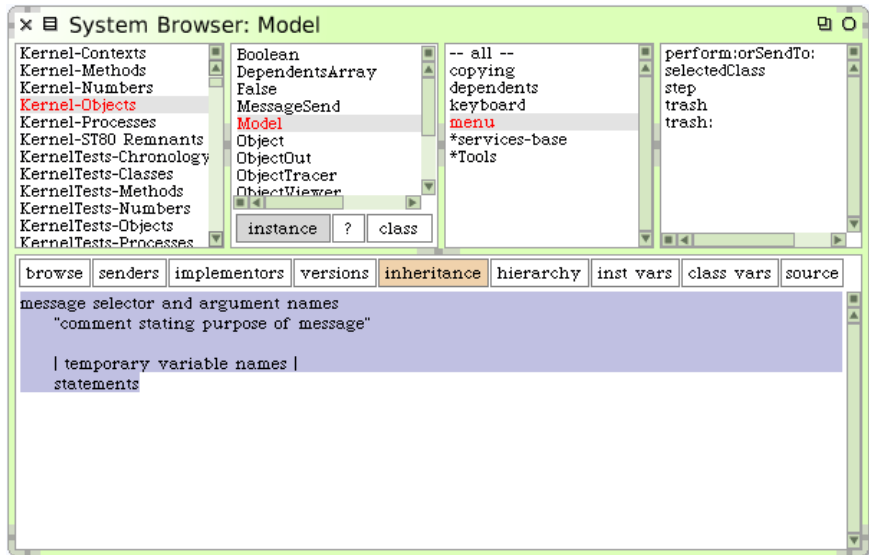



Figure 6.5: System Browser showing the method-creation template

Browsing Code

The **browse** button opens a new system browser on the class or method that is currently selected. It's often useful to have multiple browsers open at the same time. When you are writing code you will almost certainly need at least two: one for the method that you are typing, and another to browse around the system to see what to type. You can also open a browser on a class named by any selected text using the CMD-b keyboard shortcut.

 *Try this: in a workspace window, type the name of a class (for instance ScaleMorph), select it, and then press CMD-b. This trick is often useful; it works in any text window.*

Senders and Implementors of a Message

The **senders** button will give you a list of all methods that may use the selected method. With the browser open on ScaleMorph, click on the checkExtent: method in the method pane near the top right corner of the browser; the body of checkExtent: displays in the bottom part of the browser. If you now press the **senders** button, a menu will appear with checkExtent: as the topmost item, and below it, all the messages that checkExtent: sends (see Figure 6.6). Selecting an item in this menu will open a browser with

the list of all methods in the image that send the selected message.

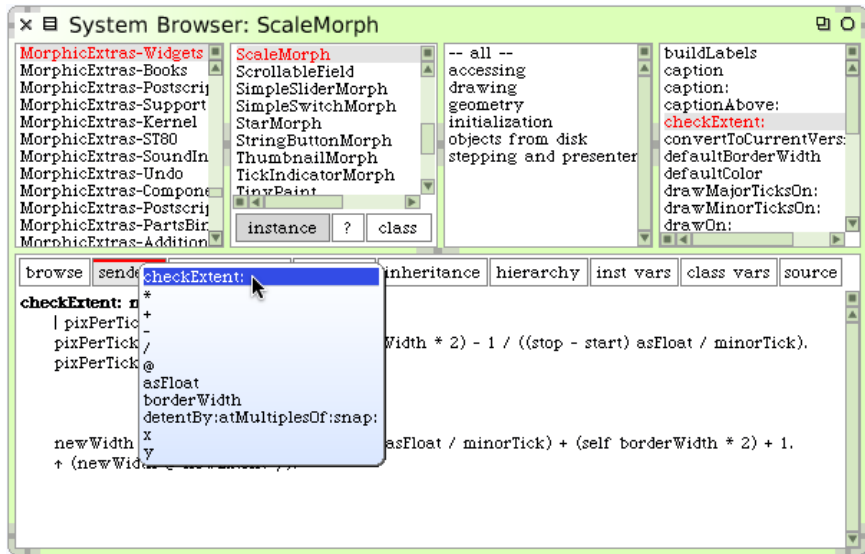


Figure 6.6: A Class Browser opened on the `ScaleMorph` class. Note the horizontal bar of buttons at the center of the browser; here we are using the `senders` button.

The `implementors` button works in a similar way, but instead of listing the senders of a message, it lists all of the classes that implement a method with the same selector. To see this, select `drawOn:` in the message pane and then bring up the “implementors of `drawOn:`” browser, either using the `implementors` button, or the yellow button menu on the method pane, or just by typing `CMD-m` (for `implementors`) in the method pane with `drawOn:` selected. You should get a method list window showing a scrolling list of 96 classes that implement a `drawOn:` method. It shouldn’t be all that surprising that so many classes implement this method: `drawOn:` is the message that is understood by every object that is capable of drawing itself on the screen. While viewing any one of these method, try browsing the senders of the `drawOn:` message: we found 63 methods that send this message. You can also bring up an implementors browser at any time by selecting a message (including the arguments if it is a keyword message) and typing `CMD-m`.

If you look at the send of `drawOn:` in `AtomMorph»drawOn:`, you will see that it is a super send. So we know that the method that will be executed will be in `AtomMorph`’s superclass. What class is that? Click the `hierarchy` button and you will see that it is `EllipseMorph`.

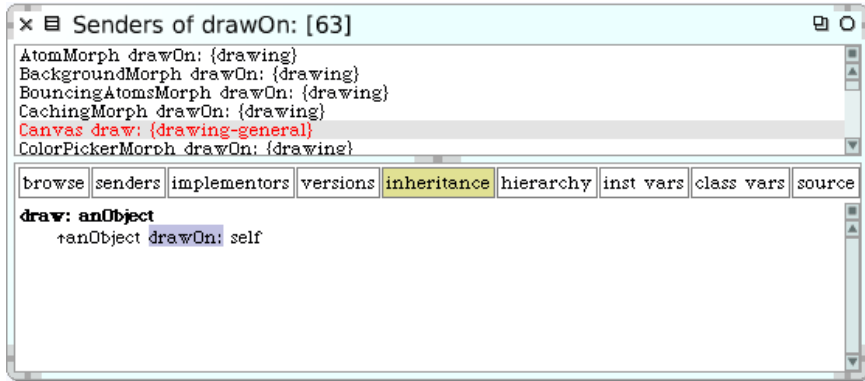


Figure 6.7: The Senders Browser showing that the Canvas»draw method sends the drawOn: message to its argument.

Now look at the fifth sender in the list, Canvas»draw, shown in Figure 6.7. You can see that this method sends drawOn: to whatever object is passed to it as argument, which could potentially be an instance of any class at all. Dataflow analysis can help figure out the class of the receiver of some messages, but in general, there is no simple way for the browser to know which message-sends might cause which methods to be executed. For this reason, the “senders” browser shows exactly what its name suggests: all of the senders of the message with the chosen selector. The `senders` button is nevertheless extremely useful when you need to understand how you can *use* a method: it lets you navigate quickly through example uses. Since all of the methods with the same selector ought to be used in the same way, all of the uses of a given message ought to be similar.

Versions of a Method

When you save a new version of a method, the old one is not lost. Squeak keeps all of the old versions, and allows you to compare different versions and to go back (“revert”) to an old version. The `versions` button gives access to the successive modifications made to the selected method. In Figure 6.8 we can see the versions of the mouseUp: method that one of the authors created while writing the Quinto game described in Chapter 2.

The top pane displays one line for each version of the method, listing the initials of the programmer who wrote it, the date and time at which it was saved, the names of the class and the method, and the protocol in which it was defined. The current (active) version is at the top of the list; whichever

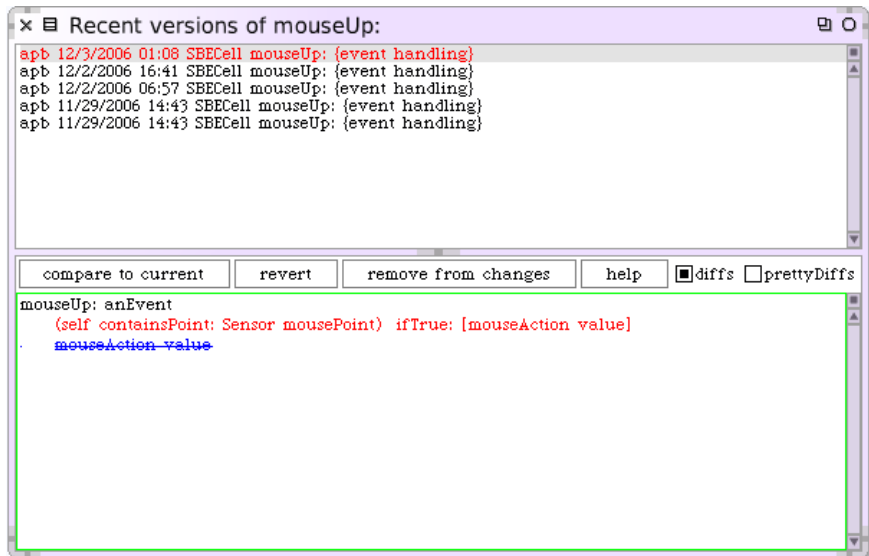


Figure 6.8: The versions browser showing several versions of the SBECcell» mouseUp: method

version is selected is displayed in the bottom pane. If the `diffs` checkbox is selected, as it is in Figure 6.8, the display also shows the differences between the selected version and the one immediately older. Buttons are also provided for displaying the differences between the selected method and the current version, and for reverting to the selected version. The `prettyDiffs` checkbox is useful if there have been changes to layout: it pretty-prints both versions before differencing, so that the differences that are displayed exclude formatting changes.

The existence of the versions browser means that you never have to worry about preserving code that you think might no longer be needed: just delete it. If you find that you *do* need it, you can always revert to the old version, or copy the needed code fragment out of the old version and paste it into a another method. Get into the habit of using versions; “commenting out” code that is no longer needed is a bad practice because it makes the current code harder to read. Smalltalkers rate code readability extremely highly.

HINT What if you delete a method entirely, and then decide that you want it back? You can find the deletion in a change set, where you can ask to see versions with the yellow button menu. The change set browser is described in Section 6.8

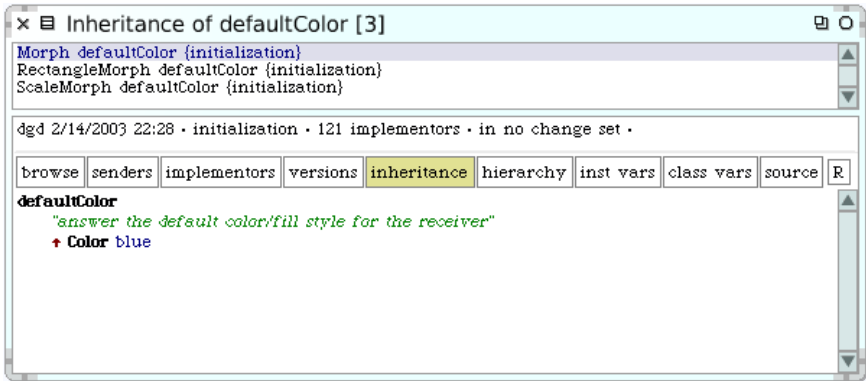


Figure 6.9: ScaleMorph»defaultColor and the methods that it overrides, in inheritance order. The `inheritance` button is gold because the displayed method is overridden in a subclass

Method Overridings

The `inheritance` button opens a specialized browser that displays all the methods overridden by the displayed method. To see how it works, display the ScaleMorph»defaultColor method and click `inheritance`. This method definition overrides RectangleMorph»defaultColor, which itself overrides Morph»defaultColor, as shown in Figure 6.9. The colour of the `inheritance` button depends on how the overriding occurs. The colours are explained in a help balloon:

- pink:* the displayed method overrides another method but doesn't use it;
- green:* the displayed method overrides another method and uses it via super;
- gold:* the displayed method is itself overridden in a subclass;
- salmon:* the displayed method overrides another method, and it itself overridden;
- violet:* the displayed method overrides, is overridden, and makes a super-send.

Note that there are currently two versions of the inheritance browser. If you are using the version of the system browser based on the OmniBrowser framework, the `inheritance` button does not change colour, and the inheritance browser looks different. It also displays more information: it shows not only the methods on the inheritance chain, but also their siblings. This version of the inheritance browser is shown in Figure 6.10.

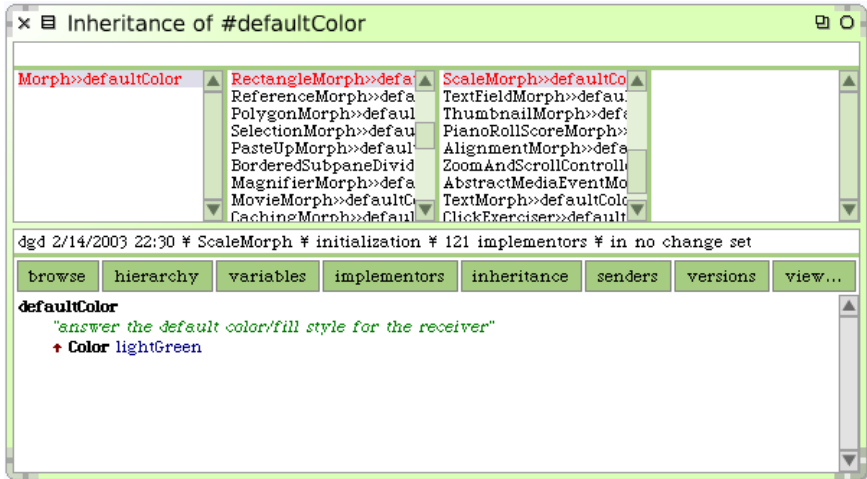


Figure 6.10: ScaleMorph>>defaultColor and the methods that it overrides, as shown by the new OmniBrowser based inheritance browser. The siblings of the selected methods are shown in the scrolling lists.

The Hierarchy Browser

The **hierarchy** button opens a hierarchy browser on the current class; this browser can also be opened by using the **browse hierarchy** menu item in the class pane. The hierarchy browser is similar to the system browser, but instead of displaying the system categories and the classes in each category, it shows a single list of classes, indented to represent inheritance. The category of the selected class is displayed in the small annotation pane at the top of the browser. The hierarchy browser is designed to make it easy to navigate through the inheritance hierarchy, but does not show all of the classes in the system: only the superclasses and subclasses of the initial class are shown. In Figure 6.11, the hierarchy browser reveals that the direct superclass of ScaleMorph is RectangleMorph.

Finding Variable References

The **inst vars** and **class vars** buttons help you find where an instance variable or a class variable is used; the same information is accessible from the yellow button menu items **inst var refs** and **class var refs** in the class pane. The menu also includes **inst var defs**, which shows the subset of the instance variable references that assign to the variable. Once you click on the button




Figure 6.11: A hierarchy browser open on ScaleMorph.

or select the menu item, you will be presented with a dialog that invites you to choose a variable from all of the variables defined in the current class, and all of the variables that it inherits. The list is in inheritance order; it can often be useful to bring up this list just to remind yourself of the name of an instance variable. If you click outside the list, it will go away and no variable browser will be created.

Also available from the yellow button menu on the class pane is **class vars**, which opens an inspector showing the class variables of the current class *and their values*, and **class refs (N)** which displays a list of all of the methods that directly reference the current class.

Source

The **source** button brings up the “what to show” menu, which allows you to choose what the browser shows in the source pane. Options include the **source** code, **prettyPrinted** source code, **byteCodes** and source code **decompiled** from the byte codes. The label on the button changes if you select one of the other modes. There are other options too; if you let the mouse linger over the names, a help balloon will appear. Try some of them.

Note that selecting `prettyPrint` in the “what to show” menu is *not* the same as prettyPrinting a method before you save it. The menu controls only what the browser displays, and has no effect on the code stored in the system. You can verify this by opening two browsers, and selecting `prettyPrint` in one and `source` in the other. In fact, focussing two browsers on the same method and selecting `byteCodes` in one and `decompile` in another is a good way to learn about the Squeak virtual machine’s byte-coded instruction set.

Refactoring

Did you notice the small `R` at the end of the button bar?¹ Although unobtrusive, this button gives access to one of the most powerful and important features of the Smalltalk environment. Clicking on `R` gives you access to a hierarchy of menus for refactoring your code. The same refactoring engine is also available in several other ways, for example, through the yellow button menu in the class, method and code panes. Refactoring was formerly available only in a special browser called the refactoring browser, but it can now be accessed from any browser.


The Browser Menus

Many additional functions are available from the browser’s yellow button menu. Since yellow button menus are context-sensitive, each pane in the browser has its own menu. Even if the labels on the menu items are the same, their *meaning* is context dependent. For example, the category pane, the class pane, the protocol pane and the messages pane all have a `file out` menu item. However, they do different things: the category pane’s `file out` menu files out the whole category, the class pane’s `file out` menu files-out the whole class, the protocol pane’s `file out` menu files out the whole protocol, and the method pane’s `file out` menu files-out just the displayed method. Although this may seem obvious, it can be a source of confusion for beginners.


Possibly the most useful menu item is `find class... (f)` in the category pane. Although the categories are useful for the code that we are actively developing, most of us do not know the categorization of the whole system, and it is much faster to type `CMD-f` followed by the first few characters of the name of a class than to guess which category it might be in. `recent classes... (r)` can also help you quickly go back to a class that you have browsed recently, even if you have forgotten its name.

¹That is, you will notice this button if you first load the AST and RefactoringBrowser packages from SqueakSource.

In the class pane, there are two menu items `find method` and `find method wildcard...` that can be useful if you want to browse a particular method. However, unless the list of methods is very long, it is often quicker to browse the `--all--` protocol (which is the default), place the mouse in the method pane, and type the first letter of the name of the method that you are looking for. This will usually scroll the pane so that the sought-for method name is visible.

 Try both ways of navigating to `OrderedCollection»removeAt:`

There are many other options available in the menus. It pays to spend a few minutes working with the browser and seeing what is there.

 Compare the result of `Browse Protocol`, `Browse Hierarchy`, and `Show Hierarchy` in the class pane menu.

Other Class Browsers

At the beginning of this section we mentioned another class browser: the *package pane browser*. This can be opened from the world menu: `World ▷ open... ▷ package pane browser`. It's basically the same as the class browser, but it knows about the naming convention for system categories. You will have noticed that the names of categories have two parts. For example, the `ScaleMorph` class belongs to the *Morphic-Widgets* category. The package browser assumes that the part before the hyphen, *Morphic* is the name of a "package", and adds a fifth pane that allows you to browse only those categories in a particular package. However, if you select no package at all, then all the categories are available, just as with the ordinary four-pane browser.


Unfortunately, the meaning of the term package has changed since the package pane browser was developed. "Package" now has a more precise meaning, related to the *Monticello* packaging tool, as we will discuss in the next section. At present there is no tool that lets you browse packages as defined by Monticello, however there is one currently being developed.

The Squeak community is in the process of developing a whole new family of browsers based on a new, highly customizable framework called the *OmniBrowser*. The implementation of the *OmniBrowser* is worth looking at as a good example of object-oriented design, but from the outside, most of the *OmniBrowser*-based tools look very similar to the ones that we have just described. The main enhancement that you will notice in the *Omni System Browser* is the addition of *virtual protocols*. In addition to the traditional programmer-defined protocols, each class has a number of virtual protocols defined on it by definitional rules. For example, the

--*supersend*-- protocol includes all of the methods that send to super, while the --*required*-- protocol lists all of the messages that are sent by the methods in the current class but not defined by it or by its superclasses.

Browsing Programmatically

The class `SystemNavigation` provides a number of utility methods that are useful for navigating around the system. Many of the functions offered by the classic browser are implemented by `SystemNavigation`.

 Open a workspace and **do it** the following code to browse the senders of `checkExtent::`

```
SystemNavigation default browseAllCallsOn: #checkExtent: .
```

To restrict the search for senders to the methods of a specific class:

```
SystemNavigation default browseAllCallsOn: #drawOn: from: ScaleMorph .
```

Because the development tools are objects, they are completely accessible from programs and you can develop your own tools or adapt the existing tools to your needs.

The programmatic equivalent to the `implementors` button is:

```
SystemNavigation default browseAllImplementorsOf: #checkExtent: .
```

To learn more about what is available, explore the class `SystemNavigation` with the browser. Further examples can be found in Chapter A.

Summary

As you have seen, there are many ways to navigate around Smalltalk code. You may find this confusing at first, in which case you can always fall back to the traditional system browser. However, we usually find that once beginners gain more experience with Squeak, the availability of different browsers becomes one of its most valued features, because they provide many ways to help you to understand and organize your code. The problem of understanding code is one of the greatest challenges of large-scale software development.

6.3 Monticello

We gave you a quick overview of Monticello, Squeak's packaging tool, in Section 2.9. However, Monticello has many more features than were discussed there. Because Monticello manages *Packages*, before telling you more about Monticello, it's important that we first explain exactly what a package is.

Packages: Declarative Categorization of Squeak Code

The package system is a simple, lightweight way of organizing Smalltalk source code. It leverages the long-used naming convention mentioned above (Section 6.2), but adds to it in an important way.

Let's explain this using an example. Suppose that you are developing a framework named to facilitate the use of relational databases from Squeak. You have decided to call your framework *SqueakLink*, and have created a series of system categories to contain all of the classes that you have written, *e.g.*,

Category 'SqueakLink-Connections' contains *OracleConnection MySQLConnection PostgresConnection*

Category 'SqueakLink-Model' contains *DBTable DBRow DBQuery*

and so on. However, not all of your code will reside in these classes. For example, you may also have a series of methods to convert objects into an SQL-friendly format:

```
Object»asSQL  
String»asSQL  
Date»asSQL
```

These methods belong in the same package as the classes in the categories *SqueakLink-Connections* and *SqueakLink-Model*. But clearly the whole of class *Object* does not belong in your package! So you need a way of putting certain *methods* in a one package, even though the rest of the class is in another package.


The way that you do this is by placing those methods in a protocol (of *Object*, *String*, *Date*, and so on) named **squeaklink* (note the initial star, and the lower-case name). The combination of the *SqueakLink-...* categories and the **squeaklink* protocols form a package named *SqueakLink*. To be precise, the rules for what goes in a package are as follows.

A package named *Foo* contains:

1. all class definitions of classes in the category *Foo*, or in categories with names starting with *Foo-*, and
2. all method definitions in any class in a protocol named **foo* or whose name starts with **foo-* (when performing this name comparison, the case of the letters in the names is ignored), and
3. all methods in classes in the category *Foo*, or in a category whose name starts with *Foo-*, *except* for those methods in protocols whose names start with ***.

A consequence of these rules is that each class definition and each method belongs to exactly one package. The *except* in the last rule has to be there because those methods must belong to other packages. The reason for ignoring case in rule 2 is that, by convention, protocol names are all lower case (and may include spaces), while category names use CamelCase (and don't include spaces).

The class `PackageInfo` implements these rules, and one way to get a feel for them is to experiment with this class.


 *Try this in your image, which should contain the classes `PackageInfo` and `RefactoringBrowser`.*

The `Refactoring Browser` code uses these package naming conventions, with `RefactoringEngine` as the package name. In a workspace, create a model of this package with

```
refactory := PackageInfo named: 'RefactoringEngine'.
```

It is now possible to introspect on this package. For example, `refactory` classes will return the long list of classes that make up the `Refactoring Engine` and the `Refactoring Browser`. `refactory coreMethods` will return a list of `MethodReferences` for all of the methods in those classes. `refactory extensionMethods` is perhaps one of the most interesting queries: it will return a list of all methods contained in the `RefactoringEngine` package but not contained within a `RefactoringEngine` class. This includes, for example, `ClassDescription>chooseThisClassInstVarThenDo:` and `SharedPool class>keys`.

Packages are a relatively new addition to Squeak, but since the package naming conventions were based on those used already in use, it is possible to use `PackageInfo` to analyze older code that has not been explicitly adapted to work with it.

 *Evaluate `(PackageInfo named: 'Collections') externalSubclasses`; this expression will answer a list of all subclasses of `Collection` that are not in the `Collections` package.*




Figure 6.12: The Monticello browser.

You can send `fileOut` to an instance of `PackageInfo` to get a change set of the entire package. For more sophisticated versioning of packages, we use Monticello.

Basic Monticello

Monticello is named after the mountaintop home of Thomas Jefferson, third president of the United States and author of the Statute of Virginia for Religious Freedom. The name means “little mountain” in Italian, and so it is always pronounced with an Italian “c”, which sounds like the “ch” in chair: Mont-y’-che-llo.

When you open the Monticello browser, you will see two list panes and a row of buttons, as shown in Figure 6.12. The left-hand pane lists all of the packages that have been loaded into the image that you are running; the particular version of the package is shown in parentheses after the name.

The right-hand pane lists all of the source-code repositories that Monticello knows about, usually because it has loaded code from them. If you select a package in the left pane, the right pane is filtered to show only those repositories that contain versions of the selected package.

One of the repositories is a directory named *package-cache*, which is a

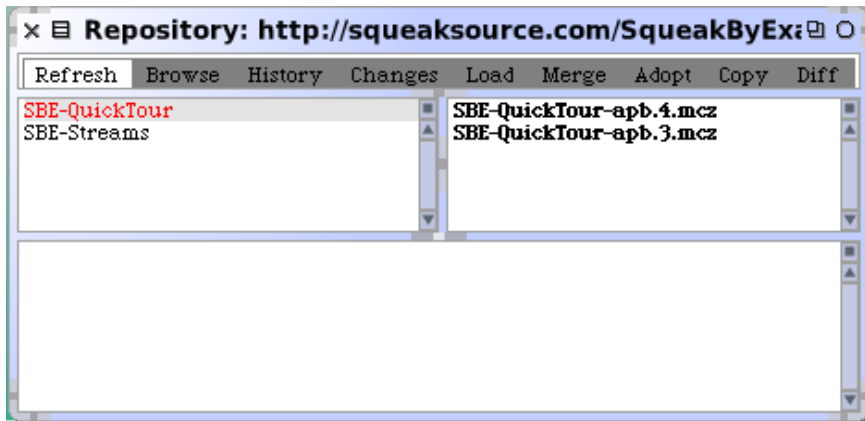



Figure 6.13: A Repository browser.

sub-directory of the directory in which your image is running. When you load code from or write code to a remote repository, a copy is also saved in the package cache. This can be useful if the network is not available and you need to access a package. Also, if you are given a Monticello (.mcz) file directly, for example as an email attachment, the most convenient way to access it is to place it in the package-cache directory.

To add a new repository to the list, click the `+Repository`, and choose the kind of repository from the pop-up menu. Let's add an HTTP repository.

 *Open Monticello, click on `+Repository`, and select `HTTP`. Edit the dialog to read:*

```

MCHttpRepository
  location: 'http://squeaksource.com/SqueakByExample'
  user: ""
  password: ""
  
```

Then click on `Open` to open a repository browser on this repository. You should see something like Figure 6.13. On the left is a list of all of the packages in the repository; if you select one, then the pane on the right will show all of the versions of the selected package in this repository.

If you select one of the versions, you can `Browse` it (without loading it into your image), `Load` it, or look at the `Changes` that will be made to your image by loading the selected version. You can also make a `Copy` of a version of a package, which you can then write to another repository.





Figure 6.14: Two classes in the “SBE” package.

As you can see, the names of versions contain the name of the package, the initials of the author of the version, and a version number. The version name is also the name of the file in the repository. Never change these names; correct operation of Monticello depends on them! Monticello version files are just zip archives, and if you are curious you can unpack them with a zip tool, but the best way to look at their contents is using Monticello itself.

To create a package with Monticello, you have to do two things: write some code, and tell Monticello about it.

 *Create a category called SBE-Monticello, and put a couple of classes in it, as shown in Figure 6.14. Also, create a method in an existing class, and put it in the same package as your classes, using the rules from page 125 — see Figure 6.15.*

To tell Monticello about your package, click on **+Package**, and type the name of the package, in this case “SBE”. Monticello will add SBE to its list of packages; the package entry will be marked with an asterisk to show that the version in the image has not yet been written to any repository.

Initially, the only repository associated with this package will be your package cache, as shown in Figure 6.16. That’s OK, you can still save the code, which will write it to the cache. Just click **Save** and you will be invited to provide a log message for the version of the package that you are about to save, as shown in Figure 6.17; when you accept the message, Monticello will save your package, and the asterisk decorating the name in

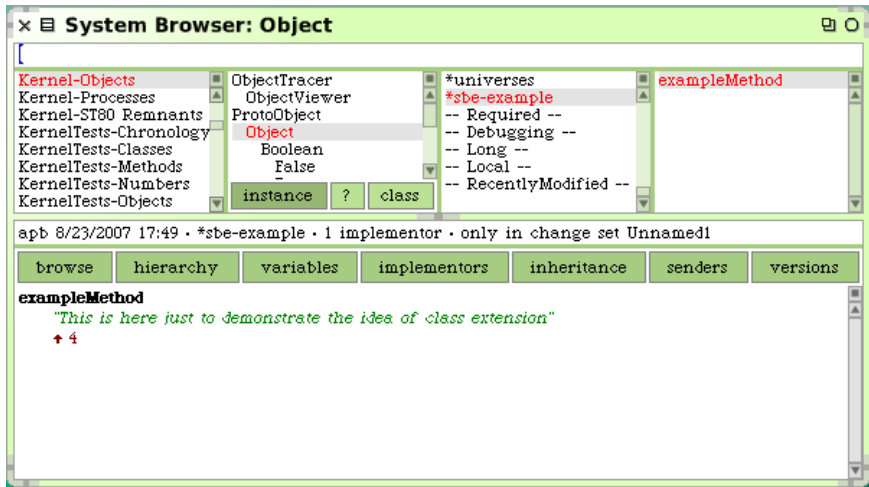


Figure 6.15: An extension method that will also be in the “SBE” package.

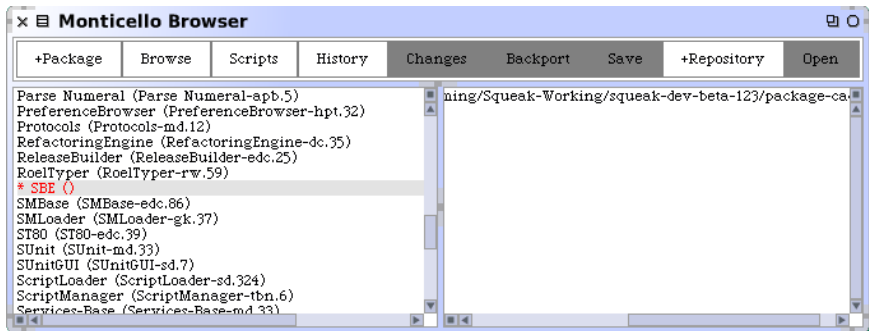


Figure 6.16: The as-yet-unsaved SBE package in Monticello.

Monticello’s package pane will be replaced by the version name.

If you then make a change to the package — say by adding a method to one of the classes — the asterisk will re-appear, showing that you have unsaved changes. If you open a repository browser on the package cache, you can select the saved version, and use **Changes** and the other buttons. You can of course save the new version to the repository too; once you **Refresh** the repository view, it should look like Figure 6.18.

To save our new package to another repository (other than the package cache), you need to first make sure that Monticello knows about the reposi-

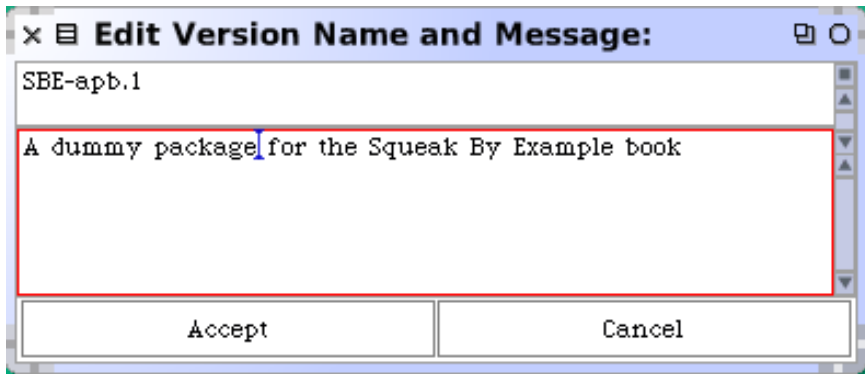


Figure 6.17: Providing a log message for a new version of a package.

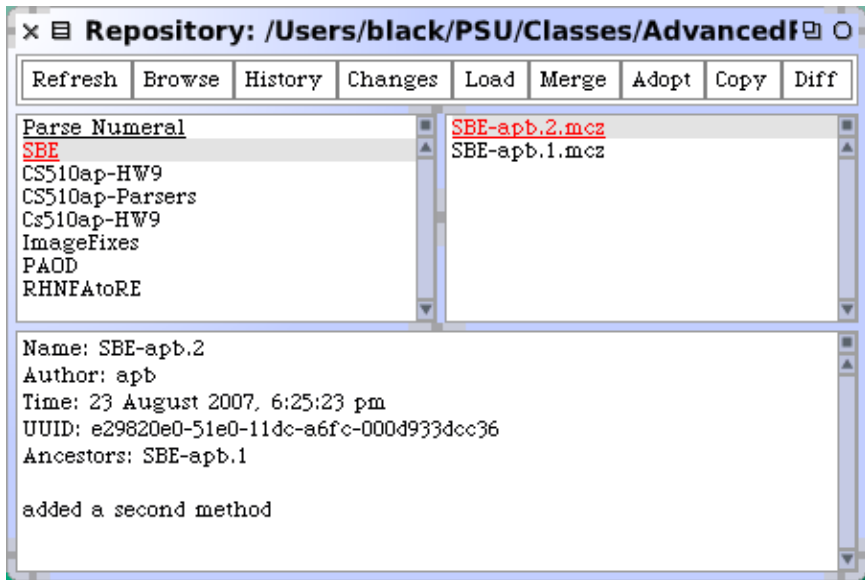


Figure 6.18: Two versions of our package are now in the package cache.

tory, adding it if necessary. Then you can use the **Copy** in the package-cache repository browser, and select the repository to which the package should be copied. You can also associate the desired repository with the package by using the yellow button menu item **add to package...** on the repository, as shown in Figure 6.19. Once the package knows about a repository, you can

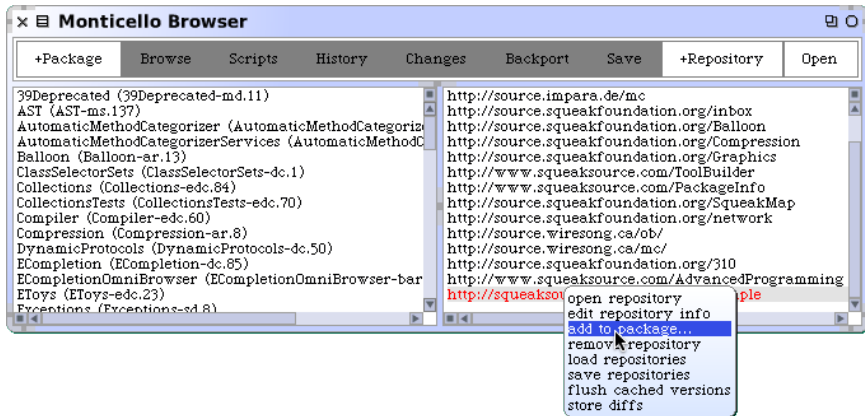


Figure 6.19: Adding a repository to the set of repositories associated with a package.

save a new version by selecting the repository and the package in the Monticello Browser, and clicking **Save**. Of course, you must have permission to write to a repository. The SqueakByExample repository on *SqueakSource* is world readable but not world writable, so if you try and save there, you will see an error message. However, you can create your own repository on *SqueakSource* by using the web interface at <http://www.squeaksource.com>, and use this to save your work. This is especially useful as a mechanism to share your code with friends, or if you use multiple computers.

If you do try and save to a repository where you don't have write permission, a version will nevertheless be written to the package-cache. So you can recover by editing the repository information (yellow button menu in the Monticello Browser) or choosing a different repository, and then using **Copy** from the package-cache browser.

6.4 The Inspector and Explorer

One of the things that makes Smalltalk so different from many other programming environments is that it provides you with a window onto a world of live objects, not a world of static code. Any of those objects can be examined by the programmer, and even changed — although some care is necessary when changing the basic objects that support the system. By all means experiment, but save your image first!





Figure 6.20: Inspecting TimeStamp now

The Inspector

 As an illustration of what you can do with an inspector, type `TimeStamp now` in a workspace, and then choose `inspect it` from the yellow button menu.

(It's not necessary to select the text before using the menu; if no text is selected, the menu operations work on the whole of the current line. You can also type `CMD-i` for `inspect it`.)

A window like that shown in Figure 6.20 will appear. This is an inspector, and can be thought of as a window onto the internals of a particular object—in this case, the particular instance of `TimeStamp` that was created when you evaluated the expression `TimeStamp now`. The title bar of the window shows the *class* of the object that is being inspected. If you select `self` at the top of the left pane, the right pane will show the `printstring` of the object. If you select `all inst vars` in the left pane, the right pane will show a list of the instance variables in the object, and the `printstring` for each one. The remaining items in the left pane represent the instance variables; this makes it easy to examine them one at a time, and also to change them.

The horizontal pane at the bottom of the inspector is a small workspace window. It is useful because in this window, the pseudo-variable `self` is bound to the very object that you are inspecting. So, if you `inspect it` on

`self` – `TimeStamp` today

in the workspace pane, the result will be a `Duration` object that represents the time interval between midnight today and the instant at which you evaluated `TimeStamp now` and created the `TimeStamp` object that you are in-




Figure 6.21: Exploring TimeStamp now

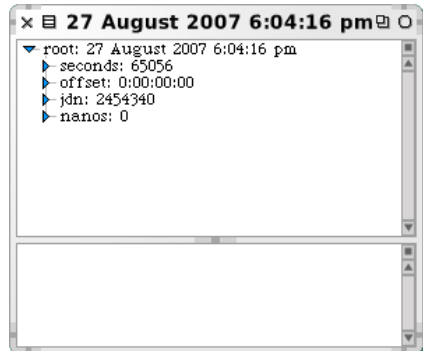


Figure 6.22: Exploring the instance variables

specting. You can also try evaluating `TimeStamp now - self`; this will tell you how long you have spent reading this section of this book!


In addition to `self`, all the instance variables of the object are in scope in the workspace pane, so you can use them in expressions or even assign to them. For example, if you evaluate `jdn := jdn - 1` in the workspace pane, you will see that the value of the `jdn` instance variable will indeed change, and the value of `TimeStamp now - self` will increase by one day.

You can change instance variables directly by selecting them, replacing the old value in the right-hand pane by a Squeak expression, and accepting. Squeak will evaluate the expression and assign the result to the instance variable.

There are special variants of the inspector for Dictionaries, OrderedCollections, CompiledMethods and a few other classes that make it easier to examine the contents of these special objects.

The Object Explorer

The *object explorer* is conceptually similar to the inspector, but presents its information in a different way. To see the difference, we'll *explore* the same object that we were just inspecting.



 Select `self` in the inspector's left-hand pane, and choose `explore (l)` from the yellow button menu.

The explorer window looks like Figure 6.21. If you click on the small triangle next to `root`, the view will change to Figure 6.22, which shows the

instance variables of object that you are exploring. Click on the triangle next to offset, and you will see *its* instance variables. The explorer is really useful when you need to explore a complex hierarchic structure—hence the name.

The workspace pane of the object explorer works slightly differently to that of the inspector. `self` is not bound to the root object, but rather to the object that is currently selected; the instance variables of the selected object are also in scope. Maybe explain what "self" refers too in the inspector (always the object) and explorer (the selected item) (p. 154)

To see the value of the explorer, let's use it to explore a deeply-nested structure of objects.

 Open a browser, and blue-click twice on the method pane to bring-up the Morphic halo on the `PluggableListMorph` that is used to represent the list of messages. Red-click on the debug handle  and select `explore morph` from the menu that appears. This will open an Explorer on the `PluggableListMorph` object that represents the method list on the screen. Open the root object (by clicking in its triangle), open its submorphs, and continue exploring the structure of the objects that underlie this Morph, as shown in Figure 6.23.

6.5 The Debugger

The debugger is arguably the most powerful tool in the Squeak tool suite. It is used not just for debugging, but also for writing new code. To demonstrate the debugger, let's start by writing a bug!

 Using the browser, add the following method to the class `String`:

Method 6.1: A buggy method

suffix

```
"assumes that I'm a file name, and answers my suffix, the part after the last dot"
| dot dotPosition |
dot := FileDirectory dot.
dotPosition := (self size to: 1 by: -1) detect: [ :i | (self at: i) = dot ].
↑ self copyFrom: dotPosition to: self size
```

Of course, we are sure that such a trivial method will work, so instead of writing an SUnit test, we just type 'readme.txt' suffix in a workspace and print it (p). What a surprise! Instead of getting the expected answer 'txt', a `PreDebugWindow` pops up, as shown in Figure 6.24.

The `PreDebugWindow` has a title-bar that tells us what error occurred, and shows us a *stack trace* of the messages that led up to the error. Starting

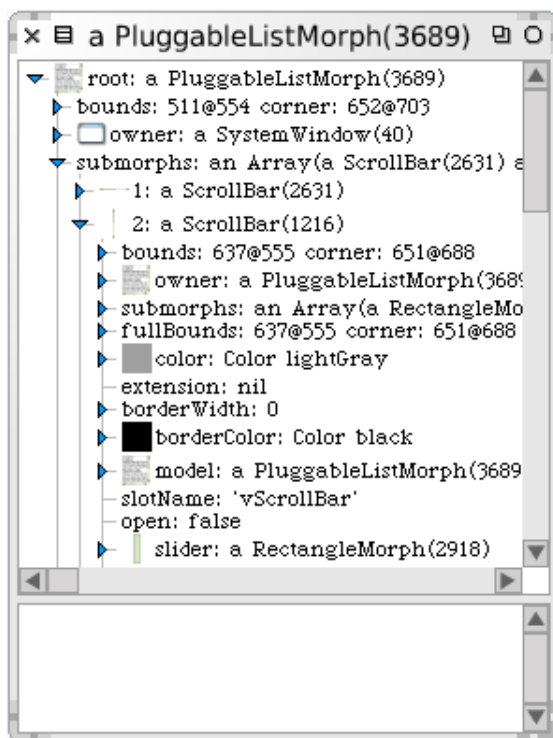


Figure 6.23: Exploring a PluggableListMorph

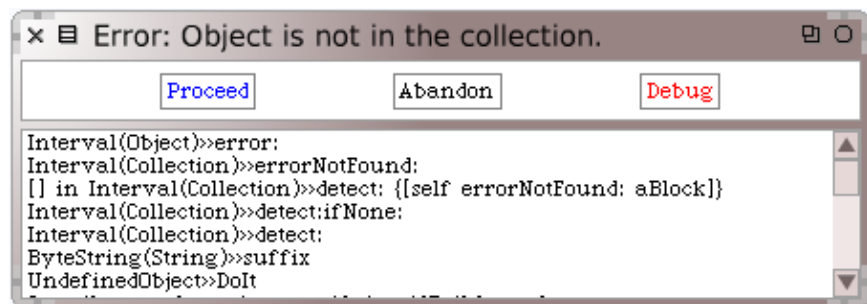



Figure 6.24: A PreDebugWindow notifies us of a bug.



Figure 6.25: The debugger.

from the bottom of the trace, `UndefinedObject>>DoIt` represents the code that was compiled and run when we selected 'readme.txt' suffix in the workspace and asked Squeak to `print it`. This code, of course, sent the message `suffix` to a `ByteString` object ('readme.txt'). This caused the inherited suffix method in class `String` to execute; all this information is encoded in the next line of the stack trace, `ByteString(String)>>suffix`. Working up the stack, we can see that `suffix` sent `detect:...` and eventually `detect;ifNone` sent `errorNotFound`.

To find out *why* the dot was not found, we need the debugger itself, so click on `Debug`.

 You can also open the debugger by red-clicking on any of the lines on the stack trace. If you do this, the debugger will open already focussed on the corresponding method.


The debugger is shown in Figure 6.25; it looks intimidating at first, but it is quite easy to use. The title-bar and the top pane are very similar to those that we saw in the `PreDebugWindow`. However, the debugger combines the stack trace with a method browser, so when you select a line in the stack

trace, the corresponding method is shown in the pane below. It's important to realize that the execution that caused the error is still in your image, but in a suspended state. Each line of the stack trace represents a frame on the execution stack that contains all of the information necessary to continue the execution. This includes all of the objects involved in the computation, with their instance variables, and all of the temporary variables of the executing methods.

In Figure 6.25 we have selected the `detect:ifNone:` method in the top pane. The method body is displayed in the center pane; the blue highlight around the message value shows that the current method has sent the message value and is waiting for an answer.


The four panes at the bottom of the debugger are really two mini-inspectors (without workspace panes). The inspector on the left shows the current object, that is, the object named `self` in the center pane. As you select different stack frames, the identity of `self` may change, and so will the contents of the `self`-inspector. If you click on `self` in the bottom-left pane, you will see that `self` is the interval (10 to: 1 by -1), which is what we expect. The workspace panes are not needed in the debugger's mini-inspectors because all of the variables are also in scope in the method pane; you should feel free to type or select expressions in this pane and evaluate them. You can always **cancel (l)** your changes using the menu or `CMD-L`.

The inspector on the right shows the temporary variables of the current context. In Figure 6.25, value was sent to the parameter `exceptionBlock`.

 *To see the current value of this parameter, click on `exceptionBlock` in the context inspector. This will tell you that `exceptionBlock` is `[self errorNotFound: ...]`. So, it is not surprising that we see the corresponding error message.*

Incidentally, if you want to open a full inspector window on one of the variables shown in the mini-inspectors, just double-click on the name of the variable, or select the name of the variable and choose **inspect (i)** or **explore (l)** from the yellow button menu. This can be useful if you want to want to watch how a variable changes while you execute other code.

Looking back at the method window, we see that we expected the penultimate line of the method to find `dot` in the string `'readme.txt'`, and that execution should never have reached the final line. Squeak does not let us run an execution backwards, but it does let us start a method again, which works very well in code such as this that does not mutate objects, but instead creates new ones.

 Click **Restart**, and you will see that the locus of execution returns to the first statement of the current method. The blue highlight shows that the next message to be sent will be `do:` (see Figure 6.26).

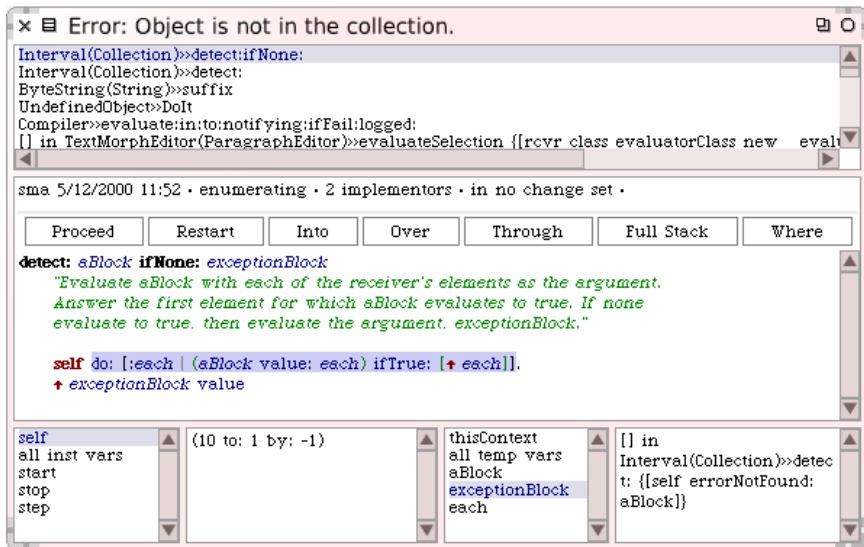





Figure 6.26: The debugger after restarting the detect: ifNone: method

The **Into** and **Over** buttons give us two different ways to step through the execution. If you click **Over**, Squeak executes the current message-send (in this case the do:) in one step, unless there is an error. So **Over** will take us to the next message-send in the current method, which is value: this is exactly where we started, and not much help. What we need to do is to find out why the do: is not finding the character that we are looking for.

 Click **Over**, and then click **Restart** to get back to the situation shown in Figure 6.26.

 Click **Into**; Squeak will go into the method corresponding to the highlighted message-send, in this case, Collection>>do:.

However, it turns out that this is not much help either: we can be fairly confident that Collection>>do: is not broken. The bug is much more likely to be in *what* we asked Squeak to do. **Through** is the appropriate button to use in this case: we want to ignore the details of the do: itself and focus on the execution of the argument block.

 Click on **Through** a few times. Select each in the context window as you do so. You should see each count down from 10 as the do: method executes.

When each is 7 we expect the ifTrue: block to be executed, but it isn't. To

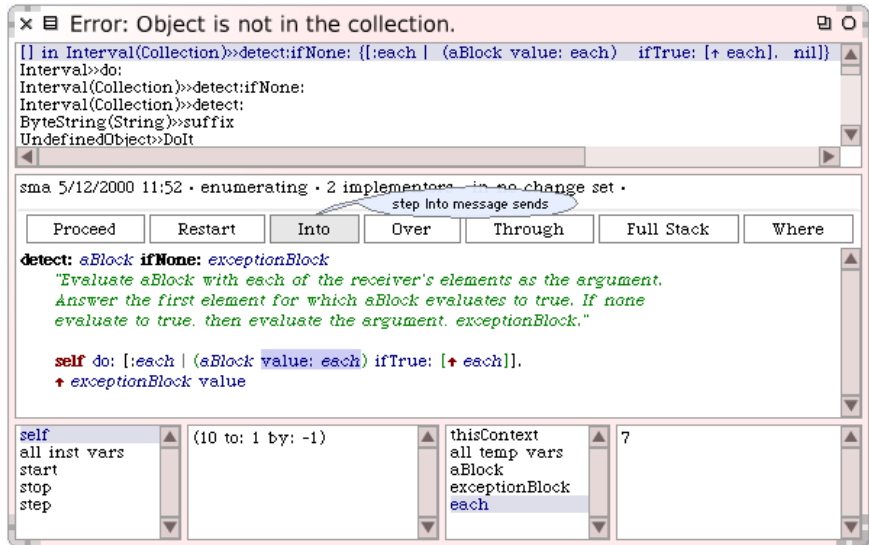



Figure 6.27: The debugger after stepping **Through** the `do:` method several times

see what is going wrong, go **Into** the execution of `value:` as illustrated in Figure 6.27.

After clicking **Into**, we find ourselves in the position shown in Figure 6.28. It looks at first that we have gone *back* to the `suffix` method, but this is because we are now executing the block that `suffix` provided as argument to `detect:`. If you select `i` in the context inspector, you can see its current value, which should be 7 if you have been following along. You can then select the corresponding element of `self` from the `self`-inspector. In Figure 6.28 you can see that element 7 of the string is character 46, which is indeed a dot. If you select `dot` in the context inspector, you will see that its value is `'.'`. And now you see why they are not equal: the seventh character of `'readme.txt'` is of course a `Character`, while `dot` is a `String`.

Now that we see the bug, the fix is obvious: we have to convert `dot` to a character before starting to search for it.

 *Change the code right in the debugger so that the assignment reads `dot := FileDirectory dot asCharacter` and **accept** the change.*

Because we are executing code inside a block that is inside a `detect:`, several stack frames will have to be abandoned in order to make this change. Squeak asks us if this is what we want (see Figure 6.29), and,

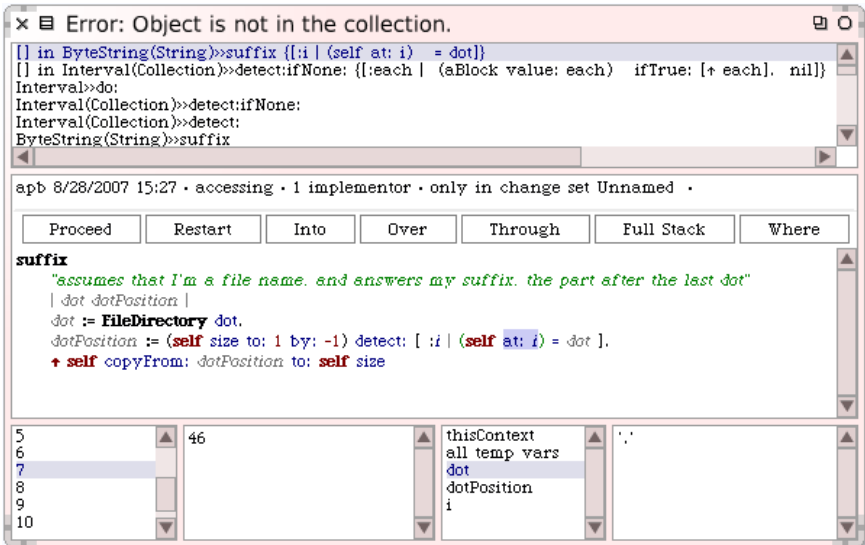


Figure 6.28: The debugger showing why 'readme.txt' at 7 is not equal to dot

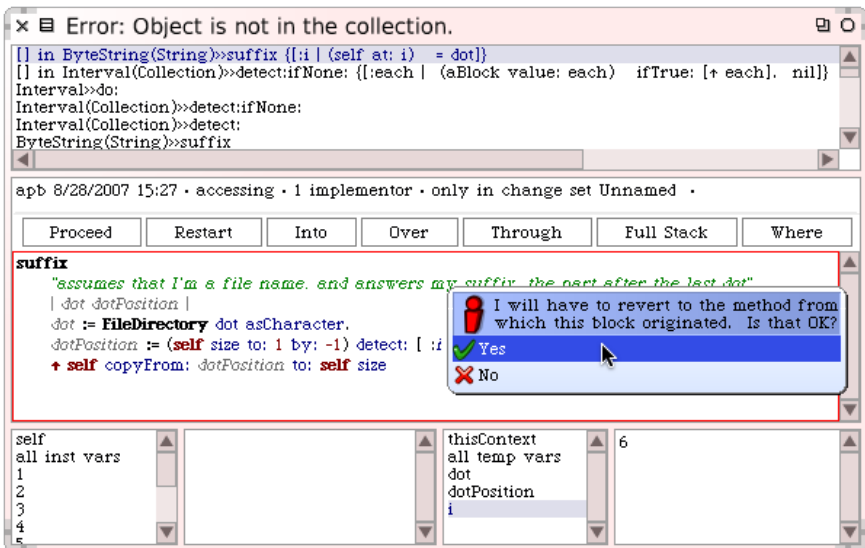



Figure 6.29: Changing the suffix method in the debugger: asking for confirmation of the exit from an inner block

assuming that we click **yes**, will save (and compile) the new method.

 Click **Restart** and then **Proceed**; the debugger window will vanish, and the evaluation of the expression 'readme.txt' suffix will complete, and print the answer 'txt'

Is the answer correct? Unfortunately, we can't say for sure. Should the suffix be .txt or txt? The method comment in suffix is not very precise. The way to avoid this sort of problem is to write an SUnit test that defines the answer.

Method 6.2: A simple test for the suffix method

```
testSuffixFound
  self assert: 'readme.txt' suffix = 'txt'
```

The effort required to do that was little more than to run the same test in the workspace, but using SUnit saves the test as executable documentation, and makes it easy for others to run. Moreover, if you add method 6.2 to the class StringTest and run that test suite with SUnit, you can very quickly get back to debugging the error. SUnit opens the debugger on the failing assertion, but you need only go back down the stack one frame, **Restart** the test and go **Into** the suffix method, and you can correct the error, as we are doing in Figure 6.30. It is then only the work of a second to click on the **Run Failures** button in the SUnit Test Runner, and confirm that the test now passes.

Here is a better test:

Method 6.3: A better test for the suffix method

```
testSuffixFound
  self assert: 'readme.txt' suffix = 'txt'.
  self assert: 'read.me.txt' suffix = 'txt'
```

Why is this test better? Because it tells the reader what the method should do if there is more than one dot in the target String.

There are a few other ways to get into the debugger in addition to catching errors and assertion failures. If you execute code that goes into an infinite loop, you can interrupt it and open a debugger on the computation by typing **CMD-.** (that's a full stop or a period, depending on where you learned English).² You can also just edit the suspect code to insert **self halt**. So, for example, we might edit the suffix method to read as follows:

²It is also useful to know that you can bring up an emergency debugger at any time by typing **CMD-SHIFT**.




Figure 6.30: Changing the suffix method in the debugger: fixing the off-by-one error after an SUnit assertion failure

Method 6.4: *Inserting a halt into the suffix method.*

suffix

```
"assumes that I'm a file name, and answers my suffix, the part after the last dot"
| dot dotPosition |
dot := FileDirectory dot asCharacter.
dotPosition := (self size to: 1 by: -1) detect: [ :i | (self at: i) = dot ].
self halt.
↑ self copyFrom: dotPosition to: self size
```

When we run this method, the execution of the `self halt` will bring up the pre-debugger, from where we can proceed, or go into the debugger and look at variables, step the computation, and edit the code.

That's all there is to the debugger, but it's not all there is to the `suffix` method. The initial bug should have made you realize that if there is no dot in the target string, the `suffix` method will raise an error. This isn't the behaviour that we want, so let's add a second test to specify what should happen in this case.

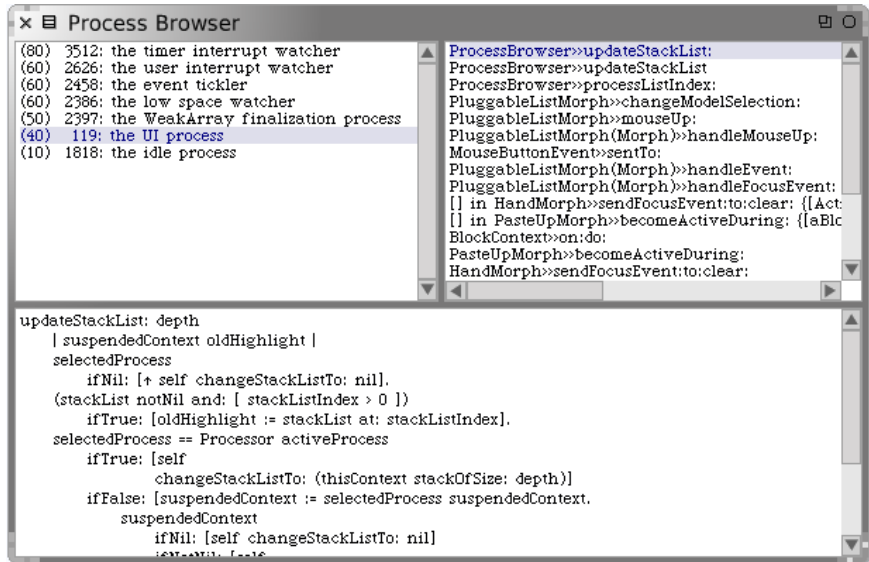



Figure 6.31: The Process Browser

Method 6.5: *A second test for the suffix method: the target has no suffix*

```
testSuffixNotFound
  self assert: 'readme' suffix = "
```

 Add method 6.5 to the test suite in class `StringTest`, and watch the test raise an error. Enter the debugger by selecting the erroneous test in `SUnit`, and edit the code so that the test passes. The easiest and clearest way to do this is to replace the `detect: message` by `detect: ifNone:`, where the second argument is a block that simply returns the empty string.

We will learn more about `SUnit` in Chapter 7.

6.6 The Process Browser

Smalltalk is a multi-threaded system: there are many lightweight processes (also known as threads) running concurrently in your image. In the future the Squeak virtual machine may take advantage of multiprocessors when they are available, but at present concurrency is implemented by time-slicing.

The process browser is a cousin of the debugger that lets you look at the various processes running inside Squeak. Figure 6.31 shows a screenshot. The top-left pane lists all of the processes in Squeak, in priority order, from the timer interrupt watcher at priority 80 to the idle process at priority 10. Of course, on a uniprocessor, the only process that can be running when you look is the UI process; all of the other will be waiting for some kind of event. By default, the display of processes is static; it can be updated using the yellow button menu, which also provides an option to turn on auto-update (a)

If you select a process in the top-left pane, its stack trace is displayed in the top-right pane, just as with the debugger. If you select a stack frame, the corresponding method is displayed in the bottom pane. The process browser is not equipped with mini-inspectors for self and thisContext, but yellow button menu items on the stack frames provide equivalent functionality.

6.7 Finding Methods

There are two tools in Squeak to help you find messages; both can be dragged out of the *Tools* flap. They differ in both interface and functionality.

The *method finder* was described at some length in Section 1.9; you can use it to find methods by name or by functionality. However, to look at the body of a method, the method finder opens a new browser. This can quickly become overwhelming.

The *message names* browser has more limited search functionality: you type a fragment of a message selector in the search box, and the browser lists all methods that contain that fragment in their names, as shown in Figure 6.32. However, it is a full-fledged browser: if you select one of the names in the left pane, all of the methods with that name are listed in the right pane, and can be browsed in the bottom pane. As with the class browser, the message names browser has a button bar that can be used to open other browsers on the selected method or its class.

6.8 Change Sets and the Change Sorter

Whenever you are working in Squeak, any changes that you make to methods and classes are recorded in a change set. This includes creating new classes, re-naming classes, changing categories, adding methods to existing classes — just about everything of significance. However, arbitrary *doits* are not included, so if, for example, you create a new global variable

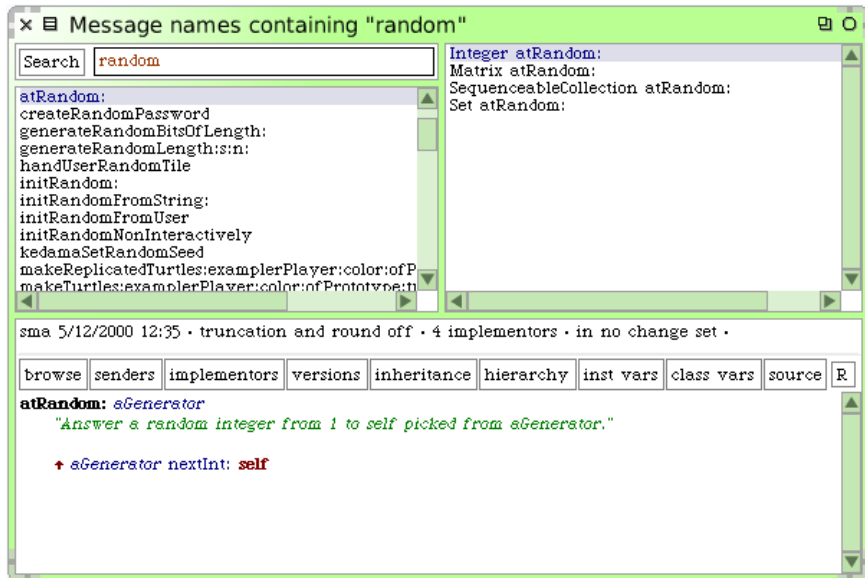


Figure 6.32: The method names browser showing all methods containing the substring random in their selectors.

by assigning to it in a workspace, the variable creation will not make it into a change set.

At any time, many change sets exist, but only one of them — `ChangeSet current` — is collecting the changes that are being made to the image. You can see which change set is current and can examine all of the change sets using the change set browser, available from `World > open... > simple change sorter` or by dragging the `Change Set` icon out of the *Tools* flap.

Figure 6.33 shows this browser. The title bar shows which change set is current, and this change set is selected when the browser opens.

Other change sets can be selected in the top-left pane; the yellow button menu allows you to make a different change set current, or to create a new change set. The top-right pane lists all of the classes affected by the selected change set (with their categories). Selecting one of the classes displays the names of those of its methods that are also in the change set (*not* all of the methods in the class) in the central pane, and selecting a method name displays the method definition in the bottom pane. Note that the browser does *not* show you whether the creation of the class itself is part of the change set, although this information is stored in the object structure that is used to represent the change set.

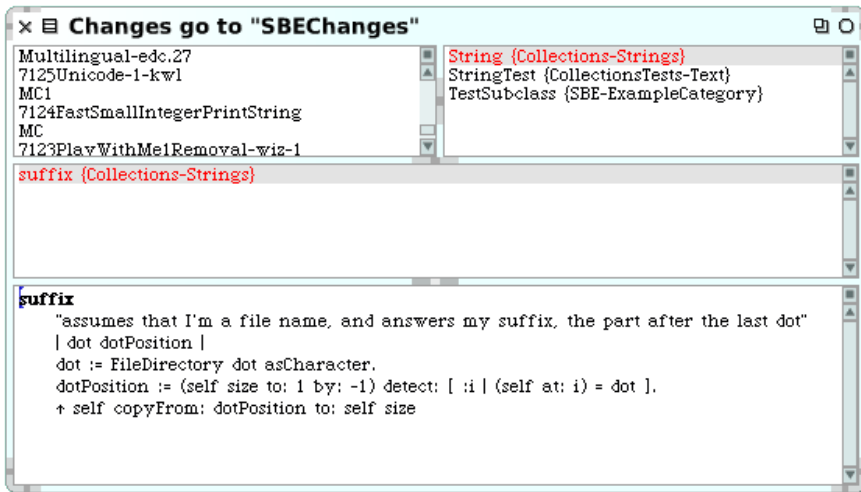


Figure 6.33: The Change Set Browser

The change set browser also lets you delete classes and methods from the change set using the yellow button menu on the corresponding items. However, for more elaborate editing of change sets, you should use a second tool, the *change sorter*, available under that name in the *Tools* flap or by selecting **World > open... > dual change sorter**, which is shown in Figure 6.34.

The change sorter is essentially two change set browsers side by side; each side can focus on a different change set, class, or method. This layout supports the change sorter's main feature, which is the ability to move or copy changes from one change set to another, as shown by the yellow button menu in Figure 6.34. It is also possible to copy individual methods from one side to the other.

You may be wondering why you should care about the composition of a change set. The answer is that change sets provide a simple mechanism for exporting code from Squeak to the file system, from where it can be imported into another Squeak image, or into another non-Squeak Smalltalk. Change set export is known as “filing-out”, and can be accomplished using the yellow button menu on any change set, class or method in either browser. Repeated file outs create new versions of the file, but change sets are not a versioning tool like Monticello: they do not keep track of dependencies.

Before the advent of Monticello, change sets were the main means for exchanging code between Squeakers. They have the advantage of simplicity

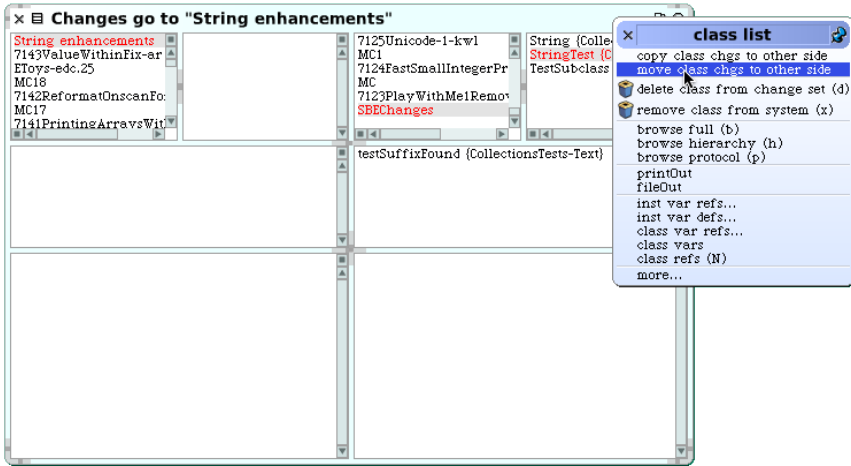


Figure 6.34: The Change Sorter

(the file out is just a text file, although we *don't* recommend that you try to edit with a text editor), and a degree of portability. It's also quite easy to create a change set that makes changes to many different, unrelated parts of the system — something that Monticello is not yet equipped to do.

The main drawback of change sets, compared to Monticello packages, is that they do not support the notion of dependencies. A filed-out change set is a set of *actions* that change any image into which it is loaded. To successfully load a change set requires that the image be in an appropriate state. For example, the change set might contain an action to add a method to a class; this can only be accomplished if the class is already defined in the image. Similarly, the change set might rename or re-categorize a class, which obviously will only work if the class is present in the image; methods may use instance variables that were declared when they were filed out, but which do not exist in the image into which they are imported. The problem is that change sets do not explicitly describe the conditions under which they can be filed in: the file in process just hopes for the best, usually resulting in a cryptic error message and a stack trace when things go wrong. Even if the file in works, one change set might silently undo a change made by another change set.

In contrast, Monticello packages represent code in a declarative fashion: they describe the state of the image should be in after they have been loaded. This permits Monticello to warn you about conflicts (when two packages require contradictory final states) and to offer to load a series of

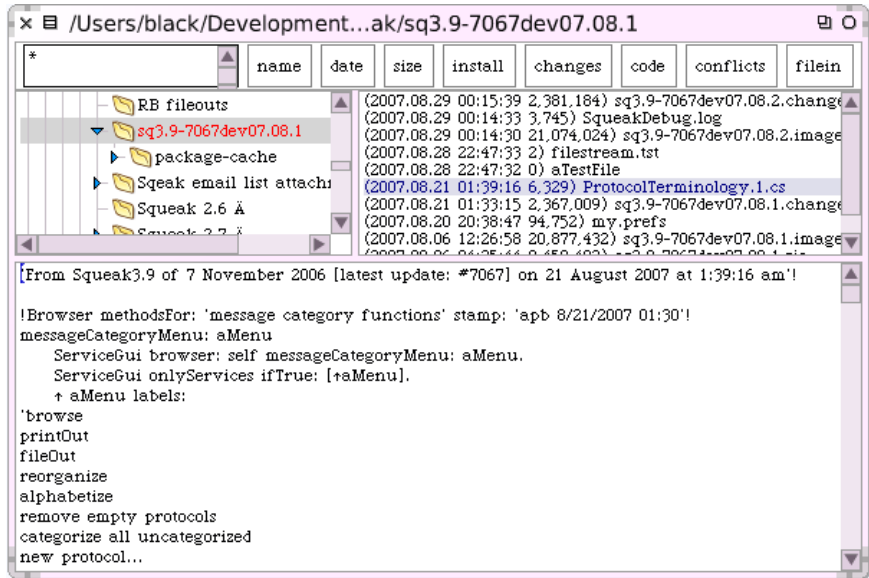


Figure 6.35: A file list browser

packages in dependency order.

In spite of these shortcomings, change sets still have their uses; in particular, you may find change sets on the Internet that you want to look at and perhaps use. So, having filed out a change set using the change sorter, we will now tell you how to file one in. This requires the use of another tool, the file list browser.

6.9 The File List Browser

The file list browser is in fact a general-purpose tool from browsing the file system (and also FTP servers) from Squeak. You can open it from the **World** ▸ **open...** ▸ **file list** menu, or drag it from the *Tools* flap. What you see of course depends on the contents of your local file system, but a typical view is shown in Figure 6.35.

When you first open a file list browser it will be focussed on the current directory, that is, the one from which you started Squeak. The title bar shows the path to this directory. The larger pane on the left-hand side can be used to navigate the file system in the conventional way. When a directory is selected, the files that it contains (but not the directories)

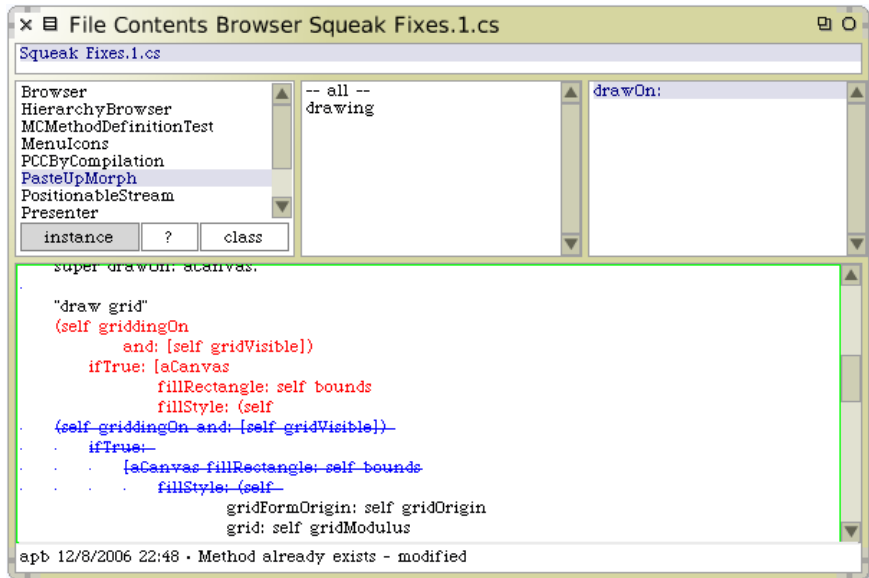


Figure 6.36: A File Contents Browser

are displayed on the right. This list of files can be filtered by entering a Unix-style pattern in the small box at the top-left of the window. Initially, this pattern is *, which matches all file names, but you can type a different string there and accept it, changing the pattern. (Note that a * is implicitly prepended and appended to the pattern that you type.) The sort order of the files can be changed using the `name`, `date` and `size` buttons. The rest of the buttons depend on the name of the file selected in the browser. In Figure 6.35, the file name has the suffix .cs, so the browser assumes that it is a change set, and provides buttons to `install` it (which *files it in* to a new change set whose name is derived from the name of the file), to browse the `changes` in the file, to examine the `code` in the file, and to `filein` the code into the *current* change set. You might think that the `conflicts` button would tell you about changes in the change set that conflicted with existing code in the image, but it doesn't. Instead it just checks for potential problems in the file that might indicate that the file cannot properly be loaded (such as the presence of linefeeds).

Because the choice of buttons to display depends on the file's *name*, and not on its contents, sometimes the button that you want won't be on the screen. However, the full set of options is always available from the yellow button `more...` menu, so you can easily work around this problem.

The `code` button is perhaps the most useful for working with change sets; it opens a browser on the contents of the change set file; an example is shown in Figure 6.36. The file contents browser is similar to the system browser except that it does not show categories, just classes, protocols and methods. For each class, the browser will tell you whether the class already exists in the system and whether it is defined in the file (but *not* whether the definitions are identical). It will show the the methods in each class, and (as shown in Figure 6.36) will show you the differences between the current version and the version in the file. Yellow-button menu items in each of the top four panes will also let you file in the whole of the change set, or the corresponding class, protocol or method.

6.10 In Smalltalk, You Can't Lose Code

It is quite possible to crash Squeak: as an experimental system, Squeak lets you change anything, including things that are vital to make Squeak work!

 *To maliciously crash Squeak, try* `Object become: nil`.

The good news is that you need never lose any work, even if you crash and go back to the last saved version of your image, which might be hours old. This is because all of the code that you executed is saved in the *.changes* file. All of it! This includes one liners that you evaluate in a workspace, as well as code that you add to a class while programming.

So here are the instructions on how to get your code back. There is no need to read this until you need it. However, when you do need it, you'll find it here waiting for you.

In the worst case, you can use a text editor on the *.changes* file, but since it is many megabytes in size, this can be slow and is not recommended. Squeak offers you better ways.

How to get your code back

Restart Squeak from the most recent snapshot, and select `World>help>useful expressions`. This will open a workspace full of useful expressions. The first three,

```
Smalltalk recover: 10000.
ChangeList browseRecentLog.
ChangeList browseRecent: 2000.
```

are most useful for recovery.

If you execute `ChangeList browseRecentLog`, you will be given the opportunity to decide how far back in history you wish to browse. Normally, it's sufficient to browse changes as far back as the last snapshot. (You can get much the same effect by editing `ChangeList browseRecent: 2000` so that the number 2000 becomes something else, using trial and error.)

One you have a *recent changes* browser, showing, say, changes back as far as your last snapshot, you will have a list of everything that you have done to Squeak during that time. You can delete items from this list using the yellow button menu. When you are satisfied, you can file-in what is left, thus incorporating the changes into your new image. It's a good idea to start a new change set, using the ordinary change set browser, before you do the file in, so that all of your recovered code will be in a new change set. You can then file out this change set.

One useful thing to do in the *recent changes* browser is to `remove doIts`. Usually, you won't want to file in (and thus re-execute) doIts. However, there is an exception. Creating a class shows up as a `dolt`. *Before you can file in the methods for a class, the class must exist.* So, if you have created any new classes, *first* file-in the class creation doIts, then `remove doIts` and file in the methods.

When I am finished with the recover, I like to file out my new change set, quit Squeak without saving the image, restart, and make sure that the new change set files back in cleanly.

6.11 Chapter Summary

In order to develop effectively with Squeak, it is important to invest some effort into learning the tools available in the environment.

- The standard *system browser* is your main interface for browsing existing class categories, classes, method protocols and methods, and for defining new ones. The system browser offers several useful buttons to directly jump to senders or implementors of a message, versions of a method, and so on.
- There exist several different class browsers (such as the *OmniBrowser* and the *Refactoring Browser*), and several specialized browsers (such as the *hierarchy browser*) which provide different views of classes and methods.
- From any of the tools, you can highlight the name of a class or a method and immediately jump to a browser by using the keyboard shortcut `CMD-b`.

- You can also browse the Smalltalk system programmatically by sending messages to `SystemNavigation` default.
- *Monticello* is a tool for exporting, importing, versioning and sharing packages of classes and methods. A Monticello package consists of a system category, subcategories, and related methods protocols in other categories.
- The *inspector* and the *explorer* are two tools that are useful for exploring and interacting with live objects in your image. You can even inspect tools by blue-clicking to bring up their morphic halo and selecting the debug handle (🔍).
- The *debugger* is a tool that not only lets you inspect the run-time stack of your program when an error is raised, but it also enables you to interact with all of the objects of your application, including the source code. In many cases you can modify your source code from the debugger and continue executing. The debugger is especially effective as a tool to support test-first development in tandem with SUnit (Chapter 7).
- The *process browser* lets you monitor, query and interact with the processes current running in your image.
- The *method finder* and the *message names browser* are two tools for locating methods. The first is more useful when you are not sure of the name, but you know the expected behavior. The second offers a more advanced browsing interface when you know at least a fragment of the name.
- *Change sets* are automatically generated logs of all changes to the source code of your image. They have largely been superseded by Monticello as a means to store and exchange versions of your source code, but are still useful, especially for recovering from catastrophic failures, however rare these may be.
- The *file list browser* is a tool for browsing the file system. It also allows you to `file in` source code from the file system.
- In case your image crashes before you could save it or backup your source code with Monticello, you can always recover your most recent changes using a *change list browser*. You can then select the changes you want to replay and file them into the most recent copy of your image.

Chapter 7

SUnit

7.1 Introduction

SUnit is a minimal yet powerful framework that supports the creation and deployment of tests. As might be guessed from its name, the design of SUnit focussed on *Unit Tests*, but in fact it can be used for integration tests and functional tests as well. SUnit was originally developed by Kent Beck and subsequently extended by Joseph Pelrine and others to incorporate the notion of a resource, which we will describe in Section 7.6.

The interest in testing and Test Driven Development is not limited to Squeak or Smalltalk. Automated testing has become a hallmark of the Agile software development movement, and any software developer concerned with improving software quality would do well to adopt it. Indeed, developers in many languages have come to appreciate the power of unit testing, and versions of *xUnit* now exist for many languages, including Java, Python, Perl, .Net and Oracle. This chapter describes SUnit 3.3 (the current version as of this writing); the official web site of SUnit is sunit.sourceforge.net, where updates can be found.

Neither testing, nor the building of test suites, is new: everybody knows that tests are a good way to catch errors. eXtreme Programming, by making testing a core practice and by emphasizing *automated* tests, has helped to make testing productive and fun, rather than a chore that programmers dislike. The Smalltalk community has a long tradition of testing because of the incremental style of development supported by its programming environment. In traditional Smalltalk development, the programmer would write tests in a workspace as soon as a method was finished. Sometimes a test would be incorporated as a comment at the head of the method that it

exercised, or tests that needed some set up would be included as example methods in the class. The problem with these practices is that tests in a workspace are not available to other programmers who modify the code; comments and example methods are better in this respect, but there is still no easy way to keep track of them and to run them automatically. Tests that are not run do not help you to find bugs! Moreover, an example method does not inform the reader of the expected result: you can run the example and see the — perhaps surprising — result, but you will not know if the observed behavior is correct.

SUnit is valuable because it allows us to write tests that are self-checking: the test itself defines what the correct result should be. It also helps us to organize tests into groups, to describe the context in which the tests must run, and to run a group of tests automatically. In less than two minutes you can write tests using SUnit, so instead of writing small code snippets in a workspace, we encourage you to use SUnit and get all the advantages of stored and automatically executable tests.

In this chapter we start by discussing why we test, and what makes a good test. We then present a series of small examples showing how to use SUnit. Finally, we look at the implementation of SUnit, so that you can understand how Smalltalk uses the power of reflection in supporting its tools.

7.2 Why Testing is Important

Unfortunately, many developers believe that tests are a waste of their time. After all, *they* do not write bugs — only *other* programmers do that. Most of us have said, at some time or other: “I would write tests if I had more time.” If you never write a bug, and if your code will never be changed in the future, then indeed tests are a waste of your time. However, this most likely also means that your application is trivial, or that it is not used by you or anyone else. Think of tests as an investment for the future: having a suite of tests will be quite useful now, but it will be *extremely* useful when your application, or the environment in which it executes, changes in the future.

Tests play several roles. First, they provide documentation of the functionality that they cover. Moreover, the documentation is active: watching the tests pass tells you that the documentation is up-to-date. Second, tests help developers to confirm that some changes that they have just made to a package have not broken anything else in the system — and to find the parts that break when that confidence turns out to be misplaced. Finally, writing tests at the same time as — or even before — programming forces

you to think about the functionality that you want to design, *and how it should appear to the client*, rather than about how to implement it. By writing the tests first—before the code—you are compelled to state the context in which your functionality will run, the way it will interact with the client code, and the expected results. Your code will improve: try it.

We cannot test all aspects of any realistic application. Covering a complete application is simply impossible and should not be the goal of testing. Even with a good test suite some bugs will still creep into the application, where they can lay dormant waiting for an opportunity to damage your system. If you find that this has happened, take advantage of it! As soon as you uncover the bug, write a test that exposes it, run the test, and watch it fail. Now you can start to fix the bug: the test will tell you when you are done.

7.3 What makes a Good Test?

Writing good tests is a skill that can be learned most easily by practicing. Let us look at the properties that tests should have to get a maximum benefit.

1. Tests should be repeatable. You should be able to run a test as often as you want, and always get the same answer.
2. Tests should run without human intervention. You should even be able to run them during the night.
3. Tests should tell a story. Each test should cover one aspect of a piece of code. A test should act as a scenario that you or some else can read to understand a piece of functionality.
4. Tests should have a change frequency lower than that of the functionality they cover: you do not want to have to change all your tests every time you modify your application. One way to achieve this is to write tests based on the public interfaces of the class that you are testing. It is OK to write a test for a private “helper” method if you feel that the method is complicated enough to need the test, but you should be aware that such a test may have to be changed, or thrown away entirely, when you think of a better implementation.

A consequence of property (3) is that the number of tests should be somewhat proportional to the number of functions to be tested: changing one aspect of the system should not break all the tests but only a limited number. This is important because having 100 tests fail should send a much


stronger message than having 10 tests fail. However, it is not always possible to achieve this ideal: in particular, if a change breaks the initialization of an object, or the set-up of a test, it is likely to cause all of the tests to fail.

eXtreme Programming advocates writing tests before writing code. This may seem to go against our deep instincts as software developers. All we can say is: go ahead and try it. We have found that writing the tests before the code helps us to know what we want to code, helps us know when we are done, and helps us conceptualize the functionality of a class and to design its interface. Moreover, test-first development gives us the courage to go fast, because we are not afraid that we will forget something important.

7.4 SUnit by Example

Before going into the details of SUnit, we will show a step by step example. We use an example that tests the class `Set`. Try entering the code as we go along.

Step 1: Create the test class

 First you should create a new subclass of `TestCase` called `ExampleSetTest`. Add two instance variables so that your new class looks like this:

Class 7.1: An Example Set Test class


```
TestCase subclass: #ExampleSetTest
  instanceVariableNames: 'full empty'
  classInstanceVariableNames: ''
  category: 'MyTest'
```

We will use the class `ExampleSetTest` to group all the tests related to the class `Set`. It defines the context in which the tests will run. Here the context is described by the two instance variables `full` and `empty` that we will use to represent a full and an empty set.

The name of the class is not critical, but by convention it should end in `Test`. If you define a class called `Pattern` and call the corresponding test class `PatternTest`, the two classes will be alphabetized together in the browser (assuming that they are in the same category). It is critical that your class be a subclass of `TestCase`.

Step 2: Initialize the test context

The method `setUp` defines the context in which the tests will run, a bit like an initialize method. `setUp` is invoked before the execution of each test method defined in the test class.

 Define the `setUp` method as follows, to initialize the empty variable to refer to an empty set and the full variable to refer to a set containing two elements.

Method 7.2: Setting up a fixture

```
ExampleSetTest»setUp
  empty := Set new.
  full := Set with: 5 with: 6
```

In testing jargon the context is called the *fixture* for the test.

Step 3: Write some test methods

Let's create some tests by defining some methods in the class `ExampleSetTest`. Each method represents one test; the name of the method should start with the string 'test' so that SUnit will collect them into test suites. Test methods take no arguments.

 Define the following test methods.

The first test, named `testIncludes`, tests the `includes:` method of `Set`. The test says that sending the message `includes: 5` to a set containing 5 should return `true`. Clearly, this test relies on the fact that the `setUp` method has already run.

Method 7.3: Testing set membership

```
ExampleSetTest»testIncludes
  self assert: (full includes: 5).
  self assert: (full includes: 6)
```

The second test, named `testOccurrences`, verifies that the number of occurrences of 5 in full set is equal to one, even if we add another element 5 to the set.

Method 7.4: *Testing occurrences*

```
ExampleSetTest»testOccurrences
  self assert: (empty occurrencesOf: 0) = 0.
  self assert: (full occurrencesOf: 5) = 1.
  full add: 5.
  self assert: (full occurrencesOf: 5) = 1
```

Finally, we test that the set no longer contains the element 5 after we have removed it.


Method 7.5: *Testing removal*

```
ExampleSetTest»testRemove
  full remove: 5.
  self assert: (full includes: 6).
  self deny: (full includes: 5)
```

Note the use of the method `deny:` to assert something that should not be true. `aTest deny: anExpression` is equivalent to `aTest assert: anExpression not`, but is much more readable.

Step 4: Run the tests

The easiest way to execute the tests is with the SUnit *Test Runner*, which you can open from the **World ▸ open...** menu, or by dragging the **TestRunner** from the *Tools* flap. The *TestRunner*, shown in Figure 7.1, is designed to make it easy to execute groups of tests. The left-most pane lists all of the system categories that contain test classes (*i.e.*, subclasses of `TestCase`); when some of these categories are selected, the test classes that they contain appear in the pane to the right. Abstract classes are italicized, and the test class hierarchy is shown by indentation, so subclasses of `ClassTestCase` are indented more than subclasses of `TestCase`.

 Open a *Test Runner*, select the category **MyTest**, and click the **Run Selected** button.

You can also run your test by executing a **print it** on the following code: (ExampleSetTest selector: `#testRemove`) run. This expression is equivalent to the shorter `ExampleSetTest run: #testRemove`. We usually include an executable comment in our test methods that allows us to run them with a **do it** from the browser, as shown in method 7.6.

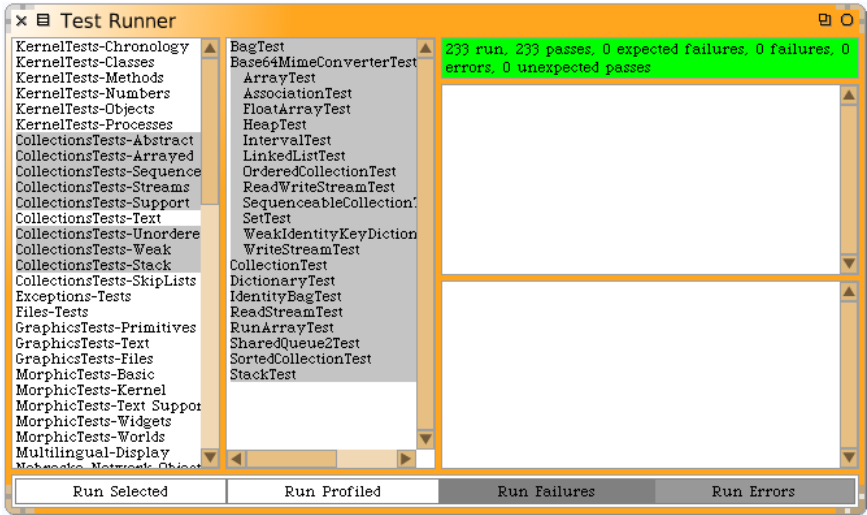



Figure 7.1: The Squeak SUnit Test Runner

Method 7.6: Executable comments in test methods

```
ExampleSetTest>testRemove
"self run: #testRemove"
full remove: 5.
self assert: (full includes: 6).
self deny: (full includes: 5)
```

 Introduce a bug in ExampleSetTest>testRemove and run the tests again. For example, change 5 to 4.

The tests that did not pass (if any) are listed in the right-hand panes of the *Test Runner*; if you want to debug one, to see why it failed, just click on the name. Alternatively, you can execute the following expressions:

```
(ExampleSetTest selector: #testRemove) debug
```


or

```
ExampleSetTest debug: #testRemove
```

Step 5: Interpret the Results

The method `assert:`, which is defined in the class `TestCase`, expects a boolean argument, usually the value of a tested expression. When the argument is true, the test passes; when the argument is false, the test fails.

There are actually three possible outcomes of a test. The outcome that we hope for is that all of the assertions in the test are true, in which case the test passes. In the test runner, when all of the tests pass, the bar at the top turns green. However, there are also two kinds of thing that can go wrong when you run a test. Most obviously, one of the assertions can be false, causing the test to *fail*. However, it is also possible that some kind of error occurs during the execution of the test, such as a *message not understood* error or an *index out of bounds* error. If an error occurs, the assertions in the test method may not have been executed at all, so we can't say that the test has failed; nevertheless, something is clearly wrong! In the test runner, failing tests cause the bar at the top to turn yellow, and are listed in the middle pane on the right, whereas erroneous tests cause the bar to turn red, and are listed in the bottom pane on the right.

 *Modify your tests to provoke both errors and failures.*

7.5 The SUnit Cook Book

This section will give you more details on how to use SUnit. If you have used another testing framework such as JUnit¹, much of this will be familiar, since all these frameworks have their roots in SUnit. Normally you will use SUnit's GUI to run tests, but there are situations where you may not want to use it.

Other assertions


In addition to `assert:` and `deny:`, there are several other methods that can be used to make assertions.

First, `assert:description:` and `deny:description:` take a second argument which is a message string that can be used to describe the reason for the failure, if it is not obvious from the test itself. These methods are described in Section 7.7.

Next, SUnit provides two additional methods, `should:raise:` and `shouldnt:raise:` for testing exception propagation. For example, you would

¹ junit.org

use (self should: aBlock raise: anException) to test that a particular exception is raised during the execution of aBlock. Method 7.7 illustrates the use of should:raise:.

 Try running this test.

Note that the first argument of the should: and shouldnt: methods is a *block* that *contains* the expression to be evaluated.


Method 7.7: Testing error raising

```
ExampleSetTest>testIllegal
  self should: [empty at: 5] raise: Error.
  self should: [empty at: 5 put: #zork] raise: Error
```

SUnit is portable: it can be used from all dialects of Smalltalk. To make SUnit portable, its developers factored-out the dialect-dependent aspects. The class method TestResult class>>error answers the system's error class in a dialect-independent fashion. You can take advantage of this: if you want to write tests that will work in any dialect of Smalltalk, instead of method 7.7 you would write:

Method 7.8: Portable error handling

```
ExampleSetTest>testIllegal
  self should: [empty at: 5] raise: TestResult error.
  self should: [empty at: 5 put: #zork] raise: TestResult error
```

 Give it a try.

Running a Single Test

Normally, you will run your tests using the Test Runner. If you don't want to launch the Test Runner from the `open...` menu or from the *Tools* flap, you can execute TestRunner open as a `print it`.

You can run a single test as follows.

```
ExampleSetTest run: #testRemove    —>    1 run, 1 passed, 0 failed, 0 errors
```

Running all the Tests in a Test Class

Any subclass of TestCase responds to the message suite, which will build a test suite that contains all the methods in the class whose names start with the string "test". To run the tests in the suite, send it the message run. For example:

ExampleSetTest suite run → 5 run, 5 passed, 0 failed, 0 errors

Must I Subclass TestCase?

In JUnit you can build a `TestSuite` from an arbitrary class containing `test*` methods. In Smalltalk you can do the same but you will then have to create a suite by hand and your class will have to implement all the essential `TestCase` methods like `assert:`. We recommend that you not try to do this. The framework is there: use it.

7.6 The SUnit Framework

SUnit consists of four main classes: `TestCase`, `TestSuite`, `TestResult`, and `TestResource`, as shown in Figure 7.2. The notion of a *test resource* was introduced in SUnit 3.1 to represent a resource that is expensive to set-up but which can be used by a whole series of tests. A `TestResource` specifies a `setUp` method that is executed just once before a suite of tests; this is in distinction to the `TestCase`»`setUp` method, which is executed before each test.




Figure 7.2: The four classes representing the core of SUnit

TestCase

`TestCase` is an abstract class that is designed to be subclassed; each of its subclasses represents a group of tests that share a common context (that

is, a test suite). Each test is run by creating a new instance of a subclass of `TestCase`, running `setUp`, running the test method itself, and then running `tearDown`.

The context is specified by instance variables of the subclass and by the specialization of the method `setUp`, which initializes those instance variables. Subclasses of `TestCase` can also override method `tearDown`, which is invoked after the execution of each test, and can be used to release any objects allocated during `setUp`.

TestSuite

Instances of the class `TestSuite` contain a collection of test cases. An instance of `TestSuite` contains tests, and other test suites. That is, a test suite contains sub-instances of `TestCase` and `TestSuite`. Both individual `TestCases` and `TestSuites` understand the same protocol, so they can be treated in the same way; for example, both can be run. This is in fact an application of the composite pattern in which `TestSuite` is the composite and the `TestCases` are the leaves — see *Design Patterns* for more information on this pattern².

TestResult

The class `TestResult` represents the results of a `TestSuite` execution. It records the number of tests passed, the number of tests failed, and the number of errors raised.

TestResource

One of the important features of a suite of tests is that they should be independent of each other: the failure of one test should not cause an avalanche of failures of other tests that depend upon it, nor should the order in which the tests are run matter. Performing `setUp` before each test and `tearDown` afterwards helps to reinforce this independence.

However, there are occasions where setting up the necessary context is just too time-consuming for it to be practical to do it once before the execution of each test. Moreover, if it is known that the test cases do not disrupt the resources used by the tests, then it is wasteful to set them up afresh for each test; it is sufficient to set them up once for each suite of tests. Suppose, for example, that a suite of tests need to query a database, or do some analysis on some compiled code. In such cases, it may make sense to

²Erich Gamma et al., *Design Patterns: Elements of Reusable Object-Oriented Software*. Reading, Mass.: Addison Wesley, 1995, ISBN 0-201-63361-2-(3).

set up the database and open a connection to it, or to compile some source code, before any of the tests start to run.

Where should we cache these resources, so that they can be shared by a suite of tests? The instance variables of a particular `TestCase` sub-instance won't do, because such an instance persists only for the duration of a single test. A global variable would work, but using too many global variables pollutes the name space, and the binding between the global and the tests that depend on it will not be explicit. A better solution is to put the necessary resources in a singleton object of some class. The class `TestResource` exists to be subclassed by such resource classes. Each subclass of `TestResource` understands the message `current`, which will answer a singleton instance of that subclass. Methods `setUp` and `tearDown` should be overridden in the subclass to ensure that the resource is initialized and finalized.

One thing remains: somehow, SUnit has to be told which resources are associated with which test suite. A resource is associated with a particular subclass of `TestCase` by overriding the `class` method `resources`. By default, the resources of a `TestSuite` are the union of the resources of the `TestCases` that it contains.

Here is an example. We define a subclass of `TestResource` called `MyTestResource` and we associate it with `MyTestCase` by specializing the class method `resources` to return an array of the test classes that it will use.

Class 7.9: An example of a `TestResource` subclass

```
TestResource subclass: #MyTestResource
    instanceVariableNames: "
```

```
MyTestCase class»resources
    "associate the resource with this class of test cases"
    ↑{ MyTestResource }
```

7.7 Advanced Features of SUnit

In addition to `TestResource`, the current version of SUnit contains assertion description strings, logging support, and resumable test failures.

Assertion Description Strings

The `TestCase` assertion protocol includes a number of methods that allow the programmer to supply a description of the assertion. The description is

a String; if the test case fails, this string will be displayed by the test runner. Of course, this string can be constructed dynamically.

```
| e |
e := 42.
self assert: e = 23
  description: 'expected 23, got ', e printString
```

The relevant methods in TestCase are:

```
#assert:description:
#deny:description:
#should:description:
#shouldnt:description:
```

Logging Support

The description strings described above may also be logged to a Stream such as the Transcript, or a file stream. You can choose whether to log by overriding TestCase»isLogging in your test class; you must also choose where to log by overriding TestCase»failureLog to answer an appropriate stream.

Continuing after a Failure

SUnit also allows us to specify whether or not a test should continue after a failure. This is a really powerful feature that uses the exception mechanisms offered by Smalltalk. To see what this can be used for, let's look at an example. Consider the following test expression:

```
aCollection do: [ :each | self assert: each even]
```

In this case, as soon as the test finds the first element of the collection that isn't even, the test stops. However, we would usually like to continue, and see both how many elements, and which elements, aren't even, and maybe also log this information. You can do this as follows:

```
aCollection do:
  [:each |
  self
    assert: each even
    description: each printString , ' is not even'
    resumable: true]
```

This will print out a message on your logging stream for each element that fails. It doesn't accumulate failures, *i.e.*, if the assertion fails 10 times in your

test method, you'll still only see one failure. All the other assertion methods that we have seen are not resumable; `assert: p description: s` is equivalent to `assert: p description: s resumable: false`.

7.8 The Implementation of SUnit

The implementation of SUnit makes an interesting case study of a Smalltalk framework. Let's look at some key aspects of the implementation by following the execution of a test.

Running One Test

To execute one test, we evaluate the expression `(aTestClass selector: aSymbol) run`.

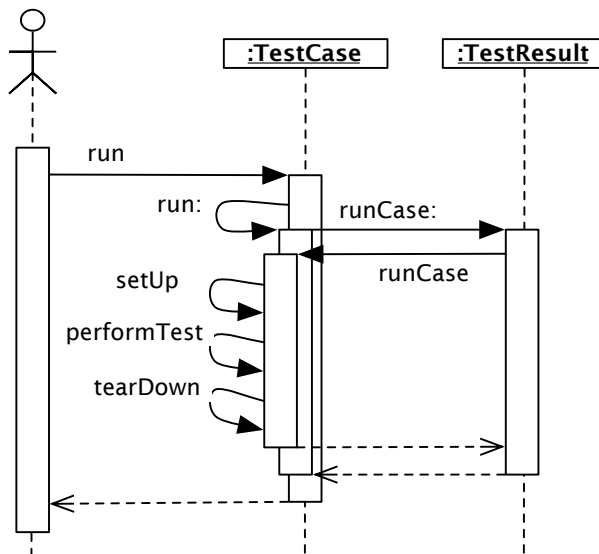


Figure 7.3: Running one test

The method `TestCase>run` creates an instance of `TestResult` that will accumulate the results of the tests, then it sends itself the message `run:`. (See Figure 7.3.)

Method 7.10: *Running a test case*

```

TestCase>>run
| result |
result := TestResult new.
self run: result.
↑result

```

The method `TestCase>>run`: sends the message `runCase`: to the test result:

Method 7.11: *Passing the test case to the test result*

```

TestCase>>run: aResult
aResult runCase: self

```

The method `TestResult>>runCase`: sends the message `runCase` to an individual test, to execute the test. `TestResult>>runCase` deals with any exceptions that may be raised during the execution of a test, runs a `TestCase` by sending it the message `runCase`, and counts the errors, failures and passes.

Method 7.12: *Catching test case errors and failures*

```

TestResult>>runCase: aTestCase
| testCasePassed |
testCasePassed := true.
[[aTestCase runCase]
on: self class failure
do:
[:signal |
failures add: aTestCase.
testCasePassed := false.
signal return: false]]
on: self class error
do:
[:signal |
errors add: aTestCase.
testCasePassed := false.
signal return: false].
testCasePassed ifTrue: [passed add: aTestCase]

```

The method `TestCase>>runCase` sends the messages `setUp` and `tearDown` as shown below.

Method 7.13: *Test case template method*

```

TestCase>>runCase
self setUp.
[self performTest] ensure: [self tearDown]

```

Running a TestSuite

To run more than one test, we send the message `run` to a `TestSuite` that contains the relevant tests. `TestCase` class provides some functionality to build a test suite from its methods. The expression `MyTestCase buildSuiteFromSelectors` returns a suite containing all the tests defined in the `MyTestCase` class. The core of this process is

Method 7.14: *Auto-building the test suite*

```
TestCase»testSelectors
  ↑self selectors asSortedCollection asOrderedCollection select: [:each |
    ('test*' match: each) and: [each numArgs isZero]]
```

The method `TestSuite»run` creates an instance of `TestResult`, verifies that all the resource are available, and then sends itself the message `run`, which runs all the tests in the suite. All the resources are then released.

Method 7.15: *Running a test suite*

```
TestSuite»run
| result |
result := TestResult new.
self areAllResourcesAvailable
  ifFalse: [↑TestResult signalErrorWith:
    'Resource could not be initialized'].
[self run: result] ensure: [self resources do:
  [:each | each reset]].
↑result
```

Method 7.16: *Passing the test suite to the test result*

```
TestSuite»run: aResult
self tests do:
  [:each |
    self sunitChanged: each.
    each run: aResult]
```

The class `TestResource` and its subclasses keep track of their currently created instances (one per class) that can be accessed and created using the class method `current`. This instance is cleared when the tests have finished running and the resources are reset.

The resource availability check makes it possible for the resource to be re-created if needed, as shown in the class method `TestResource class»isAvailable`. During the `TestResource` instance creation, it is initialized and the method `setUp` is invoked.

Method 7.17: *Test resource availability*

```
TestResource class»isAvailable
  ↑self current notNil
```

Method 7.18: *Test resource creation*

```
TestResource class»current
  current isNil ifTrue: [current := self new].
  ↑current
```

Method 7.19: *Test resource initialization*

```
TestResource»initialize
  self setUp
```

7.9 Some Advice on Testing

While the mechanics of testing are easy, writing good tests is not. Here is some advice on how to design tests.

Feathers' Rules for Unit tests. Michael Feathers, an agile process consultant and author, writes:³

A test is not a unit test if:

- *it talks to the database,*
- *it communicates across the network,*
- *it touches the file system,*
- *it can't run at the same time as any of your other unit tests,*
or
- *you have to do special things to your environment (such as editing config files) to run it.*

Tests that do these things aren't bad. Often they are worth writing, and they can be written in a unit test harness. However, it is important to be able to separate them from true unit tests so that we can keep a set of tests that we can run fast whenever we make our changes.

Never get yourself into a situation where you don't want to run your unit test suite because it takes too long.

³See <http://www.artima.com/weblogs/viewpost.jsp?thread=126923>. 9 September 2005

Unit Tests vs. Acceptance Tests. Unit tests capture one piece of functionality, and as such make it easier to identify bugs in that functionality. As far as possible try to have unit tests for each method that could possibly fail, and group them per class. However, for certain deeply recursive or complex setup situations, it is easier to write tests that represent a scenario in the larger application; these are called acceptance tests or functional tests. Tests that break Feathers' rules may make good acceptance tests. Group acceptance tests according to the functionality that they test. For example, if you are writing a compiler, you might write acceptance tests that make assertions about the code generated for each possible source language statement. Such tests might exercise many classes, and might take a long time to run because they touch the file system. You can write them using SUnit, but you won't want to run them each time you make a small change, so they should be separated from the true unit tests.

Black's Rule of Testing. For every test in the system, you should be able to identify some property for which the test increases your confidence. It's obvious that there should be no important property that you are not testing. This rule states the less obvious fact that there should be no test that does not add value to the system by increasing your confidence that a useful property holds. For example, several tests of the same property do no good. In fact they do harm: they make it harder to infer the behaviour of the class by reading the tests, and because one bug in the code might then break many tests at the same time. Have a property in mind when you write a test.

7.10 Chapter Summary

This chapter explained why tests are an important investment in the future of your code. We explained in a step-by-step fashion how to define a few tests for the class `Set`. Then we gave an overview of the core of the SUnit framework by presenting the classes `TestCase`, `TestResult`, `TestSuite` and `TestResources`. Finally we looked deep inside SUnit by following the execution of a test and a test suite.

- To maximize their potential, unit tests should be fast, repeatable, independent of any direct human interaction and cover a single unit of functionality.
- Tests for a class called `MyClass` belong in a class called `MyClassTest`, which should be introduced as a subclass of `TestCase`.

- Initialize your test data in a `setUp` method.
- Each test method should start with the word “test”.
- Use the `TestCase` methods `assert:`, `deny:` and others to make assertions.
- Run tests using the SUnit test runner tool (in the tool bar).

Chapter 8

Basic Classes

Most of the magic of Smalltalk is not in the language but in the class libraries. To program effectively with Smalltalk, you need to learn how the class libraries support the language and environment. The class libraries are entirely written in Smalltalk and can easily be extended since a package may add new functionality to a class even if it does not define this class.

Our goal here is not to present in tedious detail the whole of the Squeak class library, but rather to point out the key classes and methods that you will need to use or override to program effectively. In this chapter we cover the basic classes that you will need for nearly every application: `Object`, `Number` and its subclasses, `Character`, `String`, `Symbol` and `Boolean`.

8.1 Object

For all intents and purposes, `Object` is the root of the inheritance hierarchy. Actually, in Squeak the true root of the hierarchy is `ProtoObject`, which is used to define minimal entities that masquerade as objects, but we can ignore this point for the time being.

`Object` can be found in the *Kernel-Objects* category. Astonishingly, there are some 400 methods to be found here (including extensions). In other words, every class that you define will automatically provide these 400 methods, whether you know what they do or not. Note that some of the methods should be removed and new versions of Squeak may remove some of the superfluous methods.

The class comment for the `Object` states:

Object is the root class for almost all of the other classes in the class hierarchy. The exceptions are ProtoObject (the superclass of Object) and its subclasses. Class Object provides default behavior common to all normal objects, such as access, copying, comparison, error handling, message sending, and reflection. Also utility messages that all objects should respond to are defined here. Object has no instance variables, nor should any be added. This is due to several classes of objects that inherit from Object that have special implementations (SmallInteger and UndefinedObject for example) or the VM knows about and depends on the structure and layout of certain standard classes.

If we begin to browse the method categories on the instance side of Object we start to see some of the key behavior it provides.

Printing

Every object in Smalltalk can return a printed form of itself. You can select any expression in a workspace and select the `print it` menu: this executes the expression and asks the returned object to print itself. In fact this sends the message `printString` to the returned object. The method `printString`, which is a template method, at its core sends the message `printOn:` to its receiver. The message `printOn:` is a hook that can be specialized.

`Object>printOn:` is very likely one of the methods that you will most frequently override. This method takes as its argument a Stream on which a String representation of the object will be written. The default implementation simply write the class name preceded by “a” or “an”. `Object>printString` returns the String that is written:

For example, the class `Browser` does not redefine the method `printOn:` and sending the message `printString` to an instance executes the methods defined in `Object`.

```
Browser new printString  →  'a Browser'
```

The class `TTCFont` shows an example of `printOn:` specialization. It prints the name of the class followed by the family name, the size and the subfamily name of the font as shown by the code below which prints an instance of the class.

Method 8.1: *printOn:* redefinition.

```
TTCFont>printOn: aStream
aStream nextPutAll: 'TTCFont(';
nextPutAll: self familyName; space;
print: self pointSize; space;
```

```
nextPutAll: self subfamilyName;
nextPut: $)
```

```
TTCFont allInstances anyOne printString  →  'TTCFont(BitstreamVeraSans 6
Bold)'
```

Note that the message `printOn:` is not the same as `storeOn:`. The message `storeOn:` puts on its argument stream an expression that can be used to recreate the receiver. This expression is evaluated when the stream is read using the message `readFrom:`. `printOn:` just returns a textual version of the receiver. Of course, it may happen that this textual representation may represent the receiver as a self-evaluating expression.

A word about representation and self-evaluating representation. In functional programming, expressions return values when executed. In Smalltalk, messages (expressions) return objects (values). Some objects have the nice properties that their value is themselves. For example, the value of the object `true` is itself *i.e.*, the object `true`. We call such objects *self-evaluating objects*. You can see a *printed* version of an object value when you print the object in a workspace. Here are some examples of such self-evaluating expressions.

```
true      →  true
3@4       →  3@4
$a        →  $a
#(1 2 3)  →  #(1 2 3)
```

Note that some objects such as arrays are self-evaluating or not depending on the objects they contain. For example, an array of booleans is self-evaluating whereas an array of persons is not. In Squeak 3.9, a mechanism was introduced (via the message `isSelfEvaluating`) to print collections in their self-evaluating forms as much as possible and this especially true for brace arrays. The following examples shows that a dynamic array is self-evaluating only if its elements are:

```
{10@10 . 100@100}  →  {10@10 . 100@100}
{Browser new . 100@100} →  an Array(a Browser 100@100)
```

Remember that literal arrays can only contain literals. Hence the following array does not contain two points but rather six literal elements.

```
#(10@10 100@100)  →  #(10 #@ 10 100 #@ 100)
```

Lots of `printOn:` method specializations implement self-evaluating behavior. The implementations of `Point>printOn:` and `Interval>printOn:` are self-evaluating.

Method 8.2: *Self-evaluation of Point*

```
Point>printOn: aStream
  "The receiver prints on aStream in terms of infix notation."
  x printOn: aStream.
  aStream nextPut: $@.
  y printOn: aStream
```

Method 8.3: *Self-evaluation of Interval*

```
Interval>printOn: aStream
  aStream nextPut: $(;
  print: start;
  nextPutAll: ' to: ';
  print: stop.
  step ~ 1 ifTrue: [aStream nextPutAll: ' by: '; print: step].
  aStream nextPut: $)
```

```
1 to: 10  →  (1 to: 10)  "intervals are self-evaluating"
```

Identity and Equality

In Smalltalk, the message = tests object *equality* (i.e., whether two objects represent the same value) whereas == tests object *identity* (i.e., whether two expressions represent the same object).

The default implementation of object equality is to test for object identity:

Method 8.4:

```
Object>= anObject
  "Answer whether the receiver and the argument represent the same object.
  If = is redefined in any subclass, consider also redefining the message hash."
  ↑ self == anObject
```

This a method that you will frequently want to override. Consider the case of Complex numbers:

```
(1 + 2 i) = (1 + 2 i)  →  true   "same value"
(1 + 2 i) == (1 + 2 i) →  false  "but different objects"
```

This works because Complex overrides = as follows:

Method 8.5: Equality for complex numbers

```
Complex»= anObject
  anObject isComplex
    ifTrue: [↑ (real = anObject real) & (imaginary = anObject imaginary)]
    ifFalse: [↑ anObject adaptToComplex: self andSend: #=]
```

The default implementation of `Object»~` simply negates `Object»=`, and should not normally need to be changed.

```
(1 + 2 i) ~= (1 + 4 i)  → true
```

If you override `=`, you should consider overriding `hash`. If instances of your class are ever used as keys in a `Dictionary`, then you should make sure that instances that are considered to be equal have the same hash value:

Method 8.6: Hash must be reimplemented for complex numbers

```
Complex»hash
  "Hash is reimplemented because = is implemented."
  ↑ real hash bitXor: imaginary hash.
```

Although you should override `=` and `hash` together, you should *never* override `==`. (The semantics of object identity is the same for all classes.) `==` is a primitive method of `ProtoObject`.

Note that Squeak has some strange behavior compared to other Smalltalks: for example a symbol and a string can be equal if the string associated with the symbol is equal to the string. (We consider this behavior to be a bug, not a feature.)

```
#'lulu' = 'lulu'  → true
'lulu' = #'lulu'  → true
```

Class membership

Several methods allow you to query the class of an object.

class. You can ask any object about its class using the message `class`.

```
1 class  → SmallInteger
```

Conversely, you can ask if an object is an instance of a specific class:

```
1 isMemberOf: SmallInteger  → true  "must be precisely this class"
1 isMemberOf: Integer       → false
```

```
1 isMemberOf: Number    → false
1 isMemberOf: Object    → false
```

Since Smalltalk is written in itself, you can really navigate through its structure using the right combination of superclass and class messages (see Chapter 12).

isKindOf: Object»isKindOf: answers whether the receiver's class is either the same as, or a subclass of the argument class.

```
1 isKindOf: SmallInteger → true
1 isKindOf: Integer      → true
1 isKindOf: Number       → true
1 isKindOf: Object       → true
1 isKindOf: String       → false

1/3 isKindOf: Number     → true
1/3 isKindOf: Integer    → false
```

1/3 which is a Fraction is a kind of Number, since the class Number is a superclass of the class Fraction, but 1/3 is not a Integer.

respondsTo: Object»respondsTo: answers whether the receiver understands the message selector given as an argument.

```
1 respondsTo: #, → false
```

Normally it is a bad idea to query an object for its class, or to ask it which messages it understands. Instead of making decisions based on the class of object, you should simply send a message to the object and let it decide (*i.e.*, on the basis of its class) how it should behave.

Copying

Copying objects introduces some subtle issues. Since instance variables are accessed by reference, a *shallow copy* of an object would share its references to instance variables with the original object:

```
a1 := { { 'harry' } }.
a1 → #('harry')
a2 := a1 shallowCopy.
a2 → #('harry')
(a1 at: 1) at: 1 put: 'sally'.
a1 → #('sally')
a2 → #('sally')  "the subarray is shared"
```

`Object»shallowCopy` is a primitive method that creates a shallow copy of an object. Since `a2` is only a shallow copy of `a1`, the two arrays share a reference to the nested Array that they contain.

`Object»shallowCopy` is the “public interface” to `Object»copy` and should be overridden if instances are unique. This is the case, for example, with the classes `Boolean`, `Character`, `SmallInteger`, `Symbol` and `UndefinedObject`.

`Object»copyTwoLevel` does the obvious thing when a simple shallow copy does not suffice:

```
a1 := { { 'harry' } } .
a2 := a1 copyTwoLevel.
(a1 at: 1) at: 1 put: 'sally'.
a1  →  #(#('sally'))
a2  →  #(#('harry'))  "fully independent state"
```

`Object»deepCopy` makes an arbitrarily deep copy of an object.

```
a1 := { { { 'harry' } } } .
a2 := a1 deepCopy.
(a1 at: 1) at: 1 put: 'sally'.
a1  →  #(#('sally'))
a2  →  #(#(#('harry'))))
```

The problem with `deepCopy` is that it will not terminate when applied to a mutually recursive structure:

```
a1 := { 'harry' }.
a2 := { a1 }.
a1 at: 1 put: a2.
a1 deepCopy → ... does not terminate!
```

Although it is possible to override `deepCopy` to do the right thing, `Object»copy` offers a better solution:

Method 8.7: Copying objects as a template method

`Object»copy`

*"Answer another instance just like the receiver. Subclasses typically override
 postCopy;
 they typically do not override shallowCopy."
 ↑self shallowCopy postCopy*

You should override `postCopy` to copy any instance variables that should not be shared. `postCopy` should always do a super `postCopy`.

Debugging

The most important method here is `halt`. In order to set a breakpoint in a method, simply insert the message `send self halt` at some point in the body of the method. When this message is sent, execution will be interrupted and a debugger will open to this point in your program. (See Chapter 6 for more details about the debugger.)

The next most important message is `assert:`, which takes a block as its argument. If the block returns `true`, execution continues. Otherwise an exception will be raised. If this exception is not otherwise caught, the debugger will open to this point in the execution. `assert:` is especially useful to support *design by contract*. The most typical usage is to check non-trivial pre-conditions to public methods of objects. `Stack»pop` could easily have been implemented as follows:

Method 8.8: *Checking a pre-condition*

```
Stack»pop
  "Return the first element and remove it from the stack."
  self assert: [ self isEmpty not ].
  ↑self linkedList removeFirst element
```

Do not confuse `Object»assert:` with `TestCase»assert:`, which occurs in the `SUnit` testing framework (see Chapter 7). While the former expects a block as its argument¹, the latter expects a `Boolean`. Although both are useful for debugging, they each serve a very different intent.

Error handling

This protocol contains several methods useful for signaling run-time errors.

Sending `self deprecated: anExplanationString` signals that the current method should no longer be used, if deprecation has been turned on in the *debug* protocol of the preference browser. The `String` argument should offer an alternative.

```
1 doIfNotNil: [ :arg | arg printString, ' is not nil' ]
   → SmallInteger(Object)»doIfNotNil: has been deprecated. use ifNotNilDo:
```

`doesNotUnderstand:` is sent whenever message lookup fails. The default implementation, *i.e.*, `Object»doesNotUnderstand:` will trigger the debugger at this point. It may be useful to override `doesNotUnderstand:` to provide some other behaviour.

¹ Actually, it will take any argument that understands `value`, including a `Boolean`.

`Object»error` and `Object»error:` are generic methods that can be used to raise exceptions. (Generally it is better to raise your own custom exceptions, so you can distinguish errors arising from your code from those coming from kernel classes.)

Abstract methods in Smalltalk are implemented by convention with the body `self subclassResponsibility`. Should an abstract class be instantiated by accident, then calls to abstract methods will result in `Object»subclassResponsibility` being evaluated.

Method 8.9: *Signaling that a method is abstract*

`Object»subclassResponsibility`

"This message sets up a framework for the behavior of the class' subclasses.

Announce that the subclass should have implemented this message."

```
self error: 'My subclass should have overridden ', thisContext sender selector
printString
```

Magnitude, Number and Boolean are classical examples of abstract classes that we shall see shortly in this chapter.

`Number new + 1` \longrightarrow *Error: My subclass should have overridden #+*

`self shouldNotImplement` is sent by convention to signal that an inherited method is not appropriate for this subclass. This is generally a sign that something is not quite right with the design of the class hierarchy. Due to the limitations of single inheritance, however, sometimes it is very hard to avoid such workarounds.

A typical example is `Collection»remove:` which is inherited by `Dictionary` but flagged as not implemented. (A `Dictionary` provides `removeKey:` instead.)

Testing

The *testing* methods have nothing to do with SUnit testing! A testing method is one that lets you ask a question about the state of the receiver and returns a Boolean.

Numerous testing methods are provided by `Object`. We have already seen `isComplex`. Others include `isArray`, `isBoolean`, `isBlock`, `isCollection` and so on. Generally such methods are to be avoided since querying an object for its class is a form of violation of encapsulation. Instead of testing an object for its class, one should simply send a request and let the object decide how to handle it.

Nevertheless some of these testing methods are undeniably useful. The most useful are probably `ProtoObject»isNil` and `Object»notNil` (though the `Null`

Object² design pattern can obviate the need for even these methods).

Initialize release

A final key method that occurs not in Object but in ProtoObject is initialize.

Method 8.10: initialize *as an empty hook method*

ProtoObject»initialize

"Subclasses should redefine this method to perform initializations on instance creation"

The reason this is important is that in Squeak as of version 3.9, the default new method defined for every class in the system will send initialize to newly created instances.

Method 8.11: new *as a class-side template method*

Behavior»new

"Answer a new initialized instance of the receiver (which is a class) with no indexable variables. Fail if the class is indexable."

↑ self basicNew initialize

This means that simply by overriding the initialize hook method, new instances of your class will automatically be initialized. The initialize method should normally perform a super initialize to establish the class invariant for any inherited instance variables. Note that this is *not* standard behavior in other Smalltalks.

8.2 Numbers

Remarkably, numbers in Smalltalk are not primitive data values but true objects. Of course numbers are implemented efficiently in the virtual machine, but the Number hierarchy is as perfectly accessible and extensible as any other portion of the Smalltalk class hierarchy.

Numbers are found in the *Kernel-Numbers* category. The abstract root of this hierarchy is *Magnitude*, which represents all kinds of classes supporting comparison operators. Number adds various arithmetic and other operators as mostly abstract methods. Float and Fraction represent, respectively, floating point numbers and fractional values. Integer is also abstract,

²Bobby Woolf, Null Object. In Robert Martin, Dirk Riehle and Frank Buschmann, editors, Pattern Languages of Program Design 3. Addison Wesley, 1998.

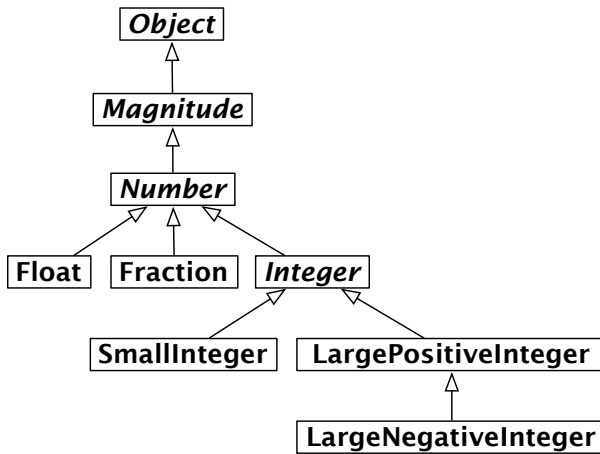


Figure 8.1: The Number Hierarchy

thus distinguishing between subclasses `SmallInteger`, `LargePositiveInteger` and `LargeNegativeInteger`. For the most part users do not need to be aware of the difference between the three `Integer` classes, as values are automatically converted as needed.

Magnitude

`Magnitude` is the parent not only of the `Number` classes, but also of other classes supporting comparison operations, such as `Character`, `Duration` and `Timespan`. (Complex numbers are not comparable, and so do not inherit from `Number`.)

Methods `<` and `=` are abstract. The remaining operators are generically defined. For example:

Method 8.12: *Abstract comparison methods*

`Magnitude`» `<` `aMagnitude`

"Answer whether the receiver is less than the argument."

↑self subclassResponsibility

`Magnitude`» `>` `aMagnitude`

"Answer whether the receiver is greater than the argument."

↑aMagnitude < self

Number

Similarly, Number defines +, -, * and / to be abstract, but all other arithmetic operators are generically defined.

All Number objects support various *converting* operators, such as asFloat and asInteger. There are also numerous *shortcut constructor methods*, such as i, which converts a Number to an instance of Complex with a zero real component, and others which generate Duration s, such as hour, day and week.

Numbers directly support common *math functions* such as sin, log, raiseTo:, squared, sqrt and so on.

Number»printOn: is implemented in terms of the abstract method Number»printOn:base:. (The default base is 10.)

Testing methods include even, odd, positive and negative. Unsurprisingly Number overrides isNumber. More interesting, isInfinite is defined to return false.


Truncation methods include floor, ceiling, integerPart, fractionPart and so on.

1 + 2.5	→	3.5	"Addition of two numbers"
3.4 * 5	→	17.0	"Multiplication of two numbers"
8 / 2	→	4	"Division of two numbers"
10 - 8.3	→	1.7	"Subtraction of two numbers"
12 = 11	→	false	"Equality between two numbers"
12 ~= 11	→	true	"Test if two numbers are different"
12 > 9	→	true	"Greater than"
12 >= 10	→	true	"Greater or equal than"
12 < 10	→	false	"Smaller than"
100@10	→	100@10	"Point creation"

The following example works surprisingly well in Smalltalk:

1000 factorial / 999 factorial	→	1000
--------------------------------	---	------

Note that 1000 factorial is really calculated which in many other languages can be quite difficult to compute. This is an excellent example of automatic coercion and exact handling of a number.

 Try to display the result of 1000 factorial. It takes more time to display it than to calculate it!

Float

Float implements the abstract Number methods for floating point numbers.

More interestingly, Float class (*i.e.*, the class-side of Float) provides methods to return the following *constants*: e, infinity, nan and pi.

Float pi	→	3.141592653589793
Float infinity	→	Infinity
Float infinity isInfinite	→	true

Fraction

Fractions are represented by instance variables for the numerator and denominator, which should be Integers. Fractions are normally created by Integer division (rather than using the constructor method Fraction» numerator:denominator:):

6/8	→	(3/4)
(6/8) class	→	Fraction

Multiplying a Fraction by an Integer or another Fraction may yield an Integer:

6/8 * 4	→	3
---------	---	---

Integer

Integer is the abstract parent of three concrete integer implementations. In addition to providing concrete implementations of many abstract Number methods, it also adds a few methods specific to integers, such as factorial, atRandom, isPrime, gcd: and many others.

SmallInteger is special in that its instances are represented compactly — instead of being stored as a reference, a SmallInteger is represented directly using the bits that would otherwise be used to hold a reference. The first bit of an object reference indicates whether the object is a SmallInteger or not.

The class methods minVal and maxVal tell us the range of a SmallInteger:

SmallInteger maxVal = ((2 raisedTo: 30) - 1)	→	true
SmallInteger minVal = (2 raisedTo: 30) negated	→	true

When a SmallInteger goes out of this range, it is automatically converted to a LargePositiveInteger or a LargeNegativeInteger, as needed:

(SmallInteger maxVal + 1) class	→	LargePositiveInteger
(SmallInteger minVal - 1) class	→	LargeNegativeInteger

Large integers are similarly converted back to small integers when appropriate.

As in most programming languages, integers can be useful for specifying iterative behavior. There is a dedicated method `timesRepeat:` for evaluating a block repeatedly. We have already seen a similar example in Chapter 3:

```
n := 2.
3 timesRepeat: [ n := n*n ].
n  → 256
```

8.3 Characters

Character is defined in the *Collections-Strings* category as a subclass of *Magnitude*. Printable characters are represented in Squeak as `$<char>`. For example:

```
$a < $b  → true
```

Non-printing characters can be generated by various class methods. `Character class>value:` takes the Unicode (or ASCII) integer value as argument and returns the corresponding character. The protocol *accessing untypeable characters* contains a number of convenience constructor methods such as `backspace`, `cr`, `escape`, `euro`, `space`, `tab`, and so on.

```
Character space = (Character value: Character space asciiValue)  → true
```

The `printOn:` method is clever enough to know which of the three ways to generate characters offers the most appropriate representation:

```
Character value: 1  → Character value: 1
Character value: 32 → Character space
Character value: 97 → $a
```

Various convenient *testing* methods are built in: `isAlphaNumeric`, `isCharacter`, `isDigit`, `isLowercase`, `isVowel`, and so on.

To convert a `Character` to the string containing just that character, send `asString`. In this case `asString` and `printString` yield different results:

```
$a asString  → 'a'
$a          → $a
$a printString → '$a'
```

Every ascii Character is a unique instance, stored in the class variable CharacterTable:

(Character value: 97) == \$a	→	true
------------------------------	---	------

Characters outside the range 0 to 255 are not unique, however:

Character characterTable size	→	256
(Character value: 500) == (Character value: 500)	→	false

8.4 Strings

The String class is also defined in the category *Collections-Strings*. A String is an indexed Collection that holds only Characters.




Figure 8.2: The String Hierarchy

In fact, String is abstract and Squeak Strings are actually instances of the concrete class ByteString.

'hello world' class	→	ByteString
---------------------	---	------------

The other important subclass of String is Symbol. The key difference is that there is only ever a single instance of Symbol with a given value.

(This is sometimes called “the unique instance property”). In contrast, two separately constructed Strings that happen to contain the same sequence of characters will often be different objects.

```
'hel','lo' == 'hello'  →  false
```

```
('hel','lo') asSymbol == #hello  →  true
```

Another important difference is that a String is mutable, whereas a Symbol is immutable.

```
'hello' at: 2 put: $u; yourself  →  'hullo'
```

```
#hello at: 2 put: $u  →  error
```

It is easy to forget that since strings are collections, they understand the same messages that other collections do:

```
#hello indexOf: $o  →  5
```

Although String does not inherit from Magnitude, it does support the usual *comparing* methods, <, = and so on. In addition, String»match: is useful for some basic glob-style pattern-matching:

```
'*or*' match: 'zorro'  →  true
```

Should you need more advanced support for regular expressions, there are a number of third party implementations available, such as Vassili Bykov’s Regex package.

Strings support rather a large number of conversion methods. Many of these are shortcut constructor methods for other classes, such as asDate, asFileName and so on. There are also a number of useful methods for converting a string to another string, such as capitalized and translateToLowercase.

For more on strings and collections, see Chapter 9.

8.5 Booleans

The class Boolean offers a fascinating insight into how much of the Smalltalk language has been pushed into the class library. Boolean is the abstract superclass of the Singleton classes True and False.

Most of the behaviour of Booleans can be understood by considering the method ifTrue:ifFalse:, which takes two Blocks as arguments.




Figure 8.3: The Boolean Hierarchy

(4 factorial > 20) ifTrue: ['bigger'] ifFalse: ['smaller'] → 'bigger'

The method is abstract in Boolean. The implementations in its concrete subclasses are both trivial:

Method 8.13: *Implementations of ifTrue:ifFalse:*

True»ifTrue: trueAlternativeBlock ifFalse: falseAlternativeBlock
 ↑trueAlternativeBlock value

False»ifTrue: trueAlternativeBlock ifFalse: falseAlternativeBlock
 ↑falseAlternativeBlock value

In fact, this is the essence of OOP: when a message is sent to an object, the object itself determines which method will be used to respond. In this case an instance of True simply evaluates the *true* alternative, while an instance of False evaluates the *false* alternative. All the abstract Boolean methods are implemented in this way for True and False. For example:


Method 8.14: *Implementing negation*

True»not
 "Negation--answer false since the receiver is true."
 ↑false

Booleans offer several useful convenience methods, such as ifTrue:, ifFalse:, ifFalse:ifTrue:. You also have the choice between eager and lazy conjunctions and disjunctions.

$(1 > 2) \ \& \ (3 < 4)$	\longrightarrow	false	<i>"must evaluate both sides"</i>
$(1 > 2) \text{ and: } [3 < 4]$	\longrightarrow	false	<i>"only evaluate receiver"</i>
$(1 > 2) \text{ and: } [(1/0) > 0]$	\longrightarrow	false	<i>"argument block is never evaluated, so no exception"</i>

In the first example, both Boolean subexpressions are evaluated, since `&` takes a Boolean argument. In the second and third examples, only the first is evaluated, since `and:` expects a Block as its argument. The Block is evaluated only if the first argument is true.

 Try to imagine how `and:` and `or:` are implemented. Check the implementations in `Boolean`, `True` and `False`.

8.6 Chapter Summary

- If you override `=` then you should override `hash` as well.
- Override `postCopy` to correctly implement copying for your objects.
- Send `self halt` to set a breakpoint.
- Return `self subclassResponsibility` to make a method abstract.
- To give an object a String representation you should override `printOn:`.
- Override the hook method `initialize` to properly initialize instances.
- Number methods automatically convert between Floats, Fractions and Integers.
- Fractions truly represent rational numbers rather than floats.
- Characters are unique instances.
- Strings are mutable; Symbols are not. Take care not to mutate string literals, however!
- Symbols are unique; Strings are not.
- Strings and Symbols are Collections and therefore support the usual Collection methods.

Chapter 9

Collections

9.1 Introduction

The collection classes form a loosely-defined group of general-purpose subclasses of `Collection` and `Stream`. The group of classes that appears in the “Blue Book”¹ contains 17 subclasses of `Collection` and 9 subclasses of `Stream`, for a total of 28 classes, and had already been redesigned several times before the Smalltalk-80 system was released. This group of classes is often considered to be a paradigmatic example of object-oriented design.

In Squeak, the abstract class `Collection` has 98 subclasses, and the abstract class `Stream` has 39 subclasses, but many of these (like `Bitmap`, `FileStream` and `CompiledMethod`) are special-purpose classes crafted for use in other parts of the system or in applications, and hence not categorized as “Collections” by the system organization. For the purposes of this chapter, we use the term “Collection Hierarchy” to mean `Collection` and its 37 subclasses that are *also* in the system categories labelled *Collections-**. We use the term “Stream Hierarchy” to mean `Stream` and its 10 subclasses that are *also* in the *Collections-Streams* system categories. The full list is shown in Figure 9.1. These 49 classes respond to 794 messages and define a total of 1236 methods!

In this chapter we focus mainly on the subset of collection classes shown in Figure 9.2. Streams will be discussed separately in Chapter 10.

¹Adele Goldberg and David Robson, *Smalltalk 80: the Language and its Implementation*. Reading, Mass.: Addison Wesley, May 1983, ISBN 0-201-13688-0.

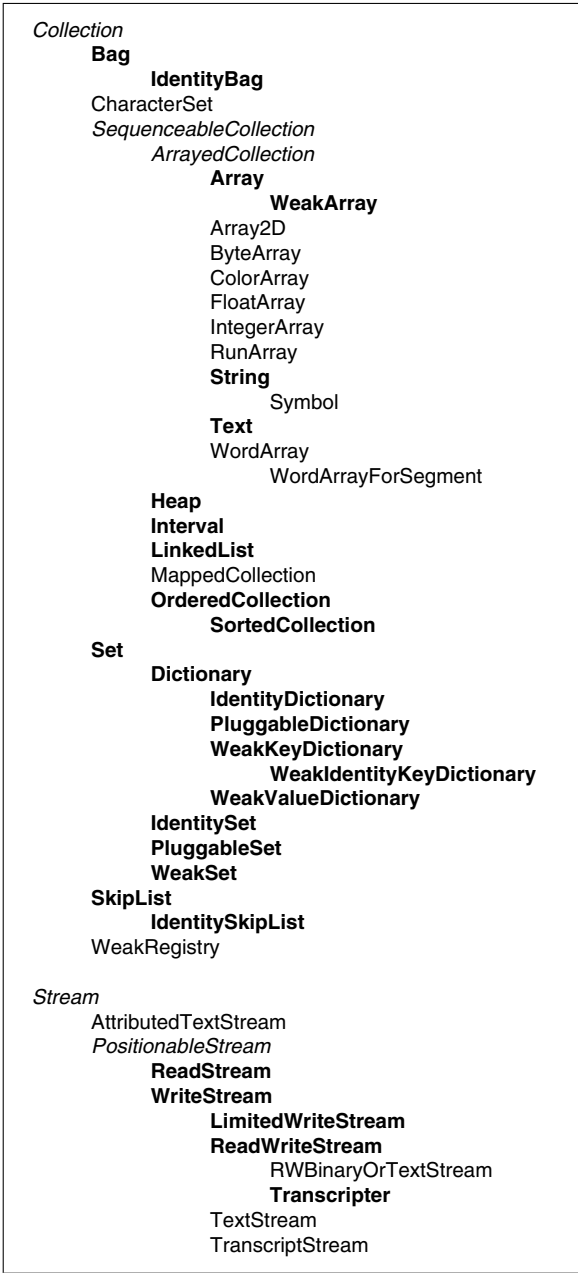


Figure 9.1: The collection classes in Squeak. Indentation indicates subclassing; *Italicized* classes are abstract.

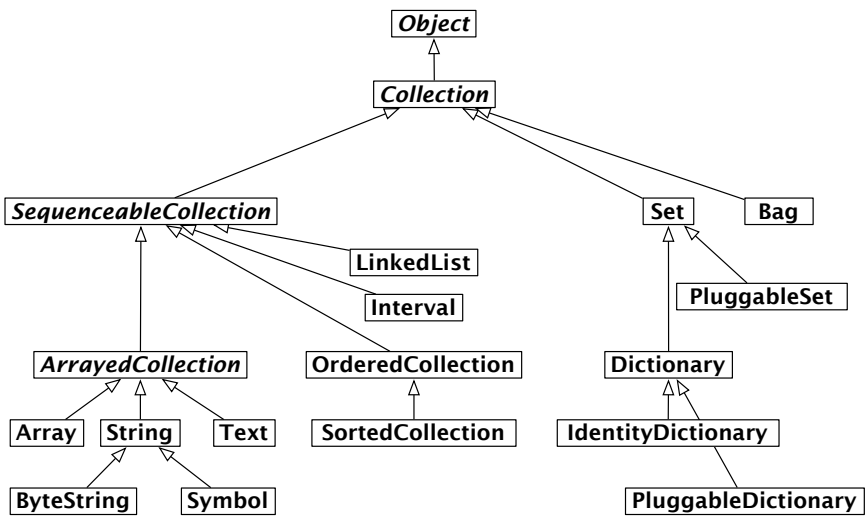


Figure 9.2: Some of the key collection classes in Squeak.

9.2 The Varieties of Collection

To make good use of the collection classes, the reader needs at least a superficial knowledge of the wide variety of collections that they implement, and their commonalities and differences.

Programming with collections rather than individual elements is an important way to raise the level of abstraction of a program. The Lisp function `map`, which applies an argument function to every element of a list and returns a new list containing the results is an early example of this style, but Smalltalk-80 adopted collection-based programming as a central tenet. Modern functional programming languages such as ML and Haskell have followed Smalltalk’s lead.

Why is this a good idea? Suppose you have a data structure containing a collection of student records, and wish to perform some action on all of the students that meet some criterion. Programmers raised to use an imperative language will immediately reach for a loop. But the Smalltalk programmer will write:

```
students select: [ :each | each gpa < threshold ]
```

which evaluates to a new collection containing precisely those elements of

Protocol	Methods
<i>accessing</i>	size, capacity, at: <i>anIndex</i> , at: <i>anIndex</i> put: <i>anElement</i>
<i>testing</i>	isEmpty, includes: <i>anElement</i> , contains: <i>aBlock</i> , occurrencesOf: <i>anElement</i>
<i>adding</i>	add: <i>anElement</i> , addAll: <i>aCollection</i>
<i>removing</i>	remove: <i>anElement</i> , remove: <i>anElement</i> ifAbsent: <i>aBlock</i> , removeAll: <i>aCollection</i>
<i>enumerating</i>	do: <i>aBlock</i> , collect: <i>aBlock</i> , select: <i>aBlock</i> , reject: <i>aBlock</i> , detect: <i>aBlock</i> , detect: <i>aBlock</i> ifNone: <i>aNoneBlock</i> , inject: <i>aValue</i> into: <i>aBinaryBlock</i>
<i>converting</i>	asBag, asSet, asOrderedCollection, asSortedCollection, asArray, asSortedCollection: <i>aBlock</i>
<i>creation</i>	with: <i>anElement</i> , with:with:, with:with:with:, with:with:with:with:, with: <i>anElement</i> All: <i>aCollection</i>

Figure 9.3: Standard Collection protocols

students for which the bracketed function returns true². The Smalltalk code has the simplicity and elegance of a domain-specific query language.

The message select: is understood by *all* collections in Smalltalk. There is no need to find out if the student data structure is an array or a linked list: the select: message is understood by both. Note that this is quite different from using a loop, where one must know whether students is an array or a linked list before the loop can be set up.

In Smalltalk, when one speaks of a collection without being more specific about the kind of collection, one means an object that supports well-defined protocols for testing membership and enumerating the elements. *All* collections understand the *testing* messages includes:, isEmpty and occurrencesOf:. *All* collections understand the *enumeration* messages do:, select:, reject: (which is the opposite of select:), collect: (which is like lisp's map), detect:ifNone:, inject:into: (which performs a left fold) and many more. It is the ubiquity of this protocol, as well as its variety, that makes it so powerful.

Figure 9.3 summarizes the standard protocols supported by most of the classes in the collection hierarchy. These methods are defined, redefined, optimized or occasionally even forbidden by subclasses of Collection.

Beyond this basic uniformity, there are many different kinds of collection either supporting different protocols, or providing different behaviour for

²The expression in brackets can be thought of as a λ -expression defining an anonymous function $\lambda x.x \text{ gpa} < \text{threshold}$.

the same requests. Let us briefly survey some of the key differences:

- **Sequenceable:** Instances of all subclasses of `SequenceableCollection` start from a first element and proceed in a well-defined order to a last element. Instances of `Set`, `Bag` and `Dictionary`, on the other hand, are not sequenceable.
- **Sortable:** A `SortedCollection` maintains its elements in sort order.
- **Indexable:** Most sequenceable collections are also indexable, that is, elements can be retrieved with `at:`. `Array` is the familiar indexable data structure with a fixed size; `anArray at: n` retrieves the n^{th} element of `anArray`, and `anArray at: n put: v` changes the n^{th} element to `v`. `LinkedLists` and `SkipLists` are sequenceable but not indexable, that is, they understand first and last, but not `at:`.
- **Keyed:** Instances of `Dictionary` and its subclasses are accessed by keys instead of indices.
- **Mutable:** Most collections are mutable, but `Intervals` and `Symbols` are not. An `Interval` is an immutable collection representing a range of `Integers`. For example, `5 to: 16 by: 2` is an interval that contains the elements 5, 7, 9, 11, 13 and 15. It is indexable with `at:`, but cannot be changed with `at:put:`.
- **Growable:** Instances of `Interval` and `Array` are always of a fixed size. Other kinds of collections (sorted collections, ordered collections, and linked lists) can grow after creation.

The class `OrderedCollection` is more general than `Array`; the size of an `OrderedCollection` grows on demand, and it has methods for `addFirst:` and `addLast:` as well as `at:` and `at:put:`.

- **Accepts duplicates:** A `Set` will filter out duplicates, but a `Bag` will not. `Dictionary`, `Set` and `Bag` use the `=` method provided by the elements; the `Identity` variants of these classes use the `==` method, which tests whether the arguments are the same object, and the `Pluggable` variants use an arbitrary equivalence relation supplied by the creator of the collection.
- **Heterogeneous:** Most collections will hold any kind of element. A `String`, `CharacterArray` or `Symbol`, however, only holds `Characters`. An `Array` will hold any mix of objects, but a `ByteArray` only holds `Bytes`, an `IntegerArray` only holds `Integers` and a `FloatArray` only holds `Floats`. A `LinkedList` is constrained to hold elements that conform to the *Link* \triangleright *accessing* protocol.

Arrayed Implementation	Ordered Implementation	Hashed Implementation	Linked Implementation	Interval Implementation
Array String Symbol	OrderedCollection SortedCollection Text Heap	Set IdentitySet PluggableSet Bag IdentityBag Dictionary IdentityDictionary PluggableDictionary	LinkedList SkipList	Interval

Figure 9.4: Some collection classes categorized by implementation technique.

9.3 Implementations of Collections

These categorizations by functionality are not our only concern; we must also consider how the collection classes are implemented. As shown in Figure 9.4, five main implementation techniques are employed.

1. Arrays store their elements in the (indexable) instance variables of the collection object itself; as a consequence, arrays must be of a fixed size, but can be created with a single memory allocation.
2. OrderedCollections and SortedCollections store their elements in an array that is referenced by one of the instance variables of the collection. Consequently, the internal array can be replaced with a larger one if the collection grows beyond its storage capacity.
3. The various kinds of set and dictionary also reference a subsidiary array for storage, but use the array as a hash table. Bags use a subsidiary Dictionary, with the elements of the bag as keys and the number of occurrences as values.
4. LinkedLists use a standard singly-linked representation.
5. Intervals are represented by three integers that record the two endpoints and the step size.

In addition to these classes, there are also “weak” variants of Array, Set and of the various kinds of dictionary. These collections hold onto their elements weakly, *i.e.*, in a way that does not prevent the elements from being garbage collected. The Squeak virtual machine is aware of these classes and handles them specially.

Readers interested in learning more about the Smalltalk collections are referred to LaLonde and Pugh's excellent book³.

9.4 Examples of Key Classes

We present now the most common or important collection classes using simple code examples. The main protocols of collections are: `at:`, `at:put:` — to access an element, `add:`, `remove:` — to add or remove an element, `size`, `isEmpty`, `include:` — to get some information about the collection, `do:`, `collect:`, `select:` — to iterate over the collection. Each collection may implement or not such protocols and when they do interpret them to fit with their semantics. We suggest you to browse the classes themselves to identify specific and more advanced protocols.

We will focus on the most common collection classes: `OrderedCollection`, `Set`, `SortedCollection`, `Dictionary`, `Interval`, and `Array`.

Common creation protocol. There are several ways to create instances of collections. The most generic ones use the methods `new:` and `with:`. `new:` `anInteger` creates a collection of size `anInteger` whose elements will all be `nil`. `with:` `anObject` creates a collection and adds `anObject` to the created collection. Different collections will realize this behavior differently.

You can create collections with initial elements using the methods `with:`, `with:with:` etc. for up to six elements.

```

Array with: 1    →  #(1)
Array with: 1 with: 2    →  #(1 2)
Array with: 1 with: 2 with: 3    →  #(1 2 3)
Array with: 1 with: 2 with: 3 with: 4    →  #(1 2 3 4)
Array with: 1 with: 2 with: 3 with: 4 with: 5    →  #(1 2 3 4 5)
Array with: 1 with: 2 with: 3 with: 4 with: 5 with: 6    →  #(1 2 3 4 5 6)

```

You can also use `addAll:` to add all elements of one kind of collection to another kind:

```

(1 to: 5) asOrderedCollection addAll: '678'; yourself →  an OrderedCollection(1 2
3 4 5 $6 $7 $8)

```

Take care that `addAll:` also returns its argument, and not the receiver!

You can also create many collections with `withAll:` or `newFrom:`

³Wilf LaLonde and John Pugh, *Inside Smalltalk: Volume 1*. Prentice Hall, 1990, ISBN 0-13-468414-1.

Array withAll: #(7 3 1 3)	→	#(7 3 1 3)
OrderedCollection withAll: #(7 3 1 3)	→	an OrderedCollection(7 3 1 3)
SortedCollection withAll: #(7 3 1 3)	→	a SortedCollection(1 3 3 7)
Set withAll: #(7 3 1 3)	→	a Set(7 1 3)
Bag withAll: #(7 3 1 3)	→	a Bag(7 1 3 3)
Dictionary withAll: #(7 3 1 3)	→	a Dictionary(1→7 2→3 3→1 4→3)

Array newFrom: #(7 3 1 3)	→	#(7 3 1 3)
OrderedCollection newFrom: #(7 3 1 3)	→	an OrderedCollection(7 3 1 3)
SortedCollection newFrom: #(7 3 1 3)	→	a SortedCollection(1 3 3 7)
Set newFrom: #(7 3 1 3)	→	a Set(7 1 3)
Bag newFrom: #(7 3 1 3)	→	a Bag(7 1 3 3)
Dictionary newFrom: {1 -> 7. 2 -> 3. 3 -> 1. 4 -> 3}	→	a Dictionary(1→7 2→3 3→1 4→3)

Note that these two methods are not identical. In particular, Dictionary class»withAll: interprets its argument as a collection of values, whereas Dictionary class»newFrom: expects a collection of associations.

Array

An Array is a fixed-sized collection of elements accessed by integer indices. Contrary to the C convention, the first element of a Smalltalk array is at position 1 and not 0. The main protocol to access array elements is the method at: and at:put:. at: anInteger returns the element at index anInteger. at: anInteger put: anObject puts anObject at index anInteger. Arrays are fixed-size collections therefore we cannot add or remove elements at the end of an array. The following code creates an array of size 5, puts values in the first 3 locations and returns the first element.

```
anArray := Array new: 5.
anArray at: 1 put: 4.
anArray at: 2 put: 3/2.
anArray at: 3 put: 'ssss'.
anArray at: 1 → 4
```

There are several ways to create instances of the class Array. We can use new:, with:, and the constructs #() and { }.

Creation with new: new: anInteger creates an array of size anInteger. Array new: 5 creates an array of size 5.

Creation with with: with: methods allows one to specify the value of the elements. The following code creates an array of three elements consisting of the number 4, the fraction 3/2 and the string 'lulu'.

```
Array with: 4 with: 3/2 with: 'lulu'  →  {4 . (3/2) . 'lulu'}
```

Literal creation with #(). #() creates literal arrays with static (or “literal”) elements that have to be known when the expression is compiled, and not when it is executed. The following code creates an array of size 2 where the first element is the (literal) number 1 and the second the (literal) string 'here'.

```
#(1 'here') size  →  2
```

Now, if you evaluate #(1+2), you do not get an array with a single element 3 but instead you get the array #(1 #+ 2) *i.e.*, with three elements: 1, the symbol #+ and the number 2.

```
#(1+2)  →  #(1 #+ 2)
```

This occurs because the construct #() causes the compiler to interpret literally the expressions contained in the array. The expression is scanned and the resulting elements are fed to a new array. Literal arrays contain numbers, nil, true, false, symbols and strings.

Dynamic creation with { }. Finally, you can create a dynamic array using the construct {}. { a . b } is equivalent to Array with: a with: b. This means in particular that the expressions enclosed by { and } are executed.

```
{ 1 + 2 }  →  #(3)
{(1/2) asFloat} at: 1  →  0.5
{10 atRandom . 1/3} at: 2  →  (1/3)
```

Element Access. Elements of all sequenceable collections can be accessed with at: and at:put:.

```
anArray := #(1 2 3 4 5 6) copy.
anArray at: 3  →  3
anArray at: 3 put: 33.
anArray at: 3  →  33
```

Be careful with code that modifies literal arrays! The compiler tries to allocate space just once for literal arrays. Unless you copy the array, the

second time you evaluate the code your “literal” array may not have the value you expect. (Without cloning, the second time around, the literal `#(1 2 3 4 5 6)` will actually be `#(1 2 33 4 5 6)`!) Dynamic arrays do not have this problem.

OrderedCollection

`OrderedCollection` is one of the collections that can grow, and to which elements can be added sequentially. It offers a variety of methods such as `add:`, `addFirst:`, `addLast:`, and `addAll:`.

```
ordCol := OrderedCollection new.
ordCol add: 'Seaside'; add: 'SqueakSource'; addFirst: 'Monticello'.
ordCol  →  an OrderedCollection('Monticello' 'Seaside' 'SqueakSource')
```

Removing Elements. The method `remove: anObject` removes the first occurrence of an object from the collection. If the collection does not include such an object, it raises an error.

```
ordCol add: 'Monticello'.
ordCol remove: 'Monticello'.
ordCol  →  an OrderedCollection('Seaside' 'SqueakSource' 'Monticello')
```

There is a variant of `remove:` named `remove:ifAbsent:` that allows one to specify as second argument a block that is executed in case the element to be removed is not in the collection.

```
res := ordCol remove: 'zork' ifAbsent: [33].
res  →  33
```

Conversion. It is possible to get an `OrderedCollection` from an `Array` (or any other collection) by sending the message `asOrderedCollection:`

```
#(1 2 3) asOrderedCollection  →  an OrderedCollection(1 2 3)
'hello' asOrderedCollection  →  an OrderedCollection($h $e $l $l $o)
```

Interval

The class `Interval` represents ranges of numbers. For example, the interval of numbers from 1 to 100 is defined as follows:

```
Interval from: 1 to: 100  →  (1 to: 100)
```

The `printString` of this interval reveals that the class `Number` provides us with a convenience method called `to: to generate intervals`:

```
(Interval from: 1 to: 100) = (1 to: 100)  →  true
```

We can use `Interval class»from:to:by:` or `Number»to:by:` to specify the step between two numbers as follow:

```
(Interval from: 1 to: 100 by: 0.5) size  →  199
(1 to: 100 by: 0.5) at: 198  →  99.5
(1/2 to: 54/7 by: 1/3) last  →  (15/2)
```

Dictionary

Dictionaries are important collections whose elements are accessed using keys. Among the most commonly used messages of dictionary you will find `at:`, `at:put:`, `at:ifAbsent:`, `keys` and `values`.

```
colors := Dictionary new.
colors at: #yellow put: Color yellow.
colors at: #blue put: Color blue.
colors at: #red put: Color red.
colors at: #yellow  →  Color yellow
colors keys        →  a Set(#blue #yellow #red)
colors values      →  {Color blue . Color yellow . Color red}
```

Dictionaries compare keys by equality. Two keys are considered to be the same if they return `true` when compared using `=`. A common and difficult to spot bug is to use as key an object whose `=` method has been redefined but not its `hash` method. Both methods are used in the implementation of dictionary and when comparing objects.

The class `Dictionary` clearly illustrates that the collection hierarchy is based on subclassing and not subtyping. Even though `Dictionary` is a subclass of `Set`, we would normally not want to use a `Dictionary` where a `Set` is expected. In its implementation, however, a `Dictionary` can clearly be seen as consisting of a set of associations (key value) created using the message `->`. We can create a `Dictionary` from a collection of associations, or we may convert a dictionary to an array of associations.

```
colors := Dictionary newFrom: { #blue->Color blue. #red->Color red. #yellow->Color
    yellow }.
colors removeKey: #blue.
colors associations  →  {#yellow->Color yellow . #red->Color red}
```

IdentityDictionary. While a dictionary uses the result of the messages = and hash to determine if two keys are the same, the class IdentityDictionary uses the identity (message ==) of the key instead of its values, *i.e.*, it considers two keys to be equal *only* if they are the same object.

Often Symbols are used as keys, in which case it is natural to use an IdentityDictionary, since a Symbol is guaranteed to be globally unique. If, on the other hand, your keys are Strings, it is better to use a plain Dictionary, or you may get into trouble:

```
a := 'foobar'.
b := a copy.
trouble := IdentityDictionary new.
trouble at: a put: 'a'; at: b put: 'b'.
trouble at: a      → 'a'
trouble at: b      → 'b'
trouble at: 'foobar' → 'a'
```

Since a and b are different objects, they are treated as different objects. Interestingly, the literal 'foobar' is allocated just once, so is really the same object as a. You don't want your code to depend on behaviour like this! A plain Dictionary would give the same value for any key equal to 'foobar'.

Use only Symbols as keys for a IdentityDictionary, and Strings (or other objects) as keys for a plain Dictionary.

Note that the global Smalltalk is an instance of SystemDictionary, a subclass of IdentityDictionary, hence all its keys are Symbols (actually, ByteSymbols, which only contain 8-bit characters).

```
Smalltalk keys collect: [ :each | each class ] → a Set(ByteSymbol)
```

Sending keys or values to a Dictionary results in a Set, which we look at next.

Set

The class Set is a collection which behaves as a mathematical set, *i.e.*, as a collection with no duplicate elements and without any order. In a Set elements are added using the message add: and they cannot be accessed using the message at:. Objects put in a set should implement the methods hash and =.

```
s := Set new.
s add: 4/2; add: 4; add:2.
s size → 2
```

You can also create sets using Set class»newFrom: or the conversion message Collection»asSet:

```
(Set newFrom: #( 1 2 3 1 4 )) = #(1 2 3 4 3 2 1) asSet    →    true
```

asSet offers us a convenient way to eliminate duplicates from a collection:

```
{ Color black. Color white. (Color red + Color blue + Color green) } asSet size    →    2
```

Note that red + blue + green = white.

A Bag is much like a Set except that it does allow duplicates:

```
{ Color black. Color white. (Color red + Color blue + Color green) } asBag size    →    3
```

The set operations *union*, *intersection* and *membership* testing are implemented by the Collection messages union:, intersection: and includes:. The receiver is first converted to a Set, so these operations work for all kinds of collections!

```
(1 to: 6) union: (4 to: 10)    →    a Set(1 2 3 4 5 6 7 8 9 10)
'hello' intersection: 'there'  →    'he'
#Smalltalk includes: $k        →    true
```

As we explain below set elements are accessed using iterators (see Section 9.5).

SortedCollection

In contrast to an OrderedCollection, a SortedCollection maintains its elements in sort order. By default, a sorted collection uses the message <= to establish sort order, *i.e.*, it can sort instances of subclasses of the abstract class Magnitude, which defines the protocol of comparable objects (<, =, >, >=, between:and:...). (See Chapter 8.)

You can create a SortedCollection by creating a new instance and adding elements to it:

```
SortedCollection new add: 5; add: 2; add: 50; add: -10; yourself.    →    a
    SortedCollection(-10 2 5 50)
```

More usually, though, one will send the conversion message asSortedCollection to an existing collection:

```
#(5 2 50 -10) asSortedCollection  →  a SortedCollection(-10 2 5 50)
```

This example answers the following FAQ:

FAQ: How do you sort a collection?
 ANSWER: Send the message `asSortedCollection` to it.

```
'hello' asSortedCollection  →  a SortedCollection($e $h $l $l $o)
```

How do you get a `String` back from this result? `asString` unfortunately returns the `printString` representation, which is not what we want:

```
'hello' asSortedCollection asString  →  'a SortedCollection($e $h $l $l $o)'
```

The correct answer is to either use `String class>>newFrom:`, `String class>>withAll:` or `Object>>as::`:

```
'hello' asSortedCollection as: String      →  'ehllo'
String newFrom: ('hello' asSortedCollection) →  'ehllo'
String withAll: ('hello' asSortedCollection) →  'ehllo'
```

It is possible to have different kinds of elements in a `SortedCollection` as long as they are all comparable. For example we can mix different kinds of numbers such as integers, floats and fractions:

```
{ 5. 2/-3. 5.21 } asSortedCollection  →  a SortedCollection((-2/3) 5 5.21)
```

Imagine you want to sort objects that do not define the method `<=` or that you would like to have a different sorting criterion. You can do that by specifying a two argument block. For example, the class `Color` is not a `Magnitude` and it does not implement the method `<=`, but we can specify a block stating that the colors should be sorted according to their luminance (a measure of brightness).

```
col := SortedCollection sortBlock: [:c1 :c2 | c1 luminance <= c2 luminance].
col addAll: { Color red. Color yellow. Color white. Color black }.
col  →  a SortedCollection(Color black Color red Color yellow Color white)
```

String

A Smalltalk `String` represents a collection of `Characters`. It is sequenceable, indexable, mutable and homogeneous, containing only `Character` instances.

Like Arrays, Strings have a dedicated syntax, and are normally created by directly specifying a String literal within single quotes, but the usual collection creation methods will work as well.

'Hello'	→	'Hello'
String with: \$A	→	'A'
String with: \$h with: \$i with: \$!	→	'hi!'
String newFrom: #(\$h \$e \$! \$! \$o)	→	'hello'

In actual fact, String is abstract. When we instantiate a String we actually get either an 8-bit ByteString or a 32-bit WideString. To keep things simple, we usually ignore the difference and just talk about instances of String.

Two instances of String can be concatenated with a comma.

s := 'no', ' ', 'worries'.	
s	→ 'no worries'

Since a string is a mutable collection we can also change it using the method at:put:.

s at: 4 put: \$h; at: 5 put: \$u.	
s	→ 'no hurries'

Note that the comma method is defined by Collection, so it will work for any kind of collection!

(1 to: 3) , '45'	→	#(1 2 3 \$4 \$5)
------------------	---	------------------

We can also modify an existing string using replaceAll:with: or replaceFrom:to:with: as shown below. Note that the number of characters and the interval should have the same size.

s replaceAll: \$n with: \$N.	
s	→ 'No hurries'
s replaceFrom: 4 to: 5 with: 'wo'.	
s	→ 'No worries'

Instead, with copyReplaceAll: a new string is created. (Curiously, here the arguments are substrings rather than individual characters, and the sizes do not have to match.)

s copyReplaceAll: 'rries' with: 'mbats'	→	'No wombats'
---	---	--------------

A quick look at the implementation of these methods reveals that they are defined not only for Strings, but for any kind of SequenceableCollection, hence the following also works:

(1 to: 6) copyReplaceAll: (3 to: 5) with: { 'three'. 'etc.' }	→	#(1 2 'three' 'etc.' 6)
---	---	-------------------------

String matching It is possible to ask whether a string matches a pattern by sending the match: message. The pattern can specify * to match an arbitrary series of characters and # to match a single character. Note that match: is sent to the pattern and not the string to be matched.

'Linux #' match: 'Linux mag'	→	true
'GNU/Linux #ag' match: 'GNU/Linux tag'	→	true

Another useful method is findString:.

'GNU/Linux mag' findString: 'Linux'	→	5
'GNU/Linux mag' findString: 'linux' startingAt: 1 caseSensitive: false	→	5

More advanced pattern matching facilities offering the capabilities of Perl are also available, but they are not included in the standard image⁴.

Some tests on strings. The following examples illustrate the use of isEmpty, includes: and anySatisfy: which are further messages defined not only on Strings but more generally on collections.

'Hello' isEmpty.	
'Hello' includes: \$a	→ false
'JOE' anySatisfy: [:c c isLowercase]	→ false
'Joe' anySatisfy: [:c c isLowercase]	→ true

String templating. There are three messages that are useful to manage string templating: format:, expandMacros and expandMacrosWith:.

'{1} is {2}' format: {'Squeak' . 'cool'}	→	'Squeak is cool'
--	---	------------------

The messages of the expandMacros family offer variable substitution, using <n> for carriage return, <t> for tabulation, <1s>, <2s>, <3s> for arguments (<1p>, <2p>, surrounds the string with single quotes), and <1?value1:value2> for conditional.

'look-<t>-here' expandMacros	→	'look- -here'
'<1s> is <2s>' expandMacrosWith: 'Squeak' with: 'cool'	→	'Squeak is cool'
'<2s> is <1s>' expandMacrosWith: 'Squeak' with: 'cool'	→	'cool is Squeak'
'<1p> or <1s>' expandMacrosWith: 'Squeak' with: 'cool'	→	'"Squeak" or Squeak'
'<1?Quentin:Thibaut> plays' expandMacrosWith: true	→	'Quentin plays'
'<1?Quentin:Thibaut> plays' expandMacrosWith: false	→	'Thibaut plays'

⁴We strongly recommend Vassili Bykov's regular expression package, available at www.squeaksource.com/Regex.html.

Some other utility methods. The class `String` offers numerous other utilities including the messages `asLowercase`, `asUppercase` and `capitalized`.

```
'XYZ' asLowercase  → 'xyz'
'xyz' asUppercase  → 'XYZ'
'hilaire' capitalized → 'Hilaire'
'1.54' asNumber    → 1.54
'this sentence is without a doubt far too long' contractTo: 20 → 'this sent...too long'
```

Note that there is generally a difference between asking an object its string representation by sending the message `printString` and converting it to a string by sending the message `asString`. Here is an example of the difference.

```
#ASymbol printString → '#ASymbol'
#ASymbol asString    → 'ASymbol'
```

A symbol is similar to a string but is guaranteed to be globally unique. For this reason symbols are preferred to strings as keys for dictionaries, in particular for instances of `IdentityDictionary`. See also Chapter 8 for more about `String` and `Symbol`.

9.5 Collection Iterators

In Smalltalk loops and conditionals are simply messages sent to collections or other objects such as integers or block (see also Chapter 3). In addition to low-level messages such as `to:do:` which evaluates a block with an argument ranging from an initial to a final number, the Smalltalk collection hierarchy offers various high-level iterators. Using such iterators will make your code more robust and compact.

Iterating (do:)

The method `do:` is the basic collection iterator. It applies its argument (a block taking a single argument) to each element of the receiver. The following example prints all the strings contained in the receiver to the transcript.

```
#('bob' 'joe' 'toto') do: [:each | Transcript show: each; cr].
```

Variants. There are a lot of variants of `do:`, such as `do:without:`, `doWithIndex:` and `reverseDo:`. For the indexed collections (`Array`, `OrderedCollection`, `SortedCollection`) the method `doWithIndex:` gives also access to the current index. This method is related to `to:do:` which is defined on `Number`.

```
#('bob' 'joe' 'toto') doWithIndex: [:each :i | (each = 'joe') ifTrue: [ ↑ i ]] → 2
```

For ordered collections, `reverseDo:` walks the collection in the reverse order.

The following code shows an interesting message: `do:separatedBy:` which executes the second block only in between two elements.

```
res := "".
#('bob' 'joe' 'toto') do: [:e | res := res, e ] separatedBy: [res := res, '.'].
res → 'bob.joe.toto'
```

Note that this code is not especially efficient since it creates intermediate strings and it would be better to use a write stream to buffer the result (see Chapter 10):

```
String streamContents: [:stream | #('bob' 'joe' 'toto') asStringOn: stream delimiter: '.' ]
→ 'bob.joe.toto'
```

Dictionaries. When the method `do:` is sent to a dictionary, the elements taken into account are the values, not the associations. The proper methods to use are `keysDo:`, `valuesDo:`, and `associationsDo:`, which iterate respectively on keys, values or associations.

```
colors := Dictionary newFrom: { #yellow -> Color yellow. #blue -> Color blue. #red ->
    Color red }.
colors keysDo: [:key | Transcript show: key; cr].           "displays the keys"
colors valuesDo: [:value | Transcript show: value; cr].     "displays the values"
colors associationsDo: [:value | Transcript show: value; cr]. "displays the associations"
```

Collecting results (collect:)

If you want to process the elements of a collection and produce a new collection as a result, rather than using `do:`, you are probably better off using `collect:`, or one of the other iterator methods. Most of these can be found in the *enumerating* protocol of `Collection` and its subclasses.

Imagine that we want a collection containing the doubles of the elements in another collection. Using the method `do:` we must write the following:

```
double := OrderedCollection new.
#(1 2 3 4 5 6) do: [:e | double add: 2 * e].
double   →   an OrderedCollection(2 4 6 8 10 12)
```

The method `collect:` executes its argument block for each element and returns a new collection containing the results. Using `collect:` instead, the code is much simpler:

```
#(1 2 3 4 5 6) collect: [:e | 2 * e]   →   #(2 4 6 8 10 12)
```

The advantages of `collect:` over `do:` are even more dramatic in the following example, where we take a collection of integers and generate as a result a collection of absolute values of these integers:

```
aCol := #( 2 -3 4 -35 4 -11).
result := aCol species new: aCol size.
1 to: aCol size do: [ :each | result at: each put: (aCol at: each) abs].
result   →   #(2 3 4 35 4 11)
```

Contrast the above with the much simpler following expression:

```
#( 2 -3 4 -35 4 -11) collect: [:each | each abs ]   →   #(2 3 4 35 4 11)
```

A further advantage of the second solution is that it will also work for sets and bags.

Generally you should avoid using `do:`, unless you want to send messages to each of the elements of a collection.

Note that sending the message `collect:` returns the same kind of collection as the receiver. For this reason the following code fails. (A String cannot hold integer values.)

```
'abc' collect: [:ea | ea asciiValue ]   "error!"
```

Instead we must first convert the string to an Array or an OrderedCollection:

```
'abc' asArray collect: [:ea | ea asciiValue ]   →   #(97 98 99)
```

Actually `collect:` is not guaranteed to return a collection of exactly the same class as the receiver, but only the same “*species*”. In the case of an Interval, the species is actually an Array!

```
(1 to: 5) collect: [ :ea | ea * 2 ]   →   #(2 4 6 8 10)
```

Selecting and Rejecting Elements

select: returns the elements of the receiver that satisfy a particular condition:

```
(2 to: 20) select: [:each | each isPrime]    →  #(2 3 5 7 11 13 17 19)
```

reject: does the opposite:

```
(2 to: 20) reject: [:each | each isPrime]   →  #(4 6 8 9 10 12 14 15 16 18 20)
```

Identifying an element with detect:

The method **detect:** returns the first element of the receiver that matches block argument.

```
'through' detect: [:each | each isVowel]    →  $o
```

The method **detect:ifNone:** is a variant of the method **detect:.** Its second block is evaluated when there is no element matching the block.

```
Smalltalk allClasses detect: [:each | 'java*' match: each asString] ifNone: [ nil ]
→ nil
```

Accumulating results with inject:into:

Functional programming languages often provide a higher-order function called *fold* or *reduce* to accumulate a result by applying some binary operator iteratively over all elements of a collection. In Squeak this is done by **Collection>>inject:into:.**

The first argument is an initial value, and the second argument is a two-argument block which is applied to the result this far, and each element in turn.

A trivial application of **inject:into:** is to produce the sum of a collection of numbers. Following Gauss, in Squeak we could write this expression to sum the first 100 integers:

```
(1 to: 100) inject: 0 into: [:sum :each | sum + each ]    →  5050
```

Another example is the following one-argument block which computes factorials:

```
factorial := [:n | (1 to: n) inject: 1 into: [:product :each | product * each ] ].
factorial value: 10    →  3628800
```

Other messages

count: The message count: returns the number of elements matching a block.

```
Smalltalk allClasses count: [:each | '*Collection*' match: each asString ]    →    14
```

includes: The message includes: checks whether the argument is contained in the collection.

```
colors := {Color white . Color yellow. Color red . Color blue . Color orange}.
colors includes: Color blue.    →    true
```

anySatisfy: The message anySatisfy: returns whether at least one element satisfies a condition.

```
colors anySatisfy: [:c | c red > 0.5]    →    true
```

9.6 Some Hints for using Collections

A common mistake with add: The following error is one of the most frequent Smalltalk mistakes.

```
collection := OrderedCollection new add: 1; add: 2.
collection    →    2
```

Here the variable collection does not hold the newly created collection but rather the last number added. Indeed the method add: returns the element added and not the receiver.

The following code yields the expected result:

```
collection := OrderedCollection new.
collection add: 1; add: 2.
collection    →    an OrderedCollection(1 2)
```

You can also use the message yourself to return the receiver of a cascade of messages:

```
collection := OrderedCollection new add: 1; add: 2; yourself    →    an
    OrderedCollection(1 2)
```

Removing an element of the collection you are iterating on. Another mistake you may make is to remove an element from a collection you are currently iterating over. `remove`:

```
range := (2 to: 20) asOrderedCollection.
range do: [:aNumber | aNumber isPrime ifFalse: [ range remove: aNumber ] ].
range  →  an OrderedCollection(2 3 5 7 9 11 13 15 17 19)
```

This result is clearly incorrect since 9 and 15 should have been filtered out!

The solution is to copy the collection before going over it.

```
range := (2 to: 20) asOrderedCollection.
range copy do: [:aNumber | aNumber isPrime ifFalse: [ range remove: aNumber ] ].
range  →  an OrderedCollection(2 3 5 7 11 13 17 19)
```

Redefining both = and hash. A difficult error to spot is when you redefine = but not hash. The symptoms are that you will lose elements that you put in sets or other strange behavior. One solution proposed by Kent Beck is to using `xor`: to redefine hash. Suppose that we want two books to be considered equal if their titles and authors are the same. Then we would redefine not only = but also hash as follows:

Method 9.1: *Redefining = and hash.*

```
Book»= aBook
self class = aBook class ifFalse: [↑ false].
↑ title = aBook title and: [ authors = aBook authors]
```

```
Book»hash
↑ title hash xor: authors hash
```

Another nasty problem that can arise is to use mutable objects, *i.e.*, that may change their hash value over time, as elements of a Set or as keys of a Dictionary. Don't do this unless you love debugging!

9.7 Chapter Summary

The Smalltalk collection hierarchy provides a common vocabulary for uniformly manipulating a variety of different kinds of collections.

- A key distinction is between `SequenceableCollections`, which maintain their elements in a given order, `Dictionary` its subclasses, which maintain key-to-value associations, and `Sets` and `Bags`, which are unordered.

- You can convert most collections to another kind of collection by sending them the messages `asArray`, `asOrderedCollection` etc..
- To sort a collection, send it the message `asSortedCollection`.
- Literal Arrays are created with the special syntax `#(...)`. Dynamic Arrays are created with the syntax `{ ... }`.
- A Dictionary compares keys by equality. It is most useful when keys are instances of `String`. An `IdentityDictionary` instead uses object identity to compare keys. It is more suitable when `Symbols` are used as keys, or when mapping object references to values.
- Strings also understand the usual collection messages. In addition, a `String` supports a simple form of pattern-matching. For more advanced application, look instead at the `RegEx` package.
- The basic iteration message is `do:`. It is useful for imperative code, such as modifying each element of a collection, or sending each element a message.
- Instead of using `do:`, it is more common to use `collect:`, `select:`, `reject:`, `includes:`, `inject:into:` and other higher-level messages to process collections in a uniform way.
- Never remove an element from a collection you are iterating over. If you must modify it, iterate over a copy instead.
- If you override `=`, remember to override `hash` as well!

Chapter 10

Streams

Stream Streams are used to iterate over sequences of elements such as sequenced collections, files, and network streams. Streams may be either readable, or writeable, or both. Reading or writing is always relative to the current position in the stream. Streams can easily be converted to collections, and vice versa.

10.1 Two Sequences of Elements

A good metaphor to understand a stream is the following: A stream can be represented as two sequences of elements: a past element sequence and a future element sequence. The stream is positioned between the two sequences. Understanding this model is important since all stream operations in Smalltalk rely on it. For this reason, most of the Stream classes are subclasses of PositionableStream. Figure 10.1 presents a stream which contains five characters. This stream is in its original position, *i.e.*, there is no element in the past. You can go back to this position using the message `reset`.




Figure 10.1: A stream positioned at its beginning.

Reading an element conceptually means removing the first element of

the future element sequence and putting it after the last element in the past element sequence. After having read one element using the message `next`, the state of your stream is that shown in Figure 10.2.




Figure 10.2: The same stream after the execution of the method `next`: the character `a` is “in the past” whereas `b`, `c`, `d` and `e` are “in the future”.

Writing an element means replacing the first element of the future sequence by the new one and moving it to the past. Figure 10.3 shows the state of the same stream after having written an `x` using the message `nextPut: anElement`.




Figure 10.3: The same stream after having written an `x`.

10.2 Streams vs. Collections

The collection protocol supports the storage, removal and enumeration of the elements of a collection, but does not allow these operations to be intermingled. For example, if the elements of an `OrderedCollection` are processed by a `do:` method, it is not possible to add or remove elements from inside the `do:` block. Nor does the collection protocol offer ways to iterate over two collections at the same time, choosing which collection goes forward and which does not. Procedures like these require that a traversal index or position reference is maintained outside of the collection itself: this is exactly the role of `ReadStream`, `WriteStream` and `ReadWriteStream`.

These three classes are defined to *stream over* some collection. For example, the following snippet creates a stream on an interval, then it reads two elements.

```

r := ReadStream on: (1 to: 1000).
r next.    → 1
r next.    → 2
r atEnd.   → false

```

WriteStreams can write data to the collection:

```

w := WriteStream on: (String new: 5).
w nextPut: $a.
w nextPut: $b.
w contents. → 'ab'

```

It is also possible to create ReadWriteStreams that support both the reading and writing protocols.

The main problem with WriteStream and ReadWriteStream is that they only support arrays and strings in Squeak. This is currently being changed by the development of a new library named Nile, but for now if you try to stream over another kind of collection, you will get an error:

```

w := WriteStream on: (OrderedCollection new: 20).
w nextPut: 12. → raises an error

```

Streams are not only meant for collections, they can be used for files or sockets too. The following example creates a file named test.txt, writes two strings to it, separated by a carriage return, and closes the file.

```

StandardFileStream
  fileName: 'test.txt'
  do: [:str | str
    nextPutAll: '123';
    cr;
    nextPutAll: 'abcd'].

```

The following sections present the protocols in more depth.

10.3 Streaming over Collections

Streams are really useful when dealing with collections of elements. They can be used for reading and writing elements in collections. We will now explore the stream features for the collections.

Reading Collections

This section presents features used for reading collections. Using a stream to read a collection essentially provides you a pointer into the collection. That pointer will move forward on reading and you can place it wherever you want. The class `ReadStream` should be used to read elements from collections.

Methods `next` and `next:` are used to retrieve one or more elements from the collection.

```
stream := ReadStream on: #(1 (a b c) false).
stream next.    → 1
stream next.    → #(a #b #c)
stream next.    → false
```

```
stream := ReadStream on: 'abcdef'.
stream next: 0. → ""
stream next: 1. → 'a'
stream next: 3. → 'bcd'
stream next: 2. → 'ef'
```

The message `peek` is used when you want to know what is the next element in the stream without going forward.

```
stream := ReadStream on: '-143'.
negative := (stream peek = $-).  "look at the first element without reading it"
negative. → true
negative ifTrue: [stream next].  "ignores the minus character"
number := stream upToEnd.
number. → '143'
```

This code sets the boolean variable `negative` according to the sign of the number in the stream and `number` to its absolute value. The method `upToEnd` returns everything from the current position to the end of the stream and sets the stream to its end. This code can be simplified using `peekFor:`, which moves forward if the following element equals the parameter and doesn't move otherwise.

```
stream := '-143' readStream.
(stream peekFor: $-) → true
stream upToEnd      → '143'
```

`peekFor:` also returns a boolean indicating if the parameter equals the element.

You might have noticed a new way of constructing a stream in the above example: one can simply send `readStream` to a sequenceable collection to get a reading stream on that particular collection.

Positioning. There are methods to position the stream pointer. If you have the index, you can go directly to it using `position::`. You can request the current position using `position`. Please remember that a stream is not positioned on an element, but between two elements. The index corresponding to the beginning of the stream is 0.

You can obtain the state of the stream depicted in Figure 10.4 with the following code:

```
stream := 'abcde' readStream.
stream position: 2.
stream peek    →    $c
```




Figure 10.4: A stream at position 2

If you want to go to the beginning or at the end is what you want, you can use `reset` or `setToEnd`. `skip:` and `skipTo:` are used to go forward to a location relative to the current position: `skip:` accepts a number as argument and skips that number of elements whereas `skipTo:` skips all elements in the stream until it finds an element equals to its parameter. Note that it positions the stream after the matched element.

```
stream := 'abcdef' readStream.
stream next.      →    $a    "stream is now positioned just after the a"
stream skip: 3.   →          "stream is now after the d"
stream position.  →    4
stream skip: -2.  →          "stream is after the b"
stream position.  →    2
stream reset.
stream position.  →    0
stream skipTo: $e. →          "stream is just after the e now"
stream next.      →    $f
stream contents.  →    'abcdef'
```

As you can see, the letter `e` has been skipped.

The method `contents` always returns a copy of the entire stream.

Testing. Some methods allow you to test the state of the current stream: `atEnd` returns true if and only if no more elements can be read whereas `isEmpty` returns true if and only if there is no element at all in the collection.

Here is a possible implementation of an algorithm using `atEnd` that takes two sorted collections as parameters and merges those collections into another sorted collection:

```
stream1 := #(1 4 9 11 12 13) readStream.
stream2 := #(1 2 3 4 5 10 13 14 15) readStream.
```

```
"The variable result will contain the sorted collection."
result := OrderedCollection new.
[stream1 atEnd not & stream2 atEnd not]
  whileTrue: [stream1 peek < stream2 peek
    "Remove the smallest element from either stream and add it to the result."
    ifTrue: [result add: stream1 next]
    ifFalse: [result add: stream2 next]].
```

```
"One of the two streams might not be at its end. Copy whatever remains."
result
  addAll: stream1 upToEnd;
  addAll: stream2 upToEnd.
```

```
result.  →  an OrderedCollection(1 1 2 3 4 4 5 9 10 11 12 13 13 14 15)
```

Writing to Collections

We have already seen how to read a collection by iterating over its elements using a `ReadStream`. We'll now learn how to create collections using `WriteStreams`.

`WriteStreams` are useful for appending a lot of data to a collection at various locations. They are often used to construct strings that are based on static and dynamic parts as in this example:

```
stream := String new writeStream.
stream
  nextPutAll: 'This Smalltalk image contains: ';
  print: Smalltalk allClasses size;
  nextPutAll: ' classes.';
  cr;
  nextPutAll: 'This is really a lot.'
```

```
stream contents.  →  'This Smalltalk image contains: 2322 classes.
This is really a lot.'
```

This technique is used in the different implementations of the method `printOn:` for example. There is a simpler and more efficient way of creating streams if you are only interested in the content of the stream:

```
string := String streamContents:
  [:stream |
    stream
      print: #(1 2 3);
      space;
      nextPutAll: 'size';
      space;
      nextPut: $=;
      space;
      print: 3. ].
string.  →  '#(1 2 3) size = 3'
```

The method `streamContents:` creates a collection and a stream on that collection for you. It then executes the block you gave passing the stream as a parameter. When the block ends, `streamContents:` returns the content of the collection.

The following `WriteStream` methods are especially useful in this context:

nextPut: adds the parameter to the stream;

nextPutAll: adds each element of the collection, passed as a parameter, to the stream;

print: adds the textual representation of the parameter to the stream.

There are also methods useful for printing different kinds of characters to the stream like `space`, `tab` and `cr` (carriage return). Another useful method is `ensureASpace` which ensures that the last character in the stream is a space; if the last character isn't a space it adds one.

About Concatenation. Using `nextPut:` and `nextPutAll:` on a `WriteStream` is often the best way to concatenate characters. Using the comma concatenation operator (`,`) is far less efficient:

```
[| temp |
  temp := String new.
  (1 to: 100000)
    do: [:i | temp := temp, i asString, ' ']] timeToRun  →  115176 "(ms)"

[| temp |
  temp := WriteStream on: String new.
```

```
(1 to: 100000)
do: [:i | temp nextPutAll: i asString; space].
temp contents] timeToRun    →    1262 "(milliseconds)"
```

The reason that using a stream can be much more efficient is that comma creates a new string containing the concatenation of the receiver and the argument, so it must copy both of them. When you repeatedly concatenate onto the same receiver, it gets longer and longer each time, so that the number of characters that must be copied goes up exponentially. This also creates a lot of garbage, which must be collected. Using a stream instead of string concatenation is a well-known optimization. In fact, you can use `streamContents:` (mentioned on page 223) to help you do this:

```
String streamContents: [ :tempStream |
  (1 to: 100000)
  do: [:i | tempStream nextPutAll: i asString; space]]
```

Reading and writing at the same time

It's possible to use a stream to access a collection for reading and writing at the same time. Imagine you want to create an `History` class which will manage backward and forward buttons in a web browser. A history would react as in figures from 10.5 to 10.11.



Figure 10.5: A new history is empty. Nothing is displayed in the web browser.




Figure 10.6: The user opens to page 1.

This behavior can be implemented using a `ReadWriteStream`.



Figure 10.7: The user clicks on a link to page 2.



Figure 10.8: The user clicks on a link to page 3.



Figure 10.9: The user clicks on the back button. He is now viewing page 2 again.




Figure 10.10: The user clicks again the back button. Page 1 is now displayed.



Figure 10.11: From page 1, the user clicks on a link to page 4. The history forgets pages 2 and 3.

```
classVariableNames: "
poolDictionaries: "
category: 'SBE-Streams'
```

```
History>>initialize
  super initialize.
  stream := ReadWriteStream on: Array new.
```

Nothing really difficult here, we define a new class which contains a stream. The stream is created during the initialize method.

We need methods to go backward and forward:

```
History>>goBackward
  self canGoBackward ifFalse: [self error: 'Already on the first element'].
  ↑ stream back
```

```
History>>goForward
  self canGoForward ifFalse: [self error: 'Already on the last element'].
  ↑ stream next
```

Until then, the code was pretty straightforward. Now, we have to deal with the `goTo:` method which should be activated when the user clicks on a link. A possible solution is:

```
History>>goTo: aPage
  stream nextPut: aPage.
```

This version is incomplete however. This is because when the user clicks on the link, there should be no more future pages to go to, *i.e.*, the forward button must be deactivated. To do this, the simplest solution is to write `nil` just after to indicate the history end:

```
History>>goTo: anObject
  stream nextPut: anObject.
  stream nextPut: nil.
  stream back.
```

Now, only methods `canGoBackward` and `canGoForward` have to be implemented.

A stream is always positioned between two elements. To go backward, there must be two pages before the current position: one page is the current page, and the other one is the page we want to go to.

```
History>>canGoBackward
  ↑ stream position > 1
```

```
History>>canGoForward
```

```
↑ stream atEnd not and: [stream peek notNil]
```

Let us add a method to peek at the contents of the stream:

```
History>>contents
```

```
↑ stream contents
```

And the history works as advertised:

```
History new
```

```
  goTo: #page1;
```

```
  goTo: #page2;
```

```
  goTo: #page3;
```

```
  goBackward;
```

```
  goBackward;
```

```
  goTo: #page4;
```

```
  contents  →  (#page1 #page4 nil nil)
```

10.4 Using Streams for File Access

You have already seen how to stream over collections of elements. It's also possible to stream over files on your hard disk. Once created, a stream on a file is really like a stream on a collection: you will be able to use the same protocol to read, write or position the stream. The main difference appears in the creation of the stream. There are several different ways to create file streams, as we shall now see.

Creating File Streams

To create file streams, you will have to use one of the following instance creation methods offered by the class `FileStream`:

fileNamed: Open a file with the given name for reading and writing. If the file already exists, its prior contents may be modified or replaced, but the file will not be truncated on close. If the name has no directory part, then the file will be created in the default directory.

newFileNamed: Create a new file with the given name, and answer a stream opened for writing on that file. If the file already exists, ask the user what to do.

forceNewFileNamed: Create a new file with the given name, and answer a stream opened for writing on that file. If the file already exists, delete it without asking before creating the new file.

oldFileNamed: Open an existing file with the given name for reading and writing. If the file already exists, its prior contents may be modified or replaced, but the file will not be truncated on close. If the name has no directory part, then the file will be created in the default directory.

readOnlyFileNamed: Open an existing file with the given name for reading.

You have to remember that each time you open a stream on a file, you have to close it too. This is done through the close method.

```
stream := FileStream forceNewFileNamed: 'test.txt'.
stream
  nextPutAll: 'This text is written in a file named ';
  print: stream localName.
stream close.
```

```
stream := FileStream readOnlyFileNamed: 'test.txt'.
stream contents.    → 'This text is written in a file named "test.txt"'
stream close.
```

The method `localName` answers the last component of the name of the file. You can also access the full path name using the method `fullName`.

You will soon notice that manually closing the file stream is painful and error-prone. That's why `FileStream` offers a message called `forceNewFileNamed:do:` to automatically close a new stream after evaluating a block that sets its contents.

```
FileStream
  forceNewFileNamed: 'test.txt'
  do: [:stream |
    stream
      nextPutAll: 'This text is written in a file named ';
      print: stream localName].
string := FileStream
  readOnlyFileNamed: 'test.txt'
  do: [:stream | stream contents].
string → 'This text is written in a file named "test.txt"'
```

The stream-creation methods that take a block as an argument first create a stream on a file, then execute the block with the stream as an argument, and finally close the stream. These methods return what is returned by the block, which is to say, the value of the last expression in

the block. This is used in the previous example to get the content of the file and put it in the variable string.

Binary Streams

By default, created streams are text-based which means you will read and write characters. If your stream must be binary, you have to send the message `binary` to your stream.

When your stream is in binary mode, you can only write numbers from 0 to 255 (1 Byte). If you want to use `nextPutAll:` to write more than one number at a time, you have to pass a `ByteArray` as argument.

```
FileStream
  forceNewFileNamed: 'test.bin'
  do: [:stream |
    stream
      binary;
      nextPutAll: #(145 250 139 98) asByteArray].

FileStream
  readOnlyFileNamed: 'test.bin'
  do: [:stream |
    stream binary.
    stream size.      → 4
    stream next.      → 145
    stream upToEnd.   → a ByteArray(250 139 98)
  ].
```

Here is another example which creates a picture in a file named “test.pgm”. You can open this file with your favorite drawing program.

```
FileStream
  forceNewFileNamed: 'test.pgm'
  do: [:stream |
    stream
      nextPutAll: 'P5'; cr;
      nextPutAll: '4 4'; cr;
      nextPutAll: '255'; cr;
      binary;
      nextPutAll: #(255 0 255 0) asByteArray;
      nextPutAll: #(0 255 0 255) asByteArray;
      nextPutAll: #(255 0 255 0) asByteArray;
      nextPutAll: #(0 255 0 255) asByteArray
  ]
```

This creates a 4x4 checkerboard as shown in Figure 10.12.




Figure 10.12: A 4x4 checkerboard you can draw using binary streams.

10.5 Chapter Summary

Streams offer a better way than collections to incrementally read and write a sequence of elements. There are easy ways to convert back and forth between streams and collections.

- Streams may be either readable, writeable or both readable and writeable.
- To convert a collection to a stream, define a stream “on” a collection, *e.g.*, `ReadStream on: (1 to: 1000)`, or send the messages `readStream`, etc. to the collection.
- To convert a stream to a collection, send the message `contents`.
- To concatenate large collections, instead of using the comma operator, it is more efficient to create a stream, append the collections to the stream with `nextPutAll:`, and extract the result by sending `contents`.
- File streams are by default character-based. Send `binary` to explicitly make them binary.


Chapter 11

Morphic

Morphic is the basic component of the graphical interface in Squeak. It is written in Smalltalk so it is fully portable between operating systems. The main particularity of Morphic is that the user can assemble and disassemble graphical interface elements on the fly.¹

11.1 The History of Morphic

Morphic was developed in the 80's for the Self language, which was largely inspired by Smalltalk; Morphic was then ported to Squeak. The basic idea in Morphic is that each object should be directly represented and manipulated as a graphical shape: the morph. For instance, if you bring up the world menu, you can then detach a menu item to keep it as a button: blue-click once on the world menu to bring up its morphic halo, then blue-click again on the menu item you want to detach to bring up its halo, and move elsewhere it by grabbing it by the black handle.

 *Let's take another example; to transform an object into a morph, execute the following code in a workspace:*

```
'Morph' asMorph openInWorld
```

You can also try this:


```
'Morph' asMorph openViewerForArgument
```

¹We thank Hilaire Fernandes for permission to translate this chapter from his original article in French.



Figure 11.1: Detaching a morph, here the `new morph` menu item, to make it an independent button.

You should obtain a graphical element that you can manipulate. Of course it is possible to define morphs that are better graphical representations. The method `asMorph` has a default implementation in the `Object` class that just creates a string-like morph. For instance, `Color tan asMorph` returns a string-like morph labeled with the result of `Color tan printString`; it would be more interesting if we got a colored rectangle morph instead. This is easy:

 Open the `Color` class and add the following method to it:

Method 11.1: Getting a morph for an instance of `Color`

`Color»asMorph`

↑ `Morph new color: self`

Now execute `Color orange asMorph openInWorld` in a workspace. Instead of the string-like morph, you get a nice orange rectangle!

11.2 Manipulating morphs

Morphs are objects, so we can manipulate them like any other object in Smalltalk: changing their properties, creating new morph classes, etc.

As a graphical object, a morph has a position and a size — those of its bounding box actually. The `position` method returns the coordinates of the morph's upper left corner, as a point. The origin of the coordinates is the screen's upper left corner. The `extent` method also returns a point, but this one specifies the width and height of the morph rather than a location.

 Type and execute the following code:

```
joe := Morph new color: Color blue.
joe openInWorld.
```

```
bill := Morph new color: Color red .
bill openInWorld.
```

Then type `joe position` and *print it*. To move joe, execute `joe position: (joe position + (10@3))` repeatedly.

It is possible to do the same with the size: `joe extent` returns the morph's size, and to have it grow, execute `joe extent: (joe extent * 1.1)`. To change the color of a morph, send it the `color:` message with the desired color as argument. For instance, `joe color: Color orange` or, to add transparency, `joe color: (Color orange alpha: 0.5)`.

To make bill follow joe, you can repeatedly execute this code:

```
bill position: (joe position + (100@0))
```


You can move joe using the mouse then execute this code; bill will move to 100 pixels to the right of joe.

11.3 Composing Morphs

To create new morphs, there are two main techniques that you can combine:

1. by composing morphs one into another,
2. by subclassing `Morph` and overriding `drawOn:` to draw original morph shapes.

Morphs can contain other morphs. You can add a morph to another one by sending the `addMorph:` to the container morph.

 *Try adding a morph to another one:*

```
star := StarMorph new color: Color yellow.
joe addMorph: star.
star position: joe position.
```

The last line positions the star at the same coordinates as joe. The coordinates of the contained morph are still absolute: they are relative to the screen, not to the containing morph (this is a limitation of `Morphic`). To position a morph, there are plenty of methods available; browse the *geometry* protocol of the `Morph` class to see for yourself. To place the star in the middle of joe, execute `star center: joe center`.

Now if you grab the star with the mouse, you actually grab joe, and the two morphs move simultaneously: the star is anchored inside joe. It is




Figure 11.2: The star is contained inside joe, the transparent blue morph.

indeed possible to add more morphs to joe; to remove a sub-morph, do joe removeMorph: star or star delete.

11.4 Creating and Drawing Your Own Morphs

While it is possible to obtain interesting representations by composing morphs, sometimes you will really need something completely different. To do this you need to define a subclass of Morph and to override the drawOn: method. The parameter to this method is an instance of Canvas on which the morph will draw itself. Let's draw a cross-shaped morph.

 From the class browser, define a new class CrossMorph inheriting from Morph:

Class 11.2: *Defining CrossMorph*

```
Morph subclass: #CrossMorph
  instanceVariableNames: "
  classVariableNames: "
  poolDictionaries: "
  category: 'sqByExample-Morphs'
```

We will the define the drawOn: method like this:

Method 11.3: *How to draw CrossMorphs*


```
drawOn: aCanvas
  | aBounds crossHeight crossWidth |
  crossHeight := (self height / 3) asFloat.
  crossWidth := (self width / 3) asFloat.
  aBounds := self bounds top: self top + crossHeight.
  aBounds := aBounds bottom: self bottom - crossHeight.
  aCanvas fillRectangle: aBounds color: self color.
```

```

aBounds := self bounds left: self left + crossWidth.
aBounds := aBounds right: self right - crossWidth.
aCanvas fillRectangle: aBounds color: self color

```

The `bound` method returns the bounding box of a morph. This is an instance of `Rectangle` defined by two points: the upper left and the lower right corner of the morph.


 To test our new morph, execute `CrossMorph new openInWorld`.

The result should look something like Figure 11.3.



Figure 11.3: The cross-shaped morph; you can resize it as you wish.

Note however that the sensitive zone — where you can click to grab the morph — is still the bounding box.

 To limit the sensitive zone of the morph to the cross shape, you need to redefine the `containsPoint:` method:

Method 11.4: *Shaping the sensitive zone of the CrossMorph*

```

containsPoint: aPoint
| aBounds1 aBounds2 crossHeight crossWidth |
crossHeight := (self height / 3) asFloat.
crossWidth := (self width / 3) asFloat.
aBounds1 := self bounds top: self top + crossHeight.
aBounds1 := aBounds1 bottom: self bottom - crossHeight.
aBounds2 := self bounds left: self left + crossWidth.
aBounds2 := aBounds2 right: self right - crossWidth.
↑ (aBounds1 containsPoint: aPoint)
| (aBounds2 containsPoint: aPoint)


```

This leverages the `containsPoint:` method from class `Rectangle`. There is duplicated code in the `drawOn:` and `containsPoint:` methods. This code should be factored out into a third method; we leave this as an exercise for the reader.

11.5 Mouse and Keyboard Interaction

Mouse Events Besides the graphical aspect of morphs, we also need to interact with them through mouse or the keyboard. To handle mouse events, a morph must answer `true` to the `handlesMouseDown:` event: when a mouse button is pressed, `Morphic` queries each morph under the mouse pointer using this message. If a morph answers `true`, then `Morphic` immediately sends it the `mouseDown:` message, and the `mouseUp:` message once the user releases the mouse button. If all morphs answer `false`, then `Morphic` initiates a drag-and-drop operation, which is the default behavior to grab morphs. As we will show below, the `mouseDown:` and `mouseUp` methods each take an instance of `MorphicEvent` as argument.

We will now extend `CrossMorph` to handle mouse events.

 First, `CrossMorph` should accept mouse events, so add this method to the class:

Method 11.5: *Declaring that `CrossMorph` will react to mouse clicks*

```
CrossMorph>handlesMouseDown: anEvent
↑true
```

Now when the left mouse button — red in the Smalltalk terminology — is clicked, we want to change the color of the cross to red; when the middle button — yellow in Smalltalk — is clicked we will change the color to yellow:

Method 11.6: *Reacting to mouse clicks by changing the morph's color*


```
CrossMorph>mouseDown: anEvent
anEvent redButtonPressed
ifTrue: [self color: Color red].
anEvent yellowButtonPressed
ifTrue: [self color: Color yellow].
self changed
```

Once the morph handles mouse events, it you cannot grab it with the mouse and move it anymore. Instead you have to use the halos: right click on the morph to make the halos appear and grab either the brown or the black handle at the top of the morph.

The `anEvent` argument of `mouseDown:` is an instance of `MouseEvent`, which is a subclass of `MorphicEvent`. `MouseEvent` defines the `redButtonPressed` and `yellowButtonPressed` methods. Browse this class to see what methods it provides to handle mouse events.

Keyboard Events To catch keyboard events, there are several steps:


1. give the keyboard focus to a specific morph: for instance we can give focus to our morph when the mouse is over it;
2. handle the keyboard event itself with the `handleKeystroke:` method: this message is sent to our morph when the user presses a key and if our morph currently has the focus;
3. release the keyboard focus when the mouse is not over our morph anymore.

 We will extend `CrossMorph` so that it reacts to keystrokes. For a morph to be notified when the mouse pointer is over it, it should answer true to the `handlesMouseOver:` message:

Method 11.7: Declaring that `CrossMorph` will react when it is under the mouse pointer

```
CrossMorph>handlesMouseOver: anEvent
↑true
```

This message is the equivalent of `handlesMouseDown:` for the mouse position. When the mouse pointer enters or leaves the morph area, the `mouseenter:` and `mouseleave:` messages are sent to the morph, respectively.

 Define two corresponding methods so that `CrossMorph` catches and releases the keyboard focus, as well as the method to actually handle the keystrokes:

Method 11.8: Getting the keyboard focus when the mouse pointer enters the morph's area

```
CrossMorph>mouseenter: anEvent
anEvent hand newKeyboardFocus: self
```

Method 11.9: Handing back the focus when the pointer goes away

```
CrossMorph>mouseleave: anEvent
anEvent hand newKeyboardFocus: nil
```

Method 11.10: *Receiving and handling keyboard events*


```

CrossMorph>>handleKeystroke: anEvent
| keyValue |
keyValue := anEvent keyValue.
keyValue = 30    "up arrow"
    ifTrue: [self position: self position - (0 @ 1)].
keyValue = 31    "down arrow"
    ifTrue: [self position: self position + (0 @ 1)].
keyValue = 29    "right arrow"
    ifTrue: [self position: self position + (1 @ 0)].
keyValue = 28    "left arrow"
    ifTrue: [self position: self position - (1 @ 0)]

```

We have written the method so that we can move the morph using the arrow keys. Note that when the mouse is not over the morph anymore, the `handleKeystroke:` message is not sent anymore, so the morph stops responding to keyboard commands. To know the key values, you can open a Transcript window and add `Transcript show: anEvent keyValue.` in the method above. The `anEvent` argument of `handleKeystroke:` is an instance of `KeyboardEvent`, another subclass of `MorphicEvent`. Browse this class to learn more about keyboard event handling methods.

Morph Animations Morphic provides a simple animation system made of two methods: `step` is called at regular intervals of time, and `stepTime` specifies the periodicity in milliseconds at which `step` is to be called.

 To make `CrossMorph` blink, we can define these methods as follows:

Method 11.11: *Defining the animation time interval*

```

CrossMorph>>stepTime
↑ 100

```

Method 11.12: *Making a step in the animation*

```

CrossMorph>>step
(self color diff: Color black) < 0.1
    ifTrue: [self color: Color red]
    ifFalse: [self color: self color darker]

```

NB: You may need to pick up and move the `CrossMorph` to trigger the animation.

11.6 Interactors

To prompt the user for input, the `FillInTheBlank` class provides a few ready-to-use dialog boxes. For instance the `request:initialAnswer:` method returns the string entered by the user (Figure 11.4).




Figure 11.4: Dialog displayed by `FillInTheBlank request: 'What's your name?' initialAnswer: 'no name'`.





Figure 11.5: PopUp displayed by `PopupMenu>startUpWithCaption:.`

To display a pop-up menu, use the `PopupMenu` class:

```
menu := PopUpMenu
  labelArray: #'(circle' 'oval' 'square' 'rectangle' 'triangle')
  lines: #(2 4).
menu startUpWithCaption: 'Choose a shape'
```

11.7 Drag-and-drop

Morphic also supports drag-and-drop. Let's examine a simple example with two morphs, a receiver morph and a dropped morph. The former will accept a morph only if the dropped morph matches a given condition: here the morph should be blue. If it is rejected, the dropped morph decides what it should do.

 *Let's first define the receiver morph:*

Class 11.13: Defining a morph on which we can drop other morphs

```
Morph subclass: #ReceiverMorph
  instanceVariableNames: ""
  classVariableNames: ""
  poolDictionaries: ""
  category: 'sqByExample-Morphs'
```

Now define the initialization method as usual:

Method 11.14: *Initializing ReceiverMorph*

```
ReceiverMorph>initialize
  super initialize.
  color := Color red.
  bounds := 0 @ 0 extent: 200 @ 200
```

The following methods decide if the receiver morph accepts the dropped morph or not; `repelsMorph:event:` defines the basic, first-hand policy, and in most cases will ensure agreement of both the receiver and the dropped morph, via `wantsDroppedMorph:event:` and its counterpart `wantsToBeDroppedInto:`.


Method 11.15: *Changing the rejection policy for dropped morphs*

```
ReceiverMorph>repelsMorph: aMorph event: ev
  ↑ (self wantsDroppedMorph: aMorph event: ev) not
```

Method 11.16: *Accepting dropped morphs based on their color*


```
ReceiverMorph>wantsDroppedMorph: aMorph event: anEvent
  ↑ aMorph color = Color blue
```

That's all we need as far as the receiver is concerned.

 *Create instances of ReceiverMorph and RectangleMorph in a workspace:*

```
ReceiverMorph new openInWorld.
RectangleMorph new openInWorld.
```

Try to drag-and-drop the blue RectangleMorph onto the receiver. It will be rejected and sent back to its initial position.

 *To change this behavior, we create a specific subclass of Morph, named DroppedMorph:*

Class 11.17: *Defining a morph we can drag-and-drop onto ReceiverMorph*

```
Morph subclass: #DroppedMorph
  instanceVariableNames: ""
  classVariableNames: ""
  poolDictionaries: ""
  category: 'sqByExample-Morphs'
```

Method 11.18: *Initializing DroppedMorph*


```
DroppedMorph»initialize
  super initialize.
  color := Color blue.
  self position: 250@100
```

Now we can specify what the dropped morph should do when it is rejected by the receiver; here it will stay attached to the mouse pointer:

Method 11.19: *Reacting when the morph was dropped but rejected*

```
DroppedMorph»rejectDropMorphEvent: anEvent
| h |
h := anEvent hand.
WorldState
  addDeferredUIMessage: [h grabMorph: self].
anEvent wasHandled: true
```

The hand method of events returns an instance of HandMorph which represents the mouse pointer and whatever it holds. We tell the World that the mouse pointer should grab self, the rejected morph.


 Create two instances of DroppedMorph, and then drag-and-drop them onto the receiver.

```
ReceiverMorph new openInWorld.
(DroppedMorph new color: Color blue) openInWorld.
(DroppedMorph new color: Color green) openInWorld.
```

The green morph is rejected and therefore stays attached to the mouse pointer.

11.8 A Complete Example

We will design a morph to roll a die². Clicking on it will display the values of all sides of the die in a quick loop, and another click will stop the animation.

 We define the die as a subclass of BorderedMorph instead of Morph, because we will use some border attributes of the former:

²NB: One die, two dice.



Figure 11.6: The die in Morphe.

Class 11.20: *Defining the die morph*

```
BorderedMorph subclass: #DieMorph
  instanceVariableNames: 'faces dieValue isStopped'
  classVariableNames: ''
  poolDictionaries: ''
  category: 'sqByExample-Morphs'
```

The instance variables `faces`, `dieValue`, and `isStopped` define respectively the number of faces the currently displayed face, and whether the die animation is running. To create a die instance, we define the `faces:` method on the *class* side of `DieMorph`:

Method 11.21: *Creating a new die with the number of faces we like*

```
DieMorph class>faces: aNumber
  ↑ self new faces: aNumber
```

The `initialize` method, which `new` calls, is defined on the instance side:

Method 11.22: *Initializing instances of DieMorph*

```
DieMorph>>initialize
  super initialize.
  self extent: 50 @ 50.
  self useGradientFill; borderWidth: 2; useRoundedCorners.
  self setBorderStyle: #complexRaised.
  self fillStyle direction: self extent.
  self color: Color green.
  dieValue := 1.
  faces := 6.
  isStopped := false
```

We use a few methods of BorderedMorph to give a nice appearance to the die: a thick border with a raised effect, rounded corners, and a color gradient on the visible face. We define the faces: instance method as follows:

Method 11.23: *Setting the number of faces of the die*

```
DieMorph»faces: aNumber
  "Set the number of faces"
  (aNumber isInteger
   and: [aNumber > 0]
   and: [aNumber <= 9])
  ifTrue: [faces := aNumber]
```

Our die can have up to nine faces! Be sure to understand the order in which the methods are called when a die is created, for instance with DieMorph faces: 9:

1. the class method DieMorph class»faces: calls DieMorph class»new;
2. new calls the DieMorph»initialize instance method, which sets faces to an initial value of six;
3. DieMorph class»new returns in the class method DieMorph class»faces:, which calls the instance method of the same name, DieMorph»faces:, setting the faces instance variable to nine.

Before defining drawOn:, we need a few methods to place the dots on the displayed face:

Methods 11.24: *Methods for placing points on each face*

```
DieMorph»face1
  ↑{0.5@0.5}
DieMorph»face2
  ↑{0.25@0.25 . 0.75@0.75}
DieMorph»face3
  ↑{0.25@0.25 . 0.75@0.75 . 0.5@0.5}
DieMorph»face4
  ↑{0.25@0.25 . 0.75@0.25 . 0.75@0.75 . 0.25@0.75}
DieMorph»face5
  ↑{0.25@0.25 . 0.75@0.25 . 0.75@0.75 . 0.25@0.75 . 0.5@0.5}
DieMorph»face6
  ↑{0.25@0.25 . 0.75@0.25 . 0.75@0.75 . 0.25@0.75 . 0.25@0.5 . 0.75@0.5}
DieMorph»face7
  ↑{0.25@0.25 . 0.75@0.25 . 0.75@0.75 . 0.25@0.75 . 0.25@0.5 . 0.75@0.5 . 0.5@0.5}
DieMorph »face8
  ↑{0.25@0.25 . 0.75@0.25 . 0.75@0.75 . 0.25@0.75 . 0.25@0.5 . 0.75@0.5 . 0.5@0.5 . 0.5@0.25}
```

DieMorph »face9

```
↑{0.25@0.25 . 0.75@0.25 . 0.75@0.75 . 0.25@0.75 . 0.25@0.5 . 0.75@0.5 . 0.5@0.5 . 0.5@0.25 . 0.5@0.75}
```

These methods define collections of the coordinates of dots for each face. The coordinates are in a square of size 1×1 ; we will simply need to scale them to place the actual dots.

The drawOn: method does two things: it draws the die background with the call to super, then it draws the dots:

Method 11.25: *Drawing the die morph*

DieMorph»drawOn: aCanvas

super drawOn: aCanvas.

(self perform: ('face' , dieValue asString) asSymbol)

do: [:aPoint | self drawDotOn: aCanvas at: aPoint]

The second part of this method uses the reflective capacities of Smalltalk. Drawing the dots of a face is a simple matter of iterating over the collection given by the faceX method for that face, sending the drawDotOn:at: message for each coordinate. To call the correct faceX method, we use the perform: method which sends a message built from a string, here ('face', dieValue asString) asSymbol. You will encounter this use of perform: quite regularly.

Method 11.26: *Drawing a single dot within a face*

DieMorph»drawDotOn: aCanvas at: aPoint

aCanvas


fillOval: (Rectangle

center: self position + (self extent * aPoint)

extent: self extent / 6)

color: Color black

Since the coordinates are normalized to the $[0:1]$ interval, we scale them to the dimensions of our die: self extent * aPoint.

 We can already create a die instance from a workspace:

```
(DieMorph faces: 6) openInWorld.
```

To change the displayed face, we create an accessor that we can use as myDie dieValue: 4:

Method 11.27: *Setting the current value of the die*

DieMorph»dieValue: aNumber


```
(aNumber isInteger
  and: [aNumber > 0]
  and: [aNumber <= faces])
ifTrue:
  [dieValue := aNumber.
  self changed]
```

Now we will use the animation system to show quickly all the faces:

Methods 11.28: *Animating the die*

```
DieMorph>>stepTime
↑ 100
```

```
DieMorph>>step
isStopped ifFalse: [self dieValue: (1 to: faces) atRandom]
```

Now the die is rolling!

To start or stop the animation by clicking, we will reuse what we saw previously about mouse events; first, activate the reception of mouse events:

Methods 11.29: *Handling mouse clicks to start and stop the animation*

```
DieMorph>>handlesMouseDown: anEvent
↑ true
```


```
DieMorph>>mouseDown: anEvent
anEvent redButtonPressed
ifTrue: [isStopped := isStopped not]
```

Now the die will roll or stop rolling when we click on it!

11.9 More About the Canvas

The `drawOn:` method has an instance of `Canvas` as its sole argument; the canvas is the area on which the morph draws itself. By using the graphics methods of the canvas you are free to give the appearance you want to the morph. While browsing the hierarchy of the `Canvas` class, you will see that it has several variants. The default variant of `Canvas` is `FormCanvas`; you will find the key graphics methods in these two classes. These methods can draw points, lines, polygons, rectangles, ellipses, text, and images with rotation and scaling. It is possible to use other canvases, to get transparent morphs, more graphics methods, antialiasing, etc. For this you should use `AlphaBlendingCanvas` or `BalloonCanvas`, but how to use them since `drawOn:`

receives an instance of FormCanvas? In fact you can transform one canvas type into another.

 To use a canvas with a 0.5 alpha-transparency in DieMorph, redefine drawOn: like this:

Method 11.30: Drawing a transparent die

```
DieMorph>>drawOn: aCanvas
| theCanvas |
theCanvas := aCanvas asAlphaBlendingCanvas: 0.5.
super drawOn: theCanvas.
(self perform: ('face' , dieValue asString) asSymbol)
do: [:aPoint | self drawDotOn: theCanvas at: aPoint]
```

That's all you need to do!

If you're curious, have a look at the asAlphaBlendingCanvas: method. You can also get antialiasing by using BalloonCanvas and transforming the die drawing methods as shown in methods 11.31.



Figure 11.7: The die displayed with alpha-transparency.

Methods 11.31: Drawing antialiased die

```
DieMorph>>drawOn: aCanvas
| theCanvas |
theCanvas := aCanvas asBalloonCanvas aaLevel: 3.
super drawOn: aCanvas.
(self perform: ('face' , dieValue asString) asSymbol)
do: [:aPoint | self drawDotOn: theCanvas at: aPoint]
```

```
DieMorph>>drawDotOn: aCanvas at: aPoint
aCanvas
drawOval: (Rectangle
center: self position + (self extent * aPoint)
extent: self extent / 6)
```

```
color: Color black  
borderWidth: 0  
borderColor: Color transparent
```

11.10 Chapter Summary

Morphic is a graphical framework in which graphical interface elements can be dynamically composed.

- You can convert an object into a morph by sending the messages `asMorph` `openInWorld` to it.
- You can manipulate a morph by right-clicking on it and using the handles that appear. (Handles have tooltip popups that explain what they are for.)
- You can compose morphs by dragging and dropping one onto another. Alternatively you can subclass an existing morph class and redefining key methods, like `drawOn:`.
- You can control how a morph reacts to mouse and keyboard events by redefining the methods `handlesMouseOver:`, etc.
- You can animate a morph by defining the methods `step` (what to do) and `stepTime` (the number of milliseconds between steps).
- Various pre-defined morphs, like `PopupMenu` and `FillInTheBlank`, are available to interact with users.

Part III

Advanced Squeak

Chapter 12

Classes and Metaclasses

As we saw in Chapter 5, in Smalltalk, everything is an object, and every object is an instance of a class. Classes are no exception: classes are objects, and class objects are instances of other classes. This object model captures the essence of object-oriented programming: it is lean, simple, elegant and uniform. However, the implications of this uniformity may confuse newcomers. The goal of this chapter is to show that there is nothing complex, “magic” or special here: just simple rules applied uniformly. By following these rules you can always understand why the situation is the way that it is.

12.1 Rules for Classes and Metaclasses

The Smalltalk object model is based on a limited number of concepts applied uniformly. Smalltalk’s designers applied Occam’s razor: any consideration leading to a model more complex than necessary was discarded.

To refresh your memory, here are the rules of the object model that we explored in Chapter 5.

Rule 1. Everything is an object.

Rule 2. Every object is an instance of a class.

Rule 3. Every class has a superclass.

Rule 4. Everything happens by message sends.

Rule 5. Method lookup follows the inheritance chain.

As we mentioned in the introduction to this chapter, a consequence of Rule 1 is that *classes are objects too*, so Rule 2 tells us that classes must also be instances of classes. The class of a class is called a *metaclass*. A metaclass is created automatically for you whenever you create a class. Most of the time you do not need to care or think about metaclasses. However, every time that you use the system browser to browse the “class side” of a class, it is helpful to recall that you are actually browsing a different class. A class and its metaclass are two separate classes, even though the former is an instance of the latter.

To properly explain classes and metaclasses, we need to extend the rules from Chapter 5 with the following additional rules.

Rule 6. Every class is an instance of a metaclass.

Rule 7. The metaclass hierarchy parallels the class hierarchy.

Rule 8. Every metaclass inherits from Class and Behavior.

Rule 9. Every metaclass is an instance of Metaclass.

Rule 10. The metaclass of Metaclass is an instance of Metaclass.

Together, these 10 rules complete Smalltalk’s object model.

We will first briefly revisit the 5 rules from Chapter 5 with a small example. Then we will take a closer look at the new rules, using the same example.

12.2 Revisiting the Smalltalk Object Model

Since everything is an object, the color blue in Smalltalk is also an object.

Color blue \longrightarrow Color blue

Every object is an instance of a class. The class of the color blue is the class Color:

Color blue class \longrightarrow Color

Interestingly, if we set the *alpha* value of a color, we get an instance of a different class, namely TranslucentColor:

(Color blue alpha: 0.4) class \longrightarrow TranslucentColor

We can create a morph and set its color to this translucent color:

EllipseMorph new color: (Color blue alpha: 0.4); openInWorld

You can see the effect in Figure 12.1.

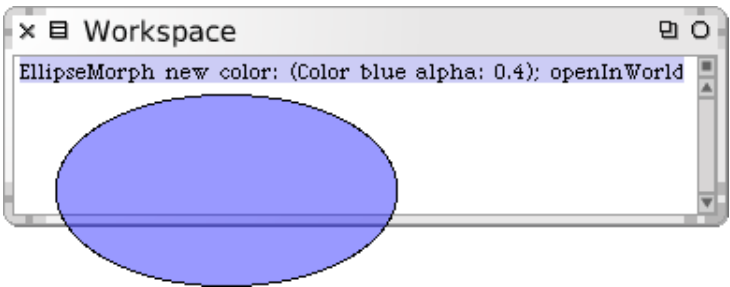


Figure 12.1: A translucent ellipse

By Rule 3, every class has a superclass. The superclass of TranslucentColor is Color, and the superclass of Color is Object:

TranslucentColor superclass	→	Color
Color superclass	→	Object

Everything happens by message sends (Rule 4), so we can deduce that blue is a message to Color, class and alpha: are messages to the color blue, openInWorld is a message to an ellipse morph, and superclass is a message to TranslucentColor and Color. The receiver in each case is an object, since everything is an object, but some of these objects are also classes.

Method lookup follows the inheritance chain (Rule 5), so when we send the message class to the result of Color blue alpha: 0.4, the message is handled when the corresponding method is found in the class Object, as shown in Figure 12.2.

The figure captures the essence of the *is-a* relationship. Our translucent blue object *is a* TranslucentColor instance, but we can also say that it *is a* Color and that it *is an* Object, since it responds to the messages defined in all of these classes. In fact, there is a message, *isKindOf:*, that you can send to any object to find out if it is in an *is a* relationship with a given class:




Figure 12.3: The metaclasses of Translucent its superclasses

The fact that classes are also objects makes it easy for us to query them by sending messages. Let’s have a look:

Color subclasses	→	{TranslucentColor}
TranslucentColor subclasses	→	#()
TranslucentColor allSuperclasses	→	an OrderedCollection(Color Object ProtoObject)
TranslucentColor instVarNames	→	#('alpha')
TranslucentColor allInstVarNames	→	#('rgb' 'cachedDepth' 'cachedBitPattern' 'alpha')
TranslucentColor selectors	→	an IdentitySet(#alpha: #asNontranslucentColor #privateAlpha #pixelValueForDepth: #isOpaque #isTranslucentColor #storeOn: #pixelWordForDepth: #scaledPixelValue32 #alpha #bitPatternForDepth: #hash #convertToCurrentVersion:refStream: #isTransparent #isTranslucent #setRgb:alpha: #balancedPatternForDepth: #storeArrayValuesOn:)

12.4 The Metaclass Hierarchy Parallels the Class Hierarchy

Rule 7 says that the superclass of a metaclass cannot be an arbitrary class: it is constrained to be the metaclass of the superclass of the metaclass’s unique instance.

TranslucentColor class superclass	→	Color class
TranslucentColor superclass class	→	Color class

This is what we mean by the metaclass hierarchy being parallel to the class hierarchy; Figure 12.4 shows how this works in the TranslucentColor hierarchy.




Figure 12.4: The metaclass hierarchy parallels the class hierarchy.

TranslucentColor class	→	TranslucentColor class
TranslucentColor class superclass	→	Color class
TranslucentColor class superclass superclass	→	Object class

Uniformity between Classes and Objects. It is interesting to step back a moment and realize that there is no difference between sending a message to an object and to a class. In both cases the search for the corresponding method starts in the class of the receiver, and proceeds up the inheritance chain.

Thus, messages sent to classes must follow the metaclass inheritance chain. Consider, for example, the method `blue`, which is implemented on the class side of `Color`. If we send the message `blue` to `TranslucentColor`, then it will be looked-up the same way as any other message. The lookup starts in `TranslucentColor` class, and proceeds up the metaclass hierarchy until it is found in `Color` class (see Figure 12.5).

TranslucentColor blue	→	Color blue
-----------------------	---	------------

Note that we get as a result an ordinary `Color blue`, and not a translucent one — there is no magic!

Thus we see that there is one uniform kind of method lookup in Smalltalk. Classes are just objects, and behave like any other objects. Classes have the power to create new instances only because classes happen to respond to the message `new`, and because the method for `new` knows how

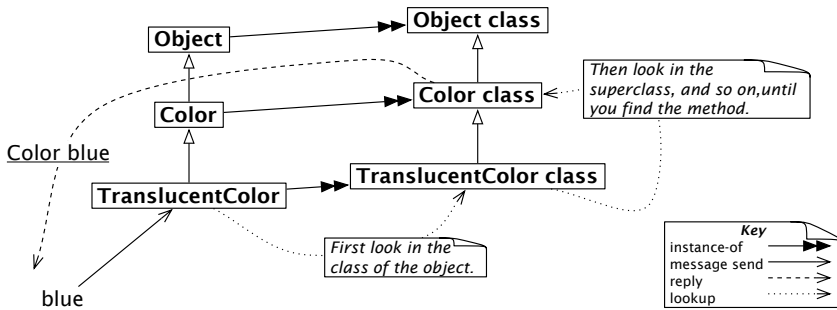


Figure 12.5: Message lookup for classes is the same as for ordinary objects.

to create new instances. Normally, non-class objects do not understand this message, but if you have a good reason to do so, there is nothing stopping you from adding a new method to a non-metaclass.

Since classes are objects, we can also inspect them.

 *Inspect Color blue and Color.*

Notice that in one case you are inspecting an instance of *Color* and in the other case the *Color* class itself. This can be a bit confusing, because the title bar of the inspector names the *class* of the object being inspected.

The inspector on *Color* allows you to see the superclass, instance variables, method dictionary, and so on, of the *Color* class, as shown in Figure 12.6.

12.5 Every Metaclass Inherits from Class and Behavior

Every metaclass *is-a* class, hence inherits from *Class*. *Class* in turn inherits from its superclasses, *ClassDescription* and *Behavior*. Since everything in *Smalltalk* *is-an* object, these classes all inherit eventually from *Object*. We can see the complete picture in Figure 12.7.

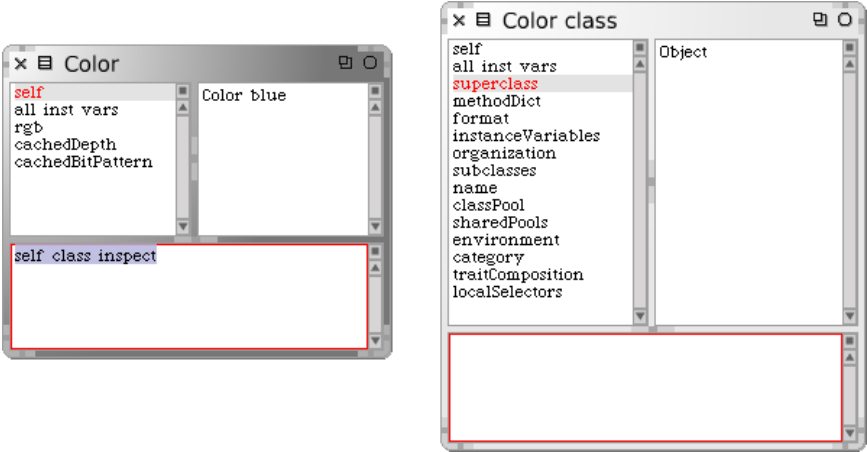


Figure 12.6: Classes are objects too.

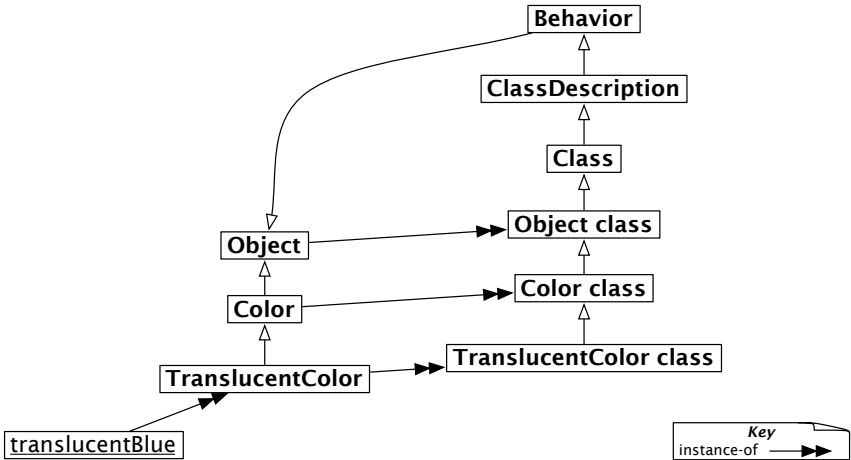


Figure 12.7: Metaclasses inherit from Class and Behavior

Where is new defined? To understand the importance of the fact that metaclasses inherit from Class and Behavior, it helps to ask where `new` is defined and how it is found. When the message `new` is sent to a class it is looked up in its metaclass chain and ultimately in its superclasses Class, ClassDescription and Behavior as shown in Figure 12.8.

The question “Where `new` is defined?” is crucial. `new` is first defined in the class Behavior, and it can be redefined in its subclasses, including any of the metaclass of the classes we define, when this is necessary. Now when a

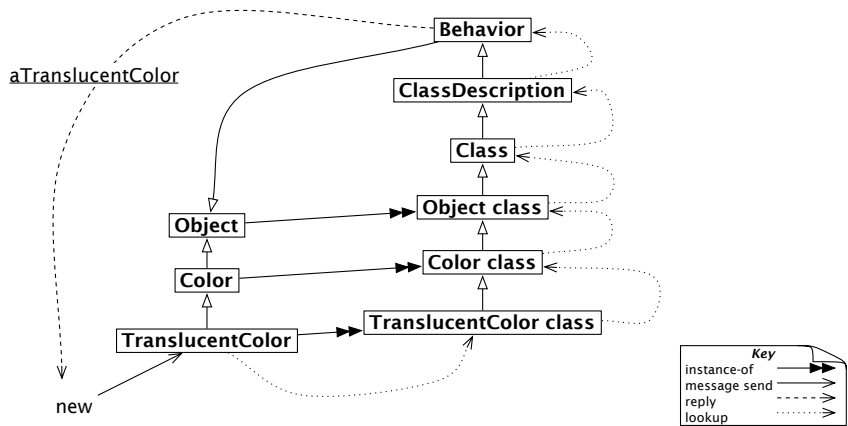


Figure 12.8: `new` is an ordinary message looked up in the metaclass chain.

message `new` is sent to a class it is looked up, as usual, in the metaclass of this class, continuing up the superclass chain right up to the class `Behavior`, if it has not been redefined along the way.

Note that the result of sending `TranslucentColor new` is an instance of `TranslucentColor` and *not* of `Behavior`, even though the method is looked-up in the class `Behavior`! `new` always returns an instance of *self*, the class that receives the message, even if it is implemented in another class.

`TranslucentColor new class` \longrightarrow `TranslucentColor` "*not Behavior*"

A common mistake is to look for `new` in the superclass of the receiving class. The same holds for `new`; the standard message to create an object of a given size. For example, `Array new: 4` creates an array of 4 elements. You will not find this method defined in `Array` or any of its superclasses. Instead you should look in `Array class` and its superclasses, since that is where the lookup will start.

Responsibilities of Behavior, ClassDescription and Class. `Behavior` provides the minimum state necessary for objects that have instances: this includes a superclass link, a method dictionary, and a description of the instances (*i.e.*, representation and number). `Behavior` inherits from `Object`, so it, and all of its subclasses, can behave like objects.

Behavior is also the basic interface to the compiler. It provides methods for creating a method dictionary, compiling methods, creating instances (*i.e.*, `new`, `basicNew`, `new:`, and `basicNew:`), manipulating the class hierarchy (*i.e.*, `superclass:`, `addSubclass:`), accessing methods (*i.e.*, `selectors`, `allSelectors`, `compiledMethodAt:`), accessing instances and variables (*i.e.*, `allInstances`, `instVarNames` ...), accessing the class hierarchy (*i.e.*, `superclass`, `subclasses`) and querying (*i.e.*, `hasMethods`, `includesSelector`, `canUnderstand:`, `inheritsFrom:`, `isVariable`).

`ClassDescription` is an abstract class that provides facilities needed by its two direct subclasses, `Class` and `Metaclass`. `ClassDescription` adds a number of facilities to the basis provided by `Behavior`: named instance variables, the categorization of methods into protocols, the notion of a name (abstract), the maintenance of change sets and the logging of changes, and most of the mechanisms needed for filing-out changes.

`Class` represents the common behavior of all classes. It provides a class name, compilation methods, method storage, and instance variables. It provides a concrete representation for class variable names and shared pool variables (`addClassVarName:`, `addSharedPool:`, `initialize`). `Class` knows how to create instances, so all metaclasses should inherit ultimately from `Class`.

12.6 Every metaclass is an instance of Metaclass

Metaclasses are objects too; they are instances of the class `Metaclass` as shown in Figure 12.9. The instances of class `Metaclass` are the anonymous metaclasses, each of which has exactly one instance, which is a class.

`Metaclass` represents common metaclass behavior. It provides methods for instance creation (`subclassOf:`) creating initialized instances of the metaclass's sole instance, initialization of class variables, metaclass instance, method compilation, and class information (inheritance links, instance variables, etc.).

12.7 The Metaclass of Metaclass is an Instance of Metaclass

The final question to be answered is: what is the class of `Metaclass` class?

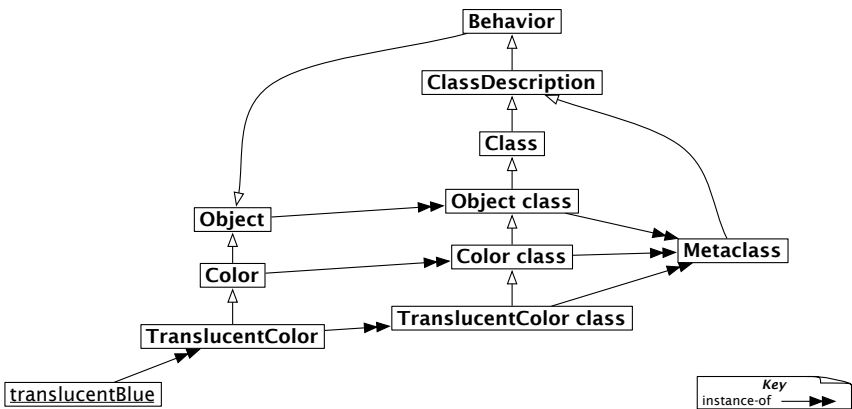


Figure 12.9: Every metaclass is a Metaclass.

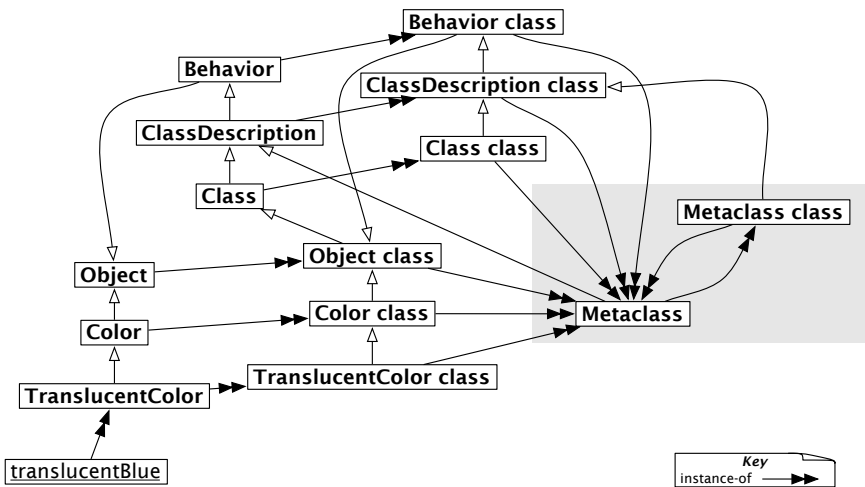


Figure 12.10: All metaclasses are instances of the class Metaclass, even the metaclass of Metaclass.

The answer is simple: it is a metaclass, so it must be an instance of Metaclass, just like all the other metaclasses in the system (see Figure 12.10).

The figure shows how all metaclasses are instances of Metaclass, including the metaclass of Metaclass itself. If you compare Figures 12.9 and 12.10 you will see how the metaclass hierarchy perfectly mirrors the class

hierarchy, all the way up to Object class.

The following examples show us how we can query the class hierarchy to demonstrate that Figure 12.10 is correct. (Actually, you will see that we told a white lie — Object class superclass → ProtoObject class, not Class. In Squeak, we must go one superclass higher to reach Class.)

Example 12.1: *The class hierarchy*

TranslucentColor superclass	→	Color
Color superclass	→	Object

Example 12.2: *The parallel metaclass hierarchy*

TranslucentColor class superclass	→	Color class
Color class superclass	→	Object class
Object class superclass superclass	→	Class "NB: skip ProtoObject class"
Class superclass	→	ClassDescription
ClassDescription superclass	→	Behavior
Behavior superclass	→	Object

Example 12.3: *Instances of Metaclass*

TranslucentColor class class	→	Metaclass
Color class class	→	Metaclass
Object class class	→	Metaclass
Behavior class class	→	Metaclass

Example 12.4: *Metaclass class is a Metaclass*

Metaclass class class	→	Metaclass
Metaclass superclass	→	ClassDescription

12.8 Chapter Summary

Now you should understand better how classes are organized and the impact of a uniform object model. If you get lost or confused, you should always remember that message passing is the key: you look for the method in the class of the receiver. This works on *any* receiver. If the method is not found in the class of the receiver, it is looked up in its superclasses.

- Every class is an instance of a metaclass. Metaclasses are implicit. A Metaclass is created automatically when you create the class that is its sole instance.

- The metaclass hierarchy parallels the class hierarchy. Method lookup for classes parallels method lookup for ordinary objects, and follows the metaclass's superclass chain.
- Every metaclass inherits from Class and Behavior. Every class *is a* Class. Since metaclasses are classes too, they must also inherit from Class. Behavior provides behavior common to all entities that have instances.
- Every metaclass is an instance of Metaclass. ClassDescription provides everything that is common to Class and Metaclass.
- The metaclass of Metaclass is an instance of Metaclass. The *instance-of* relation forms a closed loop, so Metaclass class class \longrightarrow Metaclass.

Part IV

Appendices

Appendix A

Frequently Asked Questions

A.1 Getting started

FAQ 1 *Where do I get the latest Squeak?*

Answer ftp.squeak.org/current_development

FAQ 2 *Where is the Squeak “Development Image”?*

Answer www.squeaksource.com/ImageForDevelopers This is a specially-prepared image including lots of useful packages pre-installed for developers.

A.2 Collections

FAQ 3 *How do I sort an OrderedCollection?*

Answer Send it the message `asSortedCollection`.

`#(7 2 6 1) asSortedCollection → a SortedCollection(1 2 6 7)`

FAQ 4 *How do I convert a collection of characters to a String?*

Answer

String streamContents: [:str | str nextPutAll: 'hello' asSet] → 'hlelo'

A.3 Browsing the system

FAQ 5 *How to I search for a class?*

Answer CMD-b (browse) on the class name, or CMD-f in the category pane of the class browser.

FAQ 6 *How do I find/browse all sends to super?*

Answer The second solution is much faster:

SystemNavigation default browseMethodsWithSourceString: 'super'.
SystemNavigation default browseAllSelect: [:method | method sendsToSuper].

FAQ 7 *How do I browse all super sends within a hierarchy?*

Answer

browseSuperSends := [:aClass | SystemNavigation default
 browseMessageList: (aClass withAllSubclasses gather: [:each |
 (each methodDict associations
 select: [:assoc | assoc value sendsToSuper])
 collect: [:assoc | MethodReference class: each selector: assoc key])
 name: 'Supersends of ', aClass name , ' and its subclasses'].
 browseSuperSends value: OrderedCollection.

FAQ 8 *How do I find out which are the new methods implemented in a class?*

Answer Here we ask which new methods are introduced by True:

newMethods := [:aClass| aClass methodDict keys select:
 [:aMethod | (aClass superclass canUnderstand: aMethod) not]].
newMethods value: True → an IdentitySet(#asBit)

FAQ 9 *How do I tell which methods of a class are abstract?*

Answer

```
abstractMethods :=
  [:aClass | aClass methodDict keys select:
    [:aMethod | (aClass>>aMethod) isAbstract ]].
abstractMethods value: Collection → an IdentitySet(#remove:ifAbsent: #add:
  #do:)
```

FAQ 10 *How do I generate a view of the AST of an expression?*

Answer Load AST from squeaksource.com. Then evaluate:

```
(RBPParser parseExpression: '3+4') explore
```

(Alternatively *explore it*.)

FAQ 11 *How do I find all the Traits in the system?*

Answer

```
Smalltalk allTraits
```

FAQ 12 *How do I find which classes use traits?*

Answer

```
Smalltalk allClasses select: [:each | each hasTraitComposition ]
```

A.4 Using Monticello and SqueakSource

FAQ 13 *How do I load a Squeaksource project?*

Answer

1. Find the project you want in squeaksource.com
2. Copy the registration code snippet
3. Select `open ▷ Monticello browser`
4. Select `+Repository ▷ HTTP`

5. Paste and accept the Registration code snippet; enter your password
6. Select the new repository and **Open** it
7. Select and load the latest version

FAQ 14 *How do I create a SqueakSource project?*

Answer

1. Go to squeaksource.com
2. Register yourself as a new member
3. Register a project (name = category)
4. Copy the Registration code snippet
5. **open** ▷ Monticello browser
6. **+Package** to add the category
7. Select the package
8. **+Repository** ▷ HTTP
9. Paste and accept the Registration code snippet; enter your password
10. **Save** to save the first version

FAQ 15 *How do I extend Number with Number»chf but have Monticello recognize it as being part of my Money project?*

Answer Put it in a method-category named `*Money`. Monticello gathers all methods that are in other categories named like `*package` and includes them in your package.

A.5 Tools

FAQ 16 *How do I programmatically open the SUnit TestRunner?*

Answer Evaluate `TestRunner open`.

FAQ 17 *Where can I find the Refactoring Browser?*

Answer Load AST then Refactoring Engine from squeaksource.com: www.squeaksource.com/AST www.squeaksource.com/RefactoringEngine

FAQ 18 *How do I register the browser that I want to be the default?*

Answer Click the menu icon in the top left of the Browser window.

A.6 Regular Expressions and Parsing

FAQ 19 *How can I work with regular expressions?*

Answer Load Vassili Bykov's RegEx package from: www.squeaksource.com/Regex.html

FAQ 20 *Where is the documentation for the RegEx package?*

Answer Look at the `DOCUMENTATION` protocol of RxParser class in the `VB-Regex` category.

FAQ 21 *Are there tools for writing parsers?*

Answer Use SmaCC —the Smalltalk Compiler Compiler. You should install at least SmaCC-lr.13. Load it from www.squeaksource.com/SmaccDevelopment.html. There is a nice tutorial online: www.refactory.com/Software/SmaCC/Tutorial.html

FAQ 22 *Which packages should I load from SqueakSource SmaccDevelopment to write parsers?*

Answer Load the latest version of SmaCCDev — the runtime is already there. (SmaCC-Development is for Squeak 3.8)

Bibliography

Sherman R. Alpert, Kyle Brown and Bobby Woolf: The Design Patterns Smalltalk Companion. Addison Wesley, 1998, ISBN 0-201-18462-1

Kent Beck: Smalltalk Best Practice Patterns. Prentice-Hall, 1997

Kent Beck: Test Driven Development: By Example. Addison-Wesley, 2003, ISBN 0-321-14653-0

Erich Gamma et al.: Design Patterns: Elements of Reusable Object-Oriented Software. Reading, Mass.: Addison Wesley, 1995, ISBN 0-201-63361-2-(3)

Adele Goldberg and David Robson: Smalltalk 80: the Language and its Implementation. Reading, Mass.: Addison Wesley, May 1983, ISBN 0-201-13688-0

Wilf LaLonde and John Pugh: Inside Smalltalk: Volume 1. Prentice Hall, 1990, ISBN 0-13-468414-1

Alec Sharp: Smalltalk by Example. McGraw-Hill, 1997 (URL: <http://stephane.ducasse.free.fr/FreeBooks/ByExample/>)

Bobby Woolf: Null Object. In **Robert Martin, Dirk Riehle and Frank Buschmann, editors:** Pattern Languages of Program Design 3. Addison Wesley, 1998, 5-18

Index

- **, *see* package, dirty
- :=*, *see* assignment
- ;;*, *see* cascade
- ←*, *see* assignment
- []*, *see* block
- #()*, *see* Array, literal
- #*, *see* literal symbol
- _*, *see* assignment
- .,* *see* statement separator
- ==*, *see* Object, identity
- =*, *see* Object, equality
- >>*, *see* Behavior, *>>*
- { }*, *see* Array, dynamic
- ↑*, *see* return

- abstract class, *see* class, abstract
- abstract method, *see* method, abstract
- accept it, *see* keyboard shortcut, accept
- accessing (protocol), 42, 84, 196
- accessing untypeable characters (protocol), 188
- accessor, 42, 83
- ActiveHand (global), 102
- adding (protocol), 196
- Agile software development, 155
- all (protocol), 32, 39, 112, 123
- AlphaBlendingCanvas (class), 245

- Array
 - (class), 197–200
 - at:, 200, 201
 - at:put:, 200, 201
 - copy, 201
 - dynamic, 53, 177
 - literal, 53, 177, 201

- Array class
 - new:, 200
 - with:, 201
- as yet unclassified (protocol), 39
- assignment, 53, 94
- association, *see* Object, *->*
- AST, 269
- at:, *see* Collection, at:
- at:put:, *see* Collection, at:put:
- attribute, *see* instance variable

- Bag
 - (class), 197, 198, 205
- BalloonCanvas
 - (class), 245, 246
- Beck, Kent, 97, 155
- Behavior
 - >>*, 24, 40
 - (class), 93, 252, 257, 259
 - addSubclass:, 260
 - allInstances, 260
 - allInstVarNames, 255
 - allSelectors, 260
 - allSuperclasses, 255
 - basicNew, 260
 - basicNew:, 260
 - canUnderstand:, 260
 - compiledMethodAt:, 260
 - hasMethods, 260
 - includesSelector, 260
 - inheritsFrom:, 260
 - instVarNames, 255, 260
 - isVariable, 260
 - new, 61, 184, 258, 260
 - new:, 61, 260

- selectors, 255, 260
- subclasses, 260
- superclass, 253, 260
- superclass:, 260
- Behavior Driven Development, *see* Test Driven Development
- binary message, *see* message, binary
- Bitmap
 - (class), 193
- block, 51, 53, 57, 75, 163, 182
- BlockClosure
 - fixTemps, 58
 - value, 57
 - value:, 58
 - value: value:, 58
 - valueWithArguments:, 58
 - whileFalse:, 59
 - whileTrue:, 59
- BlockContext
 - (class), 55, 58
- Blue Book, 193
- blue button, 6, 7, 231
- Boolean
 - &, 192
 - (class), 18, 19, 51, 54, 181–183, 190
 - and:, 192
 - ifFalse:, 59, 191
 - ifFalse:ifTrue, 191
 - ifTrue:, 59, 191
 - ifTrue:ifFalse:, 59, 190
- BorderedMorph
 - (class), 36, 241
 - fullPrintOn:, 99
- Browser
 - (class), 176
- browser, *see* system browser
- browsing programmatically, 124, 268
- Bykov, Vassili, 190, 208, 271
- ByteArray
 - (class), 229
- ByteString
 - (class), 137, 189, 207
- C++, 63, 67, 83, 86, 88
- camelCase, 51, 126
- Canvas
 - (class), 234, 245
 - asAlphaBlendingCanvas:, 246
- caret, *see* return
- cascade, 54, 56, 76, 213
- category, 17
 - creating, 30, 113
 - filing in, *see* file, filing in
 - filing out, *see* file, filing out
- change set, *see* file, filing out
- change set browser, 146
- changes, 4, 9, 151
- Character
 - (class), 21, 53, 91, 181, 185, 188, 206
 - asString, 188
 - asUppercase, 21
 - isAlphaNumeric, 188
 - isCharacter, 188
 - isDigit, 188
 - isLowercase, 188
 - isVowel, 188
 - printOn:, 188
 - printString, 188
- Character class
 - backspace, 188
 - cr, 105, 188
 - escape, 188
 - euro, 188
 - space, 53, 188
 - tab, 53, 188
 - value:, 188
- CharacterArray
 - (class), 197
- CharacterTable (class variable), 189
- Class
 - (class), 252, 257, 260
 - addClassVarName:, 260
 - addSharedPool:, 260
 - initialize, 260
 - subclasses, 255
- class
 - abstract, 90, 183, 190
 - comment, 18, 32
 - creation, 31, 114
 - filing in, *see* file, filing in
 - filing out, *see* file, filing out
 - finding, 18, 122
 - initialization, 104
 - instance variable, 84, 86

- invariant, 184
- method, 84, 86, 89
- recent, 122
- variable, 52, 101, 103, 104
- class browser, *see* system browser
- class category, *see* system category
- ClassDescription
 - (class), 257, 260
- closure, *see* block
- Closure Compiler, 58
- Collection
 - (class), 193
 - add;, 199
 - addAll;, 199
 - asOrderedCollection, 202
 - asSet, 205
 - asSortedCollection, 205, 267
 - at;, 199
 - at:put;, 199
 - collect;, 60, 196, 199, 211
 - comma operator, 25, 207, 223
 - common errors, 213
 - count;, 213
 - detect;, 60
 - detect:ifNone;, 196, 212
 - do;, 60, 168, 196, 199, 209, 211
 - do:separatedBy;, 210
 - do:without;, 210
 - include;, 199
 - includes;, 196, 205, 213
 - inject:into;, 60, 196, 212
 - intersection;, 205
 - isEmpty, 196, 199
 - iteration, 209
 - occurrencesOf;, 196
 - reject;, 60, 196, 212
 - remove;, 183, 199, 214
 - select;, 60, 196, 199, 212
 - size, 199
 - sorting, *see* Collection,
 - asSortedCollection
 - union;, 205
 - weak, 198
- Collection class
 - new;, 199
 - newFrom;, 199
 - with;, 199
 - with:with;, 199
 - withAll;, 199
- Collections-Streams (category), 193
- Collections-Strings (category), 188, 189
- Color
 - (class), 85, 103, 232, 252–254
 - alpha;, 253
 - name, 104
- Color class
 - (class), 254
 - blue, 86, 253, 256
 - colorNames, 105
 - initialize, 105
 - initializeNames, 104
 - showColorCube, 86
- ColorNames (class variable), 104
- comma, *see* Collection, comma operator
- comment, 53
- comparing (protocol), 190
- CompiledMethod
 - (class), 193
- Complex
 - =, 178
 - (class), 185, 186
 - hash, 179
- constant methods, 40
- control constructs, *see* iteration
- converting (protocol), 196
- copy, *see* Object, copy
- CR (global), 105
- creation (protocol), 90, 196
- CrossMorph
 - (class), 234
- CVS, 46
- debug (protocol), 182
- debugger, 24, 43, 111, 135, 182
- declaration, *see* variable declaration
- deep copy, *see* Object, deepCopy
- dependents (protocol), 112
- deprecation, 182
- design by contract, 182
- Dictionary
 - (class), 179, 183, 197, 199, 203
 - associationsDo;, 210
 - at;, 203
 - at:ifAbsent;, 203

- at:put, 203
- do:, 210
- keys, 203
- keysDo:, 210
- overriding = and hash, 203, 214
- removeKey:, 183
- values, 203
- valuesDo:, 210
- Dictionary class
 - newFrom:, 200
 - withAll:, 200
- DieMorph
 - (class), 242
- dirty package, *see* package, dirty
- do:, *see* Collection, do:
- download, 3, 4, 267
- Duration
 - (class), 133, 185, 186
- EllipseMorph
 - defaultColor, 95
- encapsulation boundary, 83
- enumerating (protocol), 196, 210
- enumeration, *see* iteration
- equality, *see* Object, equality
- EventSensor
 - (class), 102
- execution context, 43
- explorer, 15, 134
- exponent, 53
- extension, *see* method, extension
- extension package, *see* package, extension
- eXtreme Programming, 155, 158
- False
 - (class), 54, 190
 - ifTrue:, 191
- false (pseudo variable), 51, 54
- Feathers, Michael, 171
- field, *see* instance variable
- file
 - browsing, *see* file list browser
 - change set, 146
 - filing in, 45, 122
 - filing out, 45, 109, 122, 127, 147
- file list browser, 149
- FileStream
 - (class), 193, 227
 - binary, 229
 - close, 228
 - localName, 228
- FileStream class
 - fileNameed:, 227
 - fileNameed:do:, 219
 - forceNewFileNamed:, 228
 - forceNewFileNamed:do:, 228
 - newFileNamed:, 227
 - oldFileNamed:, 228
 - readOnlyFileNamed:, 228
- FillInTheBlank
 - (class), 239
 - request:initialAnswer:, 239
- fixture, *see* SUnit, fixture
- Float
 - (class), 184, 186
- Float class
 - e, 187
 - infinity, 187
 - nan, 187
 - pi, 187
- FloatArray
 - (class), 197
- floating point number, 53
- fold, *see* Collection»inject:into
- FormCanvas
 - (class), 245
- Fraction
 - (class), 180, 184
 - numerator:denominator:, 187
- Fractions
 - (class), 187
- full stop, *see* statement separator
- geometry (protocol), 233
- getter method, *see* accessor
- global variable, *see* variable, global
- HandMorph
 - (class), 102
- Haskell, 195
- hierarchy browser, 18, 120
- hook method, 184
- identity, *see* Object, identity
- IdentityDictionary

- (class), 204
- image, 4, 5, 9
- inheritance, 89, 95
 - canceling, 183
- inheritance browser, 119
- initialization, 33, 34, 88, 98
- initialization (protocol), 39, 40, 84
- inspector, 14, 34, 83, 133, 257
- instance variable, 34, 83, 94
- instance variable definition, 37
- Integer
 - (class), 184, 187
 - atRandom, 187
 - bitAnd:, 61
 - bitOr:, 61
 - bitShift:, 61
 - factorial, 186, 187
 - gcd:, 187
 - isPrime, 187
 - timesRepeat:, 59, 188
- IntegerArray
 - (class), 197
- Interval
 - (class), 60, 197–199, 202
 - at:, 197
 - printOn:, 177
- Interval class
 - from:to:, 202
 - from:to:by:, 203
 - printString, 203
- is-a, 253, 257
- iteration, 59, *see also* Collection, iteration
- Java, 67, 83, 86, 88, 96, 155
- Kernel-Classes (category), 90
- Kernel-Numbers (category), 184
- Kernel-Objects (category), 17, 18, 175
- keyboard events, 237
- keyboard shortcut, 18, 115
 - accept, 23, 32
 - browse it, 18, 19, 115, 268
 - cancel, 38, 138
 - do it, 13
 - explore it, 15, 135, 138
 - find ..., 19, 268
 - inspect it, 133, 138
 - print it, 13
- keyboard shortcuts, 13, 23, 27
- KeyboardEvent
 - (class), 238
- keys, *see* Dictionary, keys
- keyword message, *see* message, keyword
- Knight, Alan, x
- lambda expression, 196
- LargeNegativeInteger
 - (class), 185, 187
- LargePositiveInteger
 - (class), 185, 187
- launching Squeak, 4
- lexical closure, *see* block
- LF (global), 105
- LinkedList
 - (class), 197, 198
- Lisp, 195
- literal, 53
 - array, 53
 - character, 53
 - number, 53
 - string, 53
 - symbol, 53
- literal arrays, 201
- literal objects, 51
- loops, *see* iteration
- Mac OS X Finder, 112
- Magnitude
 - <, 91, 185, 205
 - <=, 205
 - =, 185, 205
 - >, 185, 205
 - >=, 91, 205
 - (class), 91, 183–185, 188, 190, 205
 - between:and:, 205
- Matrix
 - (class), 38, 41
 - free will, *see* Oracle
- Matrix class
 - new:tabulate:, 38
 - rows:columns:, 38
- message
 - binary, 54, 55, 63
 - evaluation order, 68
 - keyword, 54, 55, 63

- not understood, 100
- receiver, 64
- selector, 64
- send, 64, 93
- unary, 54, 55, 63
- message name finder, 110
- message names browser, 145
- message selector, 54
- message send, 253
- Metaclass
 - (class), 252, 260
- metaclass, 83, 84, 252, 254, 257
 - anonymous, 254
 - hierarchy, 252, 255, 261
 - implicit, 254
- Metaclass class
 - (class), 260
- method
 - abstract, 90
 - byte code, 121
 - categorize, 39, 42
 - creation, 32, 114
 - decompile, 121
 - dictionary, 257
 - extension, 97
 - filing in, *see* file, filing in
 - filing out, *see* file, filing out
 - finding, 20, 123
 - lookup, 95, 256
 - overriding, 97, 119
 - pretty-print, 38, 121
 - public, 84
 - returning self, 44
 - selector, 82
 - value, 45
 - version, 117
- method finder, 20, 110
- method lookup, 98
- MethodContext
 - (class), 55
- ML, 195
- Model
 - myDependents, 112
- Monticello, 29, 46, 109, 110, 123, 125, 127, 148, 269
- Morph
 - (class), 36
 - addMorph:, 233
 - bound, 235
 - center:, 233
 - color:, 233
 - composing, 233
 - defaultColor, 119
 - delete, 234
 - drawOn:, 233, 234
 - extent, 233
 - handleKeystroke:, 237, 238
 - handlesMouseDown:, 236, 237
 - handlesMouseOver:, 237
 - initString, 99
 - mouseDown:, 236
 - mouseEnter:, 237
 - mouseLeave:, 237
 - mouseUp:, 236
 - openInWorld, 96, 231, 232, 235, 253
 - position, 232
 - removeMorph:, 234
 - repelsMorph:event:, 240
 - subclassing, 233, 234
 - wantsDroppedMorph:event:, 240
 - wantsToBeDroppedInto:, 240
- Morphic, 9, 31, 102, 231
 - animation, 238, 245
 - halo, 6, 7, 10, 35, 45, 231
 - step, 238
 - stepTime, 238
- morphic halo, *see* Morhic
- MorphicEvent
 - (class), 236, 238
 - hand, 241
- mouse events, 236
- MouseEvent
 - (class), 237
 - redButtonPressed, 237
 - yellowButtonPressed, 237
- MyTestCase class
 - buildSuiteFromSelectors, 170
- .Net, 155
- new, *see* Behavior»new
- NeXTstep, 111
- nil (pseudo variable), 51, 54
- notifier, 43
- Null Object (pattern), 183

Number

- `*`, 186
- `+`, 186
- `-`, 186
- `/`, 186
- `(class)`, 180, 183–186
- `asFloat`, 186
- `asInteger`, 186
- `ceiling`, 186
- `day`, 186
- `even`, 186
- `floor`, 186
- `fractionPart`, 186
- `hour`, 186
- `i`, 186
- `integerPart`, 186
- `isInfinite`, 186
- `log`, 186
- `negative`, 186
- `odd`, 186
- `positive`, 186
- `printOn:`, 186
- `raiseTo:`, 186
- `sin`, 186
- `sqrt`, 186
- `squared`, 186
- `to:`, 203
- `to:by:`, 203
- `to:do:`, 59
- `week`, 186

Object

- `->`, 203
- `=`, 178
- `(class)`, 15, 17, 31, 90, 100, 175, 253
- `~=`, 179
- `asMorph`, 232
- `assert:`, 182
- `asString`, 209
- `at:`, 61
- `at:put:`, 61
- `class`, 179, 253
- `copy`, 181
- `copyTwoLevel`, 181
- `deepCopy`, 181
- `deprecated:`, 182
- `doesNotUnderstand:`, 101, 182

- `equality`, 178
- `error`, 182
- `error:`, 183
- `halt`, 182
- `hash`, 179
- `identity`, 178
- `initialization`, *see* initialization
- `isArray`, 183
- `isBlock`, 183
- `isBoolean`, 183
- `isCollection`, 183
- `isComplex`, 183
- `isKindOf:`, 180, 253
- `isMemberOf:`, 179
- `isSelfEvaluating`, 177
- `notNil`, 183
- `perform:`, 244
- `postCopy`, 181
- `printOn:`, 176
- `printString`, 18, 176, 209
- `respondsTo:`, 180
- `shallow copy`, 180
- `shallowCopy`, 181
- `shouldNotImplement`, 183
- `storeOn:`, 177
- `subclassResponsibility`, 90, 91, 183
- `yourself`, 213

Object class

- `(class)`, 254

OmniBrowser, 119, 123

- on the fly variable definition, 37

Oracle, 155

OrderedCollection

- `(class)`, 197, 199, 202, 218, 267
- `add:`, 202, 213
- `addAll:`, 202
- `addFirst:`, 197, 202
- `addLast:`, 197, 202
- `anySatisfy:`, 213
- `at:`, 197
- `at:put:`, 197
- `detect:`, 212
- `do:`, 218
- `remove:`, 202
- `remove:ifAbsent:`, 202
- `removeAt:`, 123
- `reverseDo:`, 210

- OrderedCollections
 - (class), 198
- overriding, *see* method, overriding
- package, 15, 29, 46, 123, 125
 - cache, 128, 130
 - dirty, 47
 - extension, 125
- package browser, *see* Monticello
- package cache, 47
- package pane browser, 123
- parentheses, 63, 68, 70
- PasteUpMorph
 - (class), 102
- Pelrine, Joseph, 94, 155
- period, *see* statement separator
- Perl, 155
- Pluggable collections, 197
- PluggableListMorph
 - (class), 135
- Point
 - (class), 34
 - dist:, 83
 - printOn:, 177
- pool dictionary, *see* variable, pool
- PopUpMenu
 - (class), 239
- PositionableStream
 - (class), 217
 - atEnd, 222
 - contents, 221
 - isEmpty, 222
 - peek, 220
 - peekFor:, 220
 - position, 221
 - position:, 221
 - reset, 217, 221
 - setToEnd, 221
 - skip:, 221
 - skipTo:, 221
- PositionableStream class
 - on:, 220
- pre-debugger, 143
- PreDebugWindow
 - (class), 43, 135
- preference browser, 7, 182
- pretty-print, *see* method
- primitive, 60, 94
- primitive., 54
- printing (protocol), 18
- private (protocol), 84
- process
 - browser, 145
 - interrupting, 111, 142
- process browser, 110
- protocol, 18, 39
- ProtoObject
 - ==, 204
 - (class), 90, 100, 175, 179
 - initialize, 184
 - isNil, 183
- ProtoObject class
 - (class), 262
- pseudo-variable, *see* variable, pseudo
- Python, 155
- Quinto, 29
- radix notation, 53
- RBPParser
 - (class), 269
- ReadStream
 - (class), 218, 220
 - next, 220
 - next:, 220
 - upToEnd, 220
- ReadWriteStream
 - (class), 218, 224
- ReceiverMorph
 - (class), 240
- Rectangle
 - (class), 34, 235
 - containsPoint:, 236
- RectangleMorph
 - (class), 120, 240
 - defaultColor, 119
- red button, 6, 7, 26, 135
- refactoring, 32, 126
- Refactoring Browser, 270
- reflection, 83, 156, 176, 244
- regular expression package, 190, 208, 271
- removing (protocol), 196
- required (protocol), 124
- resource, *see* test, resource

- restore display, 86
- return, 45, 54, 57, 94, 96
 - implicit, 57
- saving code, *see* categories
- SBECCell
 - (class), 31
 - initialize, 33
 - mouseAction:, 41
 - mouseUp:, 42
- SBEGame
 - (class), 36
 - cellsPerSide, 40
 - initialize, 36, 44
 - newCellAt:at:, 40, 45
 - toggleNeighboursOfCellAt:at:, 41
- ScaleMorph
 - (class), 120
 - checkExtent:, 115
 - defaultColor, 119
- Self, 231
- self
 - send, 98
- self (pseudo variable), 34, 38, 51, 52, 54, 57, 95
- self-evaluating objects, 177
- Sensor
 - (class), 102
- SequenceableCollection
 - (class), 197
 - doWithIndex:, 210
 - first, 197
 - last, 197
 - readStream, 221
- SequenceableCollection class
 - streamContents:, 223, 224
- Set
 - (class), 197, 199, 203, 204
 - add:, 204
 - intersection, *see* Collection,
 - intersection:
 - membership, *see* Collection,
 - includes:
 - union, *see* Collection, union:
- Set class
 - newFrom:, 205
- setter method, *see* accessor
- shallow copy, *see* Object, shallowCopy
- Sharp, Alex, ix
- shortcut constructor methods, 186, 190
- SimpleSwitchMorph
 - (class), 31
- Singleton (pattern), 190
- Singleton pattern, 88
- SkipList
 - (class), 197
- slot, *see* instance variable
- SmaCC, 271
- SmaCCDev, 271
- SmallInteger
 - +, 61
 - (class), 14, 181, 185, 187
 - maxVal, 187
 - minVal, 187
- Smalltalk (global), 101, 103, 204
- Sokoban, 16
- SortedCollection
 - (class), 197, 199, 205
- SortedCollection class
 - sortBlock:, 206
- SortedCollections
 - (class), 198
- SourceForge, 48
- sources, 4
- Squeak-dev, 4
- SqueakMap, 15
- SqueakSource, 48, 132
- Squeaksource, 269
- Stack
 - pop, 182
- stack trace, 135
- StandardFileStream
 - fullName, 228
- statement, 56
 - separator, 54, 56, 76
- Stream
 - (class), 176, 193, 217
 - nextPut:, 218
 - print:, 223
- String
 - (class), 21, 22, 25, 57, 176, 189, 197, 206, 209, 267
 - anySatisfy:, 208
 - asDate, 190

- asFileName, 190
- asLowercase, 209
- asUppercase, 21, 209
- at:put:, 207
- capitalized, 190, 209
- comma, *see* Collection, comma operator
- concatenation, *see* Collection, comma operator
- copyReplaceAll:, 207
- expandMacros, 208
- expandMacrosWith:, 208
- format:, 208
- includes:, 208
- isEmpty, 208
- lineCount, 57
- match:, 190, 208
- pattern matching, 208
- replaceAll:with:, 207
- replaceFrom:to:with:, 207
- templating, 208
- translateToLowercase, 190
- StringTest
 - (class), 22, 144
 - testShout, 24
- Subversion, 46
- SUnit, 22, 23, 110, 142, 155, 270
 - fixture, 159
 - set up method, 159
- super
 - initialize, 98
 - send, 98, 116, 268
- super (pseudo variable), 51, 54, 95
- superclass, 89, 95
- supersend (protocol), 124
- Symbol
 - (class), 114, 181, 189, 197, 204, 209
- symbol, 31
- syntax, 51
- system browser, 17, 18, 30, 84, 110, 111
 - browse button, 115
 - button bar, 114
 - class side, 84, 86, 88, 105, 252, 256
 - class vars button, 120
 - defining a class, 31, 114
 - defining a method, 32, 114
 - finding classes, *see* class, finding finding methods, *see* method, finding
 - hierarchy button, 116, 120
 - implementors button, 116
 - inheritance button, 119
 - inst vars button, 120
 - instance side, 84, 85, 88
 - refactor button, 122
 - senders button, 115, 117
 - source button, 121
 - versions button, 117
- system category, *see* category
- SystemDictionary
 - (class), 101, 204
- SystemNavigation
 - (class), 268
- SystemNavigation (global), 124
- SystemOrganization (global), 102
- SystemOrganizer
 - (class), 102
- template method, 176
- Test Driven Development, 22, 155
- Test Runner, 110, 160
- TestCase
 - (class), 158, 164
 - assert:, 162, 182
 - assert:description:, 162, 167
 - deny:, 160
 - deny:description:, 162, 167
 - failureLog, 167
 - isLogging, 167
 - run, 168
 - run:, 168, 169
 - runCase, 169
 - setUp, 159, 165, 169
 - should:description:, 167
 - should:raise:, 162
 - shouldnt:description:, 167
 - shouldnt:raise:, 162
 - tearDown, 165, 169
- testing, 22, *see* SUnit
- testing (protocol), 183, 196
- TestResource
 - (class), 164, 166, 170
 - setUp, 170
- TestResource class

- current, 170
- isAvailable, 170
- TestResult
 - (class), 164, 165, 168
 - runCase:, 169
- TestResult class
 - error, 163
- TestRunner, 23, 270
- TestSuite
 - (class), 164, 165
 - run, 170
 - run:, 170
- Text
 - (class), 105
- thisContext (pseudo variable), 51, 54
- three button mouse, 6
- Timespan
 - (class), 185
- TimeStamp
 - (class), 133
- Tools flap, 11, 17, 20, 23, 26, 110, 111, 145, 146, 149, 163
- Trait
 - (class), 92
- trait, 90, 91
- Transcript (global), 56, 101, 110, 238
- transcript, 11
- TranscriptStream
 - (class), 101
- TranslucentColor
 - (class), 104, 252, 253, 255, 259
- True
 - (class), 54, 190
 - ifTrue:, 191
 - not, 191
- true (pseudo variable), 51, 54, 192
- TTCFont
 - printOn:, 176
- unary message, *see* message, unary
- Undeclared (global), 102
- UndefinedObject
 - (class), 55, 137, 181
- value, *see* BlockClosure
- values, *see* Dictionary, values
- variable
 - class, *see* class, variable
 - class instance, *see* class, instance
 - variable
 - declaration, 53, 58, 94
 - global, 52, 101
 - instance, *see* instance variable
 - pool, 52, 101, 105
 - pseudo, 52, 54, 96
 - shared, 101
- versions browser, 118
- virtual machine, 3, 10, 54, 60, 94, 95, 100
- weak collections, 198
- WebServer
 - (class), 88
- WideString
 - (class), 207
- workspace, 11, 110
- World (global), 102
- world menu, 5, 7
- WriteStream
 - (class), 218, 222
 - cr, 223
 - ensureASpace, 223
 - nextPut:, 223
 - nextPutAll:, 223
 - space, 223
 - tab, 223
- xUnit, 155
- yellow button, 6, 7, 32, 39, 113, 116, 118, 120, 122, 131–134, 145–147, 150, 152

